

Computer Operating Systems, Practice Session 5

Semaphore Operations in Unix

Mustafa Ersen (ersenm@itu.edu.tr)

Istanbul Technical University
34469 Maslak, Istanbul

12 March 2014

Today

Computer Operating Systems, PS 5

Semaphore Operations

Signal Mechanism in Linux

Examples

Semaphore Creation

- ▶ Header files in Unix to be used in semaphore operations:

- ▶ `sys/ipc.h`
- ▶ `sys/sem.h`
- ▶ `sys/types.h`

- ▶ Semaphore Creation:

```
int semget(key_t key, int nsems, int semflg);
```

If successful, a nonnegative integer is returned as the semaphore set identifier, otherwise -1 is returned, with `errno` of the error.

`semflg`: `IPC_CREAT|0700` (Last 9 bits: permission flags)

A semaphore set including `nsems` semaphores is created and associated with `key`:

- ▶ `IPC_PRIVATE`
- ▶ `IPC_CREAT` & (No semaphore set exists associated with `key` value)

Semaphore Operations

- ▶ `int semop(int semid, struct sembuf *sops, unsigned nsops);`
 - ▶ `semop` operates on semaphores selected from semaphore set associated with `semid`
 - ▶ Each of the `nsops` elements, pointed by `sops`, determines operation on a specific semaphore (each element is of type: `sembuf`)
- ▶ `struct sembuf{`
 - `unsigned short sem_num; // semaphore number starts with 0`
 - `short sem_op; // semaphore operation`
 - `short sem_flg; // operation flags``};`
- ▶ The operations contained in `sops` are performed in array order **atomically** (i.e., the operations are performed either as a complete unit, or not at all)
- ▶ `sem_flg`
 - ▶ `SEM_UNDO`: Allows individual operations in the array to be automatically undone when the process exits.
 - ▶ `IPC_NOWAIT`: (Do not allow to wait) If you can not decrease, give error message and return
- ▶ `sem_op`
 - ▶ `== 0`: wait for it to be 0 (Must have read permission)
 - ▶ `!= 0`: value is added to the semaphore value (The process must have alter permission on the semaphore set)

Semaphore Control

- ▶ Control of the Value

```
int semctl(int semid, int semnum, int cmd, arg);
```

- ▶ cmd

- ▶ IPC_RMID : Remove the semaphore set, awakening all processes blocked
- GETVAL : Return the value of `semval` for the corresponding semaphore
- SETVAL : Set the value of `semval` of the corresponding semaphore to `arg.val`
- SETALL : Set `semval` values for all semaphores of the set using `arg.array`
- GETALL : Return all of the `semval` values for all semaphores of the set into `arg.array`

Basic Semaphore Operations: Increment

```
void sem_signal(int semid, int val)
{
    struct sembuf semaphore;
    semaphore.sem_num=0;
    semaphore.sem_op=val;
    semaphore.sem_flg=1; // relative: add sem_op to value
    semop(semid, &semaphore,1);
}
```

Basic Semaphore Operations: Decrement

```
void sem_wait(int semid, int val)
{
    struct sembuf semaphore;
    semaphore.sem_num=0;
    semaphore.sem_op=(-1*val);
    semaphore.sem_flg=1; // relative: add sem_op to value
    semop(semid, &semaphore,1);
}
```

Handling Signals

- Necessary header files for handling signals:
 - signal.h
 - sys/types.h

```
// signal-handling function
void mysignal(int signum){
    printf("Received signal with num=%d\n", signum);
}

void mysigset(int num){
    struct sigaction mysigaction;
    mysigaction.sa_handler=(void *)mysignal;
    // using the signal-catching function identified by sa_handler
    mysigaction.sa_flags=0;
    // sigaction() system call is used to change the action taken by a
    // process on receipt of a specific signal (specified with num)
    sigaction(num,&mysigaction,NULL);
}
```


Handling Signals

- ▶ Sending a signal (specified with `num=sig`) from a process to another process (with given `pid`):

```
int kill(pid_t pid, int sig);
```

- ▶ Waiting for a signal:

```
int pause(void);
```

Example 1

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <sys/wait.h>
5 #include <sys/ipc.h>
6 #include <sys/sem.h>
7 #include <sys/types.h>
8 #include <signal.h> // sigaction
9
10 #define SEMKEY 8
11 int sem_id;
12
13 // increment operation
14 void sem_signal(int semid, int val){
15     struct sembuf semaphore;
16     semaphore.sem_num=0;
17     semaphore.sem_op=val;
18     semaphore.sem_flg=1; // relative: add sem_op to value
19     semop(semid, &semaphore, 1);
20 }
```

Example 1

```
22 // decrement operation
23 void sem_wait(int semid, int val){
24     struct sembuf semaphore;
25     semaphore.sem_num=0;
26     semaphore.sem_op=(-1*val);
27     semaphore.sem_flg=1; // relative: add sem_op to value
28     semop(semid, &semaphore, 1);
29 }
30
31 // signal-handling function
32 void mysignal(int signum){
33     printf("Received signal with num=%d\n", signum);
34 }
35 void mysigset(int num){
36     struct sigaction mysigaction;
37     mysigaction.sa_handler=(void *)mysignal;
38     // using the signal-catching function identified by sa_handler
39     mysigaction.sa_flags=0;
40     // sigaction() system call is used to change the action taken by a
41     // process on receipt of a specific signal (specified with num)
42     sigaction(num,&mysigaction,NULL);
43 }
```

Example 1

```
45 int main(void){  
46     // signal handler with num=12  
47     mysigset(12);  
48     int f=1, i, children[10];  
49     // creating 10 child processes  
50     for(i=0; i<10; i++){  
51         if (f>0)  
52             f=fork();  
53         if (f==-1){  
54             printf("fork error...\n");  
55             exit(1);  
56         }  
57         if (f==0)  
58             break;  
59         else  
60             children[i]=f; // get pid of each child process  
61     }
```

Example 1

```
62 // parent process
63 if(f>0){
64     // creating a semaphore with key=SEMKEY
65     sem_id = semget(SEMKEY, 1, 0700|IPC_CREAT);
66     // setting value of the 0th semaphore of the set identified with sem_id to 0
67     semctl(sem_id, 0, SETVAL, 0);
68     // waiting for a second
69     sleep(1);
70     // sending the signal 12 to all child processes
71     for (i=0; i<10; i++)
72         kill(children[i], 12);
73     // decrease semaphore value by 10 (i.e., wait for all childs to increase semaphore value)
74     sem_wait(sem_id, 10);
75     printf("ALL CHILDREN HAS Finished ...\n");
76     // remove the semaphore set identified with sem_id
77     semctl(sem_id, 0, IPC_RMID, 0);
78     exit(0);
79 }
```

Example 1

```
80 // child process
81 else{
82     // wait for a signal
83     pause();
84     // returning the sem_id associated with SEMKEY
85     sem_id = semget(SEMKEY, 1, 0);
86     printf("I am the CHILD Process created in %d th order. My PROCESS ID: %d\n", i, getpid());
87     // getting value of the 0th semaphore of the set identified with sem_id
88     printf("SEMAPHORE VALUE: %d\n", semctl(sem_id, 0, GETVAL, 0));
89     // increase semaphore value by 1
90     sem_signal(sem_id, 1);
91 }
92
93 return 0;
94 }
```

Output of Example 1

Received signal with num=12

I am the CHILD Process created in 5 th order. My PROCESS ID: 2367

SEMAPHORE VALUE: 0

Received signal with num=12

I am the CHILD Process created in 2 th order. My PROCESS ID: 2364

SEMAPHORE VALUE: 1

Received signal with num=12

I am the CHILD Process created in 3 th order. My PROCESS ID: 2365

SEMAPHORE VALUE: 2

Received signal with num=12

I am the CHILD Process created in 1 th order. My PROCESS ID: 2363

SEMAPHORE VALUE: 3

Received signal with num=12

Received signal with num=12

Received signal with num=12

Output of Example 1 (Continues)

```
I am the CHILD Process created in 0 th order. My PROCESS ID: 2362
I am the CHILD Process created in 8 th order. My PROCESS ID: 2370
SEMAPHORE VALUE: 4
Received signal with num=12
I am the CHILD Process created in 7 th order. My PROCESS ID: 2369
SEMAPHORE VALUE: 4
SEMAPHORE VALUE: 6
I am the CHILD Process created in 9 th order. My PROCESS ID: 2371
SEMAPHORE VALUE: 6
Received signal with num=12
Received signal with num=12
I am the CHILD Process created in 4 th order. My PROCESS ID: 2366
SEMAPHORE VALUE: 8
I am the CHILD Process created in 6 th order. My PROCESS ID: 2368
SEMAPHORE VALUE: 9
ALL CHILDREN HAS Finished ...
```


Example 2 - Deadlock

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <sys/wait.h>
5 #include <sys/ipc.h>
6 #include <sys/sem.h>
7 #include <sys/types.h>
8 #include <signal.h>
9
10 #define SEMKEY_A 1
11 #define SEMKEY_B 2
12 #define SEMKEY_C 3
13
14 // increment operation
15 void sem_signal(int semid, int val){
16     struct sembuf semaphore;
17     semaphore.sem_num=0;
18     semaphore.sem_op=val;
19     semaphore.sem_flg=1; // relative: add sem_op to value
20     semop(semid, &semaphore, 1);
21 }
```

Example 2 - Deadlock

```
23 // decrement operation
24 void sem_wait(int semid, int val){
25     struct sembuf semaphore;
26     semaphore.sem_num=0;
27     semaphore.sem_op=(-1*val);
28     semaphore.sem_flg=1; // relative: add sem_op to value
29     semop(semid, &semaphore, 1);
30 }
31
32 // signal-handling function
33 void mysignal(int signum){
34     printf("Received signal with num=%d\n", signum);
35 }
36
37 void mysigset(int num){
38     struct sigaction mysigaction;
39     mysigaction.sa_handler=(void *)mysignal;
40     // using the signal-catching function identified by sa_handler
41     mysigaction.sa_flags=0;
42     // sigaction() system call is used to change the action taken by a
43     // process on receipt of a specific signal (specified with num)
44     sigaction(num,&mysigaction,NULL);
45 }
```

Example 2 - Deadlock

```
47 int main(void){  
48     // signal handler with num=12  
49     mysigset(12);  
50     int semA,semB,semC,c[2],f=1,i,myOrder;  
51     // creating 2 child processes  
52     for(i=0; i<2; i++){  
53         if (f>0)  
54             f=fork();  
55         if (f==-1){  
56             printf("fork error....\n");  
57             exit(1);  
58         }  
59         if (f==0)  
60             break;  
61         else  
62             c[i]=f; // get pid of each child process  
63     }
```

Example 2 - Deadlock

```
64 // parent process
65 if (f!=0){
66     printf("PARENT is starting to CREATE RESOURCES...\n");
67     // creating 3 semaphores and setting two of them as 1 and the other as 0
68     semA=semget(SEMKEY_A,1,0700|IPC_CREAT);
69     semctl(semA, 0, SETVAL, 1);
70     semB=semget(SEMKEY_B,1,0700|IPC_CREAT);
71     semctl(semB, 0, SETVAL, 1);
72     semC=semget(SEMKEY_C,1,0700|IPC_CREAT);
73     semctl(semC, 0, SETVAL, 0);
74     sleep(2);
75     printf("PARENT is starting CHILD Processes ..... \n");
76     // sending the signal 12 to all child processes
77     for (i=0; i<2; i++)
78         kill(c[i],12);
79     // decrease semaphore value by 2 (i.e., wait for all children)
80     sem_wait(semC,2);
81     printf("PARENT: Child processes has done, resources are removed back...\n");
82     // remove the created semaphore sets
83     semctl(semC,0,IPC_RMID,0);
84     semctl(semA,0,IPC_RMID,0);
85     semctl(semB,0,IPC_RMID,0);
86     exit(0);
87 }
```

Example 2 - Deadlock

```
88 // child process
89 else{
90     myOrder=i;
91     printf("CHILD %d: waiting permission from PARENT ....\n", myOrder);
92     // wait for a signal
93     pause();
94     // returning the sem_ids associated with SEMKEY_A, SEMKEY_B and SEMKEY_C
95     semA=semget(SEMKEY_A,1,0);
96     semB=semget(SEMKEY_B,1,0);
97     semC=semget(SEMKEY_C,1,0);
98     printf("CHILD %d has permission from PARENT, is starting ....\n", myOrder);
99     if (myOrder==0){
100         printf("CHILD %d: DECREASING sem A.\n", myOrder);
101         sem_wait(semA, 1);
102         sleep(1);
103         printf("CHILD %d: sem A is completed, DECREASING sem B.\n", myOrder);
104         sem_wait(semB, 1);
105         printf("CHILD %d: I am in the CRITICAL REGION.\n", myOrder);
106         sleep(5); /* Critical Region Operations */
107         // increase all the semaphore values by 1
108         sem_signal(semB, 1);
109         sem_signal(semA, 1);
110         sem_signal(semC, 1);
111     }
```

Example 2 - Deadlock

```
112         else if (myOrder==1){
113             printf("CHILD %d: DECREASING sem B.\n", myOrder);
114             sem_wait(semB, 1);
115             sleep(1);
116             printf("CHILD %d: sem B is completed, DECREASING sem A.\n", myOrder);
117             sem_wait(semA, 1);
118             printf("CHILD %d: I am in the CRITICAL REGION.\n", myOrder);
119             sleep(5); /* Critical Region Operations */
120             // increase all the semaphore values by 1
121             sem_signal(semA,1);
122             sem_signal(semB,1);
123             sem_signal(semC,1);
124         }
125     }
126     return 0;
127 }
```

Output of Example 2

```
PARENT is starting to CREATE RESOURCES....  
CHILD 1: waiting permission from PARENT ....  
CHILD 0: waiting permission from PARENT ....  
PARENT is starting CHILD Processes .....  
Received signal with num=12  
CHILD 1 has permission from PARENT, is starting ....  
CHILD 1: DECREASING sem B.  
Received signal with num=12  
CHILD 0 has permission from PARENT, is starting ....  
CHILD 0: DECREASING sem A.  
CHILD 1: sem B is completed, DECREASING sem A.  
CHILD 0: sem A is completed, DECREASING sem B.
```

Example 3 - Preventing Deadlock

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <sys/wait.h>
5 #include <sys/ipc.h>
6 #include <sys/sem.h>
7 #include <sys/types.h>
8 #include <signal.h>
9 #include <sys/errno.h>
10
11 #define SEMKEY_AB 5
12 #define SEMKEY_C 6
```


Example 3 - Preventing Deadlock

```
14 // increment operation
15 void sem_signal(int semid, int val){
16     struct sembuf semaphore;
17     semaphore.sem_num=0;
18     semaphore.sem_op=val;
19     semaphore.sem_flg=1; // relative: add sem_op to value
20     semop(semid, &semaphore, 1);
21 }
22
23 // increment operation using two semaphores
24 void sem_multi_signal(int semid, int val, int nsems){
25     struct sembuf semaphore[2];
26     int i;
27     for (i=0; i<nsems; i++){
28         semaphore[i].sem_num=i;
29         semaphore[i].sem_op=val;
30         semaphore[i].sem_flg=1;
31     }
32     // TWO Operations are performed on SAME SEMAPHORE SET
33     semop(semid, semaphore, 2);
34     for (i=0; i<nsems; i++){
35         printf("SIGNAL : SEM %d IS NOW: .... %d\n", i, semctl(semid,i,GETVAL,0));
36     }
37 }
```

Example 3 - Preventing Deadlock

```
39 // decrement operation
40 void sem_wait(int semid, int val){
41     struct sembuf semaphore;
42     semaphore.sem_num=0;
43     semaphore.sem_op=(-1*val);
44     semaphore.sem_flg=1; // relative: add sem_op to value
45     semop(semid, &semaphore, 1);
46 }
47
48 // decrement operation using two semaphores
49 void sem_multi_wait(int semid, int val, int nsems){
50     struct sembuf semaphore[2];
51     int i;
52     for (i=0; i<nsems; i++){
53         semaphore[i].sem_num=i;
54         semaphore[i].sem_op=(-1*val);
55         semaphore[i].sem_flg=1;
56     }
57     //TWO Operations are performed on SAME SEMAPHORE SET:
58     semop(semid, semaphore, 2);
59     for (i=0; i<nsems; i++){
60         printf("WAIT : SEM %d is NOW .... %d\n", i, semctl(semid,i,GETVAL,0));
61     }
62 }
```

Example 3 - Preventing Deadlock

```
65 void mysignal(int signal){ printf("Received signal with num=%d\n", signal);}  
66 void mysigset(int num){  
67     struct sigaction mysigaction;  
68     mysigaction.sa_handler=(void *)mysignal;  
69     // using the signal-catching function identified by sa_handler  
70     mysigaction.sa_flags=0;  
71     // sigaction() system call is used to change the action taken by a  
72     // process on receipt of a specific signal (specified with num)  
73     sigaction(num,&mysigaction,NULL);  
74 }  
75  
76 int main(void){  
77     // signal handler with num=12  
78     mysigset(12);  
79     int semAB,semC,c[2],f=1,i,myOrder;  
80     // creating 2 child processes  
81     for(i=0; i<2; i++){  
82         if (f>0)  
83             f=fork();  
84         if (f==-1){  
85             printf("fork error....\n");  
86             exit(1);  
87         }  
88         if (f==0)  
89             break;  
90         else  
91             c[i]=f; // get pid of each child process  
92     }  
93 }
```

Example 3 - Preventing Deadlock

```

96 // parent process
97 if (f!=0){
98     printf("PARENT is starting to CREATE RESOURCES....\n");
99     // creating a set of 2 semaphores and setting their values as 1
100    semAB=semget(SEMKEY_AB, 2, 0700|IPC_CREAT);
101    if(semAB == -1)
102        printf("SEMGET ERROR on SEM SET, Error Code: %d \n", errno);
103    if (semctl(semAB, 0, SETVAL, 1) == -1)
104        printf("SMCTL ERROR on SEM A, Error Code: %d \n", errno);
105    if (semctl(semAB, 1, SETVAL, 1) == -1)
106        printf("SMCTL ERROR on SEM B, Error Code: %d \n", errno);
107    printf("PARENT: SEM A is NOW .... %d\n", semctl(semAB,0,GETVAL,0));
108    printf("PARENT: SEM B is NOW .... %d\n", semctl(semAB,1,GETVAL,0));
109    //creating another semaphore and setting its value as 0
110    semC=semget(SEMKEY_C,1,0700|IPC_CREAT);
111    semctl(semC, 0, SETVAL, 0);
112    printf("PARENT: SEM C is NOW .... %d\n", semctl(semC,0,GETVAL,0));
113    sleep(2);
114    printf("PARENT is starting CHILD Processes ..... \n");
115    for (i=0; i<2; i++)
116        kill(c[i],12);
117    sleep(5);
118    // decrease semaphore value by 2 (i.e., wait for all children)
119    sem_wait(semC,2);
120    printf("PARENT: SEM C is NOW .... %d\n", semctl(semC,0,GETVAL,0));
121    printf("PARENT: Child processes has done, resources are removed back...\n");
122    semctl(semC,0,IPC_RMID,0);
123    semctl(semAB,0,IPC_RMID,0);
124    exit(0);
125 }

```

Example 3 - Preventing Deadlock

```
126 // child process
127 else{
128     myOrder=i;
129     printf("CHILD %d: waiting permission from PARENT ....\n", myOrder);
130     // wait for a signal
131     pause();
132     // returning the sem_ids associated with SEMKEY_AB and SEMKEY_C
133     semAB=semget(SEMKEY_AB,2,0);
134     semC=semget(SEMKEY_C,1,0);
135     printf("CHILD %d has permission from PARENT, is starting ....\n", myOrder);
136     printf("CHILD %d: DECREASING sem AB.\n", myOrder);
137     // decrease two semaphores in the set specified by semAB by 1
138     sem_multi_wait(semAB,1,2);
139     printf("CHILD %d: I am in the CRITICAL REGION.\n", myOrder);
140     sleep(5);
141     // increase two semaphores in the set specified by semAB by 1
142     sem_multi_signal(semAB,1,2);
143     // increase the third semaphore by 1
144     sem_signal(semC,1);
145 }
146 return 0;
147 }
```

Output of Example 3

```
PARENT is starting to CREATE RESOURCES....
PARENT: SEM A is NOW .... 1
PARENT: SEM B is NOW .... 1
PARENT: SEM C is NOW .... 0
CHILD 1: waiting permission from PARENT ....
CHILD 0: waiting permission from PARENT ....
PARENT is starting CHILD Processes .....
Received signal with num=12
CHILD 1 has permission from PARENT, is starting ....
CHILD 1: DECREASING sem AB.
WAIT : SEM 0 is NOW .... 0
WAIT : SEM 1 is NOW .... 0
CHILD 1: I am in the CRITICAL REGION.
Received signal with num=12
CHILD 0 has permission from PARENT, is starting ....
CHILD 0: DECREASING sem AB.
SIGNAL : SEM 0 IS NOW: .... 0
SIGNAL : SEM 1 IS NOW: .... 0
WAIT : SEM 0 is NOW .... 0
WAIT : SEM 1 is NOW .... 0
CHILD 0: I am in the CRITICAL REGION.
SIGNAL : SEM 0 IS NOW: .... 1
SIGNAL : SEM 1 IS NOW: .... 1
PARENT: SEM C is NOW .... 0
PARENT: Child processes has done, resources are removed back...
```