

Interprocess Communication

Types of Interaction

- between concurrent processes
 - resource sharing
 - communication
 - synchronization

Levels of Interaction

- interaction between processes on three levels
 - processes not aware of each other (competing)
 - using system resources (moderated by operating system)
 - processes indirectly aware of each other (sharing)
 - resource sharing (through mutual exclusion and synchronization)
 - processes directly aware of each other (communicating)

Resource Sharing


- mutual exclusion
 - two types of resources
 - can be used by more than one process at a time (e.g. reading from a file)
 - can be used by only one process at a time
 - due to physical constraints (e.g. some I/O units)
 - if the actions of one process interfered with those of another (e.g. writing to a shared memory location)
- synchronization
 - a process needs to proceed after another process completes some actions

Example

- 2 processes: Observer and Reporter
- `counter` shared variable

```
observer:          reporter:
while TRUE {      while TRUE {
  observe;        print_counter;
  counter ++;     counter=0;
}
```

Example – Possible Errors



Example – Possible Errors

```
counter++  LOAD    ACC, COUNTER
           INC     ACC
           SAVE    COUNTER, ACC
```

Race:

- when processes access a shared variable
 - outcome depends on order and running speed of processes
 - may be different for different runs

Example – Possible Errors

P1: `k=0 (initial value)`
`while TRUE`
`k=k+1;`
 what about the values of k depending on the order of P1 and P2 executions?

P2: `while TRUE`
`k=k+1;`
SOLUTION: mutual exclusion

Sharing

- two types of sharing:
 - READ (no need for mutual exclusion)
 - WRITE (mutual exclusion needed)
- for consistency
 - mutual exclusion
 - synchronization

Synchronization

- programs should not be dependent on running order of processes
- programs working together may need to be synchronized at some points
 - e.g. a program uses output calculated by another program

Mutual Exclusion

critical section (CS): Part of code in a process in which operations on shared resources are performed.

mutual exclusion: only one process can execute a CS for a resource at a time

Example


```
P1:
while TRUE {
  <non-CS>
  mx_begin
  <CS ops>
  mx_end
  <non-CS>
}

P2:
while TRUE {
  <non-CS>
  mx_begin
  <CS ops>
  mx_end
  <non-CS>
}
```

Mutual Exclusion – Possible Problems

- deadlock
 - more than one process requires the same resources
 - each process does not release the resource required by the other

Example: 3 processes and 3 resources



$P1()$ $P2()$ $P3()$
 $\text{req}(R1);$ $\text{req}(R1); \text{req}(R2); \text{req}(R3);$
 $\text{req}(R2); \text{req}(R3); \text{req}(R1);$

Mutual Exclusion

- `mx_begin`
 - are there any processes in their CS which have not yet executed `mx_end`?
 - if NOT
 - allow process to proceed into CS
 - leave mark for other processes
- `mx_end`
 - allow any process waiting to go into CS to proceed
 - if not leave mark (empty)

Mutual Exclusion Implementation

- only one process may be in its CS
- if a process wants to enter its CS and if there are no others executing their CS, it shouldn't wait
- any process not executing its CS should not prevent another process from entering its own CS
- no assumptions should be made about the order and speed of execution of processes
- no process should stay in its CS indefinitely
- no process should wait to enter its own CS indefinitely

Mutual Exclusion Solutions

- software based solutions
- hardware based solutions
- software and hardware based solutions

A Software Based Solution

- use a flag that shows whether a process is in its CS or not: *busy*
 - `busy ← TRUE` : process in CS
 - `busy ← FALSE` : no process in CS
- `mx_begin`: `while (busy);`
 `busy = TRUE;`
 – wait until process in CS is finished
 – enter CS
- `mx_end`: `busy = FALSE;`

A Software Based Solution

- a possible error
 - *busy* is also a shared variable!
 - Example:
 - P1 checks and finds `busy=FALSE`
 - P1 interrupted
 - P2 checks and finds `busy=FALSE`
 - both P1 and P2 enter CS

Solutions Requiring *Busy Waiting*

```

global variable turn = 1;

Process 1:      Process 2:
local variables  local variables
my_turn=1;       my_turn=2;
others_turn=2;   others_turn=1;

mx_begin: while (turn != my_turn);
mx_end  : turn = others_turn;

```

Solutions Requiring *Busy Waiting*

- use up CPU time
- works properly but has limitations:
 - processes enter their CS in turn
 - depends on speed of process execution
 - depends on number of processes

Solutions Requiring *Busy Waiting*

- first correct solution: Dekker algorithm
- Peterson algorithm (1981)
 - similar approach
 - simpler

Peterson Algorithm

- shared variables:


```

req_1, req_2: bool and initialized to FALSE
turn: integer and initialized to "P1" or "P2"

```
- P1:**

```

mx_begin:
  req_1 = TRUE;
  turn = P2;
  while (req_2 && turn==P2);
    < CS >

mx_end: req_1 = FALSE;

```

Peterson Algorithm

- different scenarios:
 - P1 is active, P2 is passive


```

req_1=TRUE and turn=P2
req_2=FALSE so P1 proceeds after while loop

```
 - P1 in CS, P2 wants to enter CS


```

req_2=TRUE and turn=P1;
req_1=TRUE so P2 waits in while loop
P2 continues after P1 executes  mx_end

```

Peterson Algorithm

- (*different scenarios cntd.*):
 - P1 and P2 want to enter CS at the same time


```

P1:      P2:
req_1=TRUE;  req_2=TRUE;
turn=P2;     turn=P1;

```
- ⇒ order depends on which process assigns value to the turn variable first.

Hardware Based Solutions

- with uninterruptable machine code instructions completed in one machine cycle
 - e.g.: `test_and_set`
 - busy waiting used
 - when a process exits CS, no mechanism to determine which other process enters next
 - indefinite waiting possible
- disabling interrupts
 - interferes with scheduling algorithm of operating system

Hardware Based Solutions

- `test_and_set(a):`

```
cc ← a
a ← TRUE
```

 - with one machine instruction, contents of "a" copied into condition code register and "a" is assigned TRUE
- ```
mx_begin: test_and_set(busy);
 while (cc) {
 test_and_set(busy);
 }
mx_end: busy=FALSE;
```

busy: shared variable  
cc: local condition code

## Semaphores

- hardware and software based solution
- no busy waiting
- does not waste CPU time
- **semaphore** is a special variable
  - only access through using two special operations
  - special operations cannot be interrupted
  - operating system carries out special operations

## Semaphores

- `s`: semaphore variable
  - special operations:
    - P (wait): when entering CS: `mutex_begin`
    - V (signal): when leaving CS: `mutex_end`
- ```
P(s):
if (s > 0)
    s=s-1;
else
    wait_on_s;

V(s):
if (anyone_waiting_on_s)
    activate_next_in_line;
else
    s=s+1;
```

Semaphores

- take on integer values (≥ 0)
- created through a special system call
- is assigned an initial value
- binary semaphore:
 - can be 0/1
 - used for CS
- counting semaphore:
 - can be integers ≥ 0

Example: Observer – Reporter

```
global variables:
counter: integer;
sem: semaphore;

process P1:
observe;
P(sem);
counter++;
V(sem);
....

process P2:
...
P(sem);
print(counter);
counter=0;
V(sem);

main_program:
sem=1; counter=0;
activate(P1);
activate(P2);
```

Example: Observer – Reporter

sample run:

```
P1: P(sem) ... sem=0;
P2: P(sem) ... sem=0 so P2 is suspended
P1: V(sem) ... P2 is waiting for sem; activate P2
P2: V(sem) ... no one waiting; sem=1
```

Synchronization with Semaphores

- a process may require an event to proceed
 - process is suspended
 - e.g. process waiting for input
 - another process detecting the occurrence of event wakes up suspended process
- ⇒ “suspend – wake-up” synchronization

Synchronization with Semaphores

- solution:

```
event: semaphore; event=0;
```

```
process P1:           process P2:
...                   ...
P(event);             ...
...                   V(event);
...                   ...
```

- more than two processes may be synchronized

Semaphores

Initial value for semaphore:

- =1 for mutual exclusion
- =0 for synchronization

Semaphores

- possible deadlock scenario:

```
x, y: semaphore;
x=1; y=1;
```

```
process 1:           process 2:
...                   ...
P(x);                 P(y);
...                   ...
P(y);                 P(x);
...                   ...
V(x);                 V(y);
V(y);                 V(x);
...                   ...
```

