

2. The Pipeline

In computer systems greater performance can be achieved by taking advantage of improvements in technology and material, such as faster circuits.

In addition, **organizational enhancements** to the systems can improve performance, such as pipelining, using multiple ALUs or use of cache memories (which we will see in next chapters).

In **pipelining** multiple tasks (for example instructions) are executed in parallel.

To use the pipelining approach efficiently

- We must have tasks that are repeated many times on different data
- Tasks must be divided into small pieces (operations or actions) which can be performed in parallel.

As an example for a pipeline we can consider an automobile assembly line.

The task is the construction of a car.

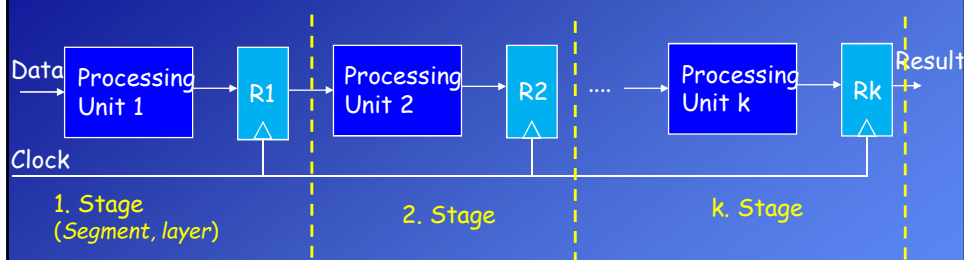
This task is repeated many times for different cars.

The task consists of some steps (operations), such as assembling the doors, assembling the tires.

Each step operates in parallel with other steps but on a different car.

For example, while a worker is assembling the doors of the i^{th} car, another worker is assembling the tires of the $(i+1)^{\text{th}}$ car at the same time.

2.1 The general structure of a pipeline:



Each processing unit performs a fixed operation.

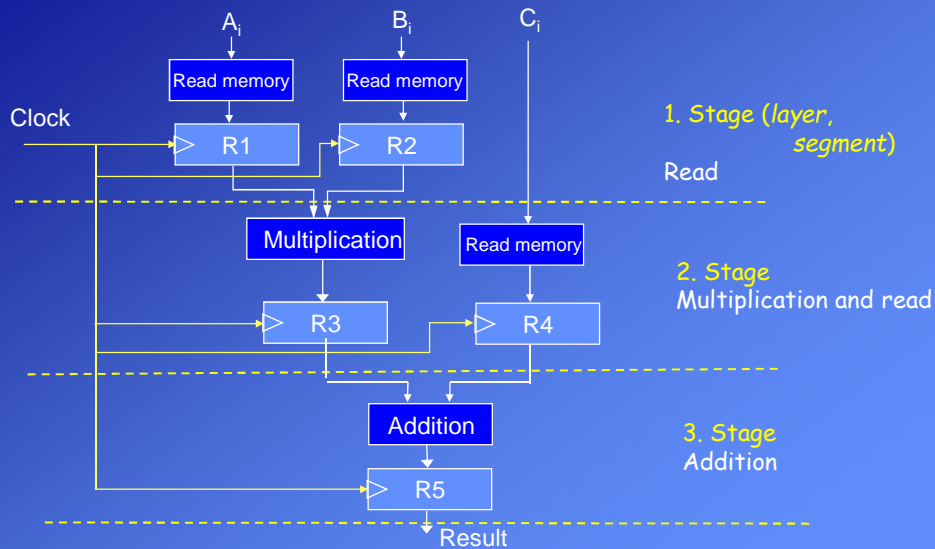
On different clock cycles the operation is performed on different data.

Registers (R1, R2, Rk) keep the intermediate results.

New inputs are accepted at one end before previously accepted inputs appear as outputs at the other end.

When all stages of the pipeline is full, on each clock cycle a new result is obtained at the output.

Example: The elements of the arrays A, B and C will be first read from the memory and then the following operation will be performed: $A_i * B_i + C_i \quad i=1,2,3,\dots$



In this example the task is decomposed as 3 operations: Reading, multiplication and addition

Functioning of the pipeline with three stages:

Clock cycle	1. Stage (Read)		2. Stage (Multiply)		3. Stage (Add)
	R1	R2	R3	R4	R5
1	A_1	B_1	-	-	-
2	A_2	B_2	$A_1 * B_1$	C_1	-
3	A_3	B_3	$A_2 * B_2$	C_2	$A_1 * B_1 + C_1$
4	A_4	B_4	$A_3 * B_3$	C_3	$A_2 * B_2 + C_2$
5	A_5	B_5	$A_4 * B_4$	C_4	$A_3 * B_3 + C_3$

Note: Under the assumption that the data is ready or the access time of the memory is very shorter than the durations of the other operations then reading is not handled as a separate operation.

In this case the pipeline could be designed with two stages which perform only arithmetical operations: multiplication and addition

2.2 Space-Time Diagram of a pipeline with 4 stages

Space-Time Diagrams (or timing diagrams) show which task is currently processed in which stage of the pipeline.

In the diagram below, clock cycles (steps) are written on columns, stages on the rows and task numbers in the cells of the table.

		Time → Clock Cycles						
		1	2	3	4	5	6	7
Stages	S1	T1	T2	T3	T4	T5	T6	
	S2		T1	T2	T3	T4	T5	T6
	S3			T1	T2	T3	T4	T5
	S4				T1	T2	T3	T4

1st task (T1) is completed in 4 clock cycles (number of stage $k=4$).

After k^{th} cycle a new task is completed in each clock cycle

Four tasks have been completed in 7 clock cycles.

Space-Time Diagram of a pipeline with 4 stages, cont'd

The space-time diagram can be also constructed in a different way.

In the diagram below, clock cycles (steps) are written on columns, tasks on the rows and stages in the cells of the table.

		Time → Clock Cycles						
		1	2	3	4	5	6	7
Tasks	T1	S1	S2	S3	S4			
	T2		S1	S2	S3	S4		
	T3			S1	S2	S3	S4	
	T4				S1	S2	S3	S4

1st task (T1) is completed in 4 clock cycles (number of stages $k=4$)

After k^{th} cycle a new task is completed in each clock cycle

Four tasks have been completed in 7 clock cycles.

2.3 Throughput and Speedup provided by the pipeline

Because all stages proceed at the same time, the length of the period of the clock signal (cycle time) is determined by the time required for the slowest stage.

The cycle time (the period of the clock) t_p can be determined as follows:

$$t_p = \max(\tau_i) + d_r = \tau_M + d_r$$

t_p : cycle time

τ_i : time delay of the circuitry in the i^{th} stage

τ_M : maximum stage delay (the slowest stage)

d_r : time delay of the register

Speedup:

k : Number of stages in the pipeline

t_p : cycle time

n : number of tasks

A total of k cycles are required to complete the execution of the first task (T_1).

Required time: $T(1) = k \cdot t_p$

Remaining $n-1$ tasks require $(n-1)$ cycles. Time: $(n-1)t_p$

Total required time for n tasks: $(k+n-1)t_p$

t_n : Required time for a task without pipelining

$$S = \frac{n \cdot t_n}{(k+n-1) \cdot t_p}$$

If we have many tasks $n \rightarrow \infty$

$$\lim_{n \rightarrow \infty} S = \frac{t_n}{t_p}$$

Under assumption $t_n = k \cdot t_p$

$S_{\max} = k$ (Theoretic maximum speedup)

Comments on speedup:

To improve the performance of the pipeline the tasks must be divided into balanced small operations with equal (at least similar) durations.

If the durations of the operations are small then the clock cycle can be short.

Remember the slowest stage determines the clock cycle.

Advantages of increasing the number of stages of a pipeline:

- If the task can be divided into many small operations increasing the number of stages can increase the speed of the clock signal and consequently the speedup.

$$S_{\max} = k$$

But increasing the number of stages may have some disadvantages:

- The cost of the pipeline increases.
- The completion time of the first task increases. $T(1) = k \cdot t_p$
- At each stage of the pipeline, there is some overhead because of registers.
- Branch penalties in the instruction pipelines caused by control hazards increase. We will discuss branch penalties in the section "2.5 Pipeline hazards".

While designing a pipeline these advantages and disadvantages should be taken into consideration.

Effects of task partitioning on the speedup:

If the task can be partitioned into small operations with small durations then a faster clock signal can be applied.

Assume that we have task T with a total duration of 100 ns.

Assume that we can decompose this task in different ways.

Case A: We partition the task into 2 equal stages.



If the delay of registers is 5 ns then the clock cycle is $t_p = 50 + 5 = 55$ ns

Case B: We partition the task into 3 not balanced stages.



The clock cycle is $t_p = 50 + 5 = 55$ ns (slowest stage $\tau_M = 50$ ns)

Although the pipeline has more stages there is no speed improvement compared to case A.

Besides the cost of the pipeline is increased.

The completion time of the first task increases. $T(1) = k \cdot t_p$

Effects of task partitioning on the speedup: (cont'd)

Case C: We partition the task into 3 stages with similar durations.



The clock cycle is $t_p = 40 + 5 = 45$ ns (slowest stage $\tau_M = 40$ ns)

The clock signal is faster compared to cases A and B.

Conclusion:

The pipelining has advantages if a task can be partitioned into small and balanced operations.

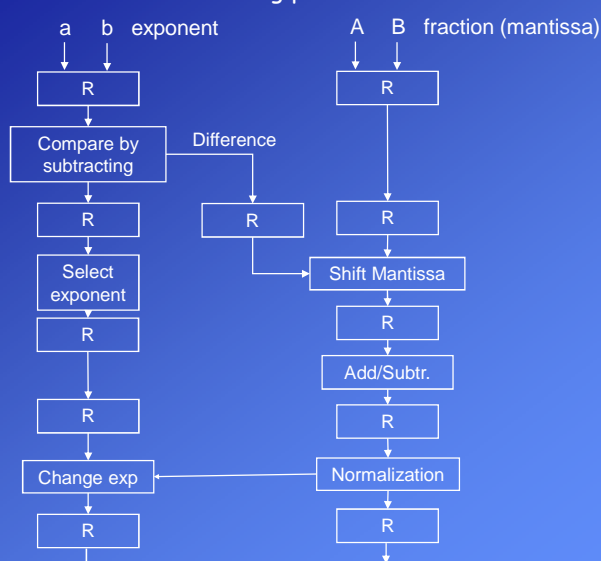
For example if we could partition the task into 5 operations each having the duration of 20ns we would have a clock cycle of 25ns.

Example: An arithmetic pipeline

Addition and subtraction of floating point numbers

$$X = A \cdot 2^a$$

$$Y = B \cdot 2^b$$



2.4 Instruction Pipeline (Instruction-Level Parallelism)

During the execution of each instruction the CPU repeats some operations.

The processing required for a single instruction is called an *instruction cycle* that includes the general stages, instruction fetch, operand fetch, execution, and, interrupt.

The simplest instruction pipeline can be constructed with two stages:

- 1) Fetch instruction 2) Execute instruction

When the main memory is not being accessed during the execution of an instruction this time can be used to fetch the next instruction in parallel with the execution of the current one.

The potential overlap among instruction is called **instruction-level** parallelism.

Remember, to gain more speedup, the pipeline must have more stages with small durations.

Instruction Pipeline (cont'd)

The instruction cycle can be decomposed into 6 operations to gain more speedup:

1. **Fetch instruction (FI):** Read the next expected instruction into a buffer.
2. **Decode instruction (DI):** Determine the opcode and the operand specifiers.
3. **Calculate addresses of operands (CO):** Calculate the effective address.
4. **Fetch operands (FO):** Fetch each operand from memory.
5. **Execute instruction (EI):** Perform the indicated operation.
6. **Write operand (WO):** Store the result in memory.

Due to several factors this decomposition may not increase the performance so much.

Problems:

- The various stages will be of different durations.
- Some instructions do not need all stages.
- Different segments can need memory access at the same time.

Therefore, some operations can be combined into same stage so that a pipeline with less stages (for example 4 or 5) is constructed.

For example, 80468 has 5 stages.

There are also processors that include instruction pipelines with more stages. For example processors of Pentium 4 family include a pipeline with 20 stages. Here internal operations are decomposed into small actions.

An instruction pipeline with 4 stages (segments)

1. FI (Fetch Instruction)
2. DA (Decode, Address)
3. FO (Fetch Operand)
4. EX (Execution): Performing the operation and updating the registers

In order to perform instruction and operand fetch operations at the same time we assume that the processor has separate instruction and data memories.

Timing Diagram for Instruction Pipeline Operation (ideal case):

Clock cycles	1	2	3	4	5	6	7	8
Instructions								
1	FI	DA	FO	EX				
2		FI	DA	FO	EX			
3			FI	DA	FO	EX		
4				FI	DA	FO	EX	
5					FI	DA	FO	EX

First instruction has been completed.
4 cycles

Just after one cycle the second instruction has been completed.

2.5 Pipeline Hazards (Conflicts)

1. Control Hazards (Branches, Interrupts):

a. Unconditional Branch

Clock cycles	1	2	3	4	5	6	7	8
Instructions								
1	FI	DA	FO	EX				
2		FI	DA	FO	EX			
Unconditional Branch			FI	DA	FO	EX		
				FI	-	-	FI	DA
								FI

After decoding, the type of the instruction is determined: branch!

The branch address is taken (absolute or relative)

Updating the PC
PC= The branch address (target)

This instruction is fetched unnecessarily. It will be discarded.

Branch penalty!
It is necessary to stall or empty the pipeline

The new instruction after branch operation (Target of branch)

b1. Conditional Branch (if the condition is not true):

Clock cycles Instructions	1	2	3	4	5	6
1	FI	DA	FO	EX		
Conditional bra. 2		FI	DA	FO	EX	
3			FI	DA	FO	EX

The previous instruction sets the conditions (flags).

Without considering the condition next instruction is fetched

No branch. PC is not changed.
No branch penalty.

b2. Conditional Branch (if the condition is true):

Clock cycles Instructions	1	2	3	4	5	6	7
1	FI	DA	FO	EX			
Conditional bra. 2		FI	DA	FO	EX		
3			FI	DA	FO	FI	DA
4				FI	DA		FI
5					FI		

The branch address is written into PC.
The pipeline must be emptied.

The target instruction of branch

The pipeline is emptied. (Branch penalty)

Pipeline Hazards (Conflicts) cont'd**2. Resource Conflict (Structural hazard):**

A resource hazard occurs when two (or more) instructions that are already in the pipeline need the same resource (memory, functional unit).

a) Memory conflict: An operand read to or write from memory cannot be performed in parallel with an instruction fetch.

Solutions:

- Instructions must be executed in serial rather than parallel for a portion of the pipeline (*stall*).
- Harvard architecture: Separate memories for instructions and data.
- Instruction queue or cache memory: There are times during the execution of an instruction when main memory is not being accessed. This time could be used to prefetch the next instruction and write it to a queue (instruction buffer).

b) Functional unit (ALU) conflict.

Solutions:

- Increasing available functional units and using multiple ALUs.
- For example different ALUs can be used address calculation and data operations.
- Fully pipelining a functional unit (for example a floating point unit FPU)

Pipeline Hazards (Conflicts) cont'd

3. Data Conflict:

There is a conflict in the access of a data location.

If the problem is not solved the program produces an incorrect result because of the use of pipelining.

a) Operand dependency:

The operand of an instruction depends on the result of another instruction

Example (68K):

Clock cycles		1	2	3	4	5
Instructions						
ADD.W D1, DATA		FI	DA	FO	EX*	
MOVE.W DATA, (A0)			FI	DA	FO	EX

DATA is updated

Operand dependency

b) Address Dependency:

Data conflict can also occur on address registers.

Example (68K):

```

ADDA.W #2, A0
MOVE.B (A0)+, D0
  
```

Previous value (not valid) of DATA is being fetched

3. Data Conflict cont'd:

There are three types of data hazards:

- **Read after write (RAW), or true dependency:** An instruction modifies a register or memory location and a succeeding instruction reads the data in that memory or register location.
A hazard occurs if the read takes place before the write operation is complete.
- **Write after read (WAR), or antidependency:** An instruction reads a register or memory location and a succeeding instruction writes to the location.
A hazard occurs if the write operation completes before the read operation takes place.
- **Write after write (WAW), or output dependency:** Two instructions both write to the same location.
A hazard occurs if the write operations take place in the reverse order of the intended sequence.

2.6 Solutions to Data Hazards:

• Operand forwarding (Bypassing):

An optional direct connection is established between the output and the inputs of the ALU.

To explain the operand forwarding the following instruction pipeline structure (that is common in RISC processors) is used:

1. FI (*Fetch Instruction*)
2. DO (*Decode, Operand (register) fetch*): It is reasonable in RISC processors
3. EX (*Execution*): Performing the operation on registers
4. WO (*Write Operand*): Writing the result to the registers

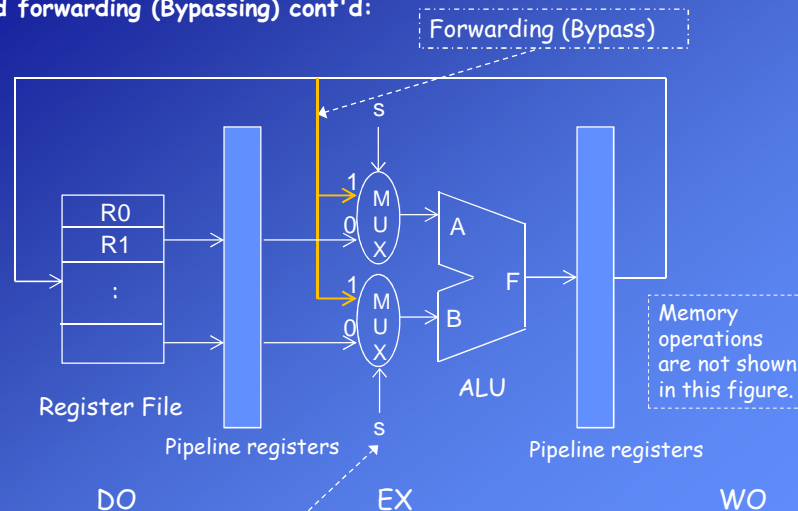
Example:

Instructions	1	2	3	4	5
ADD R1, R2, R3; $R1 \leftarrow R2 + R3$	FI	DO	EX	WO	
SUB R4, R5, R1; $R4 \leftarrow R5 - R1$		FI	DO	EX	WO

Annotations:

- R1 is updated (pointing to the WO stage of the first instruction)
- RAW dependency (pointing to the DO stage of the second instruction, which depends on the result of the first instruction)
- Previous value (not valid) of R1 is fetched (pointing to the DO stage of the second instruction, where the value of R1 is fetched before it is updated)

• Operand forwarding (Bypassing) cont'd:



s is controlled by the hazard detection unit of the pipeline. It selects the value from the register file or the forwarded result (bypass) as the ALU input.

• Operand forwarding (Bypassing) cont'd:

If the forwarding detects that the destination of the previous ALU operation is the same register as the source of the current ALU operation, control logic selects the forwarded result (bypass) as the ALU input rather than the value from the register.

Example:

	Clock cycles				
Instructions	1	2	3	4	5
ADD R1, R2, R3; $R1 \leftarrow R2 + R3$	FI	DO	EX	WO	
SUB R4, R5, R1; $R4 \leftarrow R5 - R1$		FI	DO	EX	WO

Previous value (not valid) of R1 is fetched.
This invalid value will not be used in the EX cycle.

The control unit of the pipeline selects the output of the previous ALU operation as the input, not the value from the DO.

If it is possible to solve the conflict by forwarding it is not necessary to stall the pipeline and there is not a decrease in performance.

Solutions to Data Hazards (cont'd):

• Hardware interlock:

A hardware unit tracks all instructions and stalls the pipeline when a hazard is detected.

• Compiler-based Solutions:

Delayed Load:

The compiler inserts NOP (No Operation) instructions between the instructions that cause data hazards.

The compiler changes the order of the instruction if it is possible.

We will see these solutions in the chapter RISC Pipeline

2.7 Dealing with branches:

The main problem is the conditional branch instruction because until the instruction is actually executed, it is impossible to determine whether the branch will be taken or not.

Solution mechanisms:

2.7.1 Target Instruction prefetch:

When a conditional branch is recognized, the target of the branch is prefetched, in addition to the instruction following the branch.

2.7.2 Branch prediction:**Static branch prediction strategies:**

- a) Always predict not taken: Always assumes that the branch will not be taken and fetches the next instruction in sequence.
- b) Always predict taken: Always predicts that the branch will be taken and fetches target instruction of the branch. (Performs better than a)

Studies analyzing program behavior have shown that conditional branches are taken more than 50% of the time.

Therefore; always prefetching from the branch target address should give better performance than always prefetching from the sequential path.

Dynamic branch prediction strategies

Dynamic branch strategies record the history of conditional branch instructions to predict whether the condition is true or not.

One or more bits (or counters) can be associated with each conditional branch instruction that reflect the recent history of the instruction.

These bits direct the processor to make the decision the next time the branch instruction is encountered.

1-bit prediction scheme:

Addresses of branch instructions and one prediction bit for each instruction are kept in high-speed memory location called branch history table (BHT).

The prediction bit only records whether the last execution of this instruction resulted in a branch or not.

If the branch was taken last time, it predicts to take the branch next time.

Algorithm:

Fetch the i^{th} branch instruction

If ($p_i=0$) predict not to take the branch, prefetch the next instruction in sequence

If ($p_i=1$) predict to take the branch, prefetch the target instruction of the branch

If the branch is really taken $p_i=1$

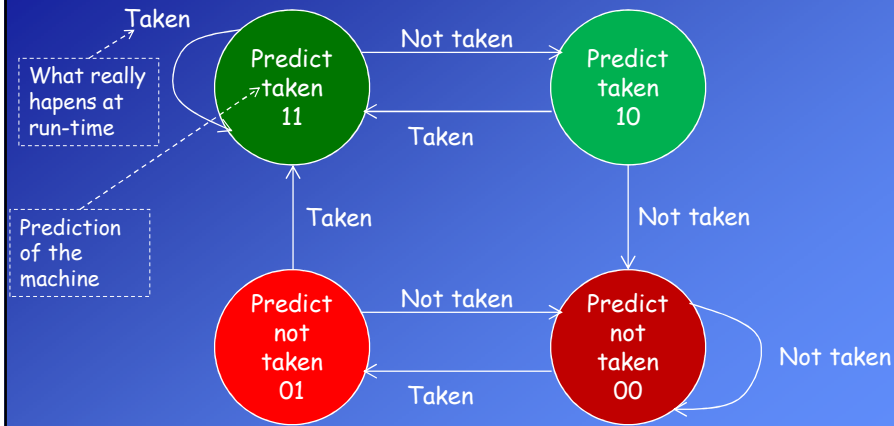
If the branch is not really taken $p_i=0$

Problem: misprediction will occur twice for each use of the loop: once on entering the loop, and once on exiting.

2-bit Branch prediction state diagram

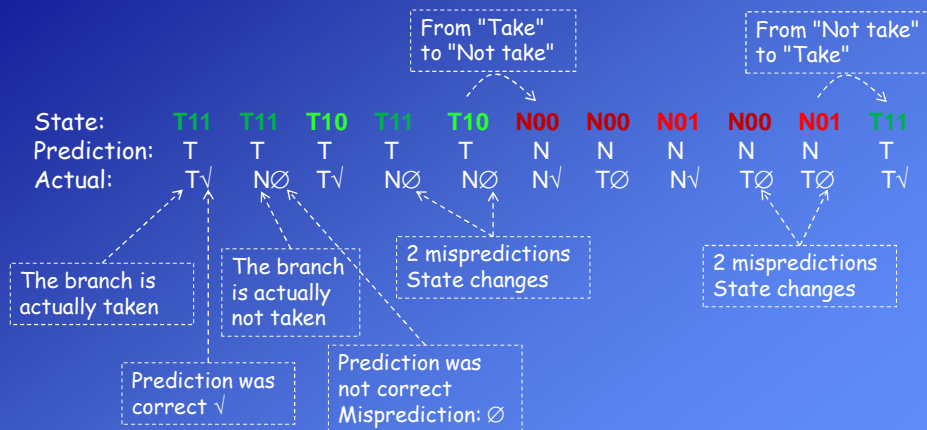
If the system is in states 11 and 10 predicts to take the branch.

If the system is in states 00 and 01 predicts not to take the branch.



In this scheme prediction is changed only if it gets misprediction twice.

Example: 2-bit Branch prediction

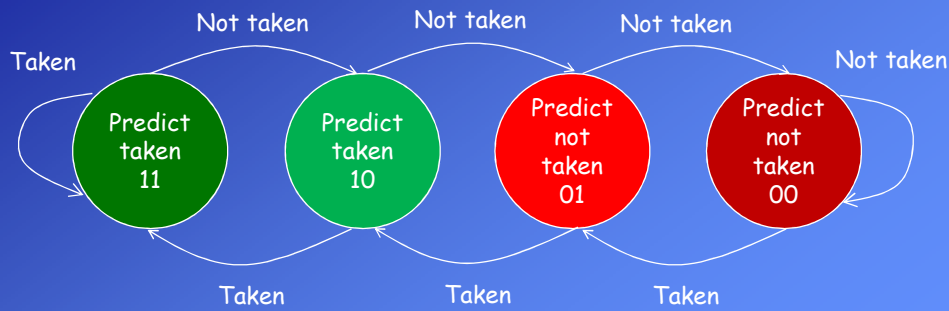


Saturating counter: Another 2-bit Branch prediction strategy

There are different ways to implement branch prediction strategies. Saturating counter is another alternative.

If the system is in states 11 and 10 predicts to take the branch.

If the system is in states 00 and 01 predicts not to take the branch.



Branch target buffer and branch history table:

In the branch history table besides the state bits also the target address of the branch instruction can be kept.

The table that keeps target address of branch instructions is called branch target buffer.

With the help of this buffer the target instruction of the branch can be prefetch without calculating the branch address.

Branch PC	Target address	State bits

2.7.3 Compiler-based solution:

Delayed branch:

- Optimized delayed branch:

The compiler changes the order of the instructions within a program, so that it is not necessary to stall or empty the pipeline because of the branch instructions.

This rearrangement of the instructions should not effect the behavior of the program.

With this method performance degradation is not encountered.

- Inserting NOOP (No Operation) instructions:

If it is not possible to change the order of the instructions the compiler inserts NOOP instructions after the branch instructions.

We will discuss these compiler based methods in the next chapter: RISC pipelining

2.8 RISC Pipelining

In RISC processors most instructions are register to register.

A pipeline with two stages would be sufficient for these instructions:

I: Instruction fetch, A: ALU operation (execution) on registers

Only for load and store operations memory to register and register to memory instructions are necessary.

For these operation an additional stage (D) to access memory is necessary.

So an instruction pipeline for a RISC processor can be designed with 3 stages:



- I: Instruction Fetch
- A: Decode, ALU Operation
- D: Data, memory access

To increase the performance there are also RISC processors that includes pipelines with more stages (4, 5 even more).

For example,

MIPS R3000 has 5 stages

MIPS R4000 has 8 stages (superpipelined)

ARM7 has 3 stages, ARM Cortex-A8 has 13 stages.

2.8.1 Example: A RISC Pipeline with 3 stages

- I (*Instruction Fetch*): Read the instruction from memory (Pointed by the PC)
- A (*Decode, ALU Operation*):
ALU is used for three different tasks:
 1. Arithmetical and logical operations on registers
 2. Memory address calculation in load/store instructions. $LDL(R5)\#10, R15$
 $PC \leftarrow PC + Y$
 3. Relative addressing
 In stage A (ALU) the operation is performed and the result is written to the destination register (R or PC) in the same clock cycle.
- D (*Data*): This stage is only used for memory access by load/store instructions.

I and D stages try to access the memory at the same time.

To solve the data hazard problem separate memories for instruction and data can be used (Harvard architecture)

Other solutions are instruction or data queues and cache memories.

Data conflict (dependency) in the RISC pipeline with 3 stages

Example:

```

100 LOAD  M[X], R1    R1 ← M[X]
104 LOAD  M[Y], R2    R2 ← M[Y]
108 ADD   R1, R2, R3   R3 ← R1 + R2
10C STORE R3, M[Z]    M[Z] ← R3
110 LOAD  M[W], R4    R4 ← M[W]
  
```

ADD instruction does not use this stage

Data dependency in the pipeline

Clock cycles	1	2	3	4	5	6	7
Instructions							
LOAD R1	I	A	D				
LOAD R2		I	A	D			
ADD R1,R2,R3			I	A	D		
STORE R3				I	A	D	
LOAD R4					I	A	D

Data conflict:

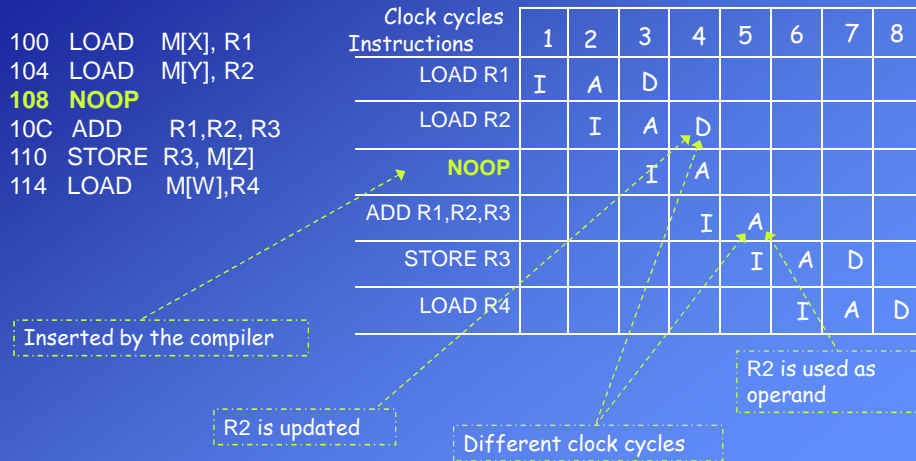
In the 4th clock cycle LOAD writes to R2, and at the same time ADD uses R2 as a source operand.

There is not a data conflict related to R3

Delayed Load

- Inserting NOOP (No Operation) instructions

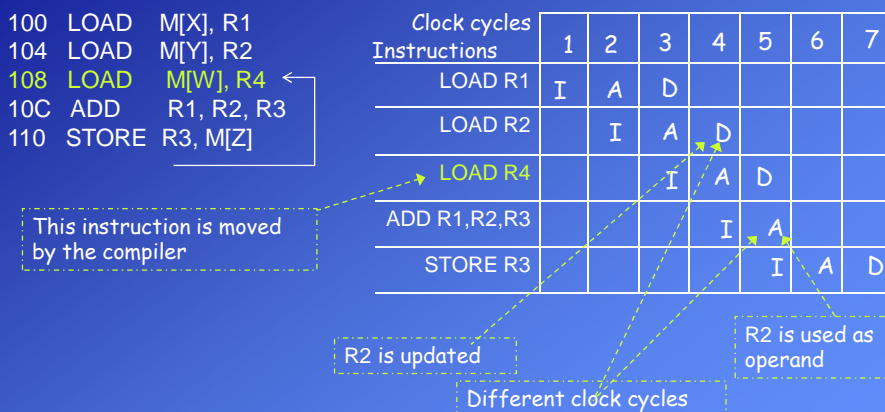
The compiler inserts NOOP instructions between LOAD and the instruction that uses the register.



Delayed Load, cont'd

- Changing the order of the instructions

The compiler rearranges the program and moves certain instructions between LOAD and instruction that uses the register.



The performance is improved: 7 clock cycles (instead of 8).

The behavior of the program is not changed.

Delayed Branch, The Problem

The branch (jump) instructions update the PC in the ALU (Execution) stage. But the next instruction in sequence (not the target of branch) is fetched at the same time that jump is executed.

In this case, either a hardware unit must empty the pipeline or compiler-based solutions (delayed branch) must be applied.

Without any precaution:**Example: Pseudo Code**

```
100 LOAD    M[X], R1
104 ADD     1, R2
108 JUMP    200
10C ADD     R1,R2
....
200 STORE   R1, M[Y]
```

Clock cycles	1	2	3	4	5	6	7
Instructions							
100 LOAD M[X], R1	I	A	D				
104 ADD 1, R2		I	A				
108 JUMP 200			I	A			
10C ADD R1,R2				I	A		
200 STORE R1,M[Y]					I	A	D

The pipeline must be emptied by hardware or a compiler-based solution must be applied.

PC is updated
PC ← 200 (Target address)

Problem: This instruction has been already fetched. Actually it should not run!

Delayed Branch, Solutions

- Inserting NOOP (No Operation) Instructions:**

To regularize the pipeline, a NOOP instruction is inserted after the branch.

Delayed branch with NOOP**Pseudo Code**

```
100 LOAD    M[X], R1
104 ADD     1, R2
108 JUMP    200
10C NOOP
110 ADD     R1,R2
....
200 STORE   R1, M[Y]
```

Clock cycles	1	2	3	4	5	6	7
Instructions							
100 LOAD M[X], R1	I	A	D				
104 ADD 1, R2		I	A				
108 JUMP 200			I	A			
10C NOOP				I	A		
200 STORE R1,M[Y]					I	A	D

Inserted by the compiler

PC is updated
PC ← 200 (Target address)

In different clock cycles

Now, the program runs in the required sequence.

But inserting NOOP instructions degrades the performance of the pipeline.

Delayed Branch, Solutions

- **Changing the order of the instructions, optimized delayed branch:**

Certain instructions (mostly from before the branch) can be placed after the branch.

Interchanging instructions improves the performance of the pipeline.

Pseudo Code

```

100 LOAD    M[X], R1
104 JUMP    200
108 ADD     1, R2
10C ADD     R1,R2
....
200 STORE   R1, M[Y]

```

Delayed branch with interchanging instructions

Clock cycles	1	2	3	4	5	6
Instructions						
100 LOAD M[X], R1	I	A	D			
108 JUMP 200		I	A			
104 ADD 1, R1			I	A		
200 STORE R1,M[Y]				I	A	D

PC is updated
PC ← 200 (Target address)

This instruction has been
already fetched.

The performance is improved: 6 clock cycles (instead of 7).

The behavior of the program is not changed.

Important points about changing the order of the instructions:

An instruction **from before** the branch can be placed just after the branch.

Branch (condition or address) must not depend on moved instruction.

This method (if it is possible) always improves the performance.

Especially, for **conditional branches**, this procedure must be applied carefully

If the condition that is tested for the branch is altered by the immediately preceding instruction, than the compiler can not move this instruction after the branch.

In this case NOOP can be inserted.

Other possibilities:

Compiler can select instructions to move

- From branch target
- Must be OK to execute moved instruction even if the branch is not taken
- Improves performance when branch is taken
- From fall through
- Must be OK to execute moved instruction even if the branch is taken
- Improves performance when branch is not taken

2.8.2 A RISC Pipeline with 4 stages

Because of the simplicity and regularity of a RISC instruction set, the design of the pipeline with three or four stages is easily accomplished.

A RISC pipeline with four stages can be designed as follows:

- I (*Instruction Fetch*): Read the instruction from memory (Pointed by the PC)
- R (*Decode, Read register file*)
- A (*ALU Operation And register write*)
- D (*Data*): Memory access by load/store instructions.

If the pipeline consists of 4 stages, in delayed load and delayed branch operations 2 NOOP instructions must be inserted or 2 instructions must be moved.

Delayed Load (In a pipeline with 4 stages)

Inserting 2 NOOP instructions:

```
100 LOAD M[X], R1
104 LOAD M[Y], R2
108 NOOP
10C NOOP
110 ADD R1, R2, R3
```

Inserted by the compiler

Clock cycles Instructions	1	2	3	4	5	6	7
LOAD R1	I	R	A	D			
LOAD R2		I	R	A	D		
NOOP			I	R	A		
NOOP				I	R	A	
ADD R1,R2,R3					I	R	A

R2 is updated

Different clock cycles

R2 is used as operand

Delayed Branch (In a pipeline with 4 stages)

Inserting 2 NOOP instructions:

Pseudo Code

```

...
108 JUMP    200
10C NOOP
110 NOOP
114 ADD     1,R1
114 ADD     R1,R2
....
200 STORE   R1, M[Y]

```

PC is updated
 $PC \leftarrow 200$ (Target address)

Delayed branch with NOOP

Clock cycles		1	2	3	4	5	6	7
Instructions								
108 JUMP 200		I	R	A				
10C NOOP			I	R	A			
110 NOOP				I	R	A		
200 STORE R1,M[Y]					I	R	A	D

Inserted by the compiler

In different clock cycles