

KERNEL ARCHITECTURE

BLG413E – System Programming, Practice Session 2

Contents

- System Calls
- Kernel Modules

1-Adding a system call

Requirements (for Ubuntu OS): linux-source, kernel_package, fakeroot, libncurses5-dev

Steps:

- define and write new system call
- modify system call table
- modify system header files
- compile kernel
- reboot to new kernel
- test new system call

Writing a system call

- **mycall.c:** under /usr/src/linux-source-x.x.x/kernel/

```
#include <linux/kernel.h>

asmlinkage int sys_mycall(int i, int j){
    return i + j;
}
```

- **Modify Makefile:** under /usr/src/linux-source-x.x.x/kernel/
 - add obj-y := mycall.o

Modifying system call table and system header files

- In /usr/src/linux-source-x.x.x/arch/x86/kernel/syscall_table_32.S:
 - append `.long sys_mycall`
- In /usr/src/linux-source-x.x.x/include/linux/syscalls.h:
 - add prototype of the new system call as

```
asmlinkage int sys_mycall(int i, int j);
```
- In /usr/src/linux-source-x.x.x/arch/x86/include/asm/unistd_32.h:
 - append `#define __NR_mycall 349`
 - The index should be the number after the last system call in the list
 - You should also increment the system call count by 1 (**NR_syscalls**)

```
#define NR_syscalls 350
```

Before compiling linux kernel

- **Configuration Options (1st one recommended):**
 1. copy config.txt file to linux-source-2.6.38 (under Ubuntu 11.04) and rename it as .config
 2. make oldconfig (only ask for new or different options)
 3. make menuconfig (menu for configuration)

Compiling linux kernel

- `make-kpkg clean` → cleans up all from previous kernel compiles
- **Compilation:** `fakeroot make-kpkg --initrd --append-to-version=-custom kernel_image kernel_headers`
- **Output:** two files in parent directory:
 - `linux-image-x.x.x-customdeb`
 - `linux-headers-x.x.x-customdeb`

Installing compiled kernel

- `sudo dpkg -i linux-image-x.x.x-customdeb`
- `sudo dpkg -i linux-headers-x.x.x-customdeb`
- **Uninstalling:**
 - `sudo dpkg -r linux-image-x.x.x-custom`
 - `sudo dpkg -r linux-headers-x.x.x-custom`

Testing new system call

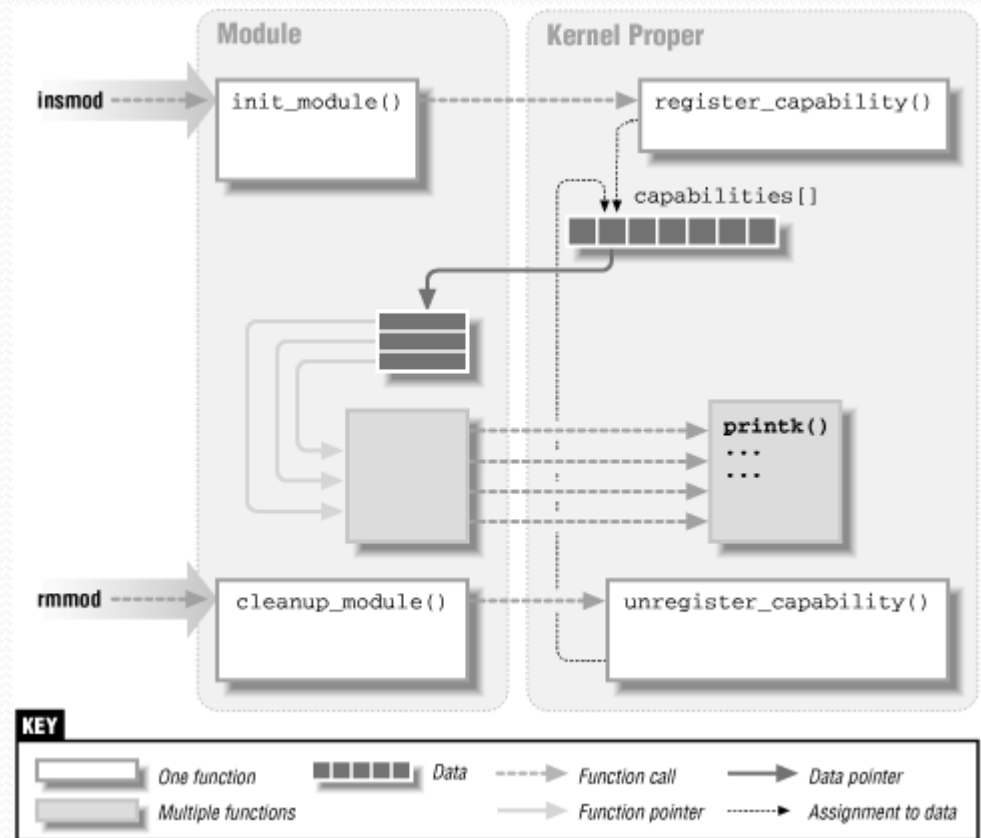
```
#include <stdio.h>
#include <sys/syscall.h>

#define __NR_mycall 349

int main(void) {
    int x1=10, x2=20, y;
    y = syscall(__NR_mycall, x1, x2);
    printf("%d\n", y);
    return 0 ;
}
```

2-Kernel modules

- A way to add new features to the kernel without rebuilding it.
- Unlike applications, modules register themselves for serving future requests.
- Applications can access the capabilities of a module through system calls.



http://www.xml.com/ldd/chapter/book/figs/ldr2_0201.gif

An example module: hello

- **hello.c:**

```
#include <linux/init.h> /* for module_init and module_exit */
#include <linux/module.h> /* needed by all modules */
MODULE_LICENSE("Dual BSD/GPL"); /* a macro to declare that this module is open source */

static int hello_init(void) /* static: invisible outside the module */
{
    /* to avoid namespace pollution */
    printk(KERN_ALERT "Hello, world\n"); /* printk: kernel print function (macros for priority) */
    return 0; /* KERN_ALERT: a situation requiring immediate action */
}

static void hello_exit(void)
{
    printk(KERN_ALERT "Goodbye, cruel world\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

- **Makefile:**

obj-m := hello.o

all:

make -C /lib/modules/\$(shell uname -r)/build M=\$(PWD) modules

M=\$(PWD) is to build external module in the working directory



Using hello module

- **Compiling:**
 - *make*
- **Loading** (check with ***dmesg*** which is used to write the kernel messages):
 - *sudo insmod ./hello.ko*
- **Unloading** (check with ***dmesg***):
 - *sudo rmmod hello*
- check with ***lsmod*** (which prints the contents of the `/proc/modules` file) before and after loading and unloading

An example module using load time parameters

- **hello.c:**

```
/* $Id: hello.c,v 1.4 2004/09/26 07:02:43 gregkh Exp $ */
#include <linux/init.h>
#include <linux/module.h>
#include <linux/moduleparam.h> /* to enable passing parameters at loadtime */
MODULE_LICENSE("Dual BSD/GPL");

/* A couple of parameters that can be passed in: how many times we say hello, and to whom */
static char *whom = "world";
static int howmany = 1;
module_param(howmany, int, S_IRUGO); /* S_IRUGO: read by the world but cannot be changed */
module_param(whom, charp, S_IRUGO);

static int hello_init(void){
    int i;
    for (i = 0; i < howmany; i++)
        printk(KERN_ALERT "(%d) Hello, %s\n", i, whom);
    return 0;
}

static void hello_exit(void){
    printk(KERN_ALERT "Goodbye, cruel world\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

Specifying module parameters

- `sudo insmod ./hellop.ko whom='Mom' howmany=4`
- `dmesg`

```
4555.764793] (0) Hello, Mom  
4555.764796] (1) Hello, Mom  
4555.764797] (2) Hello, Mom  
4555.764798] (3) Hello, Mom
```

- `sudo rmmod hellop`
- `dmesg`

```
4555.764793] (0) Hello, Mom  
4555.764796] (1) Hello, Mom  
4555.764797] (2) Hello, Mom  
4555.764798] (3) Hello, Mom  
4611.350208] Goodbye, cruel world
```