

DEVICE DRIVERS

BLG413E – System Programming, Practice Session 3

scull (Simple Character Utility for Loading Localities)

- a char driver that treats a memory area as a device
- used as an example to demonstrate and test the interface between the kernel and char drivers

Compiling scull

- **Warning:** simplified scull code under ninova is incompatible with Linux kernel versions newer than 2.6.35.
 - Use “*uname -a*” to check the version of the kernel you are currently using
- Two changes are made in the code to adapt to newer versions of the Linux kernel.

```
struct file_operations scull_fops = {
    .owner =      THIS_MODULE,
    .llseek =     scull_llseek,
    .read =       scull_read,
    .write =      scull_write,
    .unlocked_ioctl = scull_ioctl,
    .open =       scull_open,
    .release =    scull_release,
};
```

ioctl is not available after Linux kernel 2.6.36, use *unlocked_ioctl* instead

```
/* Initialize each device. */
for (i = 0; i < scull_nr_devs; i++) {
    dev = &scull_devices[i];
    dev->quantum = scull_quantum;
    dev->qset = scull_qset;
    sema_init(&dev->sem, 1);
    devno = MKDEV(scull_major, scull_minor + i);
    cdev_init(&dev->cdev, &scull_fops);
    dev->cdev.owner = THIS_MODULE;
    dev->cdev.ops = &scull_fops;
    err = cdev_add(&dev->cdev, devno, 1);
    if (err)
        printk(KERN_NOTICE "Error %d adding scull%d", err, i);
}
```

init_MUTEX is not available after Linux kernel 2.6.37, use *sema_init* instead

Compiling scull

- **Makefile:**

```
obj-m := scull.o
```

M=\$(PWD) is to build external module in the working directory

```
all:
```

```
make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
```



- **Compiling:**

```
make
```

Using scull

- when testing the scull module, it's better to become root instead of using sudo for commands:
 - `sudo su`
- **Loading:**
 - `insmod ./scull.ko`
 - `lsmod` → to see scull in the list of loaded modules
- **Getting the major number:**
 - `grep scull /proc/devices` → file displaying currently configured (and loaded) character and block devices
- **Creating the device nodes (assuming major number is 250):**
 - `mknod /dev/scull0 c 250 0`
 - `mknod /dev/scull1 c 250 1`
 - ...

Using scull

- **Writing to the device:**
 - *echo testing > /dev/scull0*
- **Reading from the device:**
 - *cat /dev/scull0*

Using scull

- Writing more than one quantum (size of the file is 58739 bytes):
 - `cp /etc/bash_completion /dev/scull0`
- Tracing the system calls:
 - `strace cp /etc/bash_completion /dev/scull0`

```
open("/etc/bash_completion", O_RDONLY|O_LARGEFILE) = 3
fstat64(3, {st_mode=S_IFREG|0644, st_size=58739, ...}) = 0
open("/dev/scull0", O_WRONLY|O_TRUNC|O_LARGEFILE) = 4
fstat64(4, {st_mode=S_IFCHR|0644, st_rdev=makedev(250, 0), ...}) = 0
read(3, "#\n#  bash_completion - programm"... 32768) = 32768
write(4, "#\n#  bash_completion - programm"... 32768) = 4000
write(4, "ular\ncomplete -f -X '!*.@(?(e)ps"... 28768) = 4000
write(4, "ulky functions in memory if we d"... 24768) = 4000
write(4, "# Default to cword unchanged\n  "... 20768) = 4000
write(4, "nt to return host:path and not o"... 16768) = 4000
write(4, " We messed up! At least return t"... 12768) = 4000
write(4, "ut\n# the bash < 4 compgen hack.\n"... 8768) = 4000
write(4, "@|}' -- \"$cur\" ) )\n      ltrim c"... 4768) = 4000
write(4, "o expand\n  # ~foo/... to /home"... 768) = 768
read(3, "es on process group IDs.\n# AIX a"... 32768) = 25971
write(4, "es on process group IDs.\n# AIX a"... 25971) = 3232
write(4, "on completes on user or user:gro"... 22739) = 4000
write(4, "word breaks. See  reassemble co"... 18739) = 4000
write(4, " /etc/ssh/ssh_config \"${HOME}/.s"... 14739) = 4000
write(4, "      #if [[ ${COMP_KNOWN_HOSTS_WI"... 10739) = 4000
write(4, "      # shift COMP_WORDS elements a"... 6739) = 4000
write(4, "cal tmp\n\n      toks=( ${toks[@]-} "... 2739) = 2739
read(3, "", 32768) = 0
close(4) = 0
close(3) = 0
```

default block size for
reading/writing

using 4000 byte sized
quantums for writing

Using scull

- **Writing more than the capacity of the device:**

- *strace cp /usr/bin/inkscape /dev/scull0*

- **Testing with quantum size 32768:**

- *rmmod scull*
- *insmod ./scull.ko scull_quantum=32768*
- *strace cp /etc/bash_completion /dev/scull0*

```
open("/etc/bash_completion", O_RDONLY|O_LARGEFILE) = 3
fstat64(3, {st_mode=S_IFREG|0644, st_size=58739, ...}) = 0
open("/dev/scull0", O_WRONLY|O_TRUNC|O_LARGEFILE) = 4
fstat64(4, {st_mode=S_IFCHR|0644, st_rdev=makedev(250, 0), ...}) = 0
read(3, "#\n#  bash_completion - programm"... , 32768) = 32768
write(4, "#\n#  bash_completion - programm"... , 32768) = 32768
read(3, "es on process group IDs.\n# AIX a"... , 32768) = 25971
write(4, "es on process group IDs.\n# AIX a"... , 25971) = 25971
read(3, "", 32768) = 0
close(4) = 0
close(3) = 0
```

Each block is written in a single write process when quantum size is the same with the block size.

References

- Corbet, J., Rubini, A., & Kroah-Hartman, G. (2005). Chapter 3: Char Drivers. In *Linux Device Drivers, Third Edition* (pp. 42-72). O'Reilly.