

Systems Programming

Kernel Development

H. Turgut Uyar Şima Uyar

2001-2009

1 / 39

Topics

Kernel

Introduction
Compiling the Kernel
System Calls
Adding a System Call

Kernel Modules

Introduction
Adding a Module

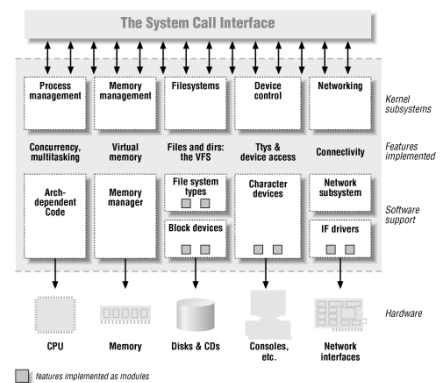
2 / 39

Kernel

- ▶ provides programs with a consistent view of the hardware
- ▶ protection against unauthorized access to resources
 - ▶ different levels of protection
 - ▶ kernel executes in supervisor-mode (kernel-space), applications execute in user-mode (user-space)
- ▶ switching to kernel-space:
 - ▶ system calls: synchronous, in the process context
 - ▶ interrupts: asynchronous

3 / 39

Kernel Subsystems



4 / 39

Kernel Subsystems

- ▶ process management
 - ▶ creating and destroying processes
 - ▶ communication between processes
 - ▶ scheduling
- ▶ memory management
 - ▶ virtual address space for each process
- ▶ filesystems
 - ▶ structured filesystem on top of unstructured hardware
- ▶ device control
- ▶ networking
 - ▶ delivering data packets across program and network interfaces
 - ▶ routing and address resolution

5 / 39

Kernel Architecture

- ▶ *monolithic*
 - ▶ all functionality in one big chunk of code
- ▶ *microkernel*
 - ▶ kernel organized as layers
 - ▶ most of the functionality in user-space
 - ▶ too much communication overhead

6 / 39

Getting Ready

- ▶ `sudo apt-get update`
- ▶ `sudo apt-get install linux-source`
- ▶ `tar xjvf /usr/src/linux-source-2.6.28.tar.bz2`
- ▶ `sudo apt-get install kernel-package`
utilities for compiling the kernel
- ▶ `sudo apt-get install fakeroot`
allows compiling the kernel without becoming root
- ▶ `sudo apt-get install libncurses5-dev`
used by script which configures the kernel

7 / 39

Configuring

- ▶ copy `config.txt` file to `linux-source-2.6.28`
- ▶ rename as `.config`
- ▶ `make oldconfig`
modifies the config file based on changes between kernel versions
- ▶ `make menuconfig`
for configuring the kernel
 - ▶ as exercise: remove the floppy disk support from the kernel
 - ▶ `[M]`: module, `[*]`: included in kernel, `[]`: disabled

8 / 39

Compiling

- ▶ `make-kpkg clean`
cleans up all from previous kernel compiles
- ▶ `fakeroot make-kpkg -initrd`
`-append-to-version=-custom`
`kernel_image kernel_headers`
- ▶ two files in parent directory:
`linux-image-2.6.28.10-custom.....deb`
`linux-headers-2.6.28.10-custom.....deb`

9 / 39

Installing

- ▶ `sudo dpkg -i linux-image-2.6.28.10-custom.....deb`
- ▶ `sudo dpkg -i linux-headers-2.6.28.10-custom.....deb`
- ▶ uninstalling:
`sudo dpkg -r linux-image-2.6.28.10-custom`
`sudo dpkg -r linux-headers-2.6.28.10-custom`

10 / 39

System Calls

- ▶ used by application programs to request service from operating system
- ▶ execute in kernel mode
- ▶ each has a unique number

11 / 39

System Call Interface

- ▶ wrapper functions usually have same name
- ▶ system call interface may change when sizes of structures or arguments change
- ▶ user programs unaware of change because of wrapper functions

12 / 39

System Calls in Linux

list of system calls

- ▶ man 2 syscalls
- ▶ man 2 unimplemented

invoking a system call

- ▶ **int** syscall (**int** no, ...);

13 / 39

Preliminaries

- ▶ new system call becomes part of kernel
- ▶ need to re-compile the kernel
- ▶ step-by-step procedure may be different for different distributions

14 / 39

Steps

- ▶ define and write new system call
- ▶ modify system header files
- ▶ modify system call table
- ▶ compile kernel
- ▶ reboot to new kernel
- ▶ test new system call

15 / 39

System Call Example

Example

- ▶ a new system call to add two parameters

16 / 39

System Call Table

- ▶ append entry to file:
arch/x86/kernel/syscall_table_32.S

Example

.long sys_mycall

17 / 39

System Calls Header File

- ▶ add system call prototype to file:
include/linux/syscalls.h

Example

asmlinkage **int** sys_mycall(**int** i, **int** j);

18 / 39

Standard Header

- ▶ append entry for system call to file:
arch/x86/include/asm/unistd_32.h

Example

```
#define __NR_mycall 333
```

19 / 39

Implement System Call

Example (mycall/mycall.c)

```
#include <linux/kernel.h>

asmlinkage int sys_mycall(int i, int j)
{
    return i + j;
}
```

20 / 39

Modify Makefile

- ▶ add your directory to Makefile
- ▶ add a Makefile to your directory

Example (Makefile)

```
core-y += .... mycall/
```

Example (mycall/Makefile)

```
obj-y := mycall.o
```

21 / 39

Sample Test Program

Example

```
#include <stdio.h>
#include <sys/syscall.h>

#define __NR_mycall 333

int main(void)
{
    int x1=10, x2=20, y;

    y = syscall(__NR_mycall, x1, x2);
    printf("%d\n", y);
    return 0;
}
```

22 / 39

Another Example

Example

- ▶ a new system call to obtain the time in seconds which has passed since 1970

23 / 39

System Call Table

- ▶ append entry to file:
arch/x86/kernel/syscall_table_32.S

Example

```
.long sys_ptime
```

24 / 39

System Calls Header File

- ▶ add system call prototype to file:
include/linux/syscalls.h

Example

```
asmlinkage int sys_ptime(struct timeval *);
```

25 / 39

Standard Header

- ▶ append entry for system call to file:
arch/x86/include/asm/unistd_32.h

Example

```
#define __NR_ptime 334
```

26 / 39

struct timeval

Example

```
struct timeval
{
    long tv_sec;           /* seconds */
    long tv_usec;         /* microseconds */
};
```

27 / 39

Implement System Call

Example (ptime/ptime.c)

```
#include <linux/kernel.h>

asmlinkage int sys_ptime(struct timeval *tm)
{
    copy_to_user(tm, &xtime,
                 sizeof(struct timeval));
    return 0;
}
```

28 / 39

Modify Makefile

- ▶ add your directory to Makefile
- ▶ add a Makefile to your directory

Example (Makefile)

```
core-y += .... ptime/
```

Example (ptime/Makefile)

```
obj-y := ptime.o
```

29 / 39

Sample Test Program

Example

```
#include <stdio.h>
#include <linux/unistd.h>
#include <sys/syscall.h>
#include <sys/time.h>

#define __NR_ptime 334
```

30 / 39

Sample Test Program

Example

```
int main(void)
{
    struct timeval utime;
    int result;

    res = syscall(__NR_ptime, &utime);
    printf("%d\n", (int) utime.tv_sec);
    sleep(2);
    res = syscall(__NR_ptime, &utime);
    printf("%d\n", (int) utime.tv_sec);
    return 0;
}
```

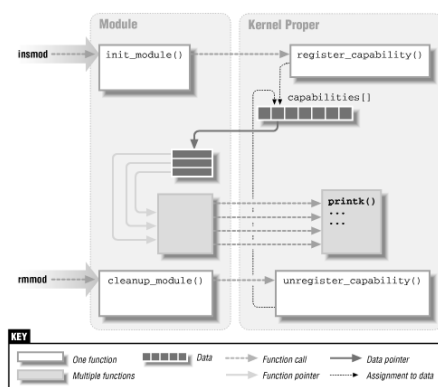
31 / 39

Modular Kernel

- ▶ monolithic architecture
- ▶ kernel extended at runtime
- ▶ easier and safer development

32 / 39

Module Registry



33 / 39

Modules vs Applications

- ▶ an application performs a task, a module registers itself in order to serve future requests
- ▶ an application runs in user-space, a module runs in kernel-space
 - ▶ an application can call external functions, a module can only call kernel functions

34 / 39

Module Example

Example (Hello world)

```
#include <linux/init.h>
#include <linux/module.h>
MODULE_LICENSE("Dual_BSD/GPL");

static int hello_init(void) { ... }

static void hello_exit() { ... }

module_init(hello_init);
module_exit(hello_exit);
```

35 / 39

Module Example

Example (Hello world)

```
static int hello_init(void)
{
    printk(KERN_ALERT "Hello ,_world!\n");
    return 0;
}

static void hello_exit()
{
    printk(KERN_ALERT "Goodbye ,_cruel_world!\n");
}
```

36 / 39

Kernel Symbol Table


- ▶ kernel symbol table contains addresses of global kernel symbols
- ▶ on loading:
 - ▶ link unresolved symbols in module to the kernel symbol table
 - ▶ exported module symbols become part of the kernel symbol table

37 / 39

Module Stacking

- ▶ modules can use symbols exported by other modules

Example



38 / 39

Reading Material

- ▶ Corbet
 - ▶ Chapter 2: **Building and Running Modules**

39 / 39