

KERNEL ARCHITECTURE

BLG413E – System Programming, Practice Session 2

Contents

- System Calls
- Kernel Modules

1-Adding a system call

Requirements (for Ubuntu OS): linux-source, kernel-package, fakeroot, libncurses5-dev (all of them are available in the provided Ubuntu 13.04 image)

Steps:

- extract linux source
- write new system call
- modify Makefiles
- modify system call table
- modify system call header file
- compile and install new kernel
- reboot to new kernel
- test new system call

Extracting linux source

- move linux source archive file (available in the provided Ubuntu 13.04 image) to desktop
 - `cd Desktop`
 - `sudo mv /usr/src/linux-source-3.8.0/linux-source-3.8.0.tar.bz2 linux-source-3.8.0.tar.bz2`
- and extract it
 - `tar -xjvf linux-source-3.8.0.tar.bz2`
- enter linux source folder
 - `cd linux-source-3.8.0`

Writing a system call

- mkdir mycall
- **mycall.c**: under /mycall

```
#include <linux/syscalls.h>
#include <linux/kernel.h>

asmlinkage int sys_mycall(int i, int j){
    return i + j;
}
```

Modifying Makefiles

- create **Makefile** under /mycall
 - write "obj-y := mycall.o" into this file
- modify Makefile under /linux-source-3.8.0 by adding "mycall/" to core-y

```
529 # Objects we will link into vmlinux / subdirs we need to visit
530 init-y      := init/
531 drivers-y   := drivers/ sound/ firmware/ ubuntu/
532 net-y       := net/
533 libs-y      := lib/
534 core-y      := usr/ mycall/
535 endif # KBUILD_EXTMOD
```

Modifying system call table and system call header files

- open arch/x86/syscalls/syscall_32.tbl
 - add "351 i386 mycall sys_mycall" to the end

```
359 350 i386      finit_module      sys_finit_module
360 351 i386      mycall             sys_mycall
```

- open include/linux/syscalls.h
 - add "asmlinkage int sys_mycall(int i, int j);" to the end of file before #endif

```
880 asmlinkage long sys_finit_module(int fd, const
881 asmlinkage int sys_mycall(int i, int j);
882 #endif
```

Before compiling linux kernel

- **Configuration Options (1st one recommended):**
 1. copy **config-3.8.13.6.txt** file to linux-source-3.8.0 and rename it as .config (**not config-3.8.13.6.config only .config without name**)
 2. make oldconfig (only ask for new or different options)
 3. make menuconfig (menu for configuration)

Compiling linux kernel

- copy **config-3.8.13.6.txt** file to linux-source-3.8.0 and rename it as .config
- make-kpkg clean → cleans up all from previous kernel compiles
- **Compilation** (**Warning: It generally takes 1-2 hours**):
fakeroot make-kpkg --initrd --append-to-version=-custom
kernel_image kernel_headers
- **Output:** two files in parent directory (i.e., Desktop):
 - linux-image-3.8.13.6-custom....deb
 - linux-headers-3.8.13.6-custom....deb

Installing compiled kernel

- `sudo dpkg -i linux-image-3.8.13.6-custom....deb`
- `sudo dpkg -i linux-headers-3.8.13.6-custom....deb`
- Then reboot to open from the new kernel:
 - `sudo reboot`

Testing new system call

- A simple C program using our new system call to add 2 numbers and printing out the result

```
#include <stdio.h>
#define NR_mycall 351


int main (void) {
    int x1=10, x2=20, y;
    y = syscall(NR_mycall, x1, x2);
    printf("%d\n", y);
    return 0 ;
}
```

Uninstalling compiled kernel

- When you need to recompile the kernel, **first boot from the original kernel** and uninstall the kernel you have compiled before by using following commands
 - `sudo dpkg -r linux-image-3.8.13.6-custom`
 - `sudo dpkg -r linux-headers-3.8.13.6-custom`

2-Kernel modules

- A way to add new features to the kernel without rebuilding it.
- Unlike applications, modules register themselves for serving future requests.
- Applications can access the capabilities of a module through system calls.



http://www.xml.com/ldd/chapter/book/figs/ldr2_0201.gif

An example module: hello

- **hello.c:**

```
#include <linux/init.h> /* for module_init and module_exit */
#include <linux/module.h> /* needed by all modules */
MODULE_LICENSE("Dual BSD/GPL"); /* a macro to declare that this module is open source */

static int hello_init(void) /* static: invisible outside the module */
{
    /* to avoid namespace pollution */
    printk(KERN_ALERT "Hello, world\n"); /* printk: kernel print function (macros for priority) */
    return 0; /* KERN_ALERT: a situation requiring immediate action */
}

static void hello_exit(void)
{
    printk(KERN_ALERT "Goodbye, cruel world\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

- **Makefile:**

obj-m := hello.o

all:

make -C /lib/modules/\$(shell uname -r)/build M=\$(PWD) modules

M=\$(PWD) is to build external module in the working directory



Using hello module

- **Compiling:**
 - *make*
- **Loading** (check with ***dmesg*** which is used to write the kernel messages):
 - *sudo insmod ./hello.ko*
- **Unloading** (check with ***dmesg***):
 - *sudo rmmod hello*
- check with ***lsmod*** (which prints the contents of the `/proc/modules` file) before and after loading and unloading

An example module using load time parameters

- **hello.c:**

```
/* $Id: hello.c,v 1.4 2004/09/26 07:02:43 gregkh Exp $ */
#include <linux/init.h>
#include <linux/module.h>
#include <linux/moduleparam.h> /* to enable passing parameters at loadtime */
MODULE_LICENSE("Dual BSD/GPL");

/* A couple of parameters that can be passed in: how many times we say hello, and to whom */
static char *whom = "world";
static int howmany = 1;
module_param(howmany, int, S_IRUGO); /* S_IRUGO: read by the world but cannot be changed */
module_param(whom, charp, S_IRUGO);

static int hello_init(void){
    int i;
    for (i = 0; i < howmany; i++)
        printk(KERN_ALERT "(%d) Hello, %s\n", i, whom);
    return 0;
}

static void hello_exit(void){
    printk(KERN_ALERT "Goodbye, cruel world\n");
}

module_init(hello_init);
module_exit(hello_exit);
```


Specifying module parameters

- `sudo insmod ./hellop.ko whom='Mom' howmany=4`
- `dmesg`

```
4555.764793] (0) Hello, Mom  
4555.764796] (1) Hello, Mom  
4555.764797] (2) Hello, Mom  
4555.764798] (3) Hello, Mom
```

- `sudo rmmod hellop`
- `dmesg`

```
4555.764793] (0) Hello, Mom  
4555.764796] (1) Hello, Mom  
4555.764797] (2) Hello, Mom  
4555.764798] (3) Hello, Mom  
4611.350208] Goodbye, cruel world
```