

Image Segmentation

Recap from last time

- Samples not squares
- Sensors are not perfect
- Quantization hurts
- Questions?

Overview

- What is image segmentation?
- Thresholding and thresholding algorithms
- Performance and the ROC curve
- Connectivity and connected components
- Region growing
- Split and merge algorithms
 - The region adjacency graph



©1994 Encyclopaedia Britannica, Inc.




Figure 2.9(a-g)

Gestalt Phenomena


- Figure-ground
- Proximity
- Similarity
- Continuation
- Closure
- Common fate
- Symmetry

What is Image Segmentation?

- It partitions an image into regions of interest.
- The first stage in many automatic image analysis systems.
- A *complete segmentation* of an image I is a finite set of regions R_1, \dots, R_N , such that

$$I = \bigcup_{i=1}^N R_i \text{ and } R_i \cap R_j = \emptyset \forall i \neq j.$$

How should I segment this?




How should I segment this?



Segmentation Quality

- The quality of a segmentation depends on what you want to do with it.
- Segmentation algorithms must be chosen and evaluated with an application in mind.


Segmentation example






Berkeley Segmentation Database and Benchmark



<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>



Berkeley Segmentation Database and Benchmark



Thresholding

- Thresholding is a simple segmentation process.
- Thresholding produces a binary image B .
- It labels each pixel **in** or **out** of the region of interest by comparison of the greylevel with a threshold T :

$$B(x, y) = \begin{cases} 1 & \text{if } I(x, y) \geq T \\ 0 & \text{if } I(x, y) < T. \end{cases}$$

Basic Thresholding Algorithm


```
for x=1:X
    for y=1:Y
        B (x, y)  =  ( I (x, y)  >=  T ) ;
    end
end
```

Don't write it like this at home!

Thresholding example



T=50



$T=100$



$T=150$



$T=200$




How do we choose T?

- Trial and error
- Compare results with ground truth
- Automatic methods

Segmentation Performance

- To use automatic analysis, we need to know the true classification of each test.
- We need to do the segmentation by hand on some example images.

Ground truth



ROC Analysis

(ROC = Receiver operating characteristic)

- An ROC curve characterizes the performance of a binary classifier.
- A binary classifier distinguishes between two different types of thing. E.g.:
 - Healthy/afflicted patients – cancer screening
 - Pregnancy tests
 - Foreground/background image pixels
 - Object detection

Classification Error

- Binary classifiers make errors.
- Two types of input to a binary classifier:
 - Positives
 - Negatives
- Four possible outcomes in any test:


True positive	False negative	} P: Total # positives
True negative	False positive	


Classified: correctly | incorrectly

Where do errors come from?

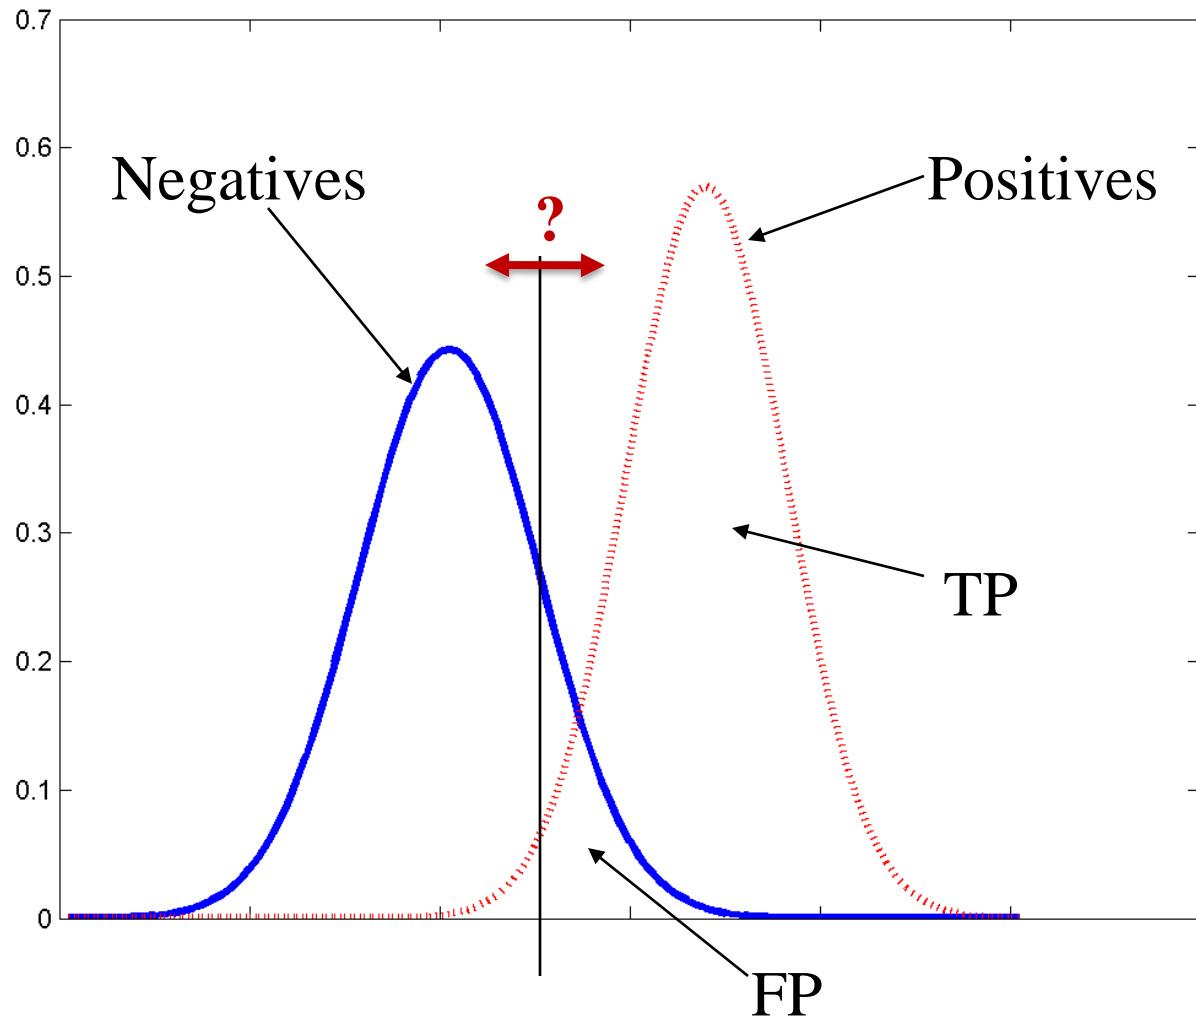
- A binary classifier thresholds a measurement made on every case.
- The measurement may be explicit
 - Example: pixel grey-level: **1 scalar**
- Or implicit
 - The degree of certainty of a doctor examining a patient:
% confidence for EACH scalar → probability distribution

Wouldn't it be nice...





Measurement Distributions



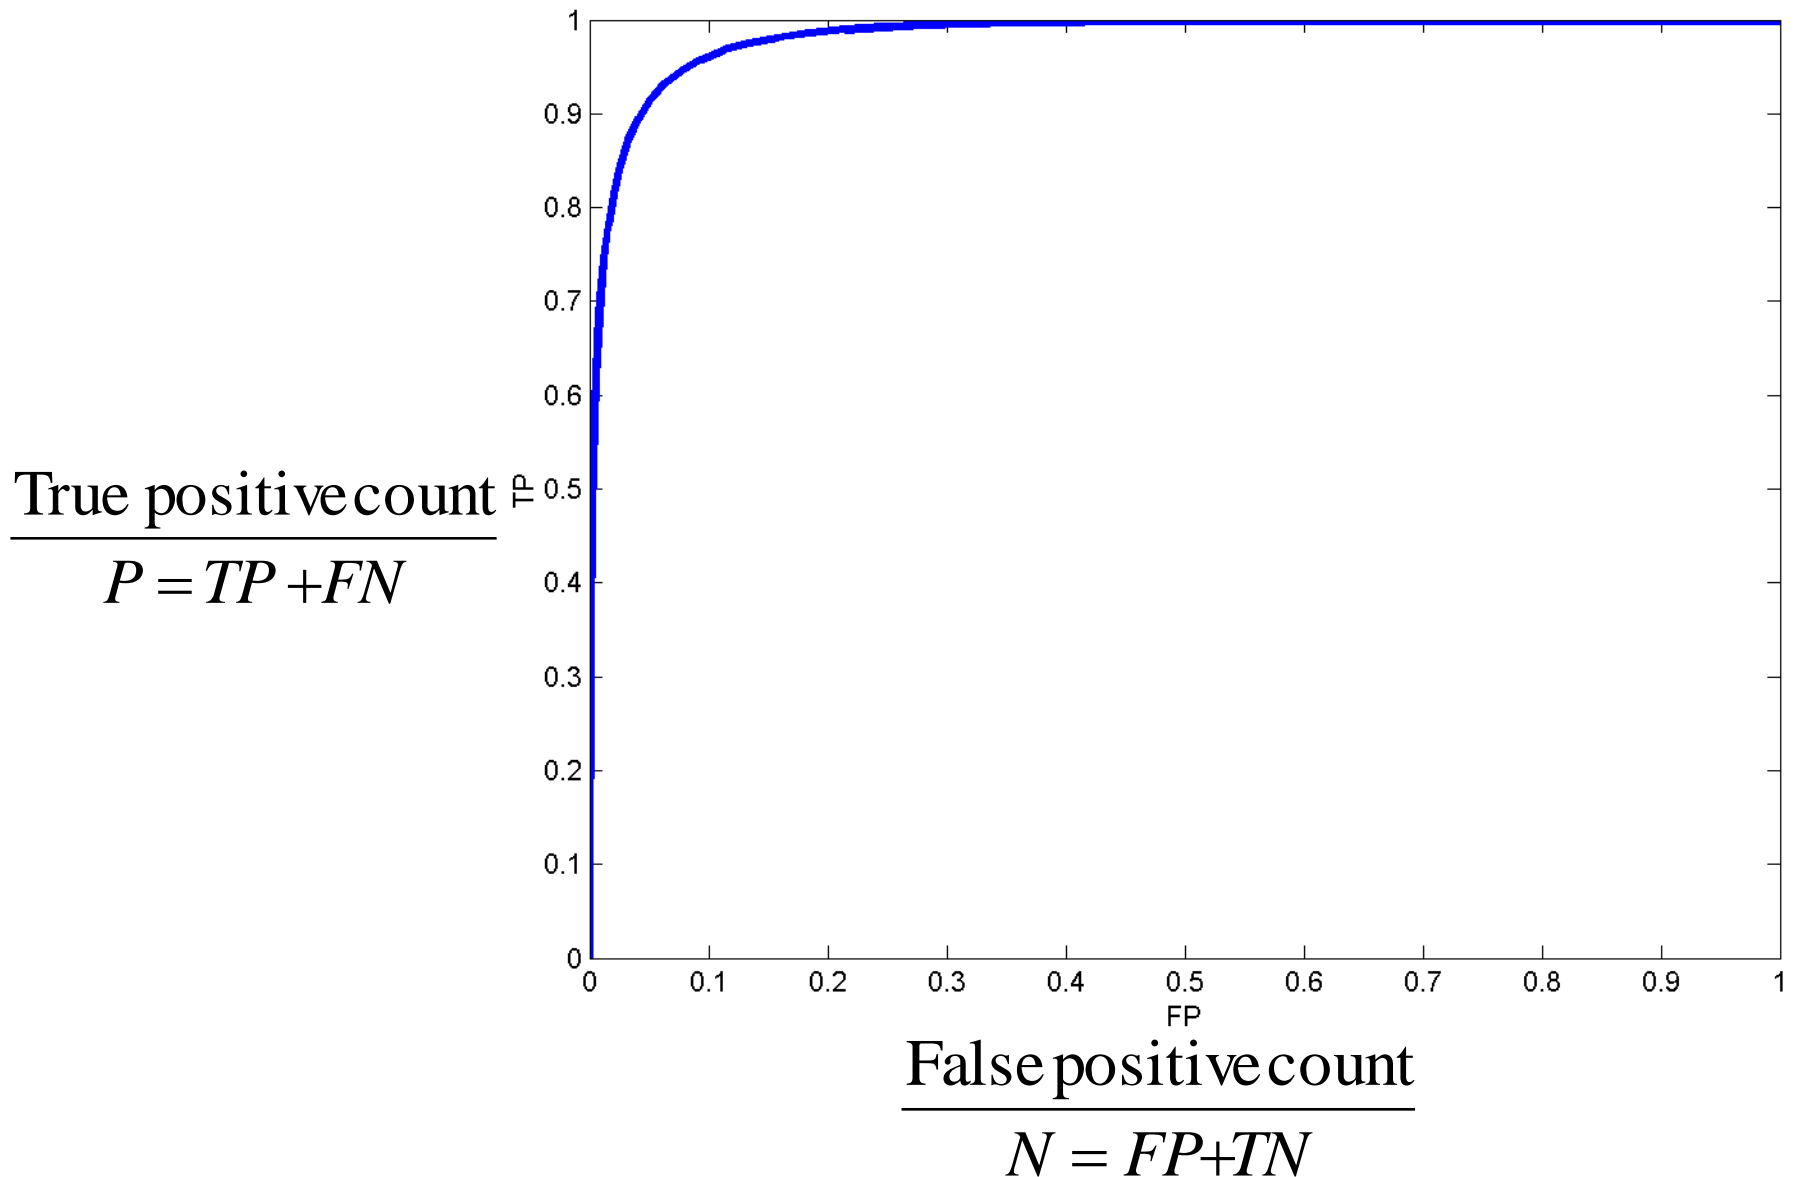
Classification outcomes

- As we change the threshold, FP and TP change.
- Notice that:
 - $FP + TN = N$ (the total number of negatives)
 - $TP + FN = P$ (total positives)
- How to evaluate performance?

The ROC curve

- Characterizes the error trade-off in binary classification tasks.
- It plots the TP fraction against FP fraction.
- TP fraction (*sensitivity*) is $\frac{\text{True positive count}}{P}$
- FP fraction (*1-specificity*) is $\frac{\text{False positive count}}{N}$

The ROC Curve



Properties of ROC curves

- An ROC curve always passes through $(0,0)$ and $(1,1)$.
- What is the ROC curve of a perfect system?
- What if the ROC curve is a straight line from $(0,0)$ to $(1,1)$?

Area under the ROC curve


- The area A under the curve measures overall classification performance.
- If the distributions of measurements on positive and negative cases are $N(\mu_P, \sigma_P)$ and $N(\mu_N, \sigma_N)$, respectively, then

$$A = Z\left(\frac{|\mu_P - \mu_N|}{(\sigma_P^2 + \sigma_N^2)^{1/2}}\right)$$

- z is the cumulative normal distribution function

$$Z(x) = \int_{-\infty}^x z(t)dt$$

Auto-select a threshold?




Operating points

- Choose an *operating point* by assigning relative costs and values to each outcome:
 - V_{TN} – value of true negative
 - V_{TP} – value of true positive
 - C_{FN} – cost of false negative
 - C_{FP} – cost of false positive
- Choose the point on the ROC curve with **gradient**

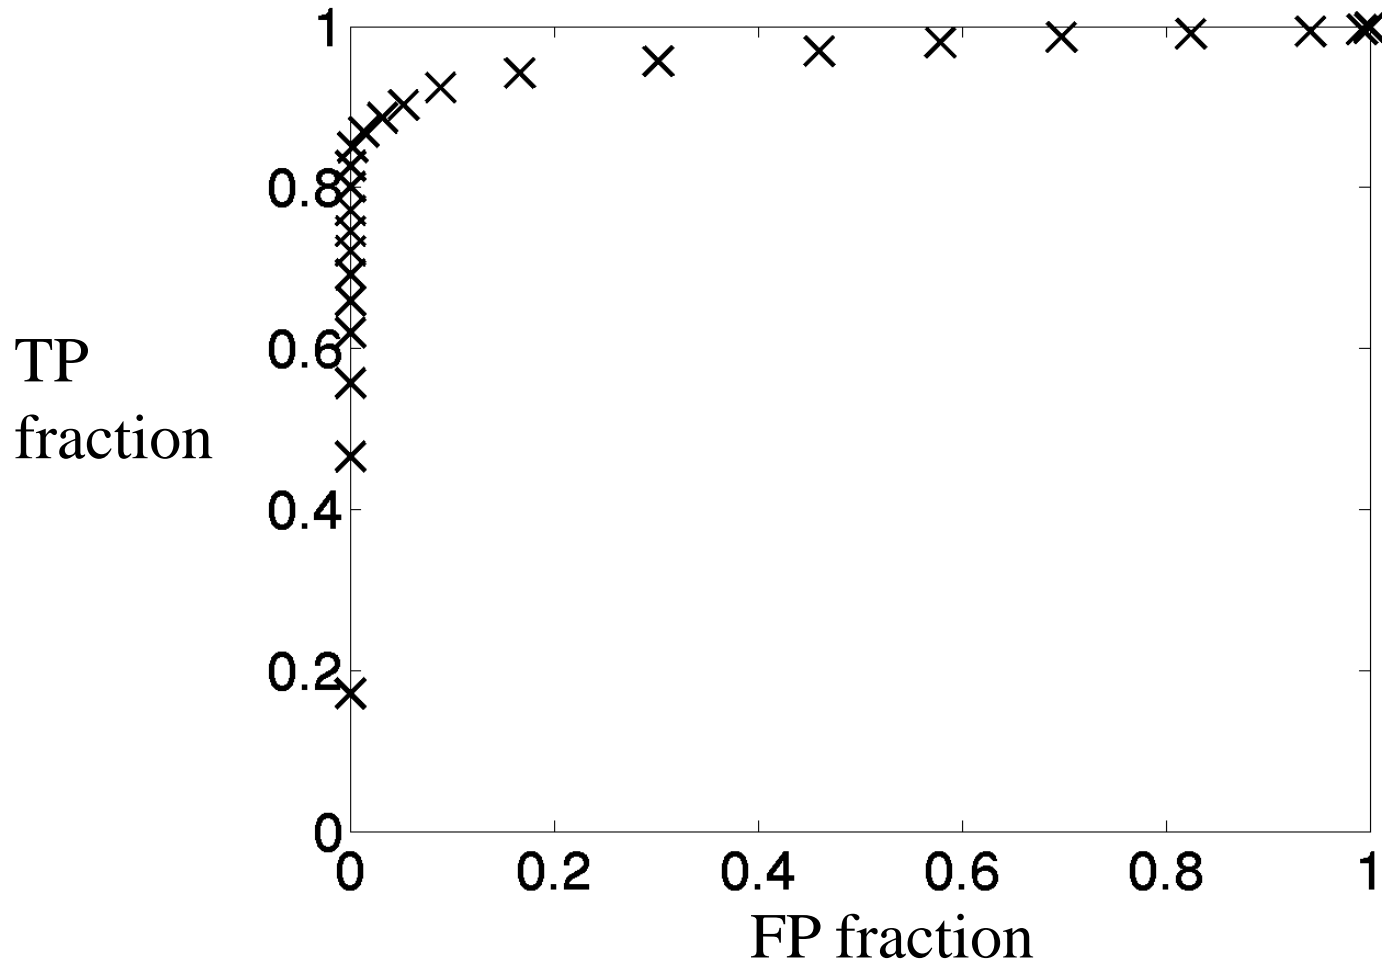
$$\beta = \frac{N V_{TN} + C_{FP}}{P V_{TP} + C_{FN}}$$

- For simplicity, we often set $V_{TN} = V_{TP} = 0$.



Classification outcomes





ROC curve



Greylevel Histograms



Positives and Negatives



Automatic Threshold Selection

- Try to find a threshold between histogram peaks
- Maybe...model positive and negative region histograms by Gaussian distributions

K-means Clustering

Initialize R_f and R_b

REPEAT

% Find the mean greylevel in each region.

$m_f = \text{mean}(I(R_f));$

$m_b = \text{mean}(I(R_b));$

% Pick threshold half way between.

$T = (m_f + m_b) / 2;$

% Find new regions.


$R_f = \{ (x, y) \mid I(x, y) \geq T \};$

$R_b = I \setminus R_f;$


UNTIL T remains unchanged.

Improvements?


- Can you extend this to use the standard deviations of the regions as well?
- What problems might you encounter?
- Can one take advantage of special situations?



Gradient.illusion.arp.jpg




(a)




(b)




(c)





(d)



(e)




(f)




(d)




(e)




(f)



(g)



(h)



(i)



(j)

Recap of Thresholding (only)


- Thresholding for binary segmentation
- Choosing the threshold can be difficult
- ROC analysis characterizes binary classification systems
- Can help choose a good threshold
- Simple algorithms can find good thresholds sometimes.

Note on Performance Assessment


- In real-life, we use two or even three separate sets of test data:
 - A *training set*, for tuning the algorithm.
 - An unseen *test set* to get a final performance score on the tuned algorithm.
 - A *validation* set for verifying the performance score obtained on the test set.

Limitations of Thresholding

- Why can we segment images much better by eye than through thresholding processes?



by [Adrian Pingstone](#), based on the [original](#) created by Edward H. Adelson



by [Adrian Pingstone](#), based on the [original](#) created by Edward H. Adelson

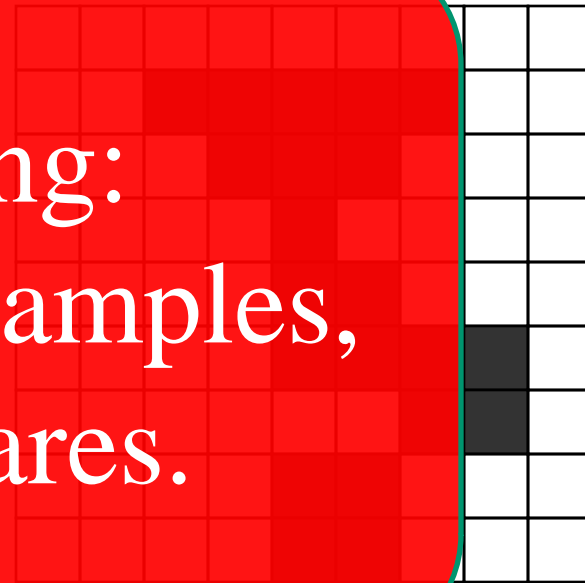
Limitations of Thresholding

- Why can we segment images much better by eye than through thresholding processes?
- We might improve results by considering image **context**: Surface Coherence

Pixel connectivity


- We need to define which pixels are neighbours.
- Are the dark pixels in this array connected?

Warning:
Pixels are samples,
not squares.




Pixel connectivity


- We need to define which pixels are neighbours.
- Are the dark pixels in this array connected?



Pixel Neighbourhoods



4-neighbourhood



8-neighbourhood

Pixel paths


- A 4-connected path between pixels p_1 and p_n is a set of pixels $\{p_1, p_2, \dots, p_n\}$ such that p_i is a 4-neighbour of p_{i+1} , $i=1, \dots, n-1$.
- In an 8-connected path, p_i is an 8-neighbour of p_{i+1} .

Connected regions

- A region is 4-connected if it contains a 4-connected path between any two of its pixels.
- A region is 8-connected if it contains an 8-connected path between any two of its pixels.


Connected regions

- Now what can we say about the dark pixels in this array?
- What about the light pixels?



Connected components labelling


- Labels each connected component of a binary image with a separate number.



1	1	1	1	1	1	1	1	1
1	1	2	2	1	3	1	1	1
1	1	1	2	2	1	1	1	1
1	1	1	1	1	1	1	1	1
1	4	1	1	5	5	1	1	1
1	1	1	1	5	5	5	1	1
6	6	1	1	1	1	5	5	1
7	6	1	1	8	8	1	1	1
7	6	1	1	8	8	1	1	1

Foreground labelling

- Only extract the connected components of the foreground



1	1	1	1	1	1	1	1	1
1	1	0	0	1	0	1	1	1
1	1	1	0	0	1	1	1	1
1	1	1	1	1	1	1	1	1
1	0	1	1	0	0	1	1	1
1	1	1	1	0	0	0	1	1
0	0	1	1	1	1	0	0	1
2	0	1	1	0	0	1	1	1
2	0	1	1	0	0	1	1	1

Connected components

% B is the binary image input. L is the labelled
% image output.

```
function L = ConnectedComponents(B)
```

```
[X,Y] = size(B);
```

```
L = zeros(X,Y);
```

```
n=1;
```

```
For each (x,y) in B
```

```
    if (B(x,y) & L(x,y)==0)
```

```
        label(x,y,n,B,L);
```

```
        n=n+1;
```

```
    end
```

```
end
```


Connected components

```
function label(x_start, y_start, n, B, L)
% Recursively give label n to this pixel
% and all its foreground neighbours.
L(x_start,y_start)=n;
For each (x,y) in N(x_start, y_start)
    if (L(x,y)==0 & B(x, y))
        label(x,y,n,B,L);
    end
end
```

Goose detector



Goose detector



Goose 4-components (26)




Goose Error




GV12/3072
Image Processing.

Goose 8-components (22)



Connected components

- What happens if we use the connected components algorithm on this image?



- How might we improve our implementation?


Region Growing

- Start from a seed point or region.
- Add neighbouring pixels that satisfy the criteria defining a region.
- Repeat until we can include no more pixels.

Region Growing

```
function B = RegionGrow(I, seed)
```

```
    [X,Y] = size(I);  
    visited = zeros(X,Y);  
    visited(seed) = 1;  
    boundary = emptyQ;  
    boundary.enQ(seed);
```




```
        while (~boundary.empty())  
            nextPoint = boundary.deQ();  
            if(include(nextPoint))  
                visited(nextPoint) = 2;  
                Foreach (x,y) in N(nextPoint)  
                    if(visited(x,y) == 0)  
                        boundary.enQ(x,y);  
                        visited(x,y) = 1;  
                    end  
                end  
            end  
        end  
    end
```

Region Growing example


- Pick a single seed pixel.
- Inclusion test is a simple grey level threshold:

```
function test = include(p)
    test = (p >= T);
```



Seed pixel

T=50



$T=100$




$T=150$



ROC Curve

- Region grower
- × Thresholding



Implementation

- The region-growing algorithm above uses a *breadth-first search*.
- The connected-components algorithm above uses a *depth-first search*.
- Both algorithms can use either search procedure.
- Breadth-first search has more robust performance.

Variations

- Seed selection
- Inclusion criteria
- Boundary constraints and snakes

Seed selection

- Point and click seed point.
- Seed region
 - By hand
 - Automatically, e.g., from a conservative thresholding.
- Multiple seeds
 - Automatically labels the regions

Inclusion criteria

- Greylevel thresholding
- Greylevel distribution model
 - Use mean μ and standard deviation σ in seed region:
 - Include if $(I(x, y) - \mu)^2 < (n\sigma)^2$. Eg: $n = 3$.
 - Can update the mean and standard deviation after every iteration.
- Colour or texture information.

Snakes

- A snake is an *active contour*.
- It is a polygon, i.e., an ordered set of points joined up by lines.
- Each point on the contour moves away from the seed while its image neighbourhood satisfies an inclusion criterion.
- Often the contour has smoothness constraints.

Snakes

- The algorithm iteratively minimizes an energy function:
- $E = E_{\text{tension}} + E_{\text{stiffness}} + E_{\text{image}}$
- See Kass, Witkin, Terzopoulos, IJCV 1988

Example



Example program...

Split and Merge Algorithms

- This is a *global* rather than *binary* segmentation algorithm

Gray Level Variance

$$\frac{1}{N - 1} \sum_{(r,c) \in Region} [I(r, c) - \bar{I}]^2$$

Split and Merge Algorithms

- This is a *global* rather than *binary* segmentation algorithm
- Recursively divide the image into homogeneous regions
- Merge adjacent regions that combine to a homogeneous region.

Splitting


```
% Splits an image I into homogeneous regions  
% labelled in L.  
  
function L = ImageSplit(I)  
  
% Initialize.  
  
[X,Y] = size(I);  
L = ones(X,Y);  
n = 2;  
  
% Call recursive splitting subroutine  
L = split(I, L, 1, X, 1, Y, n)
```

```

function L = split(I, L, xmin, xmax, ymin, ymax, n)
if(~homogeneous(I(xmin:xmax, ymin:ymax)))
    xSplit = (xmin+xmax)/2;    ySplit = (ymin+ymax)/2;
    % Top left
    L = split(I, L, xmin, xSplit, ymin, ySplit, n);
    % Top right
    L((xSplit+1):xmax, ymin:ySplit) = n;
    n = n + 1;
    L = split(I, L, (xSplit+1), xmax, ymin, ySplit,
n);
    % Bottom left
    L(xmin:xSplit, (ySplit+1):ymax) = n;
    n = n + 1;
    L = split(I, L, xmin, xSplit, (ySplit+1), ymax,
n);
    % Bottom right... similar.
end

```



Splitting Example



Merging

- Build the *Region Adjacency Graph* (RAG).
- Merge adjacent regions that are homogeneous.
- Update RAG.
- Repeat to convergence


RAG Example




Computing the RAG

```
function rag = RegionAdjacencyGraph(L)
numRegions = max(max(L)); [X,Y] = size(L);
rag = zeros(numRegions, numRegions);
For each (x1, y1) in L
    r1 = L(x1, y1);
    For each (x2, y2) in N(x1, y1)
        r2 = L(x2, y2);
        if(r1 ~= r2)
            rag(r1, r2) = rag(r2, r1) = 1;
        end
    end
end
end
```

RAG Example



274 Regions



Region Merging

```
numRegions = max(max(L));  
done = 0;  
while(~done)  
    done = 1;  
    for i=1:(numRegions-1)  
        for j=(i+1):numRegions  
            if(rag(i,j))  
                combinedR = I(find(L == i | L == j));  
                if(homogeneous(combinedR))  
                    [L, rag] = merge(i, j, rag, L);  
                    done = 0;  
                end  
            end  
        end  
    end  
end
```

... end


Region Merging

```
function [L, rag] = merge(label1, label2, rag, L)
% Merge the regions in the labelled image.
L(find(L == label2)) = label1;
% Combine the adjacency of the two regions.
rag(label1,:) = rag(label1,:) | rag(label2,:);
rag(:,label1) = rag(:,label1) | rag(:,label2);
% Make sure the combined region1 is not self-adjacent.
rag(label1, label1) = 0;
% Remove all adjacency to region2
rag(label2,:) = 0;
rag(:,label2) = 0;
```


Re-labelling

- Final step relabels L .
- Change each integer region label so that the set of labels is a contiguous set of labels starting at 1.


Merging Example - Before




274 Regions



Merging Example - After



6 Regions



Split and merge example



Split and merge example

- Regions are homogeneous if
 $\text{Max. greylevel} - \text{min. greylevel} < T$.
- $T = 20$ for split.
- $T = 100$ for merge.

Split and merge example




4087 regions

Split and merge example



135 regions

Split and merge example



Problems and Improvements

- Inefficiency
- Alternative data structures?
- Consistency.
- Homogeneity criteria
 - Statistical region similarity measures

But if this pixel is foreground...

Markov Random Fields

- In a Markov Random Field, the label in each pixel depends on the label in neighbouring pixels.
- We assign a probability to each configuration of pixels in a neighbourhood.
- We maximize the probability of the segmentation given the observed image.


Neighborhood Probability

- Simple model for 4-neighbourhood.


- $$P(x_1, x_2) = \begin{cases} 0.99 & x_1 = x_2 \\ 0.01 & x_1 \neq x_2 \end{cases}$$

- $$\begin{aligned} P(x | N(x)) &= P(x, x_1, x_2, x_3, x_4) \\ &= P(x, x_1)P(x, x_2)P(x, x_3)P(x, x_4) \end{aligned}$$


MRF Segmentation: Iter. Cond. Modes




Image




Initial segmentation




ICM Iter 1



ICM Iter 2



ICM Iter 4



ICM Iter 8

Model vs. Algorithm

- Find the set of labels that maximizes:

$$P(X | I) = \prod P(y | x)P(x)P(x | N(x))$$

- $L(X) = \sum \log P(y | x) + \log P(x) + \log P(x | N(x))$
- Initialize with thresholded image
- Iteratively replace each label if it increases $L(I)$. “Greedy” algorithm.
- MCMC, Simulated Annealing, or Graph Cuts better.

MRF Extension

- Learn the within class probabilities using an EM algorithm (see Machine Vision class).
- Extends naturally to more than two classes.

Summary

- We have looked at several segmentation algorithms:
 - Thresholding and connected components.
 - Region growing
 - Split and merge
 - Snakes and MRFs
- Used ROC analysis to compare the performance of simple systems.

Summary

- Segmentation is hard!
- But it is easier if you know what you are doing!
 - Is the segmentation task binary or global?
 - What are the regions of interest?
 - How accurately must the algorithm locate the region boundaries?
- Research problems remain!

Kanizsa Triangle

