

---

# Computer Networks Basic Protocols

Assoc. Prof. Dr. Berk CANBERK

**15 November 2017**

**-Transport Layer-**


## References:

*-Data and Computer Communications*, William Stallings, Pearson-Prentice Hall, 9<sup>th</sup> Edition, 2010.

*-Computer Networking, A Top-Down Approach Featuring the Internet*, James F.Kurose, Keith W.Ross, Pearson-Addison Wesley, 6<sup>th</sup> Edition, 2012.

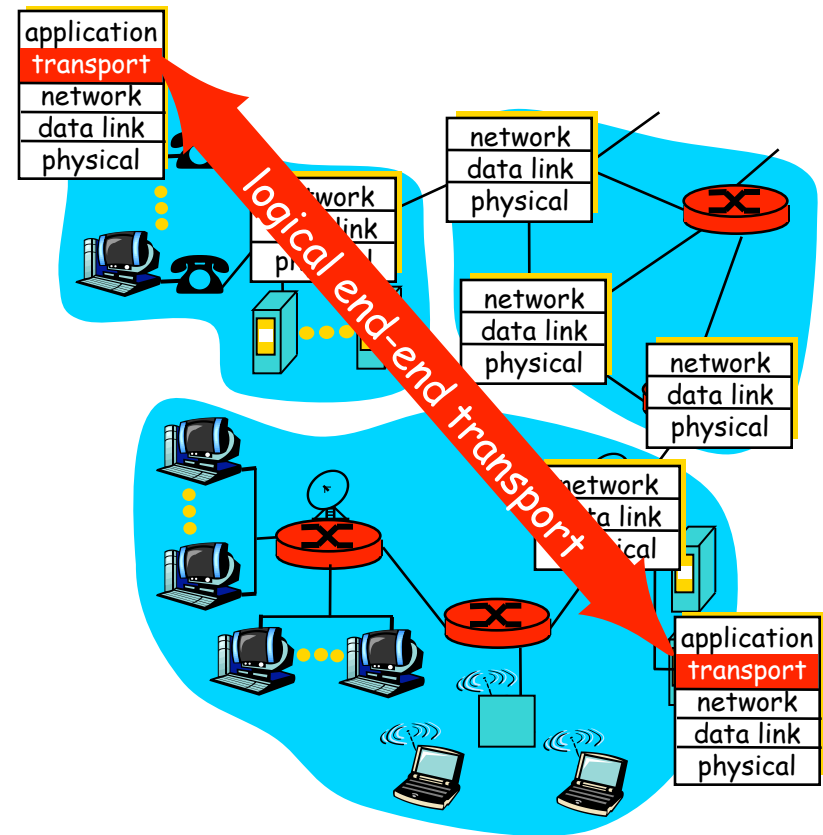
---

Transport layer provides end-to-end connectivity across the network



# Transport services and protocols

- provide *logical communication* between app processes running on different hosts
- transport protocols run in end systems
  - send side: breaks app messages into **segments**, passes to network layer
  - rcv side: reassembles segments into messages, passes to app layer
- more than one transport protocol available to apps
  - Internet: TCP and UDP



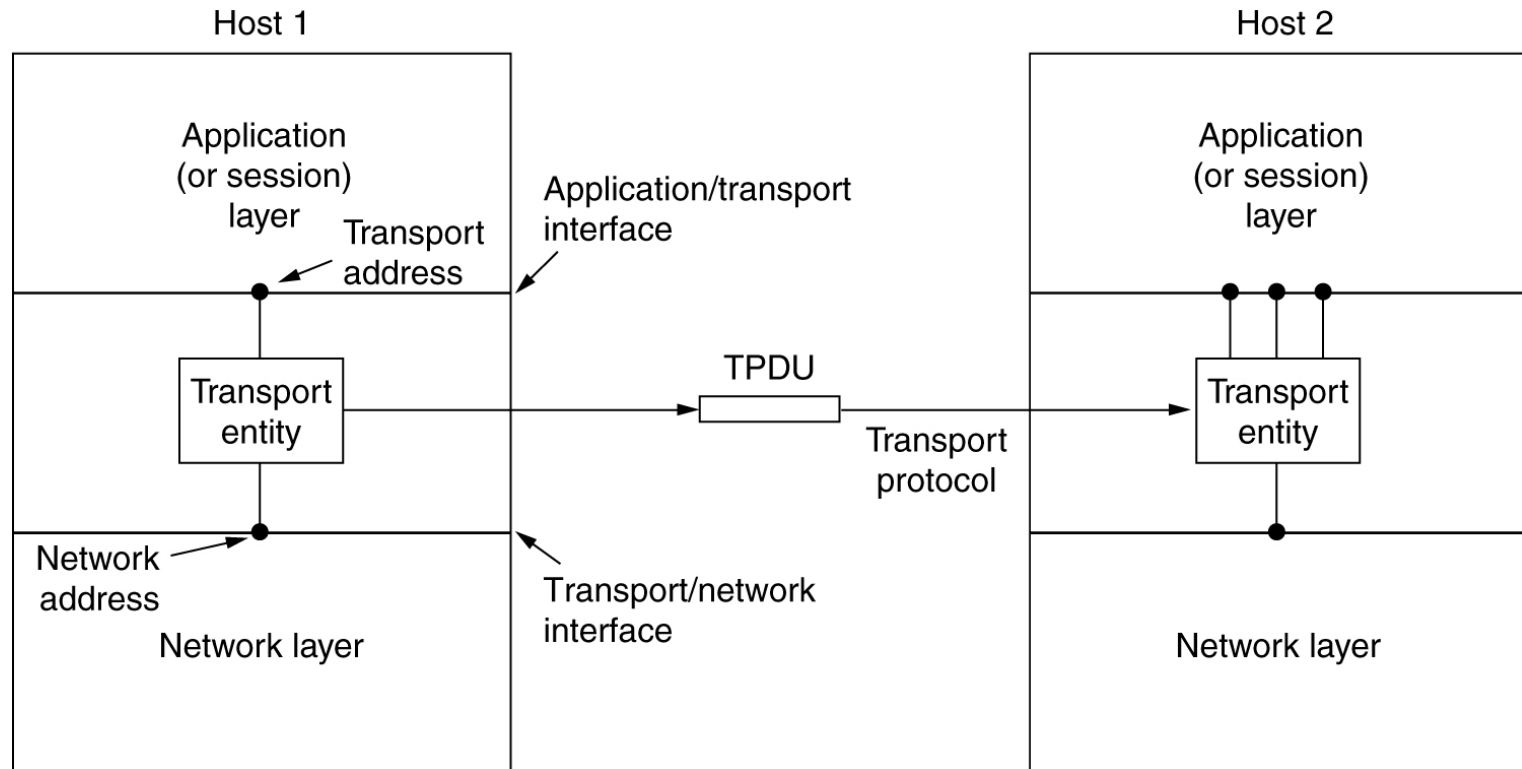
# Transport vs. network layer

---

- *Network layer:* logical communication between hosts
  - Not reliable, no control on network layer
- *Transport layer:* logical communication between processes
  - relies on, enhances, network layer services
  - Efficient reliable service to users (processes in the application layer)
  - Addressing
  - Multiplexing
  - Flow (Congestion) Control
  - Connection establishment and termination

# Services Provided to Upper Layers


The network, transport, and application layers.



# Transport Service Primitives

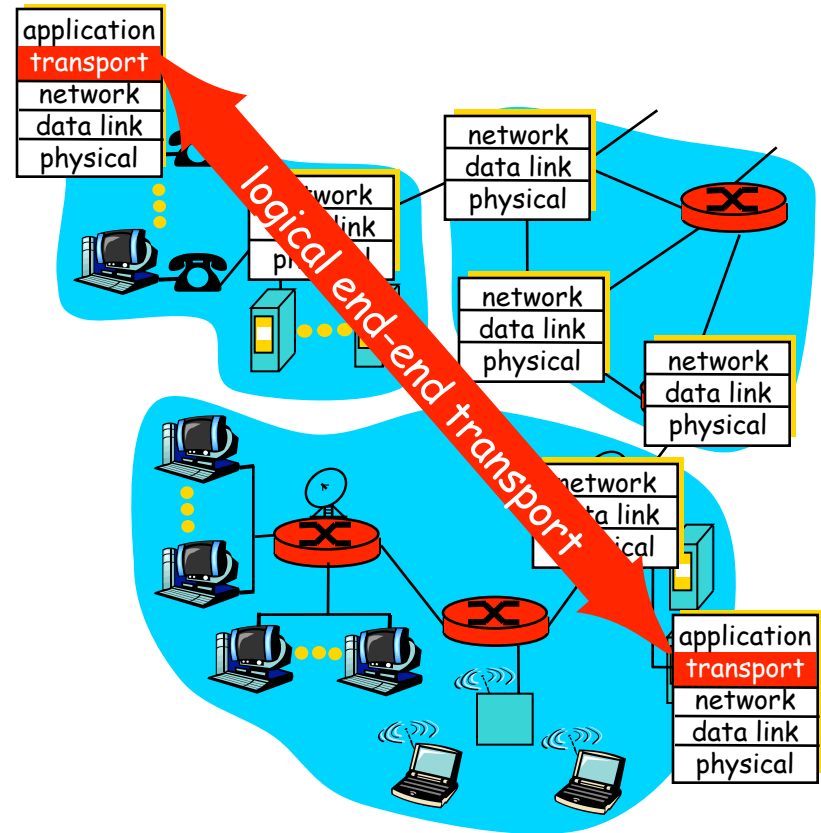
---

The nesting of TPDUs, packets, and frames.



# Internet transport-layer protocols

- reliable, in-order delivery (TCP)
  - flow/congestion control
  - connection setup
- unreliable, unordered delivery: UDP
  - extension of “best-effort” IP
- services **not** available:
  - delay guarantees
  - bandwidth guarantees



# Comparison of Internet Transport Protocols

TCP (Streams)	UDP (Datagrams)
Connections	Datagrams
Bytes are delivered once, reliably, and in order	Messages may be lost, reordered, duplicated
Arbitrary length content	Limited message size
Flow control matches sender to receiver	Can send regardless of receiver state
Congestion control matches sender to network	Can send regardless of network state



# UDP: User Datagram Protocol

## [RFC 768]

---

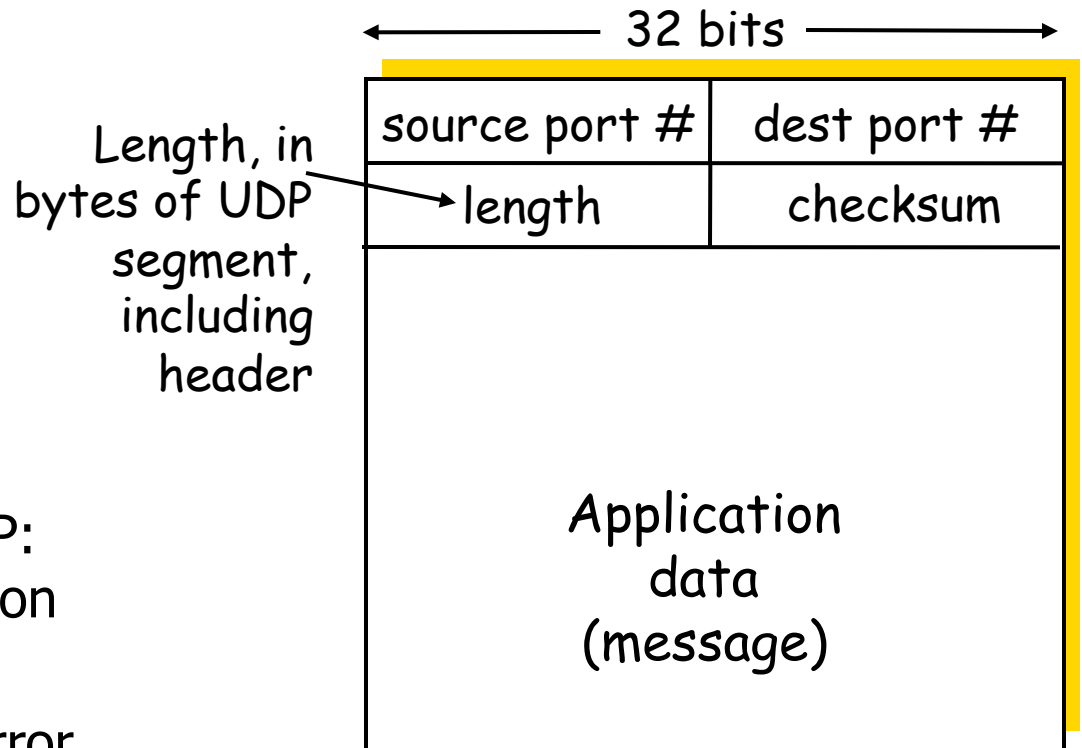
- “best effort” service, UDP segments may be:
  - lost
  - delivered out of order to app
- *connectionless*:
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others

### Why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small segment header
- no congestion control: UDP can blast away as fast as desired

# UDP

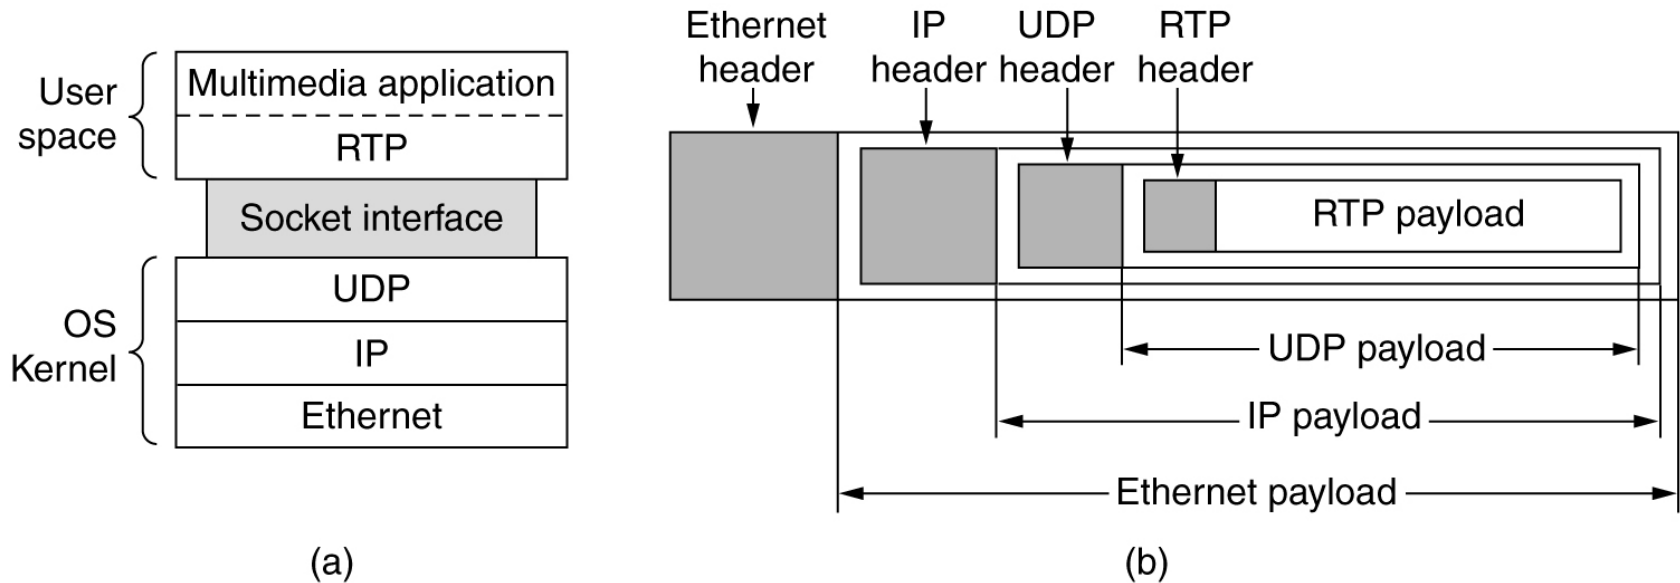
- often used for streaming multimedia apps
  - loss tolerant
  - rate sensitive
- other UDP uses
  - DNS
  - SNMP
- reliable transfer over UDP: add reliability at application layer
  - application-specific error recovery!



UDP segment format

# The Real-Time Transport Protocol

- (a) The position of RTP in the protocol stack.  
(b) Packet nesting.



# Transport Control Protocol (TCP)

RFCs: 793, 1122, 1323, 2018, 2581

## point-to-point:

one sender, one receiver

## reliable, in-order *byte stream*:

no "message boundaries"

## pipelined:

TCP congestion and flow control set window size

## *send & receive buffers*

## ■ full duplex data:


- bi-directional data flow in same connection
- MSS: maximum segment size

## ■ connection-oriented:

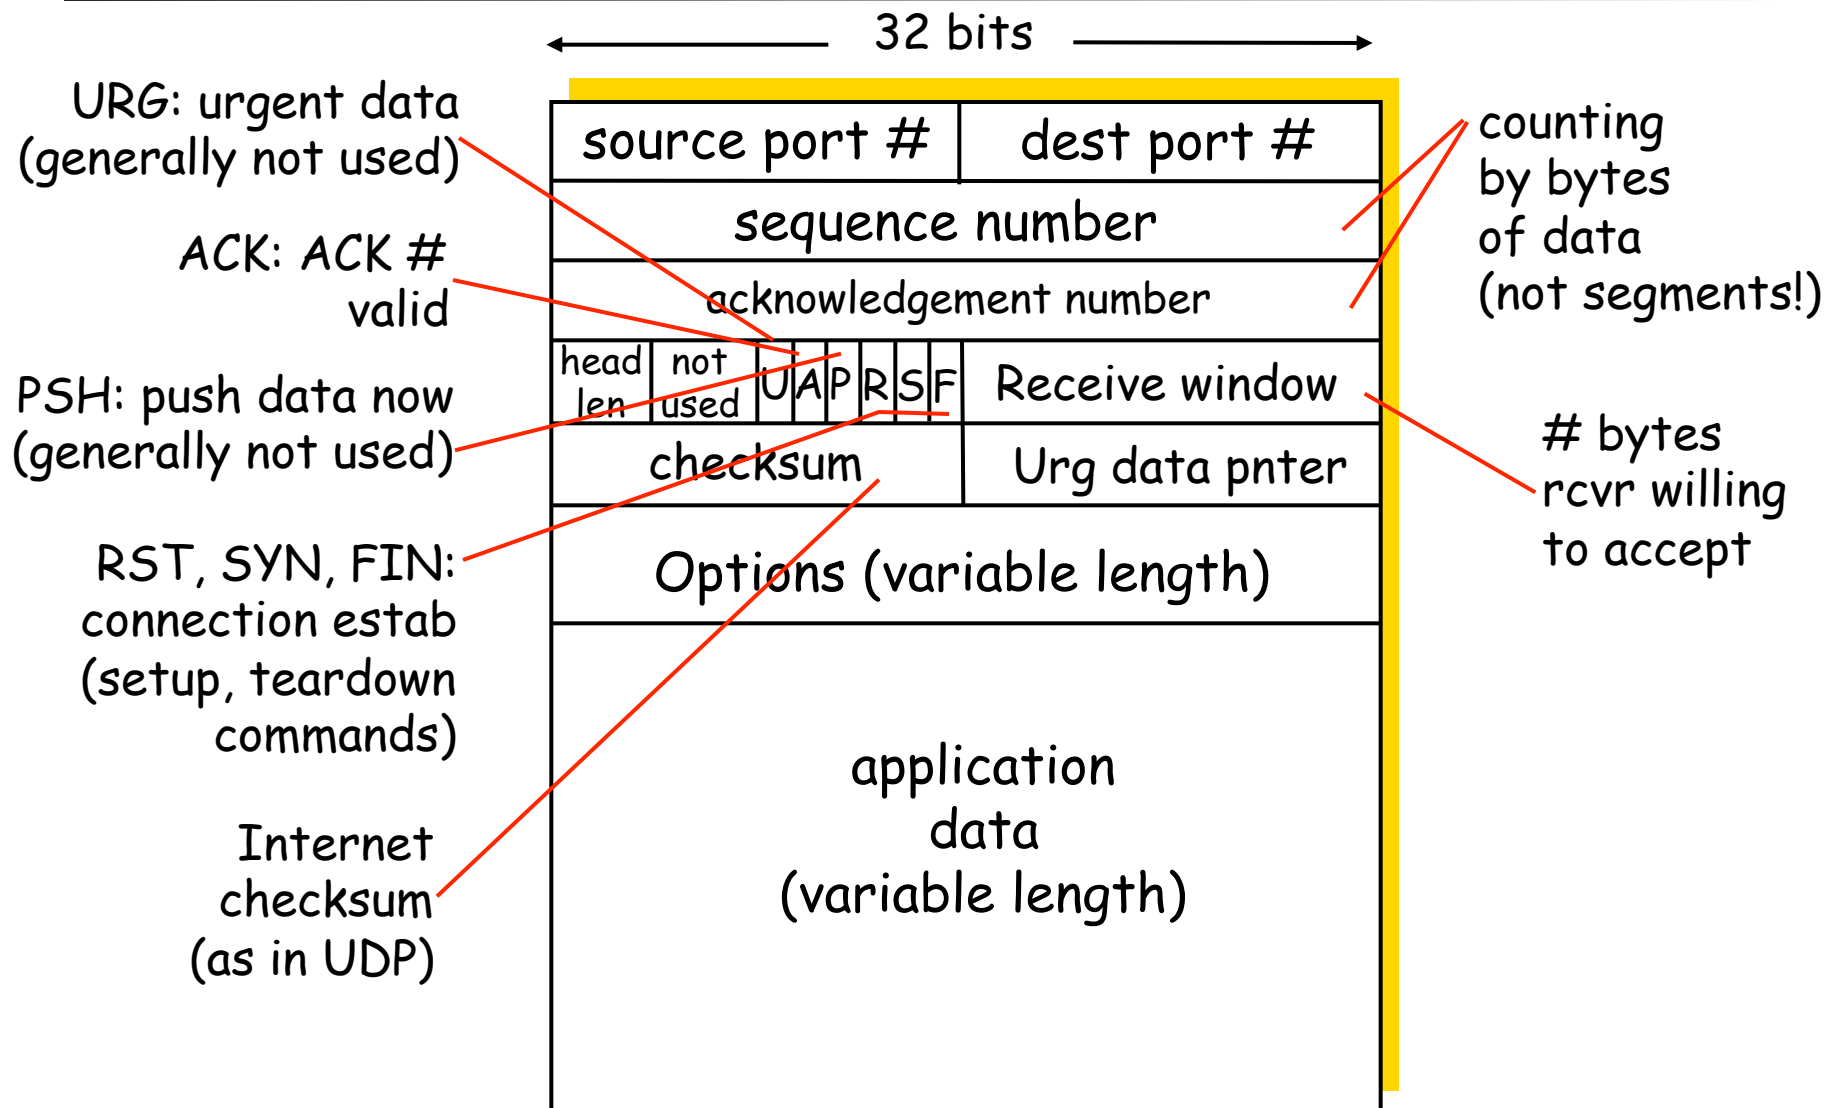
- handshaking (exchange of control msgs) init's sender, receiver state before data exchange

## ■ flow controlled:


- sender will not overwhelm receiver



# TCP segment structure



# TCP Connection Establishment



(a) TCP connection establishment in the normal case.

(b) Call collision.

# Berkeley Sockets

---

## The socket primitives for TCP

		Primitive	Meaning
Only Stream	{	SOCKET	Create a new communication end point
		BIND	Attach a local address to a socket
		LISTEN	Announce willingness to accept connections; give queue size
		ACCEPT	Block the caller until a connection attempt arrives
UDP	{	CONNECT	Actively attempt to establish a connection
		SEND	Send some data over the connection
		RECEIVE	Receive some data from the connection
		CLOSE	Release the connection

# The TCP Service Model

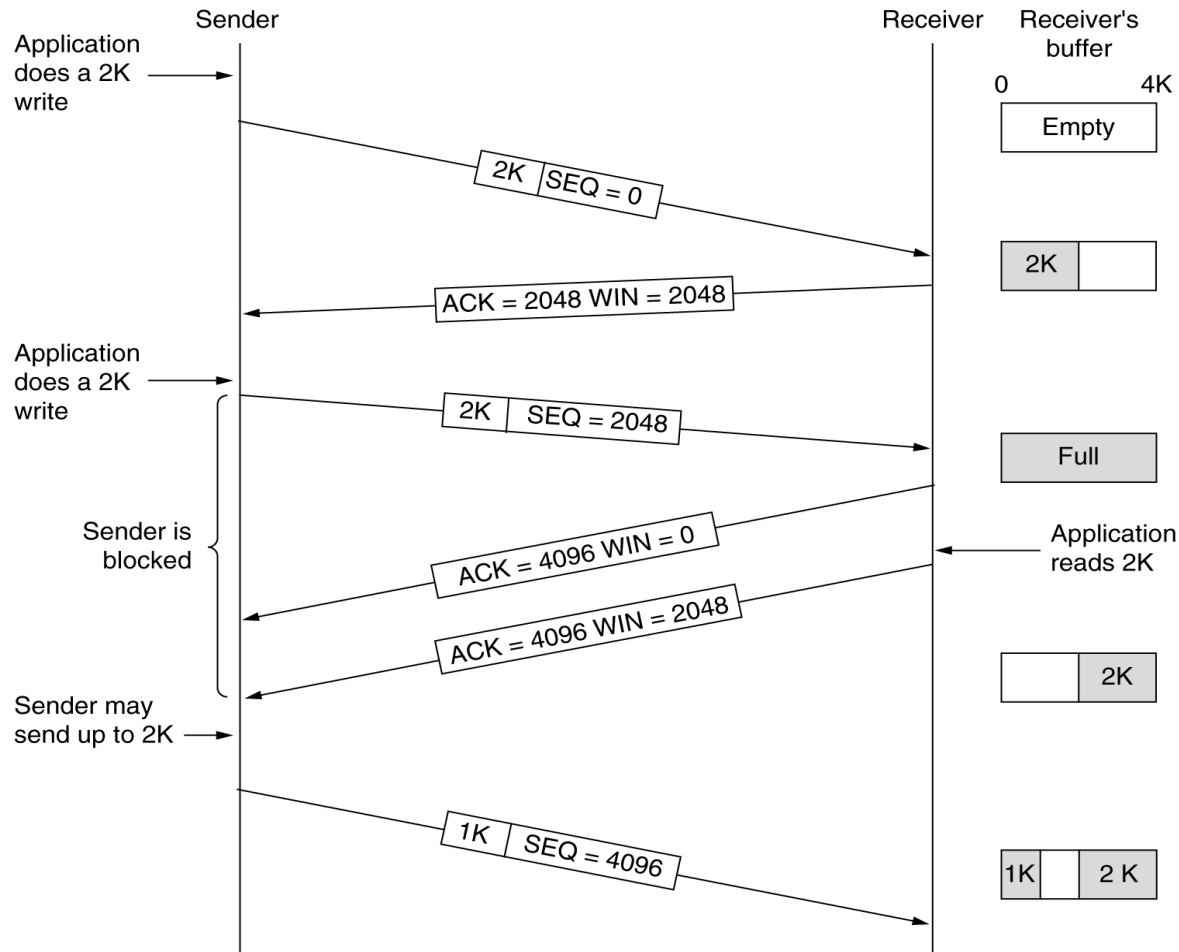
---

## Some assigned ports

Port	Protocol	Use
21	FTP	File transfer
23	Telnet	Remote login
25	SMTP	E-mail
69	TFTP	Trivial File Transfer Protocol
79	Finger	Lookup info about a user
80	HTTP	World Wide Web
110	POP-3	Remote e-mail access
119	NNTP	USENET news

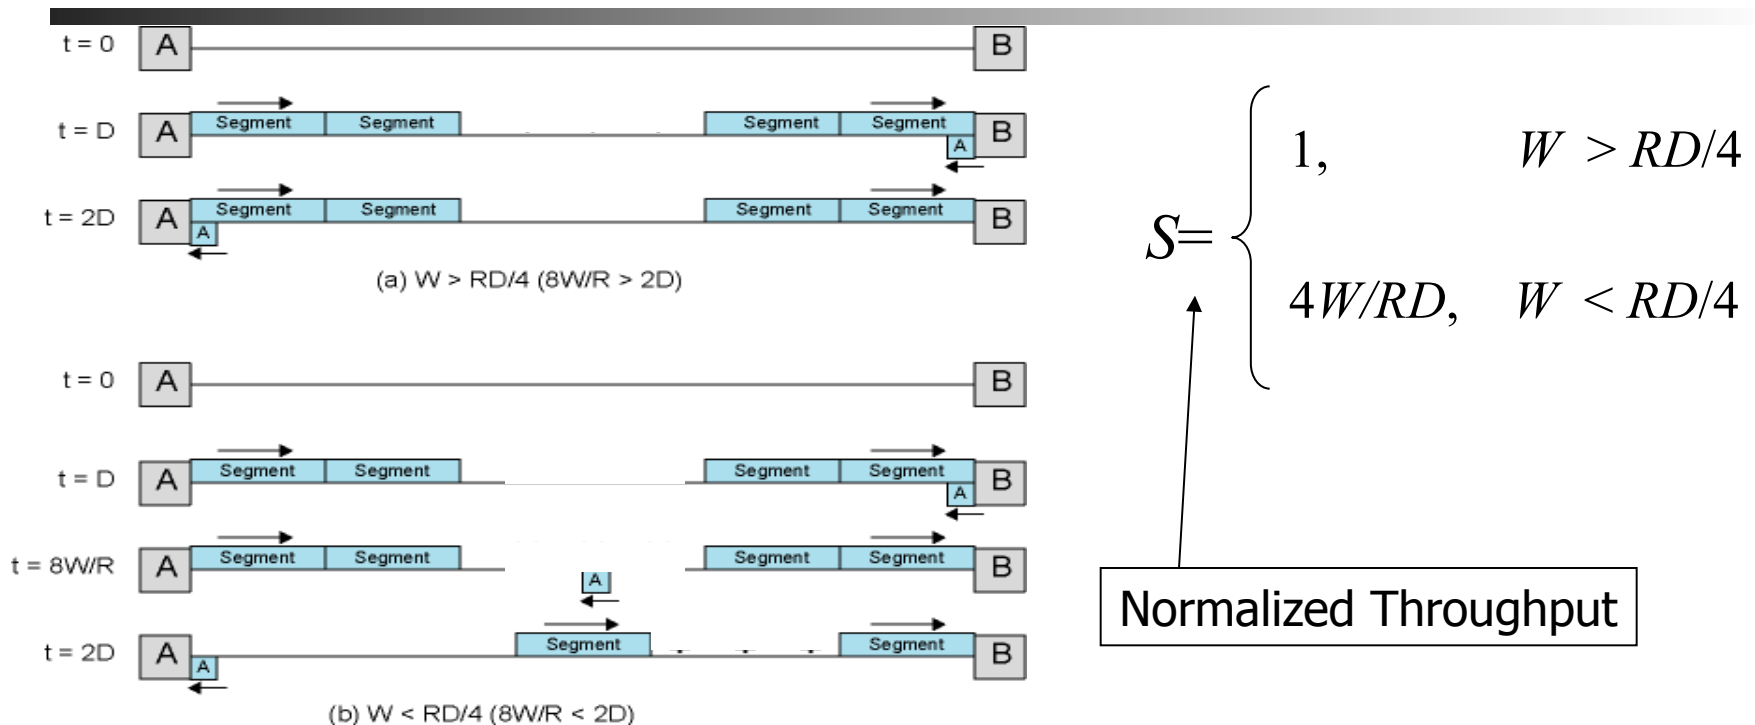


# TCP Transmission Policy



Window management in TCP.

# Effect of Window Size



- $W$  = TCP window size (octets)
- $R$  = Data rate (bps) at TCP source
- $D$  = Propagation delay (seconds)
- After TCP source begins transmitting, it takes  $D$  seconds for first octet to arrive, and  $D$  seconds for acknowledgement to return
- TCP source could transmit at most  $2RD$  bits, or  $RD/4$  octets

# TCP reliable data transfer

---

- TCP creates rdt service on top of IP's unreliable service
- Pipelined segments
- Cumulative acks
- TCP uses single retransmission timer
- Retransmissions are triggered by:
  - timeout events
  - duplicate acks

# Fast Retransmit

---

- Time-out period often relatively long:
  - long delay before resending lost packet
- Detect lost segments via **duplicate ACKs**
  - Sender often sends many segments back-to-back
  - If segment is lost, there will likely be many duplicate ACKs.
- If sender receives 3 ACKs for the same data, it supposes that segment after ACKed data was lost:
  - fast retransmit: resend segment before timer expires


# Congestion Control

(a) fast network feeding low capacity receiver


(b) Slow network feeding high-capacity receiver

## Congestion:

- informally: “too many sources sending too much data too fast for *network* to handle”
- different from flow control!
- manifestations:
  - lost packets (buffer overflow at routers)
  - long delays (queueing in router buffers)



# TCP Congestion Control



(Suppose TCP receiver discards out-of-order segments)

- spare room in buffer
- = `RcvWindow (=rcwnd)`
- = `RcvBuffer - [LastByteRcvd - LastByteRead]`

Rcvr advertises spare room by including value of

**RcvWindow** in segments

Sender limits unACKed data to **RcvWindow**

guarantees receive buffer doesn't overflow

Send `min(rcwnd, cwnd)`

# TCP Congestion Control

- end-end control (no network assistance)
- sender limits transmission:  
 $\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$
- Roughly,

$$\text{rate} = \frac{\text{cwnd}}{\text{RTT}} \text{ Bytes/sec}$$

- **cwnd** is dynamic, function of perceived network congestion

How does sender perceive congestion?

loss event = timeout *or* 3 duplicate acks

TCP sender reduces rate (**cwnd**) after loss event

three mechanisms:

AIMD


Slow Start

Congestion Avoidance

# TCP AIMD

additive increase: increase `cwnd` by 1 segment size every RTT in the absence of loss events: *probing*

multiplicative decrease: cut `cwnd` in half after loss event




Long-lived TCP connection



# TCP Slow Start

- When connection begins, **cwnd** = 1 MSS
  - Example: MSS = 500 bytes & RTT = 200 msec
  - initial rate = 20 kbps
- Available bandwidth may be  $\gg$  MSS/RTT
  - desirable to quickly ramp up to respectable rate
- When connection begins, increase rate exponentially fast until first loss event
  - double **cwnd** every RTT
  - done by incrementing **cwnd** for every ACK received
- Summary: initial rate is slow but ramps up exponentially fast



# Refinement

---

- After 3 dup ACKs:
  - `cwnd` is cut in half
  - window then grows linearly
- But after timeout event:
  - `cwnd` instead set to 1 MSS;
  - window then grows exponentially
  - to a threshold, then grows linearly


## Philosophy:

- 3 dup ACKs indicates network capable of delivering some segments
- timeout before 3 dup ACKs is “more alarming”

# Refinement (more)

**Q:** When should the exponential increase switch to linear?

**A:** When `cwnd` gets to 1/2 of its value before timeout.



## Implementation:

### Variable Threshold


At loss event, Threshold is set to 1/2 of cwnd just before loss event

# Summary: TCP Congestion Control

---

- When `cwnd` is below Threshold
  - sender in slow-start phase, `cwnd` grows exponentially
- When `cwnd` is above Threshold
  - sender is in congestion-avoidance phase, `cwnd` grows linearly
- When a triple duplicate ACK occurs
  - $cwnd = cwnd / 2$
- When timeout occurs
  - threshold set to  $cwnd / 2$  and `cwnd` is set to 1 MSS.

# TCP Congestion Control



An example of the Internet congestion algorithm.

# TCP Round Trip Time and Timeout

---

## How to set TCP timeout value?

- longer than RTT
  - but RTT varies
- too short: premature timeout
  - unnecessary retransmissions
- too long: slow reaction to segment loss

## How to estimate RTT?

- **SampleRTT**: measured time from segment transmission until ACK receipt
  - ignore retransmissions
- **SampleRTT** will vary, want estimated RTT “smoother”
  - average several recent measurements, not just current **SampleRTT**

# TCP Round Trip Time and Timeout

- **RTT**: the best current estimate of the round-trip time to the destination
- TCP measures how long an ACK took (**M**)
- Then updates RTT according to the formula

$$RTT = \alpha RTT + (1 - \alpha)M$$

**$\alpha$** : a smoothing factor that determines how much weight is given to the old value (Typically  **$\alpha = 7/8$** )

- For timeout period TCP uses  **$\beta RTT$** , but the trick is choosing  **$\beta$** .
- In the initial implementations,  **$\beta$**  was always **2**, but experience showed that a constant value was inflexible

# TCP Round Trip Time and Timeout

- Jacobson (1988) proposed  $\beta \sim \text{std. dev. of ACK arrival time pdf}$ 
  - Keep track of another smoothed variable,  $D$ , the deviation.
  - Whenever an ACK, the difference between the expected and observed values,  $|RTT - M|$ , is computed.
  - A smoothed value of this is maintained in  $D$  by the formula

$$D = \alpha D + (1 - \alpha) |RTT - M|$$

- Here,  $\alpha$  may or may not be the same value used to smooth RTT (typically  $\frac{3}{4}$ ). Most TCP implementations use this

$$\text{Timeout} = RTT + 4D$$

- When a segment times out and is sent again, it is unclear whether the ACK is for the 1<sup>st</sup> transmission or later one
- Do not update RTT on any segments that have been retransmitted.
- Timeout is doubled on each failure until the segments get through the first time.
  - This fix is called Karn's algorithm. Most TCP implementations use it.



# TCP Fairness

**Fairness goal:** if  $K$  TCP sessions share same bottleneck link of bandwidth  $R$ , each should have average rate of  $R/K$

