

## Introduction


The Art of Multiprocessor Programming  
by Maurice Herlihy & Nir Shavit

## Outline


- The Concurrency Speedup
- Java Threading Basics
- Classical Problems of Concurrency
- Basic Locks

## The Concurrency Speedup

## The Uniprocessor




## The Shared Memory Multiprocessor (SMP)

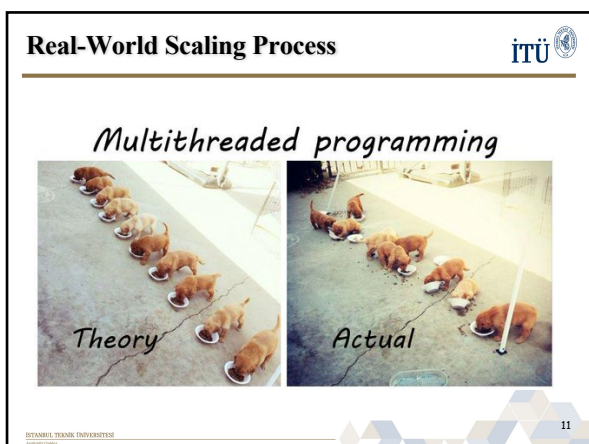
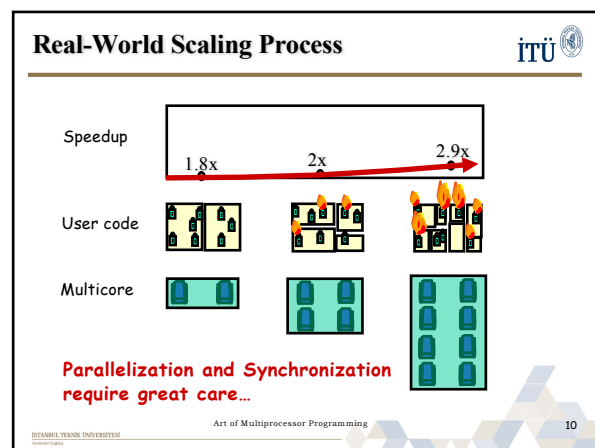
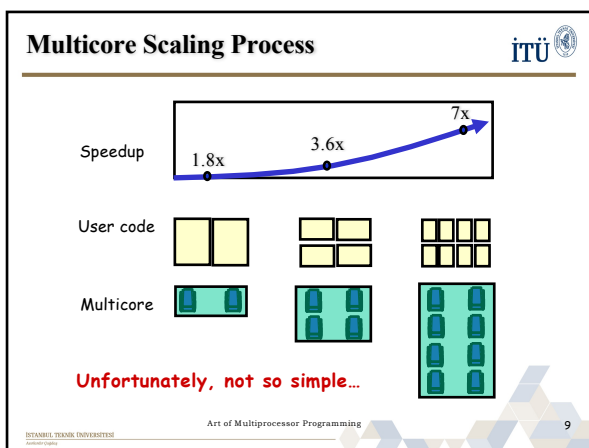
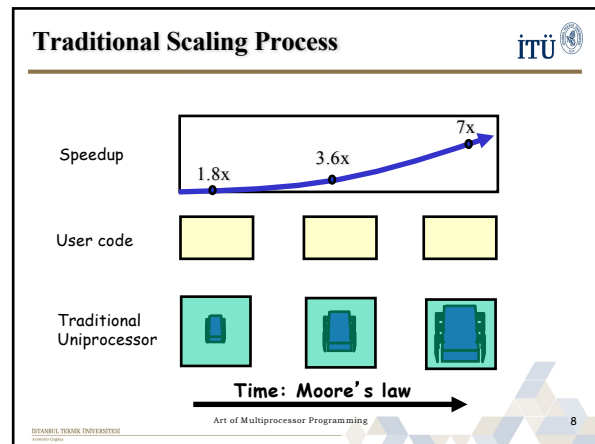



## The Multicore Processor (CMP)

All on the same chip




Sun  
T2000  
Niagara



- ### Parallel Primality Testing
- Challenge
    - Print primes from 1 to  $10^{10}$
  - Given
    - Ten-processor multiprocessor
    - One thread per processor
  - Goal
    - Get ten-fold speedup (or close)
- Art of Multiprocessor Programming
- 12

## Load Balancing



- Split the work evenly
- Each thread tests range of  $10^9$


## Procedure for Thread $i$

```
void primePrint {  
    int i = ThreadID.get(); // IDs in {0..9}  
    for (j = i*109+1, j<(i+1)*109; j++) {  
        if (isPrime(j))  
            print(j);  
    }  
}
```

## Issues

- Higher ranges have fewer primes
- Yet larger numbers harder to test
- Thread workloads
  - Uneven
  - Hard to predict

## Shared Counter



## Why do we care?

- We want as much of the code as possible to execute concurrently (in parallel)
- A larger sequential part implies reduced performance
- Amdahl's law: this relation is not linear...

## Amdahl's Law

$$\text{Speedup} = \frac{\text{OldExecutionTime}}{\text{NewExecutionTime}}$$

...of computation given  $n$  CPUs instead of 1

## Amdahl's Law



$$\text{Speedup} = \frac{1}{1 - p + \frac{p}{n}}$$

## Amdahl's Law



$$\text{Speedup} = \frac{1}{1 - p + \frac{p}{n}}$$

Parallel fraction

## Amdahl's Law



$$\text{Speedup} = \frac{1}{1 - p + \frac{p}{n}}$$

Sequential fraction

Parallel fraction

## Amdahl's Law



$$\text{Speedup} = \frac{1}{1 - p + \frac{p}{n}}$$

Sequential fraction

Parallel fraction

Number of processors

## Example



- Ten processors
- 60% concurrent, 40% sequential
- How close to 10-fold speedup?

## Example



- Ten processors
- 60% concurrent, 40% sequential
- How close to 10-fold speedup?

$$\text{Speedup} = 2.17 = \frac{1}{1 - 0.6 + \frac{0.6}{10}}$$

### Example



- Ten processors
- 80% concurrent, 20% sequential
- How close to 10-fold speedup?

### Example



- Ten processors
- 80% concurrent, 20% sequential
- How close to 10-fold speedup?

$$\text{Speedup}=3.57=\frac{1}{1-0.8+\frac{0.8}{10}}$$

### Example



- Ten processors
- 90% concurrent, 10% sequential
- How close to 10-fold speedup?

### Example



- Ten processors
- 90% concurrent, 10% sequential
- How close to 10-fold speedup?

$$\text{Speedup}=5.26=\frac{1}{1-0.9+\frac{0.9}{10}}$$

### Example



- Ten processors
- 99% concurrent, 01% sequential
- How close to 10-fold speedup?

### Example



- Ten processors
- 99% concurrent, 01% sequential
- How close to 10-fold speedup?

$$\text{Speedup}=9.17=\frac{1}{1-0.99+\frac{0.99}{10}}$$

## The Moral



- Making good use of our multiple processors (cores) means
- Finding ways to effectively parallelize our code
  - Minimize sequential parts
  - Reduce idle time in which threads wait without

Art of Multiprocessor Programming

31

## Process Concept



- All multiprogramming OSs are built around the concept of processes.
- Process – a program in execution; an instance of a running program; the entity that can be assigned to, and executed on, a processor.
- Program is a *passive entity*, process is an *active entity*.
- A process includes three segments:
  1. Program: code/text.
  2. Data: program variables and heap.
  3. Stack: for procedure calls and parameter passing.

## Process Characteristics (1)



- Unit of resource ownership – process is allocated:
  - an address space to hold the process image.
  - control of some resources (files, I/O devices...).
- Unit of dispatching – process is an execution path through one or more programs:
  - execution may be interleaved with other process.
  - the process has an execution state and a dispatching priority.

## Process Characteristics (2)




- These two characteristics are treated independently by some recent OSs:
  1. The unit of resource ownership is usually referred to as a Task or (for historical reasons) also as a Process.
  2. The unit of dispatching is usually referred to a Thread or a Light-Weight Process (LWP).
- A traditional Heavy-Weight Process (HWP) is equal to a task with a single thread.
- Several threads can exist in the same task.
  - **Multithreading** - The ability of an OS to support multiple, concurrent paths of execution within a single process.

## Processes and Threads (1)



- Process Items (shared by all threads of task):
  - address space which holds the process image
  - global variables
  - protected access to files, I/O and other resources
- Thread Items:
  - an execution state (Running, Ready, etc.)
  - program counter, register set
  - execution stack
  - some per-thread static storage for local variables
  - saved thread context when not running

## Each thread has its own stack



## Tasks/Processes and Threads (2)



### Per process items

Address space  
Global variables  
Open files  
Child processes  
Pending alarms  
Signals and signal handlers  
Accounting information

### Per thread items

Program counter  
Registers  
Stack  
State


İSTANBUL TEKNİK ÜNİVERSİTESİ  
Gönderen: Yalçın

## The Process vs. Thread Model



(a) Three processes each with one thread.

(b) One process with three threads.



İSTANBUL TEKNİK ÜNİVERSİTESİ  
Gönderen: Yalçın

## Processes vs. Threads



- Creating and managing processes is generally regarded as an expensive task (fork system call).
- Making sure all the processes peacefully co-exist on the system is not easy (as concurrency transparency comes at a price).
- Threads can be thought of as an “execution of a part of a program (in user-space)”.
- Rather than make the OS responsible for concurrency transparency, it is left to the individual application to manage the creation and scheduling of each thread.

İSTANBUL TEKNİK ÜNİVERSİTESİ  
Gönderen: Yalçın

## Thread operation latencies ( $\mu s$ )



Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840

Source: Anderson, T. et al, “Scheduler Activations: Effective Kernel Support for the User-Level Management of Parallelism”, ACM TOCS, February 1992.

İSTANBUL TEKNİK ÜNİVERSİTESİ  
Gönderen: Yalçın

## Classical Problems of Concurrency



İSTANBUL TEKNİK ÜNİVERSİTESİ  
Gönderen: Yalçın

## Classical Problems of Concurrency



- There are many of them – let’s briefly see three famous problems:
  1. Critical Section
  2. Producer-Consumer
  3. Readers and Writers

İSTANBUL TEKNİK ÜNİVERSİTESİ  
Gönderen: Yalçın

## The Critical-Section Problem



- $n$  processes competing to use some shared data.
- No assumptions may be made about speeds or the number of CPUs.
- Each process has a code segment, called *Critical Section (CS)*, in which the shared data is accessed.
- Problem – ensure that when one process is executing in its CS, no other process is allowed to execute in its CS.

İSTANBUL TEKNİK ÜNİVERSİTESİ

## What It Means



```
public class Counter {
    private long value;

    public long getAndIncrement() {
        return value++;
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

44

## What It Means



```
public class Counter {
    private long value;


    public long getAndIncrement() {
        return value++;
    }
}
```

```
temp = value;
value = value + 1;
return temp;
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming 45


## Not so good...



İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming 46

## Is this problem inherent?



If we could only glue reads and writes...

İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming 47

## An Aside: Java™



```
public class Counter {
    private long value;

    public long getAndIncrement() {
        synchronized {
            temp = value;
            value = temp + 1;
        }
        return temp;
    }
}
```

Synchronized block

İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming 48



## Solution to Critical-Section Problem



- There are 3 requirements that must stand for a correct solution:
  - 1. Mutual Exclusion**
  - 2. Progress**
  - 3. Bounded Waiting**
- We can check on all three requirements in each proposed solution, even though the non-existence of each one of them is enough for an incorrect solution.

İSTANBUL TEKNİK ÜNİVERSİTESİ

## Types of solutions to CS problem



- Software solutions – e.g. Locks/Monitors
  - algorithms whose correctness does not rely on any other assumptions.
- Hardware solutions – e.g. Atomic Instructions
  - rely on some special machine instructions.
- Operating System solutions – e.g. Semaphores/IPC
  - provide some functions and data structures to the programmer through system/library calls.

İSTANBUL TEKNİK ÜNİVERSİTESİ

## Initial Attempts to Solve Problem



- General structure of process  $P_i$  (other is  $P_j$ ) –

```
do {
    entry section
    critical section
    leave section
    remainder section
} while (TRUE);
```
- Processes may share some common variables to synchronize their actions.

İSTANBUL TEKNİK ÜNİVERSİTESİ

## Algorithm 1



- Shared variables
  - boolean flag[2];** initially **flag [0] = flag [1] = FALSE**
  - flag [i] = TRUE**  $\Rightarrow P_i$  ready to enter its critical section
- Process  $P_i$ 

```
do {
    while (flag[j]);
    flag[i] = TRUE;
    critical section
    flag [i] = FALSE;
    remainder section
} while (TRUE);
```
- Satisfies progress, but not mutual exclusion and bounded waiting requirements.

İSTANBUL TEKNİK ÜNİVERSİTESİ

## Algorithm 2



- Shared variables
  - boolean flag[2];** initially **flag [0] = flag [1] = FALSE**
  - flag [i] = TRUE**  $\Rightarrow P_i$  wants to enter its critical section
- Process  $P_i$ 

```
do {
    flag[i] = TRUE;
    while (flag[j]);
    critical section
    flag [i] = FALSE;
    remainder section
} while (TRUE);
```
- Satisfies mutual exclusion, but not progress and bounded waiting (?) requirements.

İSTANBUL TEKNİK ÜNİVERSİTESİ

## Algorithm 3



- Shared variables:
  - int turn;** initially **turn = 0**
  - turn = i**  $\Rightarrow P_i$  can enter its critical section
- Process  $P_i$ 

```
do {
    while (turn != i);
    critical section
    turn = j;
    remainder section
} while (TRUE);
```
- Satisfies mutual exclusion and bounded waiting, but not progress.

İSTANBUL TEKNİK ÜNİVERSİTESİ

#### Algorithm 4




- Combined shared variables of algorithms 1 and 2/3.
- Process  $P_i$ 

```
do {
    flag[i] = TRUE;
    turn = i;
    while (flag[j] and turn == j);
    critical section
    flag[i] = FALSE;
    remainder section
} while (TRUE);
```
- Meets all three requirements; solves the critical-section problem for two processes.

İSTANBUL TEKNİK ÜNİVERSİTESİ

#### Multiple Producers and Consumers



İSTANBUL TEKNİK ÜNİVERSİTESİ

#### Producer/Consumer (P/C) Dynamics



- A producer process produces information that is consumed by a consumer process.
- At any time, a producer activity may create some data.
- At any time, a consumer activity may want to accept some data.
- The data should be saved in a buffer until they are needed.
- If the buffer is finite, we want a producer to block if its new data would overflow the buffer.
- We also want a consumer to block if there are no data available when it wants them.

İSTANBUL TEKNİK ÜNİVERSİTESİ

#### P/C Bounded-Buffer Problem



- We need 3 semaphores:
  1. A semaphore **mutex** (initialized to 1) to have mutual exclusion on buffer access.
  2. A semaphore **full** (initialized to 0) to synchronize producer and consumer on the number of consumable items.
  3. A semaphore **empty** (initialized to n) to synchronize producer and consumer on the number of empty spaces.

İSTANBUL TEKNİK ÜNİVERSİTESİ

#### Bounded-Buffer – Semaphores



- Shared data

**semaphore full, empty, mutex;**

Initially:

**full = 0, empty = n, mutex = 1**

İSTANBUL TEKNİK ÜNİVERSİTESİ

#### Bounded-Buffer – Producer Process



```
do {
    ...
    produce an item in nextp
    ...
    wait(empty);
    wait(mutex);
    ...
    add nextp to buffer
    ...
    signal(mutex);
    signal(full);
} while (TRUE);
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

## Bounded-Buffer – Consumer Process



```
do {
    wait(full)
    wait(mutex);
    ...
    remove an item from buffer to nextc
    ...
    signal(mutex);
    signal(empty);
    ...
    consume the item in nextc
    ...
} while (TRUE);
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

## Readers-Writers Problem



- A data set/repository is shared among a number of concurrent processes:
  - Readers – only read the data set; they do **not** perform any updates.
  - Writers – can both read and write.
- Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.

İSTANBUL TEKNİK ÜNİVERSİTESİ


## Readers-Writers Dynamics



- Any number of reader activities and writer activities are running.
- At any time, a reader activity may wish to read data.
- At any time, a writer activity may want to modify the data.
- Any number of readers may access the data simultaneously.
- During the time a writer is writing, no other reader or writer may access the shared data.


İSTANBUL TEKNİK ÜNİVERSİTESİ

## Readers-Writers with active readers




İSTANBUL TEKNİK ÜNİVERSİTESİ

## Readers-Writers with an active writer



İSTANBUL TEKNİK ÜNİVERSİTESİ

## Should readers wait for waiting writer?



İSTANBUL TEKNİK ÜNİVERSİTESİ

## To post a message



W<sub>4</sub> A<sub>3</sub> S<sub>1</sub> H<sub>4</sub> T<sub>3</sub> H<sub>4</sub> E<sub>1</sub> C<sub>3</sub> A<sub>1</sub> R<sub>1</sub>



ISTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming 67

## Let's send another message



S<sub>1</sub> E<sub>1</sub> L<sub>1</sub> L<sub>1</sub> A<sub>3</sub> V<sub>4</sub> A<sub>1</sub> L<sub>1</sub> A<sub>1</sub> M<sub>3</sub> P<sub>3</sub> S<sub>1</sub>



ISTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming 68

## Uh-Oh



S<sub>1</sub> E<sub>1</sub> L<sub>1</sub> L<sub>1</sub> T<sub>1</sub> H<sub>4</sub> E<sub>1</sub> C<sub>3</sub> A<sub>1</sub> R<sub>1</sub>



ISTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming 69

## Readers-Writers problem



- There are various versions with different readers and writers preferences:
1. The **first** readers-writers problem, requires that no reader will be kept waiting unless a writer has obtained access to the shared data.
  2. The **second** readers-writers problem, requires that once a writer is ready, no new readers may start reading.
  3. In a solution to the **first** case writers may starve; In a solution to the **second** case readers may starve.

ISTANBUL TEKNİK ÜNİVERSİTESİ

## First Readers-Writers Solution (1)



- **readcount** (initialized to 0) counter keeps track of how many processes are currently reading.
- **mutex** semaphore (initialized to 1) provides mutual exclusion for updating readcount.
- **wrt** semaphore (initialized to 1) provides mutual exclusion for the writers; it is also used by the first or last reader that enters or exits the CS.

ISTANBUL TEKNİK ÜNİVERSİTESİ

## First Readers-Writers Solution (2)



- Shared data

**semaphore mutex, wrt;**  
**int readcount;**

Initially

**mutex = 1, wrt = 1, readcount = 0**

ISTANBUL TEKNİK ÜNİVERSİTESİ

## First Readers-Writers – Writer Process

```
do {
    wait(wrt);
    ...
    writing is performed
    ...
    signal(wrt);
} while(TRUE);
```

ISTANBUL TEKNİK ÜNİVERSİTESİ

## First Readers-Writers – Reader Process

```
do {
    wait(mutex);
    readcount++;
    if (readcount == 1)
        wait(wrt);
    signal(mutex);
    ...
    reading is performed
    ...
    wait(mutex);
    readcount--;
    if (readcount == 0)
        signal(wrt);
    signal(mutex);
} while(TRUE);
```

ISTANBUL TEKNİK ÜNİVERSİTESİ

## Java Threading Basics

ISTANBUL TEKNİK ÜNİVERSİTESİ

## Threads Overview

- Java supports threads
  - Threads execute within a single JVM
  - Native threads map a single Java thread to an OS thread
  - Green threads adopt the thread library approach (threads are invisible to the OS)
  - On a multiprocessor system, native threads are required to get true parallelism (but this is still implementation dependent)

ISTANBUL TEKNİK ÜNİVERSİTESİ

## Threads Overview

- There are various ways in which concurrency can be introduced by
  - an API for explicit thread creation or thread forking
  - a high-level language construct such as PAR (occam), tasks (Ada), or processes (Modula)
- Integration with OOP, various models:
  - asynchronous method calls
  - early return from methods
  - futures
  - active objects

ISTANBUL TEKNİK ÜNİVERSİTESİ

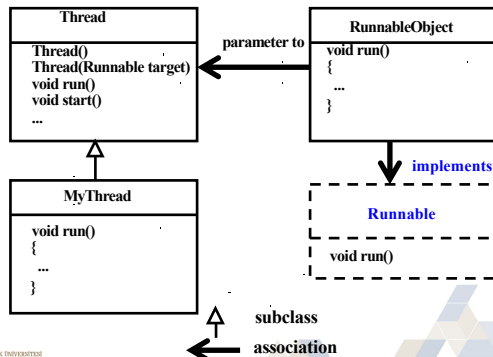
## Concurrency in Java

- Java has a predefined class `java.lang.Thread` which provides the mechanism by which threads are created
- However to avoid all threads having to be child classes of `Thread`, it also uses a standard interface

```
public interface Runnable {
    public void run();
}
```
- Hence, any class which wishes to express concurrent execution must implement this interface and provide the `run` method
- Threads do not begin their execution until the `start` method in the `Thread` class is called

ISTANBUL TEKNİK ÜNİVERSİTESİ

## Threads in Java



İSTANBUL TEKNİK ÜNİVERSİTESİ

## Thread Identification



- The identity of the currently running thread can be found using the `currentThread` method
- This has a static modifier, which means that there is only one method for all instances of `Thread` objects
- The method can always be called using the `Thread` class

```

public class Thread extends Object
    implements Runnable {
    ...
    public static Thread currentThread();
    ...
}
    
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

## Threads in Java



```

public class PrintNumbers {
    public static void printNumbers() {
        for(int i=0; i<1000; i++) {
            System.out.println(
                Thread.currentThread().getId() + ": " + i);
        }
    }
}
    
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

81

## Threads in Java



```

public class Thread1 extends Thread {
    @Override
    public void run() {
        System.out.println("Thread1 ThreadId: " +
            Thread.currentThread().getId());
        // do our thing

        PrintNumbers.printNumbers();
    }
}
    
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

82

## Threads in Java



```

static public void main(String[] args) {
    System.out.println("Main ThreadId: " +
        Thread.currentThread().getId());

    for(int i=0; i<3; i++)
        new Thread1().start(); // don't call run!

    printNumbers();
}
    
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

83

## Threads in Java



```

public class Thread2 implements Runnable {
    @Override
    public void run() {
        System.out.println("Thread2 ThreadId: " +
            Thread.currentThread().getId());

        PrintNumbers.printNumbers();
    }
}
    
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

84

## Threads in Java



```
static public void main(String[] args) {
    System.out.println("Main ThreadId: " +
        Thread.currentThread().getId());

    for(int i=0; i<3; i++)
        new Thread(new Thread2()).start();

    printNumbers();
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

85

## Threads in Java



```
static public void main(String[] args) {
    System.out.println("Main ThreadId: " +
        Thread.currentThread().getId());

    new Thread(new Runnable() {
        @Override
        public void run() {
            System.out.println("Thread3 ThreadId: " +
                Thread.currentThread().getId());

            printNumbers();
        }
    }).start();
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

86

## Threads in Java



```
public class PrintNumbers {
    public static void main(String[] args) {
        Runnable task1 = () -> {
            for(int i=0; i<1000; i++)
                System.out.println(Thread.currentThread()
                    .getId()+" : " + i);

        };

        Thread[] t = new Thread[12];
        for(int i=0; i<t.length; i++)
            t[i] = new Thread(task1);
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

87

## A Thread Terminates:



- when it completes execution of its **run** method either normally or as the result of an unhandled exception
- (DEPRECATED) via a call to its **stop** method — the **run** method is stopped and the thread class cleans up before terminating the thread (releases locks and executes any finally clauses)
  - the thread object is now eligible for garbage collection.
  - **stop** is inherently unsafe as it releases locks on objects and can leave those objects in inconsistent states; the method is now deprecated and should not be used
- (NOT IMPLEMENTED) by its **destroy** method being called — **destroy** terminates the thread without any cleanup (not provided by many JVMs, now deprecated)

İSTANBUL TEKNİK ÜNİVERSİTESİ

## Daemon Threads



- Java threads can be of two types: **user** threads or **daemon** threads
- Daemon threads are those threads which provide general services and typically never terminate
- When a new thread is created it inherits the daemon status of its parent.
- Normal thread and daemon threads differ in what happens when they exit. When the JVM halts any remaining daemon threads are abandoned: finally blocks are not executed, stacks are not unwound - JVM just exits. Due to this reason daemon threads should be used sparingly and it is dangerous to use them for tasks that might perform any sort of I/O.
- The **setDaemon** method must be called before the thread is started

İSTANBUL TEKNİK ÜNİVERSİTESİ

## Joining



- One thread can wait (with or without a timeout) for another thread (the target) to terminate by issuing the **join** method call on the target's thread object
- The **isAlive** method allows a thread to determine if the target thread has terminated

İSTANBUL TEKNİK ÜNİVERSİTESİ

## Communication in Java



- Via reading and writing to data encapsulated in shared objects protected by simple monitors
- Every object is implicitly derived from the `Object` class which defines a mutual exclusion lock
- Methods in a class can be labeled as `synchronized`, this means that they can only be executed if the lock can be acquired (this happens automatically)
- The lock can also be acquired via a `synchronized` statement which names the object
- A thread can `wait` and `notify` on a single anonymous condition variable

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ

## Synchronization



Synchronization of threads is needed in order to control threads coordination, mainly in order to **prevent simultaneous operations on data**

For simple synchronization Java provides the `synchronized` keyword

For more sophisticated **locking mechanisms**, starting from Java 5, the package `java.concurrent.locks` provides additional locking options

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ

## Synchronization



```
public class SynchronizedCounter {  
    private int c = 0;  
    public synchronized void increment() { c++; }  
    public synchronized void decrement() { c--; }  
    public synchronized int value() { return c; }  
}
```

The `synchronized` keyword on a method means that if this is already locked anywhere (on this method or elsewhere) by another thread, we need to wait till this is unlocked before entering the method

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ

## Synchronization



```
public void addName(String name) {  
    synchronized(this) {  
        lastName = name;  
        nameCount++;  
    }  
    nameList.add(name);  
}
```

When synchronizing a block, key for the locking should be supplied (usually would be this)

The advantage of not synchronizing the entire method is efficiency

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ

## Synchronization



```
public class TwoCounters {  
    private long c1 = 0, c2 = 0;  
    private Object lock1 = new Object();  
    private Object lock2 = new Object();  
    public void inc1() {  
        synchronized(lock1) {  
            c1++;  
        }  
    }  
    public void inc2() {  
        synchronized(lock2) {  
            c2++;  
        }  
    }  
}
```

You must be absolutely sure that there is no tie between c1 and c2

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ

## Synchronization



Having a static method be `synchronized` means that ALL objects of this type are locked on the method and can get in one thread at a time.

The lock is the Class object representing this class.

The performance penalty might be sometimes too high – needs careful attention!

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ



## Peterson's Algorithm



```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}
```

Art of Multiprocessor Programming

97

## Peterson's Algorithm



```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}
```

Announce I'm interested

Art of Multiprocessor Programming

98

## Peterson's Algorithm



```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}
```

Announce I'm interested

Defer to other

Art of Multiprocessor Programming

99

## Peterson's Algorithm



```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}
```

Announce I'm interested

Defer to other

Wait while other interested & I'm the victim

Art of Multiprocessor Programming

100

## Peterson's Algorithm



```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}
```

Announce I'm interested

Defer to other

Wait while other interested & I'm the victim

No longer interested

Art of Multiprocessor Programming

101

## Mutual Exclusion



```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
```

- If thread 0 in critical section,
  - flag[0]=true,
  - victim = 1

- If thread 1 in critical section,
  - flag[1]=true,
  - victim = 0

Cannot both be true

Art of Multiprocessor Programming

102

## Deadlock Free



```
public void lock() {
    ...
    while (flag[j] && victim == i) {};
```

- Thread blocked
  - only at **while** loop
  - only if it is the victim
- One or the other must not be the victim

Art of Multiprocessor Programming

103

## Starvation Free



- Thread  $i$  blocked only if  $j$  repeatedly re-enters so that

**flag[j] == true and victim == i**

- When  $j$  re-enters
  - it sets **victim** to  $j$ .
  - So  $i$  gets in

```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
```

```
public void unlock() {
    flag[i] = false;
```

Art of Multiprocessor Programming

104

## The Filter Algorithm for $n$ Threads



There are  $n-1$  “waiting rooms” called levels

- At each level
  - At least one enters level
  - At least one blocked if many try
- Only one thread makes it through



Art of Multiprocessor Programming

105

## Filter



```
class Filter implements Lock {
    int[] level; // level[i] for thread i
    int[] victim; // victim[L] for level L

    public Filter(int n) {
        level = new int[n];
        victim = new int[n];
        for (int i = 1; i < n; i++) {
            level[i] = 0;
        }
    }
    ...
}
```

Thread 2 at level 4

Art of Multiprocessor Programming

106

## Filter



```
class Filter implements Lock {
    ...
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i level[k] >= L) &&
                    victim[L] == i) {}
        }
    }
    public void unlock() {
        level[i] = 0;
    }
}
```

Art of Multiprocessor Programming

107

## Filter



```
class Filter implements Lock {
    ...
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i level[k] >= L) &&
                    victim[L] == i) {}
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}
```

One level at a time

Art of Multiprocessor Programming

108

## Filter

```
class Filter implements Lock {
    ...
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i) level[k] >= L) &&
                victim[L] == i;
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}
```

Announce  
intention to  
enter level L

## Filter

```
class Filter implements Lock {
    int level[n];
    int victim[n];
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i) level[k] >= L) &&
                victim[L] == i;
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}
```

Give priority to  
anyone but me

## Filter

Wait as long as someone else is at same or  
higher level, and I'm designated victim

```
public void lock() {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while ((∃ k != i) level[k] >= L) &&
            victim[L] == i;
    }
}
public void release(int i) {
    level[i] = 0;
}
```


## Filter

```
class Filter implements Lock {
    int level[n];
    int victim[n];
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i) level[k] >= L) &&
                victim[L] == i;
        }
    }
}
```

Thread enters level L when it completes  
the loop

## Claim

- Start at level  $L=0$
- At most  $n-L$  threads enter level  $L$
- Mutual exclusion at level  $L=n-1$



## No Starvation

- Filter Lock satisfies properties:
  - Just like Peterson Alg at any level
  - So no one starves
- But what about fairness?
  - Threads can be overtaken by others

## Bakery Algorithm



- Provides First-Come-First-Served
- How?
  - Take a “number”
  - Wait until lower numbers have been served
- Lexicographic order
  - $(a,i) > (b,j)$ 
    - If  $a > b$ , or  $a = b$  and  $i > j$

## Bakery Algorithm



```
class Bakery implements Lock {
    boolean[] flag;
    Label[] label;
    public Bakery (int n) {
        flag = new boolean[n];
        label = new Label[n];
        for (int i = 0; i < n; i++) {
            flag[i] = false; label[i] = 0;
        }
    }
    ...
}
```

## Bakery Algorithm



```
class Bakery implements Lock {
    boolean[] flag;
    Label[] label;
    public Bakery (int n) {
        flag = new boolean[n];
        label = new Label[n];
        for (int i = 0; i < n; i++) {
            flag[i] = false; label[i] = 0;
        }
    }
    ...
}
```

Diagram illustrating the Bakery Algorithm state:

0	2	6	n-1
f	f	f	f
0	0	4	0
0	0	5	0

Arrows point from the second and third rows to the label array, and a vertical arrow points from the label array to the Critical Section (CS).

## Bakery Algorithm



```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

## Bakery Algorithm



```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

Diagram illustrating the Bakery Algorithm state:

Doorway

## Bakery Algorithm



```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

Diagram illustrating the Bakery Algorithm state:

I'm interested

## Bakery Algorithm



```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

Take increasing  
label (read labels  
in some arbitrary  
order)

## Bakery Algorithm



```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

Someone is  
interested

## Bakery Algorithm



```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

Someone is  
interested

With lower (label,i)  
in lexicographic order

## Bakery Algorithm



```
class Bakery implements Lock {
    ...
    public void unlock() {
        flag[i] = false;
    }
}
```

## Bakery Algorithm



```
class Bakery implements Lock {
    ...
    public void unlock() {
        flag[i] = false;
    }
}
```

No longer  
interested

labels are always increasing

## Bakery Y2<sup>32</sup>K Bug



```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

## Bakery Y2<sup>32</sup>K Bug



```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

**Mutex breaks if label[i] overflows**

## Does Overflow Actually Matter?




- Yes
  - Y2K
  - 18 January 2038 (Unix `time_t` rollover)
  - 16-bit counters
- No
  - 64-bit counters
- Maybe
  - 32-bit counters

## Timestamps




- Label variable is really a **timestamp**
- Need ability to
  - Read others' timestamps
  - Compare them
  - Generate a **later** timestamp
- Can we do this without overflow?

## Precedence Graphs




- Timestamps form directed graph
- Edge x to y
  - Means x is later timestamp
  - We say x **dominates** y

## Unbounded Counter Precedence Graph




- Timestamping = move tokens on graph
- Atomically
  - read others' tokens
  - move mine
- Ignore tie-breaking for now

## Unbounded Counter Precedence Graph




## Two-Thread Bounded Precedence Graph



Art of Multiprocessor Programming

133


## Two-Thread Bounded Precedence Graph



Art of Multiprocessor Programming

134


## Two-Thread Bounded Precedence Graph



Art of Multiprocessor Programming

135


## Two-Thread Bounded Precedence Graph



Art of Multiprocessor Programming

136

## Two-Thread Bounded Precedence Graph $T^2$




and so on ...

Art of Multiprocessor Programming

137


## Three-Thread Bounded Precedence Graph?



Art of Multiprocessor Programming

138


### Three-Thread Bounded Precedence Graph?



Art of Multiprocessor Programming

139


### Graph Composition



Art of Multiprocessor Programming

140


### Three-Thread Bounded Precedence Graph $T^3$



Art of Multiprocessor Programming

141


### Three-Thread Bounded Precedence Graph $T^3$



Art of Multiprocessor Programming

142


### Three-Thread Bounded Precedence Graph $T^3$



Art of Multiprocessor Programming

143

### In General



Art of Multiprocessor Programming

144



## Slides adopted from



- [Java Concurrency Framework](#) by Sidartha Gracias
- [Java Threads](#) by Amir Kirsh
- [Concurrent Programming in Java](#) by Andy Wellings

... and examples from

- [Java Concurrency in Practice](#) by Brian Goetz
- [Concurrent Programming in Java](#) by Doug Lea

