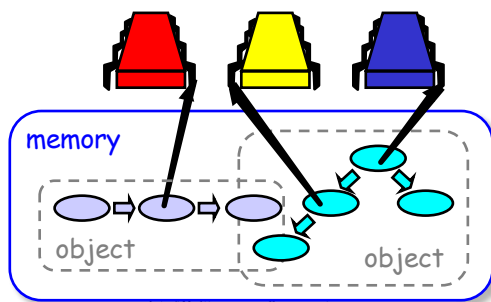


## Concurrent Objects

Companion slides for  
The Art of Multiprocessor  
Programming  
by Maurice Herlihy & Nir Shavit

## Correctness of Concurrent Objects

## Concurrent Computaton



Art of Multiprocessor Programming

3

## Objectivism

- What is a concurrent object?
  - How do we describe one?
  - How do we implement one?
  - How do we tell if we're right?

Art of Multiprocessor Programming

4

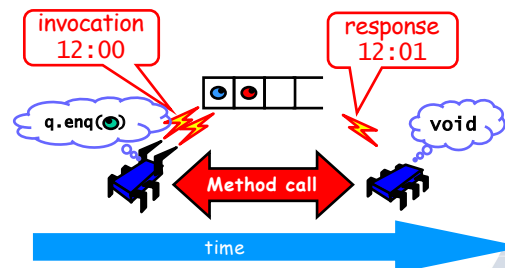
## Correctness and Progress

- In a concurrent setting, we need to specify both the safety and the liveness properties of an object
- Need a way to define
  - when an implementation is correct
  - the conditions under which it guarantees progress

Art of Multiprocessor Programming

5

## Methods Take Time



Art of Multiprocessor Programming

6

## Sequential vs Concurrent

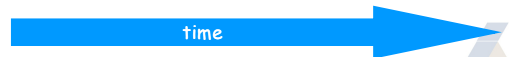
- Sequential
  - Methods take time? Who knew?
- Concurrent
  - Method call is not an event
  - Method call is an interval.

ISTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

7

## Concurrent Methods Take Overlapping Time

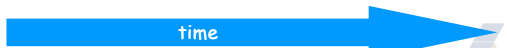


ISTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

8

## Concurrent Methods Take Overlapping Time

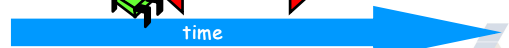
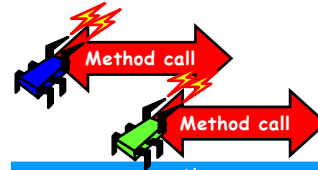


ISTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

9

## Concurrent Methods Take Overlapping Time

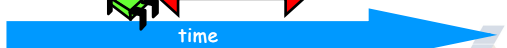
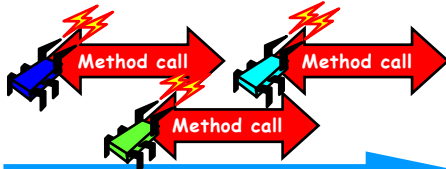


ISTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

10

## Concurrent Methods Take Overlapping Time



ISTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

11

## Sequential vs Concurrent

- Sequential:
  - Object needs meaningful state only between method calls
- Concurrent
  - Because method calls overlap, object might never be between method calls

ISTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

12

## Sequential vs Concurrent

- Sequential:
  - Each method described in isolation
- Concurrent
  - Must characterize all possible interactions with concurrent calls
    - What if two enqs overlap?
    - Two deqs? enq and deq? ...

## Sequential vs Concurrent

- Sequential:
  - Can add new methods without affecting older methods
- Concurrent:
  - Everything can potentially interact with everything else

## The Big Question

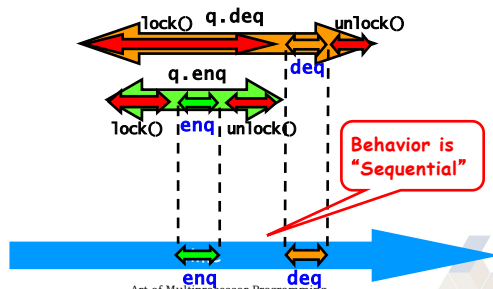
- What does it mean for a concurrent object to be correct?
  - What is a concurrent FIFO queue?
  - FIFO means strict temporal order
  - Concurrent means ambiguous temporal order

## Intuitively...

```
public T deq() throws EmptyException {
    lock.lock();
    try {
        if (tail == head)
            throw new EmptyException();
        T x = items[head % items.length];
        head++;
        return x;
    } finally {
        lock.unlock();
    }
}
```

## Intuitively

Lets capture the idea of describing the concurrent via the sequential



## Linearizability

- Each method should
  - “take effect”
  - Instantaneously
  - Between invocation and response events
- Object is correct if this “sequential” behavior is correct
- Any such concurrent object is
  - Linearizable

## Is it really about the object?

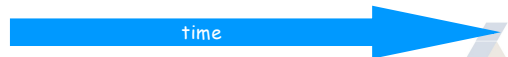
- Each method should
  - “take effect”
  - Instantaneously
  - Between invocation and response events
- Sounds like a property of an execution...
- A linearizable object: one whose all possible executions are linearizable

İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

19

## Example

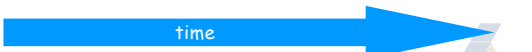
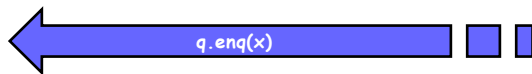


(4)

Art of Multiprocessor Programming

20

## Example

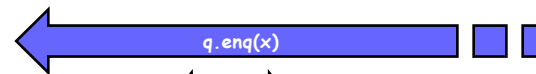


(4)

Art of Multiprocessor Programming

21

## Example

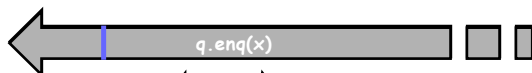


(4)

Art of Multiprocessor Programming

22

## Example

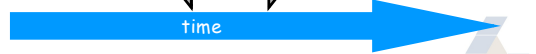
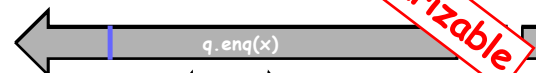


(4)

Art of Multiprocessor Programming

23

## Example



(4)

Art of Multiprocessor Programming

24

**linearizable**

Example

itü

time

(6)

Art of Multiprocessor Programming

25

Example

itü

q.enq(x)

time

(6)

Art of Multiprocessor Programming

26

Example

itü

q.enq(x)

q.enq(y)

time

(6)

Art of Multiprocessor Programming

27

Example

itü

q.enq(x)

q.enq(y)

q.deq(x)

time

(6)

Art of Multiprocessor Programming

28

Example

itü

q.enq(x)

q.enq(y)

q.deq(x)

q.deq(y)

time

(6)

Art of Multiprocessor Programming

29

Example

itü

q.enq(x)

q.enq(y)

q.deq(x)

q.deq(y)

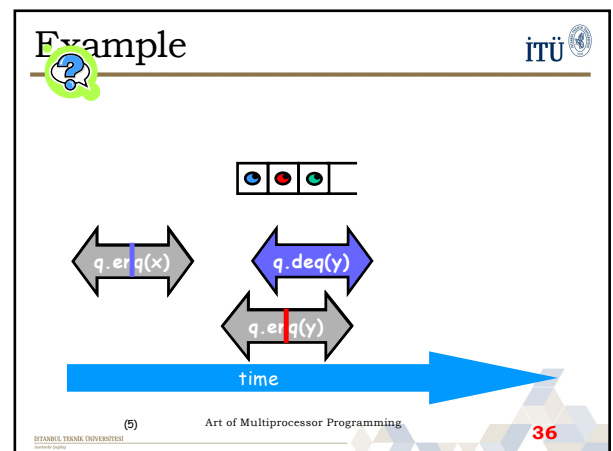
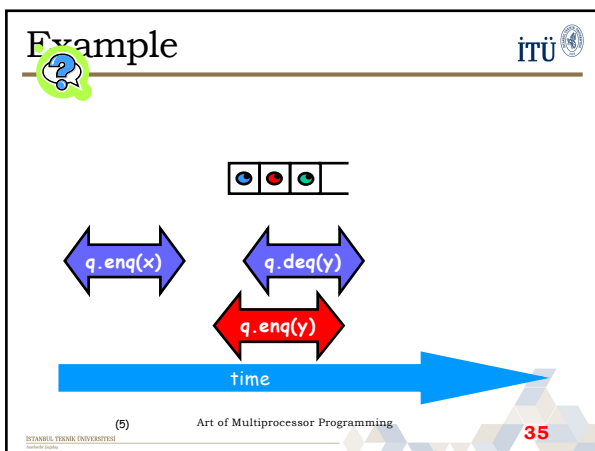
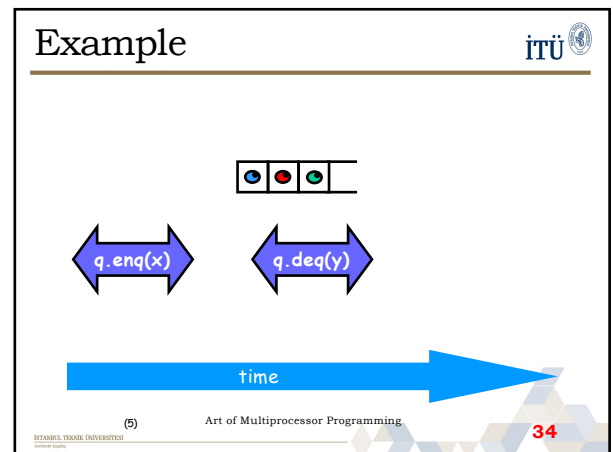
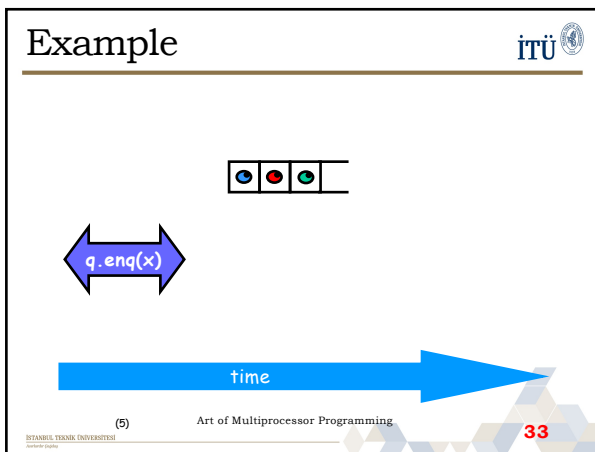
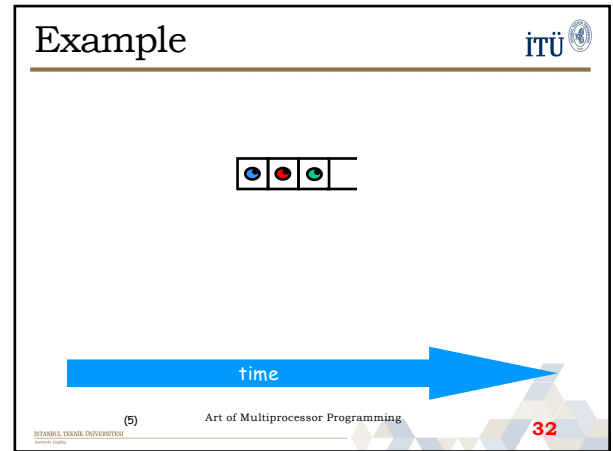
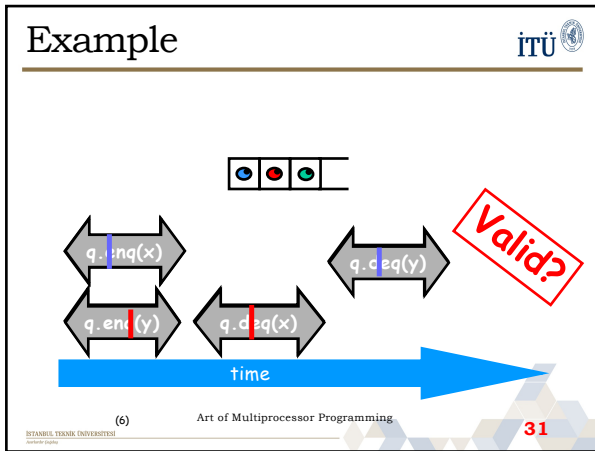
time

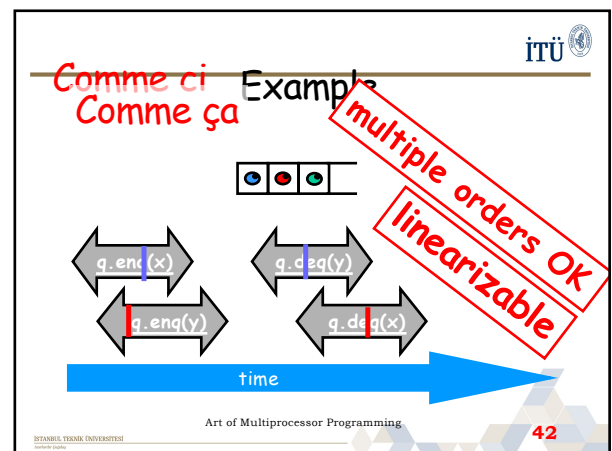
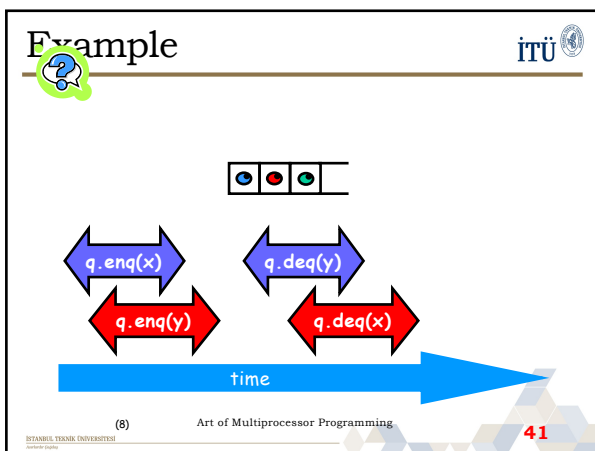
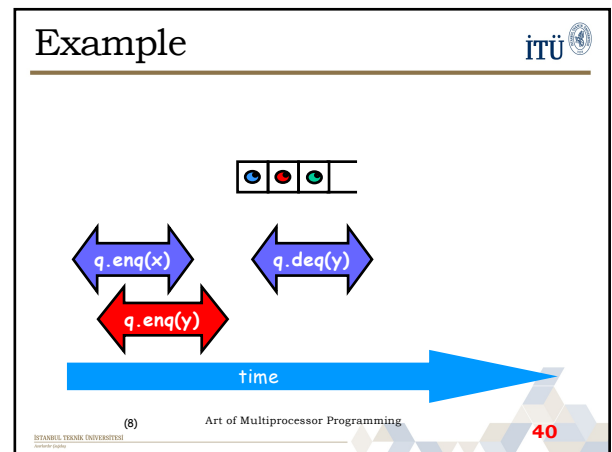
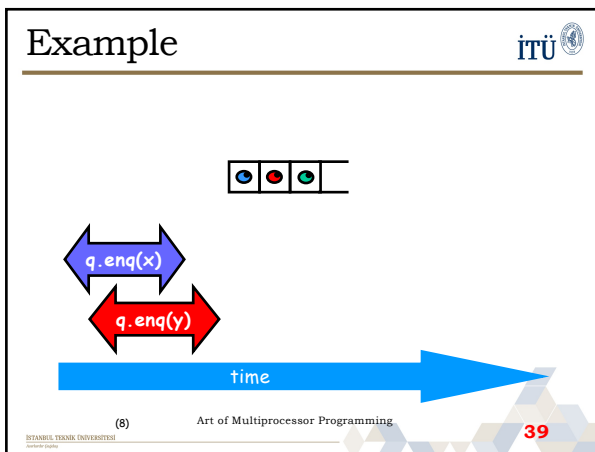
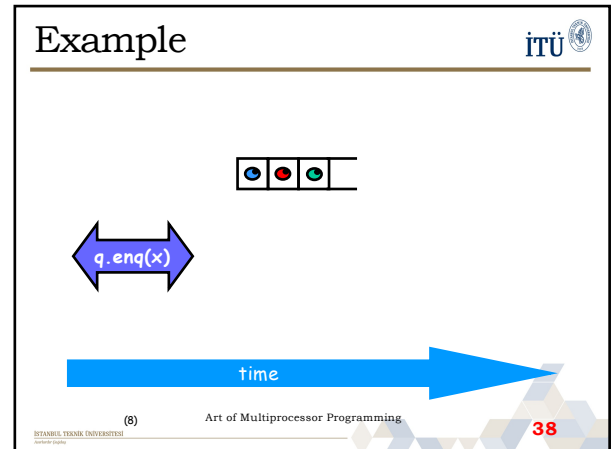
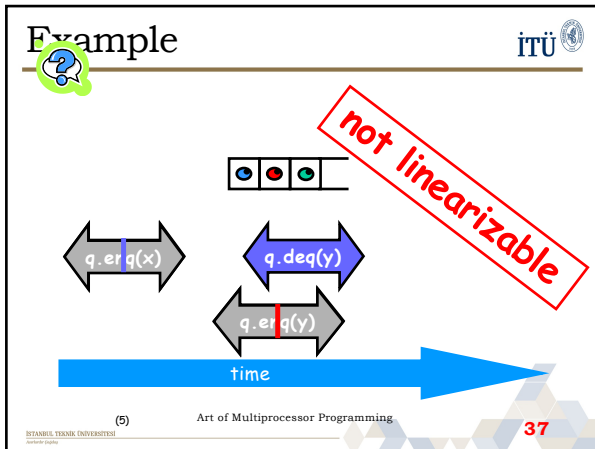
(6)

Art of Multiprocessor Programming

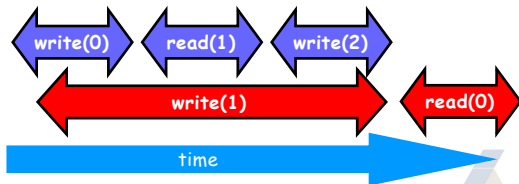
30

linearizable





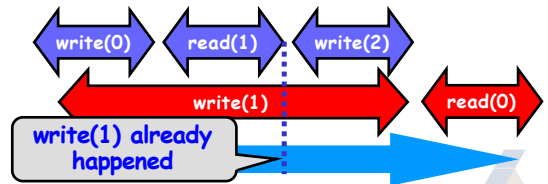
## Read/Write Register Example



(4) Art of Multiprocessor Programming

43

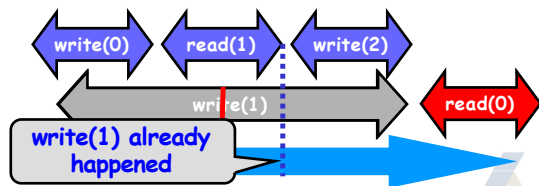
## Read/Write Register Example



(4) Art of Multiprocessor Programming

44

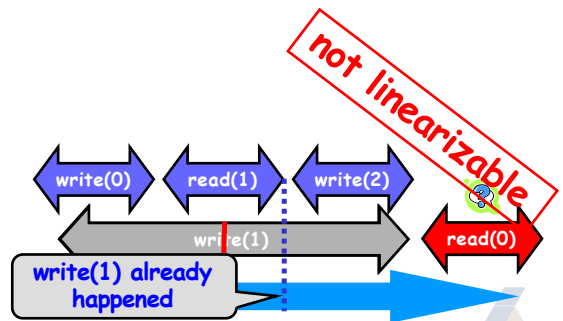
## Read/Write Register Example



(4) Art of Multiprocessor Programming

45

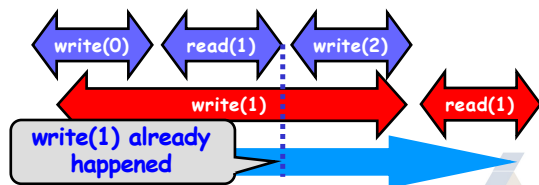
## Read/Write Register Example



(4) Art of Multiprocessor Programming

46

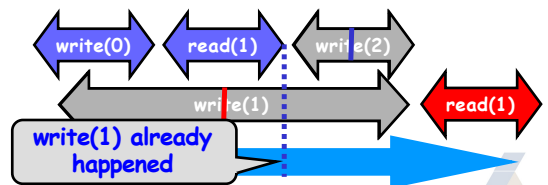
## Read/Write Register Example



(4) Art of Multiprocessor Programming

47

## Read/Write Register Example

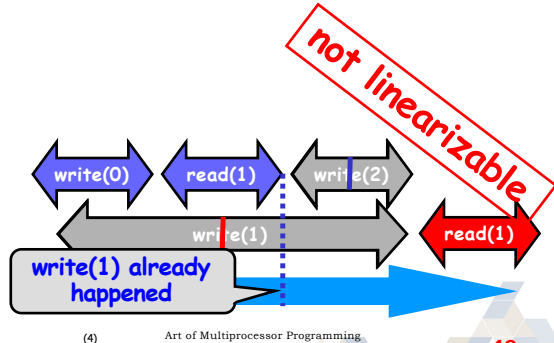


(4) Art of Multiprocessor Programming

48

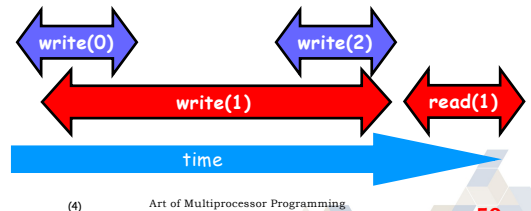


## Read/Write Register Example



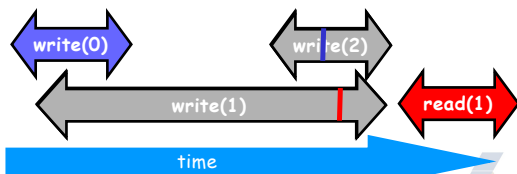
49

## Read/Write Register Example



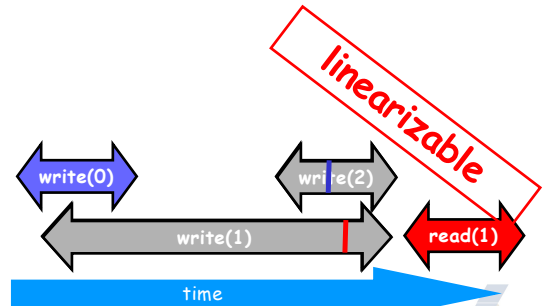
50

## Read/Write Register Example



51

## Read/Write Register Example



52

## Strategy

- Identify one atomic step where method “happens”
  - Critical section
  - Machine instruction
- Doesn't always work
  - Might need to define several different steps for a given method

Art of Multiprocessor Programming

53

## Linearizability: Summary

- Powerful specification tool for shared objects
- Allows us to capture the notion of objects being “atomic”
- Don't leave home without it

Art of Multiprocessor Programming

54

## Alternative: Sequential Consistency



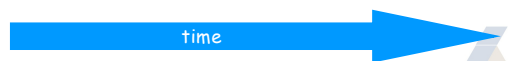
- No need to preserve real-time order
  - Cannot re-order operations done by the same thread
  - Can re-order non-overlapping operations done by different threads
- Often used to describe multiprocessor memory architectures

İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

55

## Example

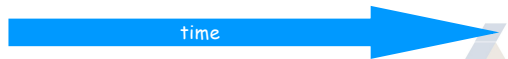
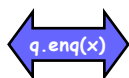


İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

56

## Example

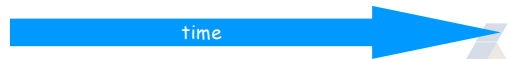
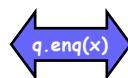


İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

57

## Example

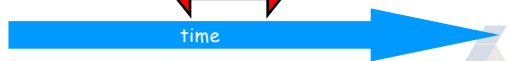
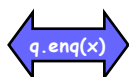


İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

58

## Example

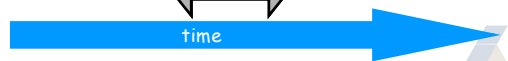


İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

59

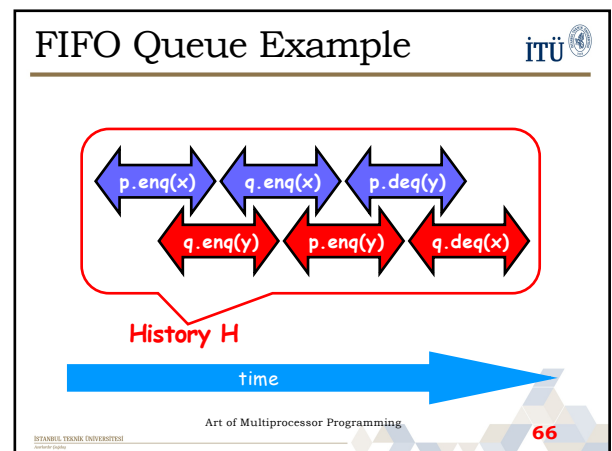
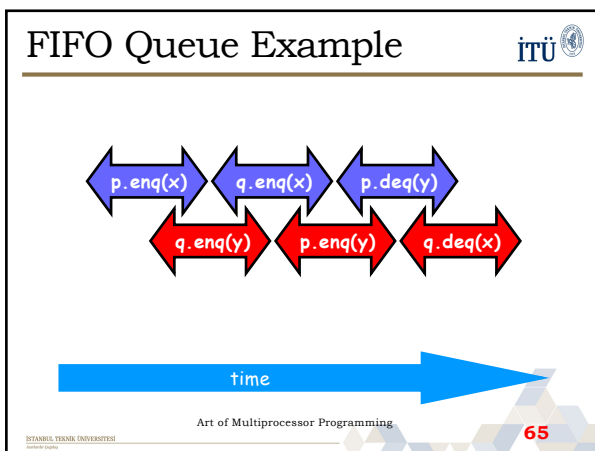
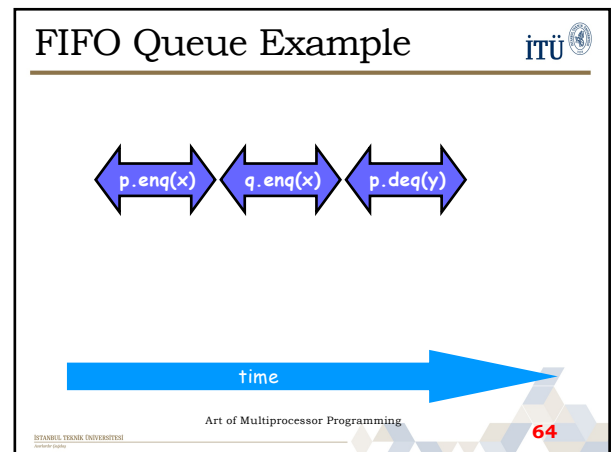
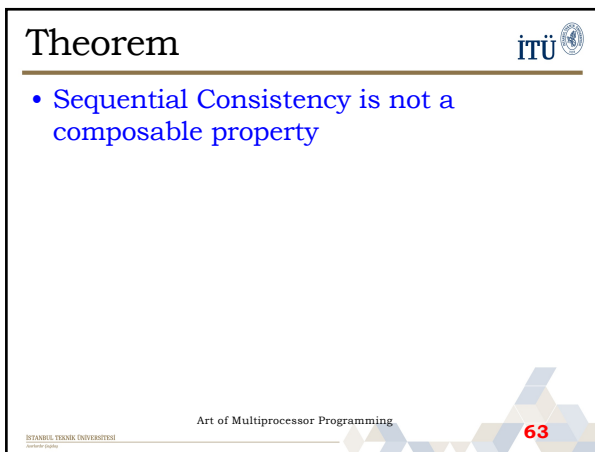
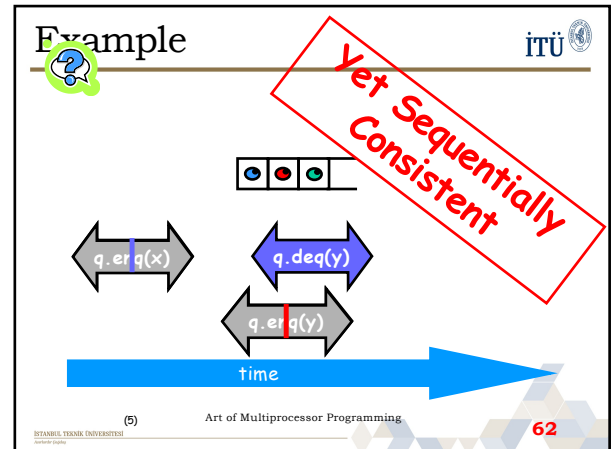
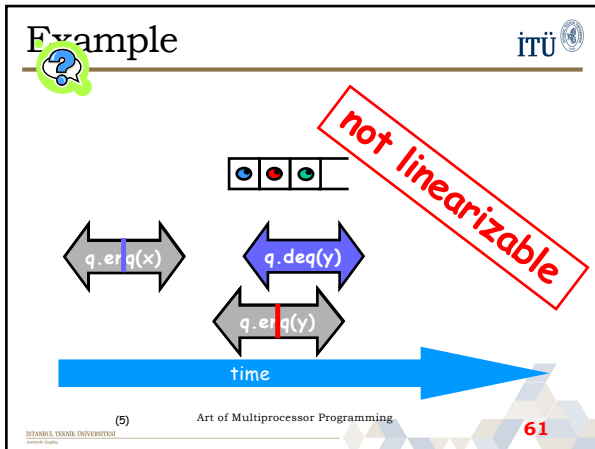
## Example



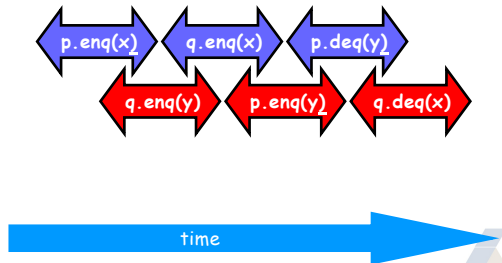
İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

60



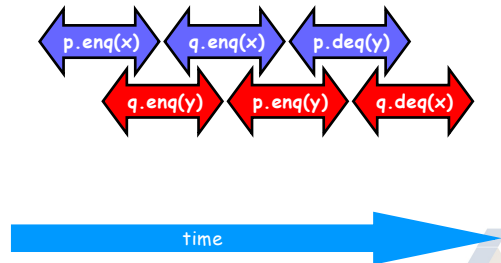
## H | p Sequentially Consistent



Art of Multiprocessor Programming

67

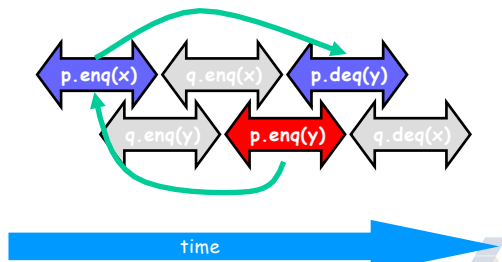
## H | q Sequentially Consistent



Art of Multiprocessor Programming

68

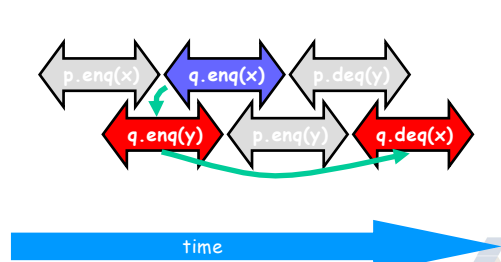
## Ordering imposed by p



Art of Multiprocessor Programming

69

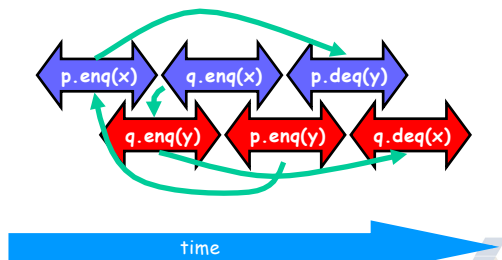
## Ordering imposed by q



Art of Multiprocessor Programming

70

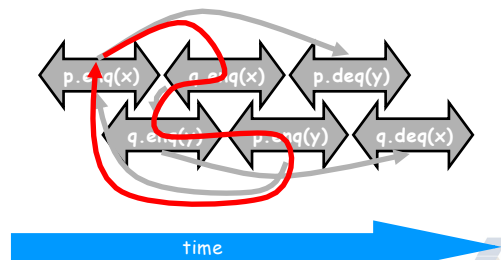
## Ordering imposed by both



Art of Multiprocessor Programming

71

## Combining orders



Art of Multiprocessor Programming

72

## Fact



- Most hardware architectures don't support sequential consistency
- Because they think it's too strong

İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

73

## Memory Operations



- To read a memory location,
  - load data into cache.
- To write a memory location
  - update cached copy,
  - Lazily write cached data back to memory

İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

74

## While Writing to Memory



- A processor can execute hundreds, or even thousands of instructions
- Why delay on every memory write?
- Instead, write back in parallel with rest of the program.

İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

75

## Explicit Synchronization



- Memory barrier instruction
  - Flush unwritten caches
  - Bring caches up to date
- Compilers often do this for you
  - Entering and leaving critical sections
- Expensive

İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

76

## Critical Sections



- Easy way to implement linearizability
  - Take sequential object
  - Make each method a critical section
- Problems
  - Blocking
  - No concurrency

İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

77

## Progress Conditions



- Deadlock-free: some thread trying to acquire the lock eventually succeeds.
- Starvation-free: every thread trying to acquire the lock eventually succeeds.
- Lock-free: some thread calling a method eventually returns.
- Wait-free: every thread calling a method eventually returns.

İSTANBUL TEKNİK ÜNİVERSİTESİ

Art of Multiprocessor Programming

78

## Progress Conditions



	Non-Blocking	Blocking
Everyone makes progress	Wait-free	Starvation-free
Someone makes progress	Lock-free	Deadlock-free

Art of Multiprocessor Programming

79

## Concurrent Registers



## Registers



```
public interface Register<T> {
    public T read();
    public void write(T v);
}
```

81

## Registers



```
public interface Register<T> {
    public T read();
    public void write(T v);
}
```

Type of register  
(usually Boolean or m-bit Integer)

82

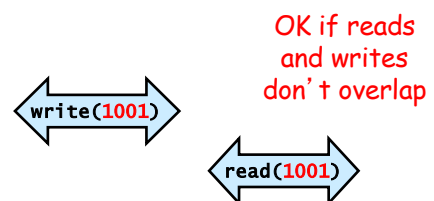
## Jargon Watch



- SRSW
  - Single-reader single-writer
- MRSW
  - Multi-reader single-writer
- MRMW
  - Multi-reader multi-writer

83

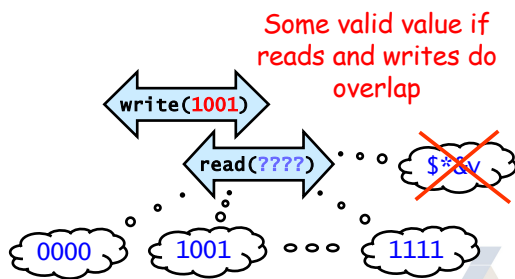
## Safe Register



(2)

84

## Safe Register



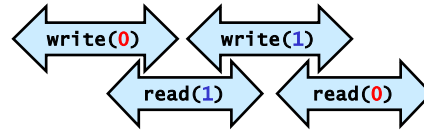
İTAMBUS, TEKNİK UNIVERSİTESİ

85

## Regular Register



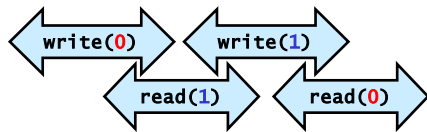
- Single Writer
- Readers return:
  - Old value if no overlap (safe)
  - Old or one of new values if overlap



İTAMBUS, TEKNİK UNIVERSİTESİ

86

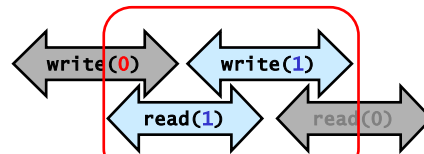
## Regular or Not?



İTAMBUS, TEKNİK UNIVERSİTESİ

87

## Regular or Not?

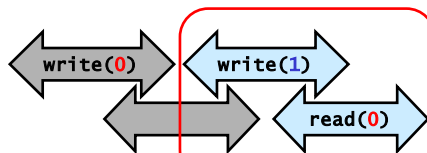


Overlap: returns new value

İTAMBUS, TEKNİK UNIVERSİTESİ

88

## Regular or Not?

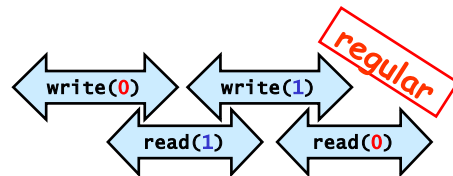


Overlap: returns old value

İTAMBUS, TEKNİK UNIVERSİTESİ

89

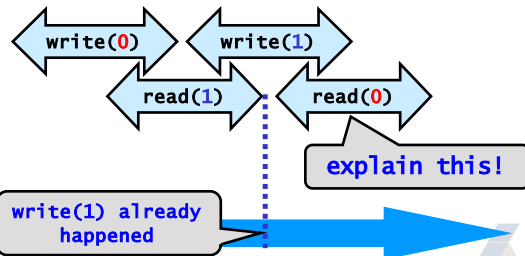
## Regular or Not?



İTAMBUS, TEKNİK UNIVERSİTESİ

90

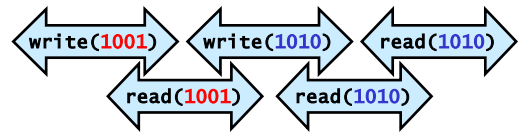
## Regular $\neq$ Atomic



İTAMBUS, TEKNİK UNIVERSİTESİ

91

## Atomic Register

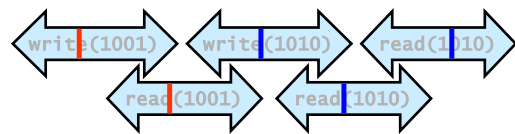


Linearizable to sequential safe register

İTAMBUS, TEKNİK UNIVERSİTESİ

92

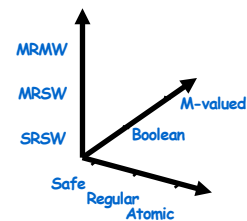
## Atomic Register



İTAMBUS, TEKNİK UNIVERSİTESİ

93

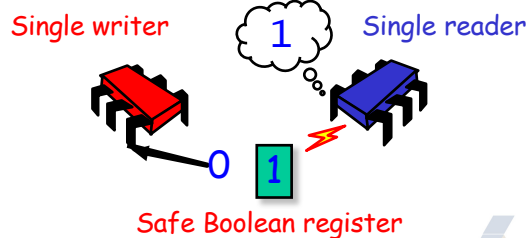
## Register Space



İTAMBUS, TEKNİK UNIVERSİTESİ

94

## Weakest Register



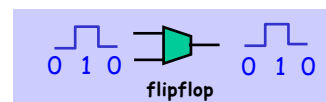
İTAMBUS, TEKNİK UNIVERSİTESİ

95

## Weakest Register



Single writer Single reader



Get correct reading if not during state transition

İTAMBUS, TEKNİK UNIVERSİTESİ

96



## Results



- From SRSW safe Boolean register
  - All the other registers
  - Mutual exclusion
- But not everything!
  - Consensus hierarchy

İTAMUL TEKNİK ÜNİVERSİTESİ

97

## Locking within Registers



- Not interesting to rely on mutual exclusion in register constructions
- We want registers to implement mutual exclusion!
- No fun to use mutual exclusion to implement itself!

İTAMUL TEKNİK ÜNİVERSİTESİ

98

## Road Map




- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic

İTAMUL TEKNİK ÜNİVERSİTESİ

99

## Road Map



- SRSW safe Boolean
- MRSW safe Boolean  Next
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic

İTAMUL TEKNİK ÜNİVERSİTESİ

100

## Register Names

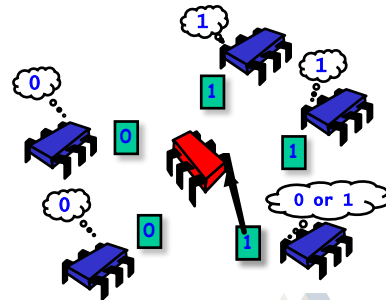


```
public class SafeBoolMRSWRegister
implements Register<Boolean> {
    public boolean read() { ... }
    public void write(boolean x) { ... }
}
```

İTAMUL TEKNİK ÜNİVERSİTESİ

101

## Safe Boolean MRSW from Safe Boolean SRSW



İTAMUL TEKNİK ÜNİVERSİTESİ

102

## Safe Boolean MRSW from Safe Boolean SRSW



```
public class SafeBoolMRSWRegister
implements Register<Boolean> {
    private SafeBoolSRSWRegister[] r =
        new SafeBoolSRSWRegister[N];
    public void write(boolean x) {
        for (int j = 0; j < N; j++)
            r[j].write(x);
    }
    public boolean read() {
        int i = ThreadID.get();
        return r[i].read();
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

103

## Safe Boolean MRSW from Safe Boolean SRSW



```
public class SafeBoolMRSWRegister
implements BooleanRegister {
    private SafeBoolSRSWRegister[] r =
        new SafeBoolSRSWRegister[N];
    public void write(boolean x) {
        for (int j = 0; j < N; j++)
            r[j].write(x);
    }
    public boolean read() {
        int i = ThreadID.get();
        return r[i].read();
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

104

Each thread has own safe SRSW register

## Road Map



- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean **Next**
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot

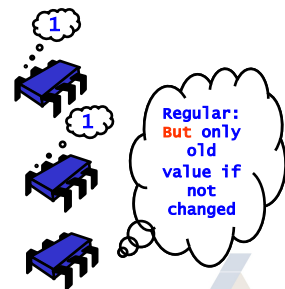
İSTANBUL TEKNİK ÜNİVERSİTESİ

105

## Regular Boolean MRSW from Safe Boolean MRSW



Safe register can return 0 or 1 even if the same value is written



İSTANBUL TEKNİK ÜNİVERSİTESİ

106

## Regular Boolean MRSW from Safe Boolean MRSW



```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    private boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

107

## Regular Boolean MRSW from Safe Boolean MRSW



```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    ThreadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

108

Last bit this thread wrote  
(OK, we're cheating here on Java syntax)

## Regular Boolean MRSW from Safe Boolean MRSW



```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    ThreadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

Actual value

İSTANBUL TEKNİK ÜNİVERSİTESİ

109

## Regular Boolean MRSW from Safe Boolean MRSW



```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    ThreadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

Is new value different from last value I wrote?

İSTANBUL TEKNİK ÜNİVERSİTESİ

110

## Regular Boolean MRSW from Safe Boolean MRSW



```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    ThreadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

If so, change it (otherwise don't!)

İSTANBUL TEKNİK ÜNİVERSİTESİ

111

## Regular Boolean MRSW from Safe Boolean MRSW



```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    ThreadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

•Overlap? No Overlap?  
•No problem  
•either Boolean value works

İSTANBUL TEKNİK ÜNİVERSİTESİ

112

## Road Map



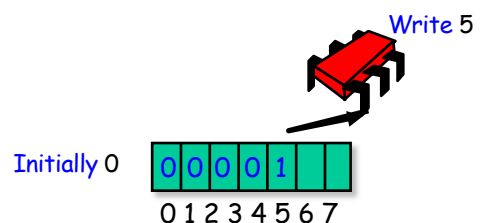
- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot

Next

İSTANBUL TEKNİK ÜNİVERSİTESİ

113

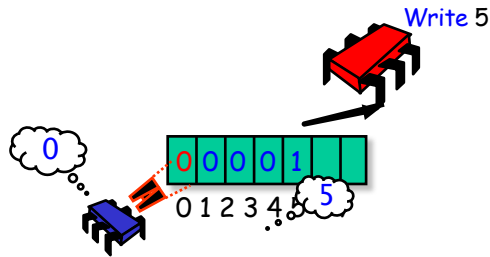
## Writing M-Valued



İSTANBUL TEKNİK ÜNİVERSİTESİ

114

## Writing M-Valued



İSTANBUL TEKNİK ÜNİVERSİTESİ

115

## MRSW Regular M-valued from MRSW Regular Boolean



```
public class RegMRSWRegister implements Register{
    RegBoolMRSWRegister[M] bit;
```

```
    public void write(int x) {
        this.bit[x].write(true);
        for (int i=x-1; i>=0; i--)
            this.bit[i].write(false);
    }
```

```
    public int read() {
        for (int i=0; i < M; i++)
            if (this.bit[i].read())
                return i;
    }
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

116

## MRSW Regular M-valued from MRSW Regular Boolean



```
public class RegMRSWRegister implements Register{
    RegBoolMRSWRegister[M] bit;
```

```
    public void write(int x) {
        this.bit[x].write(true);
        for (int i=x-1; i>=0; i--)
            this.bit[i].write(false);
    }
```

Unary representation:  
bit[i] means value i

```
    public int read() {
        for (int i=0; i < M; i++)
            if (this.bit[i].read())
                return i;
    }
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

117

## MRSW Regular M-valued from MRSW Regular Boolean



```
public class RegMRSWRegister implements Register {
    RegBoolMRSWRegister[m] bit;
```

```
    public void write(int x) {
        this.bit[x].write(true);
        for (int i=x-1; i>=0; i--)
            this.bit[i].write(false);
    }
```

Set bit x

```
    public int read() {
        for (int i=0; i < M; i++)
            if (this.bit[i].read())
                return i;
    }
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

118

## MRSW Regular M-valued from MRSW Regular Boolean



```
public class RegMRSWRegister implements Register {
    RegBoolMRSWRegister[m] bit;
```

```
    public void write(int x) {
        this.bit[x].write(true);
        for (int i=x-1; i>=0; i--)
            this.bit[i].write(false);
    }
```

Clear bits  
from higher  
to lower

```
    public int read() {
        for (int i=0; i < M; i++)
            if (this.bit[i].read())
                return i;
    }
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

119

## MRSW Regular M-valued from MRSW Regular Boolean



```
public class RegMRSWRegister implements Register {
    RegBoolMRSWRegister[m] bit;
```

```
    public void write(int x) {
        this.bit[x].write(true);
        for (int i=x-1; i>=0; i--)
            this.bit[i].write(false);
    }
```

Scan from lower  
to higher & return  
first bit set


```
    public int read() {
        for (int i=0; i < M; i++)
            if (this.bit[i].read())
                return i;
    }
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

120

## Road Map





- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic  **Next**
- MRMW atomic

İSTANBUL TEKNİK ÜNİVERSİTESİ

121

## Road Map (Slight Detour)



- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular 
- MRSW atomic  **SRSW Atomic**
- MRMW atomic

İSTANBUL TEKNİK ÜNİVERSİTESİ

122

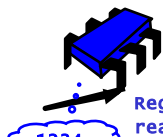
## SRSW Atomic From SRSW Regular



Regular writer



5678



Regular reader

1234

Instead of 5678...

Concurrent Reading

When is this a problem?

İSTANBUL TEKNİK ÜNİVERSİTESİ

123

## SRSW Atomic From SRSW Regular



Regular writer



5678



Regular reader

Initially 1234

Reg write(5678)

Reg read(5678)

Same as Atomic

time

İSTANBUL TEKNİK ÜNİVERSİTESİ

124

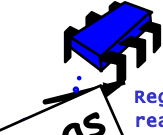
## SRSW Atomic From SRSW Regular



Regular writer



5678



Regular reader

678...

Initially 1234

Reg write(5678)

Reg read(1234)

Same as Atomic

time

İSTANBUL TEKNİK ÜNİVERSİTESİ

125

## SRSW Atomic From SRSW Regular



Regular writer



5678



Regular reader

5678...

Initially 1234

Reg write(5678)

Reg read(5678)

Reg read(1234)

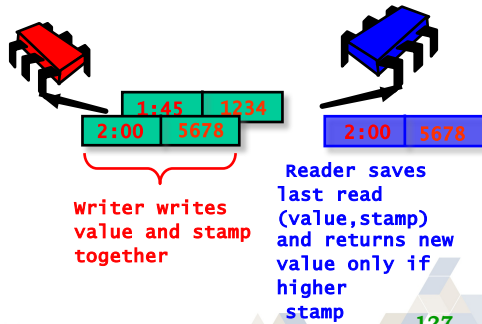
not Atomic!

Write 5678 happened

İSTANBUL TEKNİK ÜNİVERSİTESİ

126

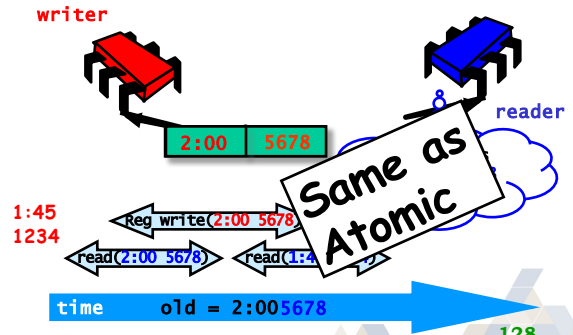
## Timestamped Values



İSTANBUL TEKNİK ÜNİVERSİTESİ

127

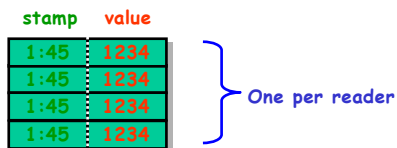
## SRSW Atomic From SRSW Regular



İSTANBUL TEKNİK ÜNİVERSİTESİ

128

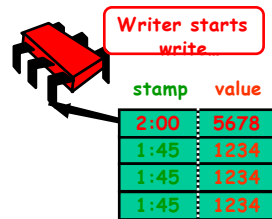
## Atomic Single Reader to Atomic Multi-Reader



İSTANBUL TEKNİK ÜNİVERSİTESİ

129

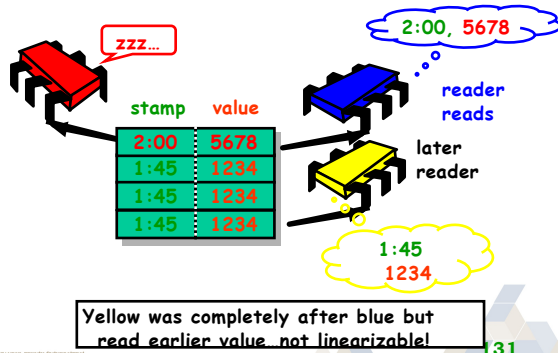
## Another Scenario



İSTANBUL TEKNİK ÜNİVERSİTESİ

130

## Another Scenario



İSTANBUL TEKNİK ÜNİVERSİTESİ

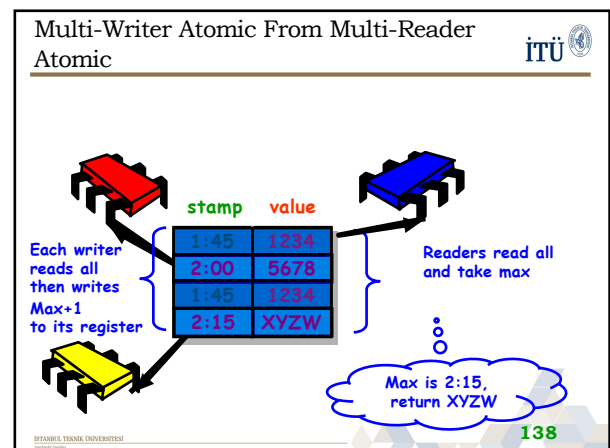
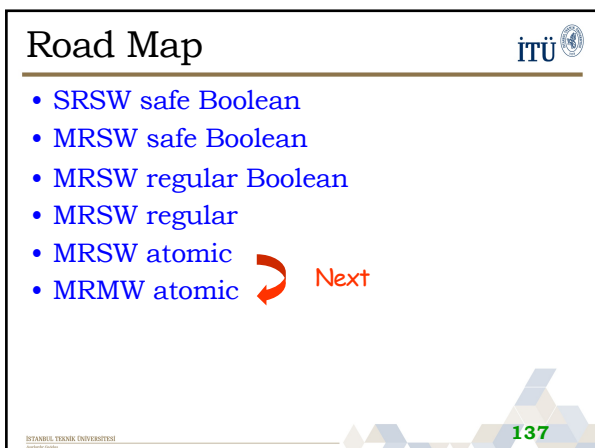
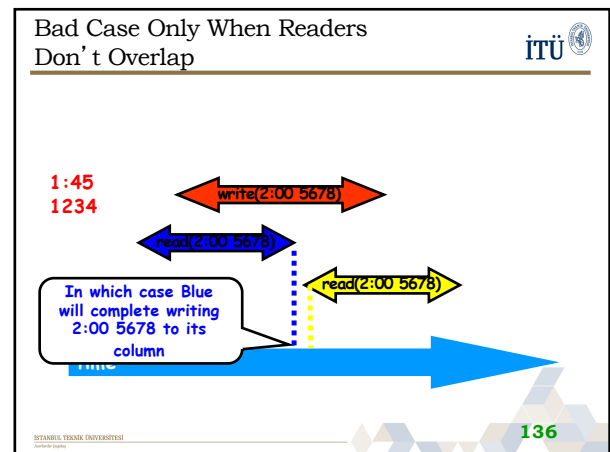
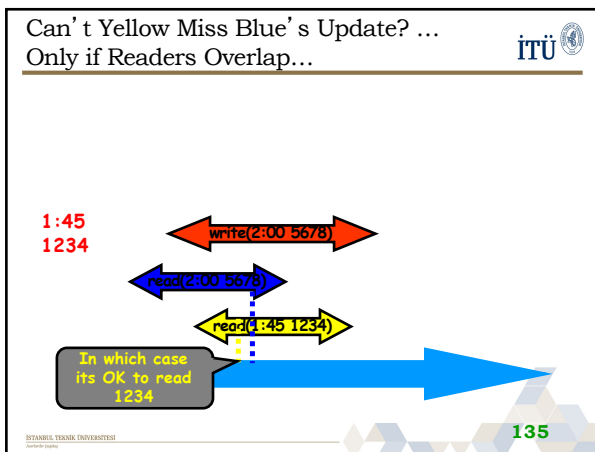
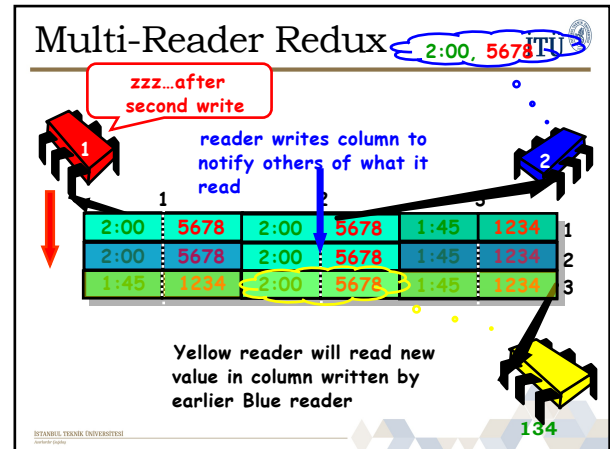
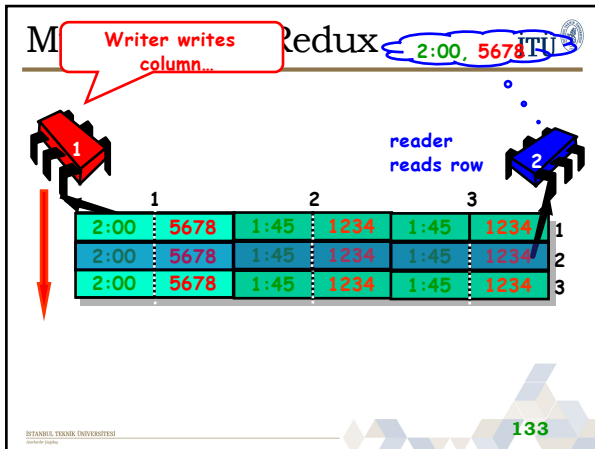
131

## Multi-Reader Redux



İSTANBUL TEKNİK ÜNİVERSİTESİ

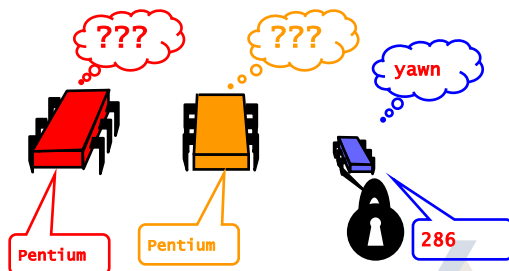
132



## Consensus Values

## Why is Mutual Exclusion so wrong?

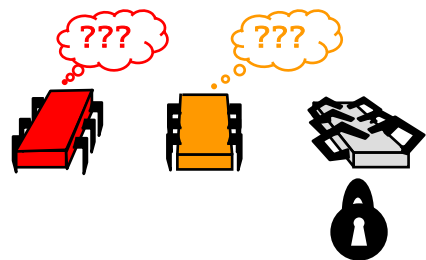
## Heterogeneous Processors



Art of Multiprocessor Programming

141

## Fault-tolerance



(2)

Art of Multiprocessor Programming

142

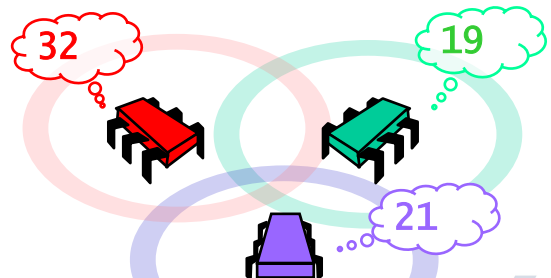
## Basic Questions

- Wait-Free synchronization might be a good idea in principle
- But how do you do it
  - Systematically?
  - Correctly?
  - Efficiently?

Art of Multiprocessor Programming

143

## Consensus: Each Thread has a Private Input

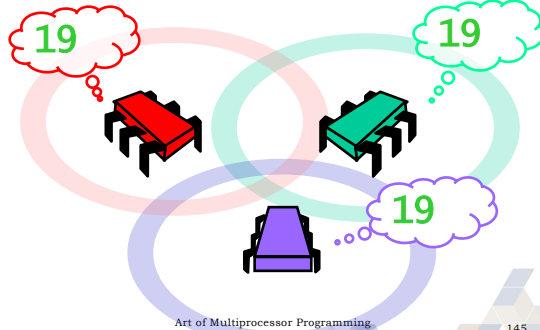


Art of Multiprocessor Programming

144



## They Agree on One Thread's Input



145

## Formally: Consensus



Consistent: all threads decide the same value  
Valid: the common decision value is some thread's input

Art of Multiprocessor Programming.

146

## Consensus Object



```
public interface Consensus {
    Object decide(Object value);
}
```

Art of Multiprocessor Programming.

147

## Java Jargon Watch



- Define Consensus protocol as an abstract class
- We implement some methods

Art of Multiprocessor Programming.

148

## Generic Consensus Protocol

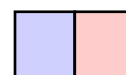


```
abstract class ConsensusProtocol implements Consensus {
    protected Object[] proposed = new Object[N];
    protected void propose(Object value) {
        proposed[ThreadID.get()] = value;
    }
    abstract public Object decide(Object value);
}
```

Art of Multiprocessor Programming.

149

## FIFO Consensus



proposed array



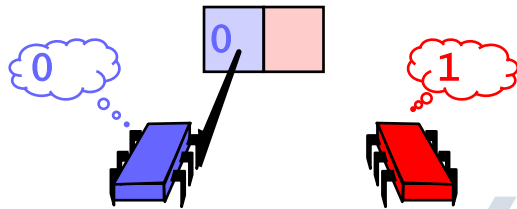
FIFO Queue  
with red and  
black balls

Coveted red ball      Dreaded black ball

Art of Multiprocessor Programming.

150

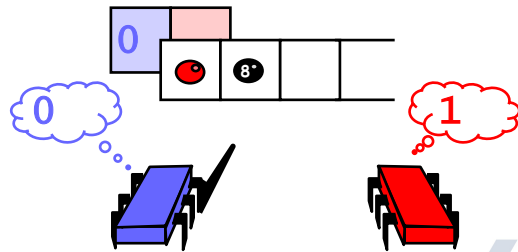
## Protocol: Write Value to Array



Art of Multiprocessor Programming

151

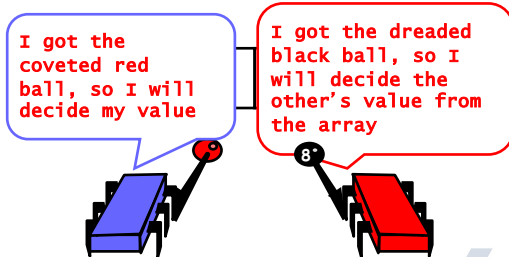
## Protocol: Take Next Item from Queue



Art of Multiprocessor Programming

152

## Protocol: Take Next Item from Queue



Art of Multiprocessor Programming

153

## Consensus Using FIFO Queue



```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;
    public QueueConsensus() {
        queue = new Queue();
        queue.enq(Ball.RED);
        queue.enq(Ball.BLACK);
    }
    ...
}
```

Art of Multiprocessor Programming

154

## Who Won?



```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;
    ...
    public decide(object value) {
        propose(value);
        Ball ball = this.queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

Art of Multiprocessor Programming

155

## Why does this Work?



- If one thread gets the red ball
- Then the other gets the black ball
- Winner decides her own value
- Loser can find winner's value in array
  - Because threads write array
  - Before dequeuing from queue

Art of Multiprocessor Programming

156

## Consensus Numbers



- An object  $X$  has **consensus number**  $n$ 
  - If it can be used to solve  $n$ -thread consensus
    - Take any number of instances of  $X$
    - together with atomic read/write registers
    - and implement  $n$ -thread consensus
  - But not  $(n+1)$ -thread consensus

## Consensus Numbers



- Theorem
  - Atomic read/write registers have consensus number 1
- Theorem
  - Multi-dequeue FIFO queues have consensus number 2

## Consensus Numbers Measure Synchronization Power



- Theorem
  - If you can implement  $X$  from  $Y$
  - And  $X$  has consensus number  $c$
  - Then  $Y$  has consensus number at least  $c$

## Synchronization Speed Limit



- Conversely
  - If  $X$  has consensus number  $c$
  - And  $Y$  has consensus number  $d < c$
  - Then there is no way to construct a wait-free implementation of  $X$  by  $Y$

## Read-Modify-Write Objects



- Method call
  - Returns object's prior value  $x$
  - Replaces  $x$  with **mumble( $x$ )**

## RMW Everywhere!



- Most synchronization instructions
  - are RMW methods
- The rest
  - Can be trivially transformed into RMW methods

## Example: getAndSet



```
public abstract class RMWRegister {
    private int value;

    public int synchronized getAndSet(int v)
    {
        int prior = this.value;
        this.value = v;
        return prior;
    }
    ...
}
```

Art of Multiprocessor Programming

163

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ

## getAndIncrement



```
public abstract class RMWRegister {
    private int value;

    public int synchronized getAndIncrement()
    {
        int prior = this.value;
        this.value = this.value + 1;
        return prior;
    }
    ...
}
```

Art of Multiprocessor Programming

164

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ

## getAndAdd



```
public abstract class RMWRegister {
    private int value;

    public int synchronized getAndAdd(int a)
    {
        int prior = this.value;
        this.value = this.value + a;
        return prior;
    }
    ...
}
```

Art of Multiprocessor Programming

165

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ

## compareAndSet



```
public abstract class RMWRegister {
    private int value;
    public boolean synchronized compareAndSet(
        int expected, int update) {

        int prior = this.value;
        if (this.value == expected) {
            this.value = update;
            return true;
        }
        return false;
    } ...
}
```

Art of Multiprocessor Programming

166

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ

## compareAndSet Has $\infty$ Consensus Number



```
public class RMWConsensus extends ConsensusProtocol {
    private AtomicInteger r = new AtomicInteger(-1);
    public Object decide(Object value) {
        propose(value);
        r.compareAndSet(-1, 1);
        return proposed[r.get()];
    }
}
```

Art of Multiprocessor Programming

167

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ

## The Consensus Hierarchy



1 Read/Write Registers, Snapshots...
2 getAndSet, getAndIncrement, ...
•
•
•
$\infty$ compareAndSet, ...

Art of Multiprocessor Programming

168

İSTANBUL TEKNİK ÜNİVERSİTESİ  
İTÜ

## Atomic variables



- `java.util.concurrent.atomic` contains classes representing scalars supporting "CAS"  
`boolean compareAndSet(expectedV, newV)`
  - Atomically set to new value if holding expected value
  - Always used in a loop
- Essential for writing efficient code on MPs
  - Nonblocking data structures, optimistic algorithms, reducing overhead and contention when updates center on a single field
- JVMs use best construct available on platform
  - Compare-and-swap, Load-linked/Store-conditional, Locks

İTAMUL YENI İNİVİRSİTİ  
İTAMUL YENI İNİVİRSİTİ

## Atomic Variables



- Java concurrency package also supplies reflection-based classes that allow CAS on given volatile fields of other classes
- Also provides methods for getting and unconditionally setting values
- Instance of classes `AtomicInteger`, `AtomicBoolean`, `AtomicLong`, `AtomicReference` provide access and update to single variables of that type

İTAMUL YENI İNİVİRSİTİ  
İTAMUL YENI İNİVİRSİTİ

## Atomic Variable Example



- Faster and less contention in programs with a single `Random` accessed by many threads
- Let's see `java.util.Random`

```
class Random { // snippets
    private AtomicLong seed;
    Random(long s) {
        seed = new AtomicLong(s);
    }
    long next(){
        for(;;) {
            long s = seed.get();
            long nexts = s * ... + ...;
            if (seed.compareAndSet(s,nexts))
                return s;
        }
    }
}
```

İTAMUL YENI İNİVİRSİTİ  
İTAMUL YENI İNİVİRSİTİ

This work is licensed under a [Creative Commons Attribution-ShareAlike 3.0 license](https://creativecommons.org/licenses/by-sa/3.0/).



- You are free:
  - to Share — to copy, distribute and transmit the work
  - to Remix — to adapt the work
- Under the following conditions:
  - Attribution. You must attribute the work to "The Art of Multiprocessor Programming" (but not in any way that suggests that the authors endorse you or your use of the work).
  - Share Alike. If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.

Art of Multiprocessor Programming

172

İTAMUL YENI İNİVİRSİTİ  
İTAMUL YENI İNİVİRSİTİ