

Futures, Scheduling, and Work Distribution

Companion slides for
The Art of Multiprocessor Programming
by Maurice Herlihy & Nir Shavit

How to write Parallel Apps?

- How to
 - split a program into parallel parts
 - In an effective way
 - Thread management

Matrix Multiplication

$$(C) = (A) \bullet (B)$$

Matrix Multiplication

$$c_{ij} = \sum_{k=0}^{N-1} a_{ik} * b_{kj}$$

Matrix Multiplication

```
class Worker extends Thread {
    int row, col;
    Worker(int row, int col) {
        this.row = row;
        this.col = col;
    }
    public void run() {
        double dotProduct = 0.0;
        for (int i = 0; i < n; i++)
            dotProduct += a[row][i] * b[i][col];
        c[row][col] = dotProduct;
    }
}
```

Matrix Multiplication

```
class Worker extends Thread {
    int row, col;
    Worker(int row, int col) {
        this.row = row; this.col = col;
    }
    public void run() {
        double dotProduct = 0.0;
        for (int i = 0; i < n; i++)
            dotProduct += a[row][i] * b[i][col];
        c[row][col] = dotProduct;
    }
}
```

a thread

Matrix Multiplication



```
class Worker extends Thread {
    int row, col;
    Worker(int row, int col) {
        this.row = row; this.col = col;
    }
    public void run() {
        double dotProduct = 0.0;
        for (int i = 0; i < n; i++)
            dotProduct += a[row][i] * b[i][col];
        c[row][col] = dotProduct;
    }
}
```

Which matrix entry
to compute

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Multiplication



```
class Worker extends Thread {
    int row, col;
    Worker(int row, int col) {
        this.row = row; this.col = col;
    }
    public void run() {
        double dotProduct = 0.0;
        for (int i = 0; i < n; i++)
            dotProduct += a[row][i] * b[i][col];
        c[row][col] = dotProduct;
    }
}
```

Actual computation

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Multiplication



```
void multiply() {
    Worker[][] worker = new Worker[n][n];
    for (int row ...)
        for (int col ...)
            worker[row][col] = new Worker(row,col);
    for (int row ...)
        for (int col ...)
            worker[row][col].start();
    for (int row ...)
        for (int col ...)
            worker[row][col].join();
}
```

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Multiplication



```
void multiply() {
    Worker[][] worker = new Worker[n][n];
    for (int row ...)
        for (int col ...)
            worker[row][col] = new Worker(row,col);
    for (int row ...)
        for (int col ...)
            worker[row][col].start();
    for (int row ...)
        for (int col ...)
            worker[row][col].join();
}
```

Create nxn
threads

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Multiplication



```
void multiply() {
    Worker[][] worker = new Worker[n][n];
    for (int row ...)
        for (int col ...)
            worker[row][col] = new Worker(row,col);
    for (int row ...)
        for (int col ...)
            worker[row][col].start();
    for (int row ...)
        for (int col ...)
            worker[row][col].join();
}
```

Start them

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Multiplication



```
void multiply() {
    Worker[][] worker = new Worker[n][n];
    for (int row ...)
        for (int col ...)
            worker[row][col] = new Worker(row,col);
    for (int row ...)
        for (int col ...)
            worker[row][col].start();
    for (int row ...)
        for (int col ...)
            worker[row][col].join();
}
```

Wait for
them to
finish

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Addition



$$\begin{pmatrix} C_{00} & C_{01} \\ C_{10} & C_{11} \end{pmatrix} = \begin{pmatrix} A_{00} + B_{00} & B_{01} + A_{01} \\ A_{10} + B_{10} & A_{11} + B_{11} \end{pmatrix}$$

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Addition



$$\begin{pmatrix} C_{00} & C_{01} \\ C_{10} & C_{11} \end{pmatrix} = \begin{pmatrix} A_{00} + B_{00} & B_{01} + A_{01} \\ A_{10} + B_{10} & A_{11} + B_{11} \end{pmatrix}$$

4 parallel additions

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Addition Task



```
class AddTask implements Runnable {
    Matrix a, b; // multiply this!
    public void run() {
        if (a.dim == 1) {
            c[0][0] = a[0][0] + b[0][0]; // base case
        } else {
            (partition a, b into half-size matrices aij and bij)
            Future<?> f00 = exec.submit(add(a00, b00));
            ...
            Future<?> f11 = exec.submit(add(a11, b11));
            f00.get(); ...; f11.get();
            ...
        }
    }
}
```

This is not real Java
Code

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Addition Task



```
class AddTask implements Runnable {
    Matrix a, b; // multiply this!
    public void run() {
        if (a.dim == 1) {
            c[0][0] = a[0][0] + b[0][0]; // base case
        } else {
            (partition a, b into half-size matrices aij and bij)
            Future<?> f00 = exec.submit(add(a00, b00));
            ...
            Future<?> f11 = exec.submit(add(a11, b11));
            f00.get(); ...; f11.get();
            ...
        }
    }
}
```

Base case: add directly

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Addition Task



```
class AddTask implements Runnable {
    Matrix a, b; // multiply this!
    public void run() {
        if (a.dim == 1) {
            c[0][0] = a[0][0] + b[0][0]; // base case
        } else {
            (partition a, b into half-size matrices aij and bij)
            Future<?> f00 = exec.submit(add(a00, b00));
            ...
            Future<?> f11 = exec.submit(add(a11, b11));
            f00.get(); ...; f11.get();
            ...
        }
    }
}
```

Constant-time operation

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Addition Task



```
class AddTask implements Runnable {
    Matrix a, b; // multiply this!
    public void run() {
        if (a.dim == 1) {
            c[0][0] = a[0][0] + b[0][0]; // base case
        } else {
            (partition a, b into half-size matrices aij and bij)
            Future<?> f00 = exec.submit(add(a00, b00));
            ...
            Future<?> f11 = exec.submit(add(a11, b11));
            f00.get(); ...; f11.get();
            ...
        }
    }
}
```

Submit 4 tasks

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Addition Task



```
class AddTask implements Runnable {
    Matrix a, b; // multiply this!
    public void run() {
        if (a.dim == 1) {
            c[0][0] = a[0][0] + b[0][0]; // base case
        } else {
            (partition a, b into half-size matrices aij and bij)
            Future<?> f00 = exec.submit(add(a00, b00));
            ...
            Future<?> f11 = exec.submit(add(a11, b11));
            f00.get(); ...; f11.get();
        }
    }
}
```

Let them finish

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Dependencies



- Matrix example is not typical
- Tasks are independent
 - Don't need results of one task ...
 - To complete another
- Often tasks are not independent

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Fibonacci



- Note
 - potential parallelism
 - Dependencies

$$F(n) = \begin{cases} 1 & \text{if } n = 0 \text{ or } 1 \\ F(n-1) + F(n-2) & \text{otherwise} \end{cases}$$

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Multithreaded Fibonacci



```
class FibTask implements Callable<Integer> {
    static ExecutorService exec = Executors.newCachedThreadPool();
    int arg;
    public FibTask(int n) {
        arg = n;
    }
    public Integer call() {
        if (arg > 2) {
            Future<Integer> left = exec.submit(new FibTask(arg-1));
            Future<Integer> right = exec.submit(new FibTask(arg-2));
            return left.get() + right.get();
        } else {
            return 1;
        }
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Multithreaded Fibonacci



```
class FibTask implements Callable<Integer> {
    static ExecutorService exec =
    Executors.newCachedThreadPool();
    int arg;
    public FibTask(int n) {
        arg = n;
    }
    public Integer call() {
        if (arg > 2) {
            Future<Integer> left = exec.submit(new FibTask(arg-1));
            Future<Integer> right = exec.submit(new FibTask(arg-2));
            return left.get() + right.get();
        } else {
            return 1;
        }
    }
}
```

Parallel calls

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Multithreaded Fibonacci



```
class FibTask implements Callable<Integer> {
    static ExecutorService exec =
    Executors.newCachedThreadPool();
    int arg;
    public FibTask(int n) {
        arg = n;
    }
    public Integer call() {
        if (arg > 2) {
            Future<Integer> left = exec.submit(new FibTask(arg-1));
            Future<Integer> right = exec.submit(new FibTask(arg-2));
            return left.get() + right.get();
        } else {
            return 1;
        }
    }
}
```

Pick up & combine results

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ


Dynamic Behavior



- Multithreaded program is
 - A directed acyclic graph (DAG)
 - That unfolds dynamically
- Each node is
 - A single unit of work


İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Fib DAG



İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Arrows Reflect Dependencies



İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ


How Parallel is That?



- Define work:
 - Total time on one processor
- Define critical-path length:
 - Longest dependency path
 - Can't beat that!


İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Fib Work




İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Fib Work




İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Fib Critical Path



İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Fib Critical Path



İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Notation Watch



- T_P = time on P processors
- T_1 = work (time on 1 processor)
- T_∞ = critical path length (time on ∞ processors)

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Simple Bounds



- $T_P \geq T_\infty \geq T_1 / P$
 - In one step, can't do more than P work
 - Can't beat inherent sequentialism

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

More Notation Watch



- Speedup on P processors
 - Ratio T_1 / T_P How much faster with P processors
- Linear speedup
 - $T_1 / T_P = \Theta(P)$
- Max speedup (average parallelism)
 - T_1 / T_∞

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Addition



$$\begin{pmatrix} C_{00} & C_{01} \\ C_{10} & C_{11} \end{pmatrix} = \begin{pmatrix} A_{00} + B_{00} & B_{01} + A_{01} \\ A_{10} + B_{10} & A_{11} + B_{11} \end{pmatrix}$$

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Addition



$$\begin{pmatrix} C_{00} & C_{01} \\ C_{10} & C_{11} \end{pmatrix} = \begin{pmatrix} A_{00} + B_{00} & B_{01} + A_{01} \\ A_{10} + B_{10} & A_{11} + B_{11} \end{pmatrix}$$

4 parallel additions

İSTANBUL TEKNİK ÜNİVERSİTESİ

Addition



- Let $A_P(n)$ be running time
 - For $n \times n$ matrix
 - on P processors
- For example
 - $A_1(n)$ is work
 - $A_\infty(n)$ is critical path length

İSTANBUL TEKNİK ÜNİVERSİTESİ

Addition



- Work is

Partition, synch, etc

$$A_1(n) = 4 A_1(n/2) + \Theta(1)$$

4 spawned additions

İSTANBUL TEKNİK ÜNİVERSİTESİ

Addition



- Work is

$$\begin{aligned} A_1(n) &= 4 A_1(n/2) + \Theta(1) \\ &= \Theta(n^2) \end{aligned}$$

Same as double-loop summation

İSTANBUL TEKNİK ÜNİVERSİTESİ

Addition



- Critical Path length is

$$A_\infty(n) = A_\infty(n/2) + \Theta(1)$$

spawned additions in parallel

Partition, synch, etc

İSTANBUL TEKNİK ÜNİVERSİTESİ

Addition



- Critical Path length is

$$\begin{aligned} A_\infty(n) &= A_\infty(n/2) + \Theta(1) \\ &= \Theta(\log n) \end{aligned}$$

İSTANBUL TEKNİK ÜNİVERSİTESİ

Matrix Multiplication Redux

$$(C) = (A) \cdot (B)$$

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Matrix Multiplication Redux

$$\begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix} = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \cdot \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix}$$

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

First Phase ...

$$\begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix} = \begin{pmatrix} A_{11}B_{11} + A_{12}B_{21} & A_{11}B_{12} + A_{12}B_{22} \\ A_{21}B_{11} + A_{22}B_{21} & A_{21}B_{12} + A_{22}B_{22} \end{pmatrix}$$

8 multiplications

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Second Phase ...

$$\begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix} = \begin{pmatrix} A_{11}B_{11} + A_{12}B_{21} & A_{11}B_{12} + A_{12}B_{22} \\ A_{21}B_{11} + A_{22}B_{21} & A_{21}B_{12} + A_{22}B_{22} \end{pmatrix}$$

4 additions

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Multiplication

- Work is 8 parallel multiplications
 - Final addition
- $$M_1(n) = 8 M_1(n/2) + 4A_1(n/2)$$
- $$= 8 M_1(n/2) + A_1(n)$$

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Multiplication

- Work is
- $$M_1(n) = 8 M_1(n/2) + \Theta(n^2)$$
- $$= \Theta(n^3)$$

Same as serial triple-nested loop

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Multiplication



- Critical path length is

$$M_{\infty}(n) = M_{\infty}(n/2) + A_{\infty}(n)$$

Final addition

Half-size parallel multiplications

İSTANBUL TEKNİK ÜNİVERSİTESİ

Multiplication



- Critical path length is

$$\begin{aligned} M_{\infty}(n) &= M_{\infty}(n/2) + A_{\infty}(n) \\ &= M_{\infty}(n/2) + \Theta(\log n) \\ &= \Theta(\log^2 n) \end{aligned}$$

İSTANBUL TEKNİK ÜNİVERSİTESİ

Parallelism



- $M_1(n) / M_{\infty}(n) = \Theta(n^3 / \log^2 n)$
- To multiply two 1000 x 1000 matrices
 - $1000^3 / 10^2 = 10^7$
- Much more than number of processors on any real machine

İSTANBUL TEKNİK ÜNİVERSİTESİ

Shared-Memory Multiprocessors



- Parallel applications
 - Do not have direct access to HW processors
- Mix of other jobs
 - All run together
 - Come & go dynamically

İSTANBUL TEKNİK ÜNİVERSİTESİ

Matrix Multiplication



```
void multiply() {
    Worker[][] worker = new Worker[n][n];
    for (int row ...)
        for (int col ...)
            worker[row][col] = new Worker(row,col);
    for (int row ...)
        for (int col ...)
            worker[row][col].join();
}
```

Start them

What's wrong with this picture?

Wait for them to finish

İSTANBUL TEKNİK ÜNİVERSİTESİ

Thread Overhead



- Threads Require resources
 - Memory for stacks
 - Setup, teardown
- Scheduler overhead
- Worse for short-lived threads

İSTANBUL TEKNİK ÜNİVERSİTESİ

Thread Pools



- More sensible to keep a pool of long-lived threads
- Threads assigned short-lived tasks
 - Runs the task
 - Rejoins pool
 - Waits for next assignment

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Thread Pool = Abstraction



- Insulate programmer from platform
 - Big machine, big pool
 - And vice-versa
- Portable code
 - Runs well on any platform
 - No need to mix algorithm/platform concerns

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

ExecutorService Interface



- In `java.util.concurrent`
 - Task = `Runnable` object
 - If no result value expected
 - Calls `run()` method.
 - Task = `Callable<T>` object
 - If result value of type `T` expected
 - Calls `T call()` method.

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Future<T>



```
Callable<T> task = ...;  
...  
Future<T> future = executor.submit(task);  
...  
T value = future.get();
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Future<T>



```
Callable<T> task = ...;  
...  
Future<T> future = executor.submit(task);  
...  
T value = future.get();
```

Submitting a `Callable<T>` task
returns a `Future<T>` object

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Future<T>



```
Callable<T> task = ...;  
...  
Future<T> future = executor.submit(task);  
...  
T value = future.get();
```

The Future's `get()` method
blocks until the value is available

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Future<?>



```
Runnable task = ...;  
...  
Future<?> future = executor.submit(task);  
...  
future.get();
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Future<?>



```
Runnable task = ...;  
...  
Future<?> future = executor.submit(task);  
...  
future.get();
```

Submitting a **Runnable** task
returns a **Future<?>** object

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Future<?>



```
Runnable task = ...;  
...  
Future<?> future = executor.submit(task);  
...  
future.get();
```

The Future's **get()** method blocks
until the computation is complete

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ


Note



- Executor Service submissions
 - Are purely advisory in nature
- The executor
 - Is free to ignore any such advice
 - And could execute tasks sequentially ...


İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Ideal Scheduling Hierarchy



İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Realistic Scheduling Hierarchy



İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

For Example



- Initially,
 - All P processors available for application
- Serial computation
 - Takes over one processor
 - Leaving P-1 for us
 - Waits for I/O
 - We get that processor back

İSTANBUL TEKNİK ÜNİVERSİTESİ

Speedup



- Map threads onto P processes
- Cannot get P-fold speedup
 - What if the kernel doesn't cooperate?
- Can try for speedup proportional to
 - time-averaged number of processors
 - the kernel gives us

İSTANBUL TEKNİK ÜNİVERSİTESİ

Scheduling Hierarchy



- User-level scheduler
 - Tells kernel which threads are ready
- Kernel-level scheduler
 - Synchronous (for analysis, not correctness!)
 - Picks p_i threads to schedule at step i
 - Processor average
 - over T steps is:

$$P_A = \frac{1}{T} \sum_{i=1}^T p_i$$

İSTANBUL TEKNİK ÜNİVERSİTESİ


Greed is Good



- Greedy scheduler
 - Schedules as much as it can
 - At each time step
- Optimal schedule is greedy
- But not every greedy schedule is optimal


İSTANBUL TEKNİK ÜNİVERSİTESİ

Work Distribution




İSTANBUL TEKNİK ÜNİVERSİTESİ

Work Dealing



İSTANBUL TEKNİK ÜNİVERSİTESİ

Work Balancing



İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Work-Balancing Thread



```
public void run() {
    int me = ThreadID.getID();
    while (true) {
        Runnable task = queue[me].deq();
        if (task != null) task.run();
        int size = queue[me].size();
        if (random.nextInt(size+1) == size) {
            int victim = random.nextInt(queue.length);
            int min = ..., max = ...;
            synchronized (queue[min]) {
                synchronized (queue[max]) {
                    balance(queue[min], queue[max]);
                }
            }
        }
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Work-Balancing Thread



```
public void run() {
    int me = ThreadID.getID();
    while (true) {
        Runnable task = queue[me].deq();
        if (task != null) task.run();
        int size = queue[me].size();
        if (random.nextInt(size+1) == size) {
            int victim = random.nextInt(queue.length);
            int min = ..., max = ...;
            synchronized (queue[min]) {
                synchronized (queue[max]) {
                    balance(queue[min], queue[max]);
                }
            }
        }
    }
}
```

Keep running tasks

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Work-Balancing Thread



```
public void run() {
    int me = ThreadID.getID();
    while (true) {
        Runnable task = queue[me].deq();
        if (task != null) task.run();
        int size = queue[me].size();
        if (random.nextInt(size+1) == size) {
            int victim = random.nextInt(queue.length);
            int min = ..., max = ...;
            synchronized (queue[min]) {
                synchronized (queue[max]) {
                    balance(queue[min], queue[max]);
                }
            }
        }
    }
}
```

With probability $1/|queue|$

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Work-Balancing Thread



```
public void run() {
    int me = ThreadID.getID();
    while (true) {
        Runnable task = queue[me].deq();
        if (task != null) task.run();
        int size = queue[me].size();
        if (random.nextInt(size+1) == size) {
            int victim = random.nextInt(queue.length);
            int min = ..., max = ...;
            synchronized (queue[min]) {
                synchronized (queue[max]) {
                    balance(queue[min], queue[max]);
                }
            }
        }
    }
}
```

Choose random victim

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Work-Balancing Thread



```
public void run() {
    int me = ThreadID.getID();
    while (true) {
        Runnable task = queue[me].deq();
        if (task != null) task.run();
        int size = queue[me].size();
        if (random.nextInt(size+1) == size) {
            int victim = random.nextInt(queue.length);
            int min = ..., max = ...;
            synchronized (queue[min]) {
                synchronized (queue[max]) {
                    balance(queue[min], queue[max]);
                }
            }
        }
    }
}
```

Lock queues in canonical order

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Work-Balancing Thread



```
public void run() {
    int me = ThreadID.getID();
    while (true) {
        Runnable task = queue[me].deq();
        if (task != null) task.run();
        int size = queue[me].size();
        if (random.nextInt(size+1) == size) {
            int victim = random.nextInt(queue.length);
            int min = ..., max = ...;
            synchronized (queue[min]) {
                synchronized (queue[max]) {
                    balance(queue[min], queue[max]);
                }
            }
        }
    }
}
```

Rebalance queues

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Starting Threads



- Executor framework

```
public interface Executor {
    void execute(Runnable command);
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Executor Example



```
class MyTask implements Runnable {
    public void run() { ... }
    public static void main(...) {
        Runnable task1 = new MyTask();
        Executor exec = Executors.newCachedThreadPool();
        exec.execute(task1);
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Provided Thread Pool Executors



- `newFixedThreadPool`
 - Bounded size
 - Replace if thread dies
- `newCachedThreadPool`
 - Demand driven variable size
- `newSingleThreadExecutor`
 - Just one thread
 - Replace if thread dies
- `newScheduledThreadPool`
 - Delayed and periodic task execution
 - Good replacement for class `Timer`

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Executor Shut-down



- Executor is a service provider
- Executor abstracts thread management
- To maintain the abstraction, the service should generally provide methods for shutting down the service
- JVM cannot shut down until threads do

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

ExecutorService



```
public interface ExecutorService extends Executor {
    void shutdown();
    List<Runnable> shutdownNow();
    boolean isShutdown();
    boolean isTerminated();
    boolean awaitTermination(long timeout, TimeUnit unit)
        throws InterruptedException;
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

ExecutorService Notes



- Implies three states:
 - Running
 - Shutting down
 - Terminated
- `shutdown()` runs tasks in queue
- `shutdownNow()` returns unstarted tasks
- `awaitTermination()` blocks until the service is in the terminated state

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

ExecutorService implementation



- `ExecutorService` is an interface
- How do you implement shut down and cancellation?

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Cancellation in Java



- *Cooperative*, not mandated
- Traditional method:
interruption

```
public class Thread {  
    public void interrupt();  
    public void isInterrupted();  
    public static boolean interrupted();  
    ...  
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Interruption



- Just a *request*!
- Sets a flag that must be explicitly cleared!
- Some methods clear the flag & throw *InterruptedException*
- Others *ignore* it, leaving other code to handle it

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

IMPORTANT!



- Your code should follow protocol when interrupted
- If *interrupted()*,
 - Throw exception
 - Reset the flag for other code
- If *InterruptedException*
 - Pass it along
 - Reset the flag for other code
- Don't swallow, unless your code handles the interruption policy

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Restoring Interrupted Status



```
catch (InterruptedException e) {  
    Thread.currentThread().interrupt();  
}  
  
// checks (AND CLEARS!) current thread  
if (Thread.interrupted()) {  
    Thread.currentThread().interrupt();  
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Handling Interruption



- Long computations “break” to check interrupted status
- If interrupted, stop computation, throw *InterruptedException* or restore the interrupted status

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Interrupting Non-Blocking Operations



- Socket read/writes do not support interruption!
 - If you detect interruption, close the socket causing read/write to throw an exception
- Locks via synchronized can not be interrupted

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Future



- Interface representing the *lifecycle* of a task

```
public interface Future<V> {  
    boolean cancel(boolean mayInterruptIfRunning);  
    boolean isCancelled();  
    boolean isDone();  
    V get() throws InterruptedException,  
        ExecutionException, CancellationException;  
    V get(long timeout, TimeUnit unit) throws InterruptedException,  
        ExecutionException, CancellationException,  
        TimeoutException;  
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

CompletionService



- Combines Executor with BlockingQueue
- *ExecutorCompletionService*

```
public interface CompletionService<V> {  
    Future<V> poll();  
    Future<V> poll(long timeout, TimeUnit unit);  
    Future<V> submit(Callable<V> task);  
    Future<V> submit(Runnable task);  
    Future<V> take();  
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Futures Again



- *ExecutorService* interface also provides

```
public interface ExecutorService {  
    ...  
    Future<T> submit(Callable<T> task);  
    Future<?> submit(Runnable task);  
    Future<T> submit(Runnable task, T result);  
}
```

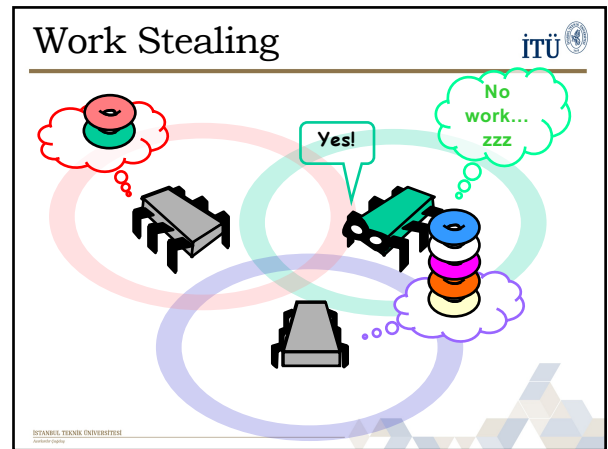
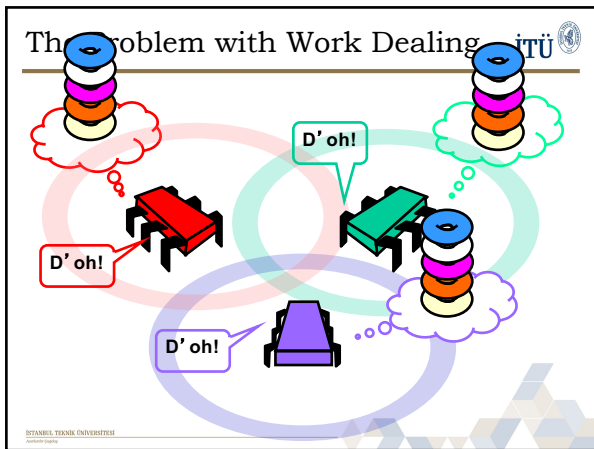
İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Executor Example



```
class WebServer {  
    Executor pool = Executors.newFixedThreadPool(7);  
    public static void main(String[] args) {  
        ServerSocket soc = new ServerSocket(80);  
        while (true) {  
            Socket conn = soc.accept();  
            Runnable r = new Runnable() {  
                public void run() {  
                    handleRequest(conn);  
                }  
            };  
            pool.execute(r);  
        }  
    }  
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

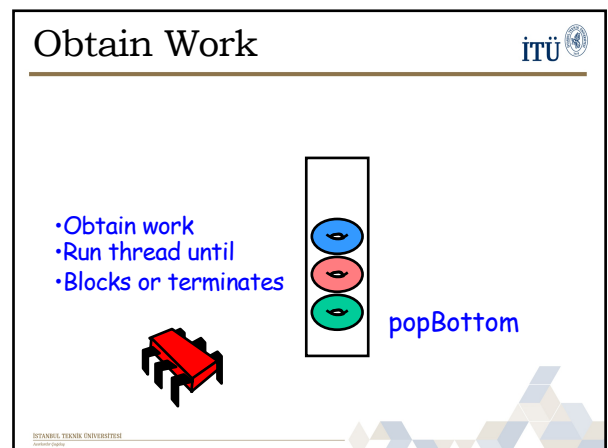
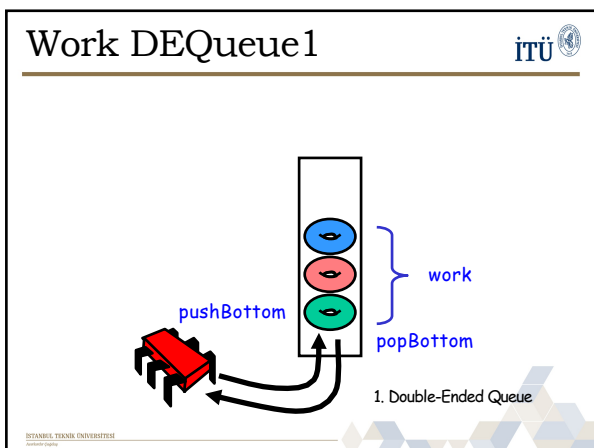
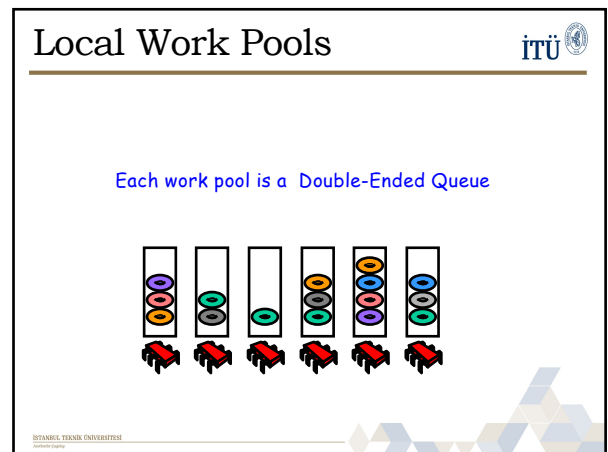


Lock-Free Work Stealing

- Each thread has a pool of ready work
- Remove work without synchronizing
- If you run out of work, steal someone else's
- Choose victim at random

İTÜ


İSTANBUL TEKNİK ÜNİVERSİTESİ



New Work




- Unblock node
- Spawn node



İSTANBUL TEKNİK ÜNİVERSİTESİ

Whatcha Gonna do When the Well Runs Dry?




İSTANBUL TEKNİK ÜNİVERSİTESİ

Steal Work from Others




Pick random guy's DEQueue



İSTANBUL TEKNİK ÜNİVERSİTESİ

Steal this Thread!



İSTANBUL TEKNİK ÜNİVERSİTESİ

Thread DEQueue



- Methods
 - pushBottom
 - popBottom
 - popTop
- } Never happen concurrently

İSTANBUL TEKNİK ÜNİVERSİTESİ

Thread DEQueue



- Methods
 - pushBottom
 - popBottom
 - popTop
- } These most common - make them fast

İSTANBUL TEKNİK ÜNİVERSİTESİ


Compromise



- Method popTop may fail if
 - Concurrent popTop succeeds, or a
 - Concurrent popBottom takes last work


İSTANBUL TEKNİK ÜNİVERSİTESİ

Dreaded ABA Problem




İSTANBUL TEKNİK ÜNİVERSİTESİ

Dreaded ABA Problem




İSTANBUL TEKNİK ÜNİVERSİTESİ

Dreaded ABA Problem




İSTANBUL TEKNİK ÜNİVERSİTESİ

Dreaded ABA Problem

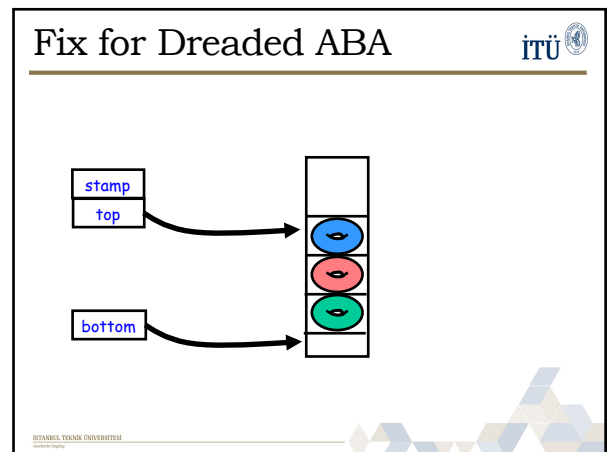
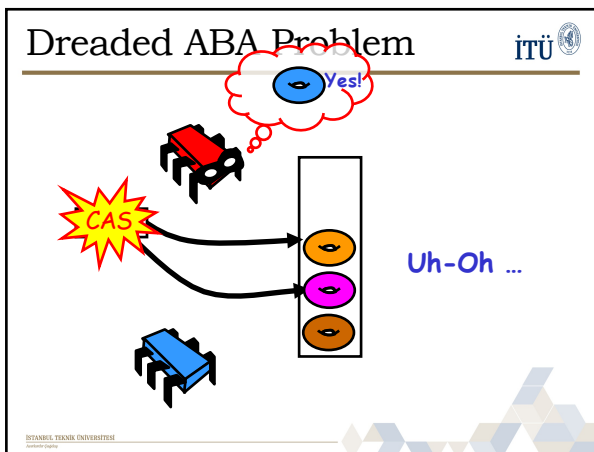
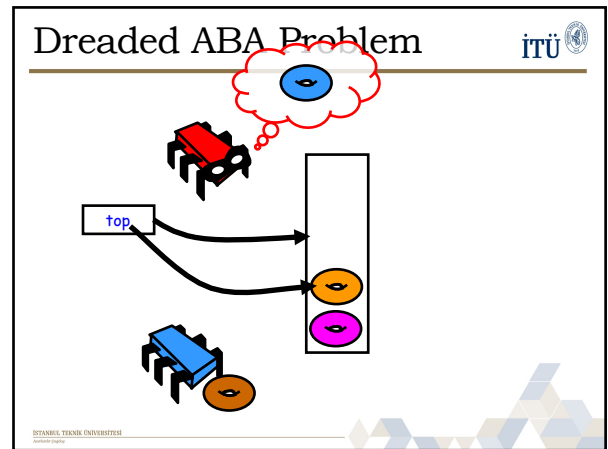
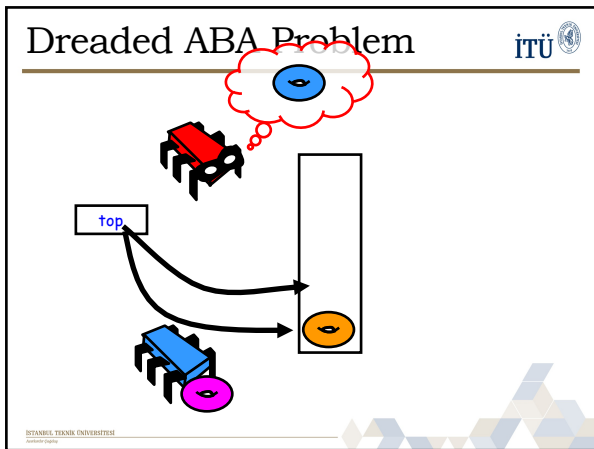


İSTANBUL TEKNİK ÜNİVERSİTESİ

Dreaded ABA Problem



İSTANBUL TEKNİK ÜNİVERSİTESİ



Bounded DEQueue

```

public class BDEQueue {
    AtomicStampedReference<Integer> top;
    volatile int bottom;
    Runnable[] tasks;
    ...
}

```

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Bounded DQueue

```

public class BDEQueue {
    AtomicStampedReference<Integer> top;
    volatile int bottom;
    Runnable[] tasks;
    ...
}

```

Index & Stamp
(synchronized)

ISTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Bounded DEQueue



```
public class BDEQueue {
    AtomicStampedReference<Integer> top;
    volatile int bottom;
    Runnable[] deq;
    ...
}
```

Index of bottom thread
(no need to synchronize
The effect of a write
must be seen - so we
need a memory barrier)

İSTANBUL TEKNİK ÜNİVERSİTESİ

Bounded DEQueue



```
public class BDEQueue {
    AtomicStampedReference<Integer> top;
    volatile int bottom;
    Runnable[] tasks;
    ...
}
```

Array holding tasks

İSTANBUL TEKNİK ÜNİVERSİTESİ

pushBottom()



```
public class BDEQueue {
    ...
    void pushBottom(Runnable r){
        tasks[bottom] = r;
        bottom++;
    }
    ...
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

pushBottom()



```
public class BDEQueue {
    ...
    void pushBottom(Runnable r){
        tasks[bottom] = r;
        bottom++;
    }
    ...
}
```

Bottom is the index to store
the new task in the array

İSTANBUL TEKNİK ÜNİVERSİTESİ

pushBottom()



```
public class BDEQueue {
    ...
    void pushBottom(Runnable r){
        tasks[bottom] = r;
        bottom++;
    }
    ...
}
```

Adjust the bottom index

İSTANBUL TEKNİK ÜNİVERSİTESİ

Steal Work



```
public Runnable popTop() {
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = oldTop + 1;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom <= oldTop)
        return null;
    Runnable r = tasks[oldTop];
    if (top.CAS(oldTop, newTop, oldStamp, newStamp))
        return r;
    return null;
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

Steal Work



```
public Runnable popTopO {
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = oldTop + 1;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom <= oldTop)
        return null;
    Runnable r = tasks[oldTop];
    if (top.CAS(oldTop, newTop, oldStamp, newStamp))
        return r;
    return null;
}
```

Read top (value & stamp)

İSTANBUL TEKNİK ÜNİVERSİTESİ

Steal Work



```
public Runnable popTopO {
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = oldTop + 1;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom <= oldTop)
        return null;
    Runnable r = tasks[oldTop];
    if (top.CAS(oldTop, newTop, oldStamp, newStamp))
        return r;
    return null;
}
```

Compute new value & stamp

İSTANBUL TEKNİK ÜNİVERSİTESİ

Steal Work



```
public Runnable popTopO {
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = oldTop + 1;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom <= oldTop)
        return null;
    Runnable r = tasks[oldTop];
    if (top.CAS(oldTop, newTop, oldStamp, newStamp))
        return r;
    return null;
}
```

Quit if queue is empty

İSTANBUL TEKNİK ÜNİVERSİTESİ

Steal Work



```
public Runnable popTopO {
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = oldTop + 1;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom <= oldTop)
        return null;
    Runnable r = tasks[oldTop];
    if (top.CAS(oldTop, newTop, oldStamp, newStamp))
        return r;
    return null;
}
```

Try to steal the task

İSTANBUL TEKNİK ÜNİVERSİTESİ

Steal Work



```
public Runnable popTopO {
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = oldTop + 1;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom <= oldTop)
        return null;
    Runnable r = tasks[oldTop];
    if (top.CAS(oldTop, newTop, oldStamp, newStamp))
        return r;
    return null;
}
```

Give up if conflict occurs

İSTANBUL TEKNİK ÜNİVERSİTESİ

Take Work



```
Runnable popBottomO {
    if (bottom == 0) return null;
    bottom--;
    Runnable r = tasks[bottom];
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = 0;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom > oldTop) return r;
    if (bottom == oldTop){
        bottom = 0;
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))
            return r;
    }
    top.set(newTop, newStamp); return null;
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

Take Work



```
Runnable popBottom() {
    if (bottom == 0) return null;
    bottom--;
    Runnable r = tasks[bottom];
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = 0;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom > oldTop) return r;
    if (bottom == oldTop){
        bottom = 0;
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))
            return r;
    }
    top.set(newTop, newStamp); return null;
}
```

Make sure queue is non-empty

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Take Work



```
Runnable popBottom() {
    if (bottom == 0) return null;
    bottom--;
    Runnable r = tasks[bottom];
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = 0;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom > oldTop) return r;
    if (bottom == oldTop){
        bottom = 0;
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))
            return r;
    }
    top.set(newTop, newStamp); return null;
}
```

Prepare to grab bottom task

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Take Work



```
Runnable popBottom() {
    if (bottom == 0) return null;
    bottom--;
    Runnable r = tasks[bottom];
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = 0;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom > oldTop) return r;
    if (bottom == oldTop){
        bottom = 0;
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))
            return r;
    }
    top.set(newTop, newStamp); return null;
}
```

Read top, & prepare new values

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Take Work



```
Runnable popBottom() {
    if (bottom == 0) return null;
    bottom--;
    Runnable r = tasks[bottom];
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = 0;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom > oldTop) return r;
    if (bottom == oldTop){
        bottom = 0;
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))
            return r;
    }
    top.set(newTop, newStamp); return null;
}
```

If top & bottom 1 or more apart, no conflict

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Take Work



```
Runnable popBottom() {
    if (bottom == 0) return null;
    bottom--;
    Runnable r = tasks[bottom];
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = 0;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom > oldTop) return r;
    if (bottom == oldTop){
        bottom = 0;
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))
            return r;
    }
    top.set(newTop, newStamp); return null;
}
```

At most one item left

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Take Work



```
Runnable popBottom() {
    if (bottom == 0) return null;
    bottom--;
    Runnable r = tasks[bottom];
    int[] stamp = new int[1];
    int oldTop = top.get(stamp), newTop = 0;
    int oldStamp = stamp[0], newStamp = oldStamp + 1;
    if (bottom > oldTop) return r;
    if (bottom == oldTop){
        bottom = 0;
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))
            return r;
    }
    top.set(newTop, newStamp); return null;
}
```

Try to steal last item.
In any case reset Bottom
because the DEQueue will be empty
even if unsuccessful (why?)

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Take Work



```

Runnable popBottom() {
    if (bottom == 0) return null;
    bottom--;
    Runnable r = tasks[bottom];
    int[] stamp = new int[1];
    int oldtop = top.get(stamp), newtop = 0;
    int oldstamp = stamp[0], newstamp = oldstamp + 1;
    if (bottom > oldtop) return r;
    if (bottom == oldtop) {
        bottom = 0;
        if (top.CAS(oldtop, newtop, oldstamp, newstamp))
            return r;
    }
    top.set(newtop, newstamp); return null;
}
    
```

I win CAS

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Take Work



```

Runnable popBottom() {
    if (bottom == 0) return null;
    bottom--;
    Runnable r = tasks[bottom];
    int[] stamp = new int[1];
    int oldtop = top.get(stamp), newtop = 0;
    int oldstamp = stamp[0], newstamp = oldstamp + 1;
    if (bottom > oldtop) return r;
    if (bottom == oldtop) {
        bottom = 0;
        if (top.CAS(oldtop, newtop, oldstamp, newstamp))
            return r;
    }
    top.set(newtop, newstamp); return null;
}
    
```

**I lose CAS
Thief must have won...**

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Take Work



```

Runnable popBottom() {
    if (bottom == 0) return null;
    bottom--;
    Runnable r = tasks[bottom];
    int[] stamp = new int[1];
    int oldtop = top.get(stamp), newtop = 0;
    int oldstamp = stamp[0], newstamp = oldstamp + 1;
    if (bottom > oldtop) return r;
    if (bottom == oldtop) {
        bottom = 0;
        if (top.CAS(oldtop, newtop, oldstamp, newstamp))
            return r;
    }
    top.set(newtop, newstamp); return null;
}
    
```

**failed to get last item
Must still reset top**

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Work Stealing & Balancing



- Clean separation between app & scheduling layer
- Works well when number of processors fluctuates.
- Works on “black-box” operating systems

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Fork/Join Framework



- The fork/join framework was presented in Java 7.
- It provides a **divide and conquer approach**.

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Fork/Join Framework



- In practice, this means that **the framework first “forks”**, recursively breaking the task into smaller independent subtasks until they are simple enough to be executed asynchronously.
- After that, **the “join” part begins**, in which results of all subtasks are recursively joined into a single result, or in the case of a task which returns void, the program simply waits until every subtask is executed.

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Fork/Join Framework



- To provide effective parallel execution, the fork/join framework uses a pool of threads called the *ForkJoinPool*, which manages worker threads of type *ForkJoinWorkerThread*.

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

ForkJoinPool



- The *ForkJoinPool* is the heart of the framework. It is an implementation of the *ExecutorService* that manages worker threads and provides us with tools to get information about the thread pool state and performance.
- This architecture is vital for balancing the thread's workload with the help of the **work-stealing algorithm**.

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

ForkJoinPool



```
ForkJoinPool commonPool = ForkJoinPool.commonPool();
```

```
public static ForkJoinPool forkJoinPool  
    = new ForkJoinPool(2);
```

```
ForkJoinPool forkJoinPool = PoolUtil.forkJoinPool;
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

ForkJoinTask<V>



- ForkJoinTask* is the base type for tasks executed inside *ForkJoinPool*. In practice, one of its two subclasses should be extended:
 - the *RecursiveAction* for void tasks and
 - the *RecursiveTask<V>* for tasks that return a value.
- They both have an abstract method *compute()* in which the task's logic is defined.

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

RecursiveAction



```
public class CustomRecursiveAction extends RecursiveAction {  
    private String workload = "";  
    private static final int THRESHOLD = 4;  
  
    private static Logger logger =  
        Logger.getAnonymousLogger();  
  
    public CustomRecursiveAction(String workload) {  
        this.workload = workload;  
    }  
  
    @Override  
    protected void compute() {  
        if (workload.length() > THRESHOLD) {  
            ForkJoinTask.invokeAll(createSubtasks());  
        } else {  
            processing(workload);  
        }  
    }  
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

RecursiveAction



```
private List<CustomRecursiveAction> createSubtasks() {  
    List<CustomRecursiveAction> subtasks = new ArrayList<>();  
  
    String partOne = workload.substring(0, workload.length() / 2);  
    String partTwo = workload.substring(workload.length() / 2, workload.length());  
  
    subtasks.add(new CustomRecursiveAction(partOne));  
    subtasks.add(new CustomRecursiveAction(partTwo));  
  
    return subtasks;  
}  
  
private void processing(String work) {  
    String result = work.toUpperCase();  
    logger.info("This result - (" + result + ") - was processed by "  
        + Thread.currentThread().getName());  
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

RecursiveTask<V>



```
public class MaxNumberCalculator extends RecursiveTask {
    public static final int THRESHOLD = 5;

    private int[] numbers;
    private int start;
    private int end;

    public MaxNumberCalculator(int[] numbers) {
        this(numbers, 0, numbers.length);
    }

    public MaxNumberCalculator(int[] numbers, int start, int end) {
        this.start = start;
        this.end = end;
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

RecursiveTask<V>



```
@Override
public Integer compute() {
    int length = end - start;
    int max = 0;
    if (length < THRESHOLD) {
        for (int x = start; x < end; x++) {
            max = numbers[x];
        }
        return max;
    } else {
        int split = length / 2;
        MaxNumberCalculator left = new MaxNumberCalculator(numbers, start,
            start + split);
        left.fork();
        MaxNumberCalculator right = new MaxNumberCalculator(numbers, start
            + split, end);
        right.fork();
        return Math.max(right.compute(), left.join());
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

ForkJoinPool



- To submit tasks to the thread pool, use the **submit()** or **execute()** method.

```
forkJoinPool.execute(customRecursiveTask); int result =
customRecursiveTask.join();
```

- The **invoke()** and **invokeAll()** method forks the task and waits for the result, and doesn't need any manual joining

```
int result = forkJoinPool.invoke(customRecursiveTask);
```

- Alternatively, you can use separate **fork()** and **join()** methods. The **fork()** method submits a task to a pool, but it doesn't trigger its execution. The **join()** method is used for this purpose.

```
customRecursiveTaskFirst.fork();
result = customRecursiveTaskLast.join();
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

ForkJoinPool



```
public static void main(String[] args) {
    final int[] numbers = new int[SIZE];
    int maxNum = 0;

    // Start sequential calculation
    long st = System.currentTimeMillis();

    for (int i = 0; i < SIZE; i++) {
        numbers[i] = (int) (Math.random() * 10000);
        if (numbers[i] > maxNum) {
            maxNum = numbers[i];
        }
    }


    System.out.println("Calculated maximum number (sequential execution): "
        + maxNum + " -- Total time: "
        + (System.currentTimeMillis() - st));

    // Start parallel calculation
    long pt = System.currentTimeMillis();

    ForkJoinPool pool = new ForkJoinPool(4);
    MaxNumberCalculator fbn = new MaxNumberCalculator(numbers);
    System.out.println("Calculated maximum number (parallel execution): "
        + pool.invoke(fbn) + " -- Total time: "
        + (System.currentTimeMillis() - pt));
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

ForkJoinPool



İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Example: WordCount



```
class Document {
    private final List<String> lines;

    Document(List<String> lines) {
        this.lines = lines;
    }

    List<String> getLines() {
        return this.lines;
    }

    static Document fromFile(File file) throws IOException {
        List<String> lines = new LinkedList<>();
        try(BufferedReader reader = new BufferedReader(new FileReader(file))) {
            String line = reader.readLine();
            while (line != null) {
                lines.add(line);
                line = reader.readLine();
            }
        }
        return new Document(lines);
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Example: WordCount



```
class Folder {
    private final List<Folder> subFolders;
    private final List<Document> documents;

    Folder(List<Folder> subFolders, List<Document> documents) {
        this.subFolders = subFolders;
        this.documents = documents;
    }

    List<Folder> getSubFolders() {
        return this.subFolders;
    }

    List<Document> getDocuments() {
        return this.documents;
    }

    static Folder fromDirectory(File dir) throws IOException {
        List<Document> documents = new LinkedList<>();
        List<Folder> subFolders = new LinkedList<>();
        for (File entry : dir.listFiles()) {
            if (entry.isDirectory()) {
                subFolders.add(Folder.fromDirectory(entry));
            } else {
                documents.add(Document.fromFile(entry));
            }
        }
        return new Folder(subFolders, documents);
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Example: WordCount



Sequential Recursive Version

```
Long countOccurrencesOnSingleThread(Folder folder, String searchedWord) {
    long count = 0;
    for (Folder subFolder : folder.getSubFolders()) {
        count = count + countOccurrencesOnSingleThread(subFolder, searchedWord);
    }
    for (Document document : folder.getDocuments()) {
        count = count + occurrencesCount(document, searchedWord);
    }
    return count;
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Example: WordCount



```
public class WordCounter {
    String[] wordsIn(String line) {
        return line.trim().split("\\s|\\p{Punct}+");
    }

    Long occurrencesCount(Document document, String searchedWord) {
        long count = 0;
        for (String line : document.getLines()) {
            for (String word : wordsIn(line)) {
                if (searchedWord.equals(word)) {
                    count = count + 1;
                }
            }
        }
        return count;
    }

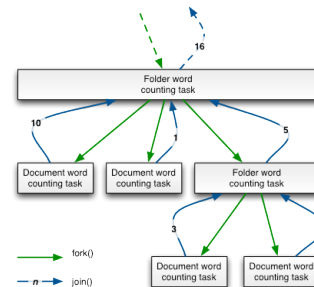
    class DocumentSearchTask extends RecursiveTask<Long> {
        private final Document document;
        private final String searchedWord;

        DocumentSearchTask(Document document, String searchedWord) {
            super();
            this.document = document;
            this.searchedWord = searchedWord;
        }

        @Override
        protected Long compute() {
            return occurrencesCount(document, searchedWord);
        }
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Example: WordCount



İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Example: WordCount



```
class FolderSearchTask extends RecursiveTask<Long> {
    private final Folder folder;
    private final String searchedWord;

    FolderSearchTask(Folder folder, String searchedWord) {
        super();
        this.folder = folder;
        this.searchedWord = searchedWord;
    }

    @Override
    protected Long compute() {
        long count = 0L;
        List<RecursiveTask<Long>> forks = new LinkedList<>();
        for (Folder subFolder : folder.getSubFolders()) {
            FolderSearchTask task = new FolderSearchTask(subFolder, searchedWord);
            forks.add(task);
            task.fork();
        }
        for (Document document : folder.getDocuments()) {
            DocumentSearchTask task = new DocumentSearchTask(document, searchedWord);
            forks.add(task);
            task.fork();
        }
        for (RecursiveTask<Long> task : forks) {
            count = count + task.join();
        }
        return count;
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Example: WordCount



```
private final ForkJoinPool forkJoinPool = new ForkJoinPool();

Long countOccurrencesInParallel(Folder folder, String searchedWord) {
    return forkJoinPool.invoke(new FolderSearchTask(folder, searchedWord));
}


public static void main(String[] args) throws IOException {
    WordCounter wordCounter = new WordCounter();
    Folder folder = Folder.fromDirectory(new File(args[0]));
    System.out.println(wordCounter.countOccurrencesOnSingleThread(folder, args[1]));
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Example: WordCount



Number of Cores	Single-Thread Execution Time (ms)	Fork/Join Execution Time (ms)	Speedup
2	18788	11026	1.704879376
4	19473	8329	2.337975747
8	18911	4208	4.494058935
12	19410	2876	6.748956885



ISTANBUL TEKNİK ÜNİVERSİTESİ
©2015-2016

This work is licensed under a [Creative Commons Attribution-ShareAlike 3.0 License](https://creativecommons.org/licenses/by-sa/3.0/).

- **You are free:**
 - **to Share** — to copy, distribute and transmit the work
 - **to Remix** — to adapt the work
- **Under the following conditions:**
 - **Attribution.** You must attribute the work to "The Art of Multiprocessor Programming" (but not in any way that suggests that the author's endorse you or your use of the work).
 - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.

ISTANBUL TEKNİK ÜNİVERSİTESİ
©2015-2016