

Concurrent Queues and Stacks

Companion slides for
The Art of Multiprocessor Programming
by Maurice Herlihy & Nir Shavit

1

The Five-Fold Path

- Coarse-grained locking
- Fine-grained locking
- Optimistic synchronization
- Lazy synchronization
- Lock-free synchronization

İSTANBUL TEKNİK ÜNİVERSİTESİ

2

Another Fundamental Problem

- We told you about
 - Sets implemented using linked lists
- Next: queues
- Next: stacks

İSTANBUL TEKNİK ÜNİVERSİTESİ

3

Queues & Stacks

- Both: pool of items
- Queue
 - enq() & deq()
 - First-in-first-out (FIFO) order
- Stack
 - push() & pop()
 - Last-in-first-out (LIFO) order

İSTANBUL TEKNİK ÜNİVERSİTESİ

4

Bounded vs Unbounded

- Bounded
 - Fixed capacity
 - Good when resources an issue
- Unbounded
 - Holds any number of objects

İSTANBUL TEKNİK ÜNİVERSİTESİ

5

Blocking vs Non-Blocking

- Problem cases:
 - Removing from empty pool
 - Adding to full (bounded) pool
- Blocking
 - Caller waits until state changes
- Non-Blocking
 - Method throws exception

İSTANBUL TEKNİK ÜNİVERSİTESİ

6

This Lecture




- Bounded, Blocking, Lock-based Queue
- Unbounded, Non-Blocking, Lock-free Queue
- Lost Wake-up problem
- Unbounded Non-Blocking Lock-free Stack
- Elimination-Backoff Stack

İSTANBUL TEKNİK ÜNİVERSİTESİ

7


Queue: Concurrency



İSTANBUL TEKNİK ÜNİVERSİTESİ

8


Concurrency



İSTANBUL TEKNİK ÜNİVERSİTESİ

9

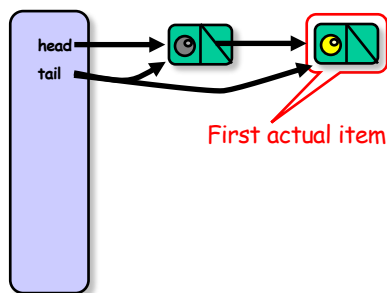
Bounded Queue



İSTANBUL TEKNİK ÜNİVERSİTESİ

10


Bounded Queue



İSTANBUL TEKNİK ÜNİVERSİTESİ

11


Bounded Queue



İSTANBUL TEKNİK ÜNİVERSİTESİ


12

Bounded Queue




13

Not Done Yet




14

Not Done Yet




15

Not Done Yet




16

Enqueuer

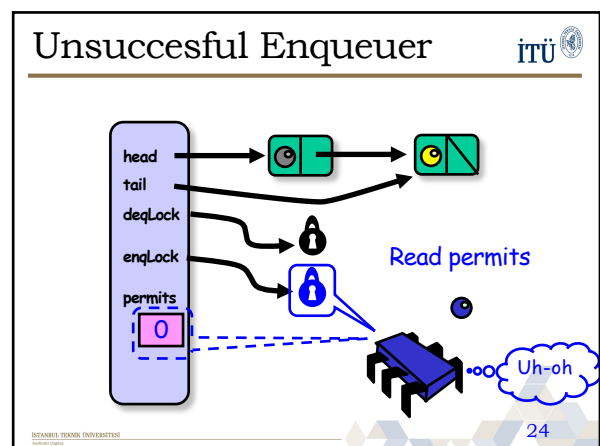
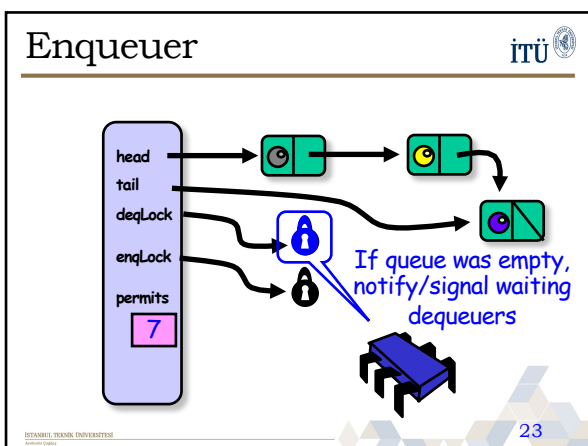
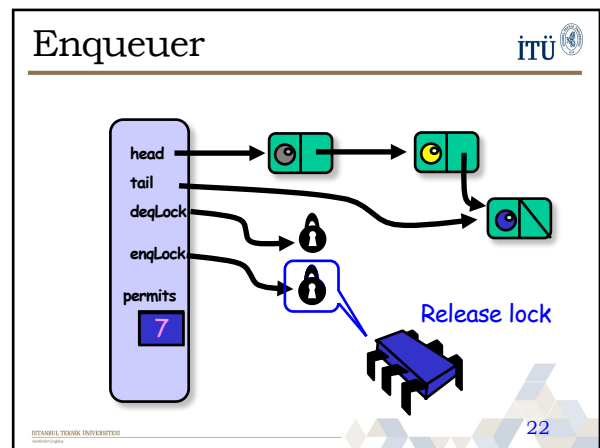
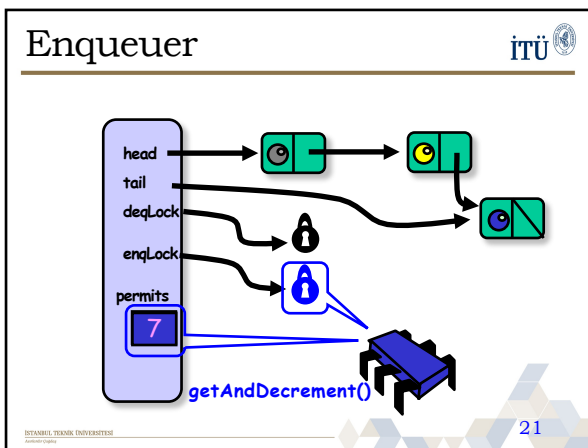
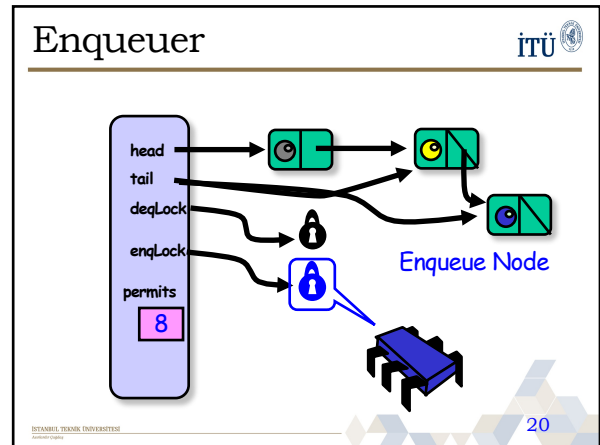
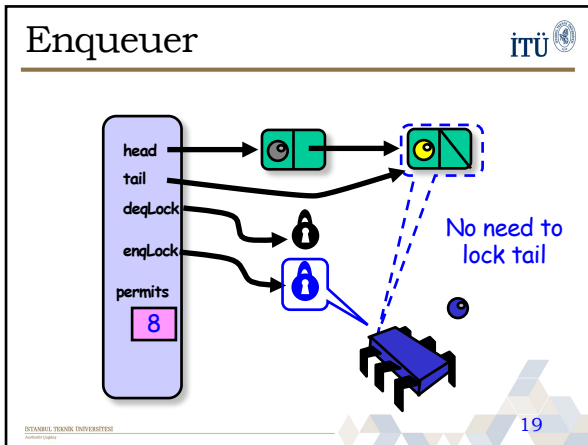


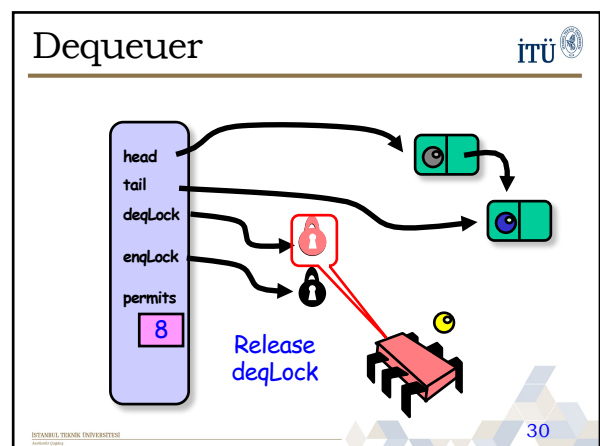
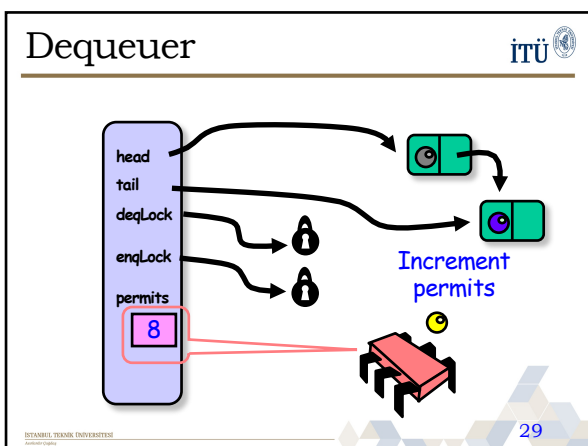
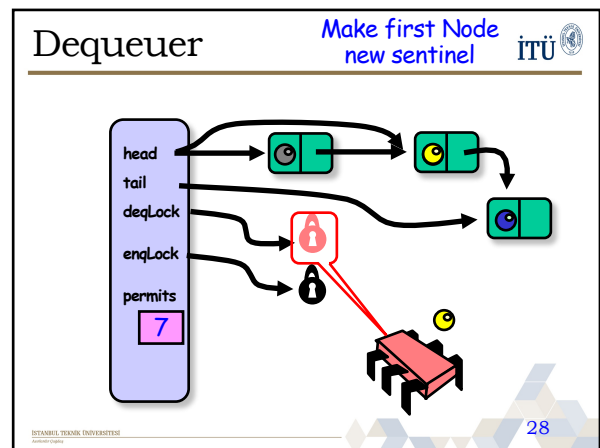
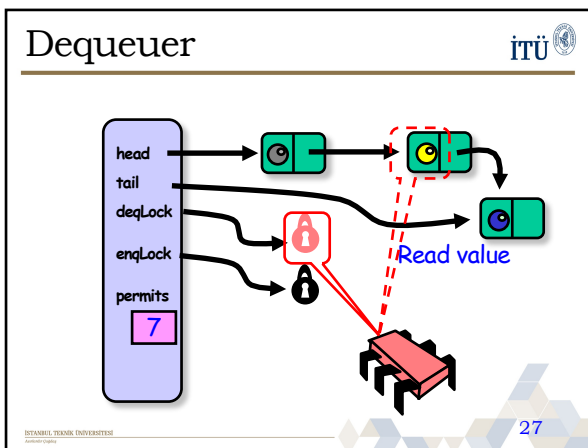
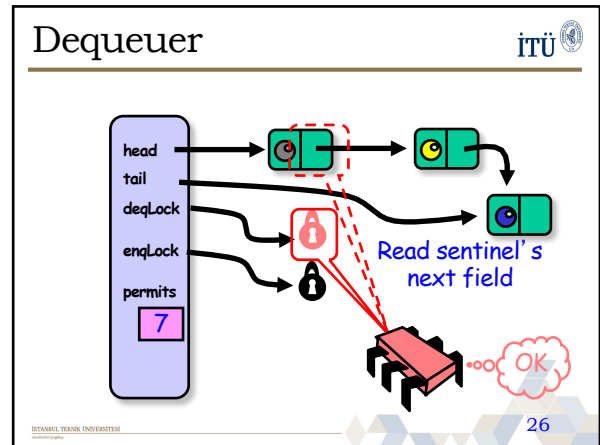
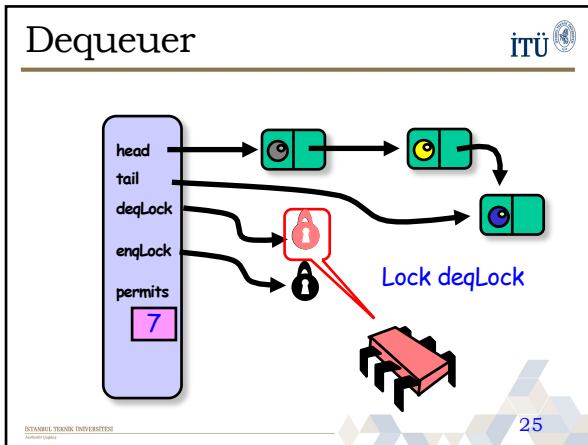
17

Enqueuer




18





Unsuccessful Dequeueer



İSTANBUL TEKNİK ÜNİVERSİTESİ

31

Bounded Queue



```
public class BoundedQueue<T> {
    ReentrantLock enqLock, deqLock;
    Condition notEmptyCondition, notFullCondition;
    AtomicInteger permits;
    Node head;
    Node tail;
    int capacity;
    enqLock = new ReentrantLock();
    notFullCondition = enqLock.newCondition();
    deqLock = new ReentrantLock();
    notEmptyCondition = deqLock.newCondition();
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

32

Bounded Queue



```
public class BoundedQueue<T> {
    ReentrantLock enqLock, deqLock;
    Condition notEmptyCondition, notFullCondition;
    AtomicInteger permits;
    Node head;
    Node tail;
    int capacity;
    enqLock = new ReentrantLock();
    notFullCondition = enqLock.newCondition();
    deqLock = new ReentrantLock();
    notEmptyCondition = deqLock.newCondition();
}
```

Enq & deq locks

İSTANBUL TEKNİK ÜNİVERSİTESİ

33

Digression: Monitor Locks



- Java Synchronized objects and Java ReentrantLocks are monitors
- Allow blocking on a condition rather than spinning
- Threads:
 - acquire and release lock
 - wait on a condition

İSTANBUL TEKNİK ÜNİVERSİTESİ

34

The Java Lock Interface



```
public interface Lock {
    void lock();
    void lockInterruptibly() throws InterruptedException;
    boolean tryLock();
    boolean tryLock(long time, TimeUnit unit);
    Condition newCondition();
    void unlock();
}
```

Acquire lock

İSTANBUL TEKNİK ÜNİVERSİTESİ

35

The Java Lock Interface



```
public interface Lock {
    void lock();
    void lockInterruptibly() throws InterruptedException;
    boolean tryLock();
    boolean tryLock(long time, TimeUnit unit);
    Condition newCondition();
    void unlock();
}
```

Release lock

İSTANBUL TEKNİK ÜNİVERSİTESİ

36

The Java Lock Interface



```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit)  
        Condition newCondition();  
    void unlock();  
}
```

Try for lock, but not too hard

İSTANBUL TEKNİK ÜNİVERSİTESİ

37

The Java Lock Interface



```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit)  
        Condition newCondition();  
    void unlock();  
}
```

Create condition to wait on

İSTANBUL TEKNİK ÜNİVERSİTESİ

38

The Java Lock Interface



```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit)  
        Condition newCondition();  
    void unlock();  
}
```

Guess what this method does?

İSTANBUL TEKNİK ÜNİVERSİTESİ

39

Lock Conditions



```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

40

Lock Conditions



```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```

Release lock and
wait on condition

İSTANBUL TEKNİK ÜNİVERSİTESİ

41

Lock Conditions



```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```

Wake up one waiting thread

İSTANBUL TEKNİK ÜNİVERSİTESİ

42

Lock Conditions



```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```

Wake up all waiting threads

43

İSTANBUL TEKNİK ÜNİVERSİTESİ

Await



- Releases lock associated with q
- Sleeps (gives up processor)
- Awakens (resumes running)
- Reacquires lock & returns

```
q.await()
```

44

İSTANBUL TEKNİK ÜNİVERSİTESİ

Signal



- Awakens one waiting thread
 - Which will reacquire lock

```
q.signal();
```

45

İSTANBUL TEKNİK ÜNİVERSİTESİ

Signal All




- Awakens all waiting threads
 - Which will each reacquire lock

```
q.signalAll();
```

46

İSTANBUL TEKNİK ÜNİVERSİTESİ


A Monitor Lock



47

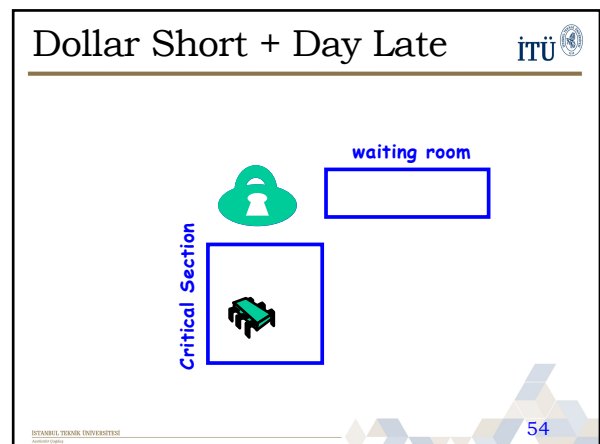
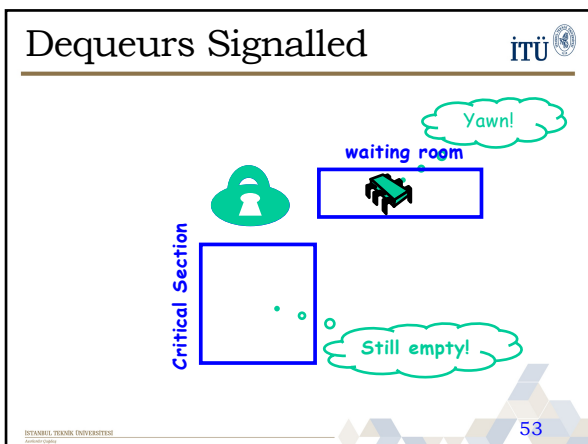

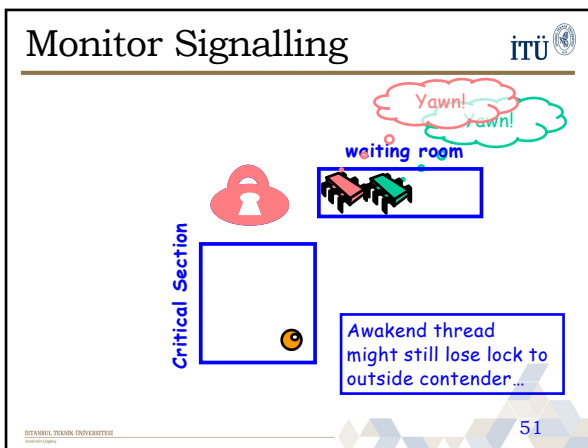
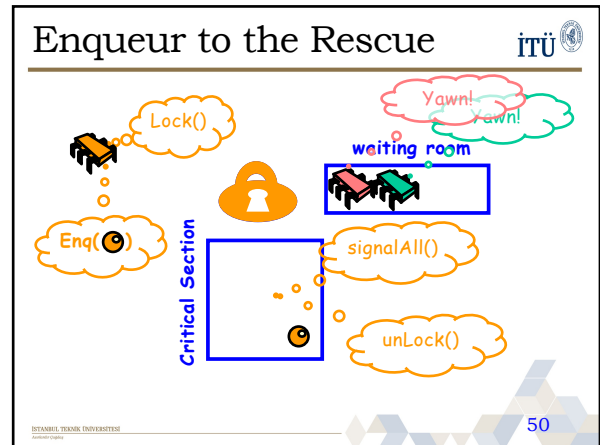
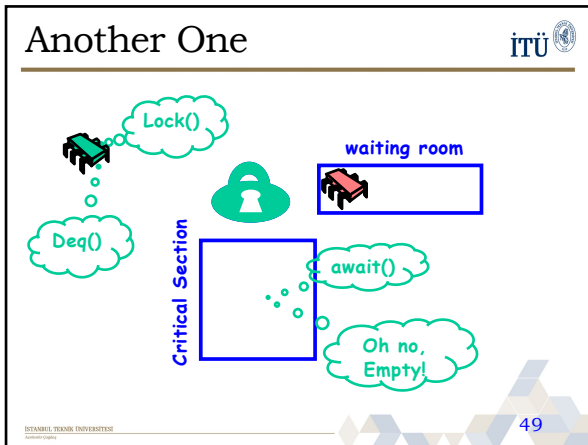
İSTANBUL TEKNİK ÜNİVERSİTESİ

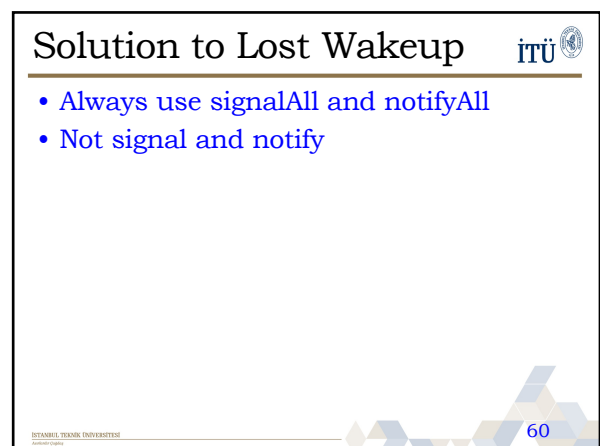
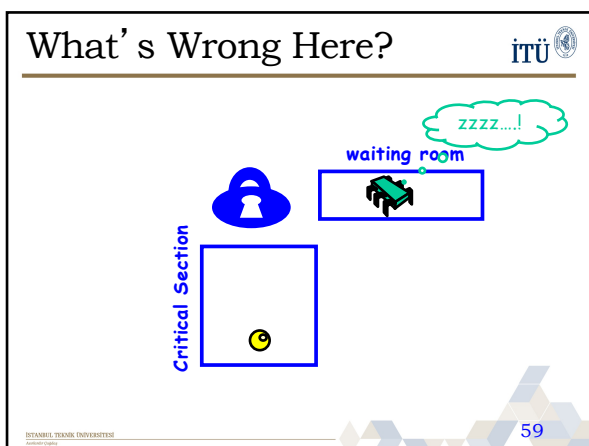
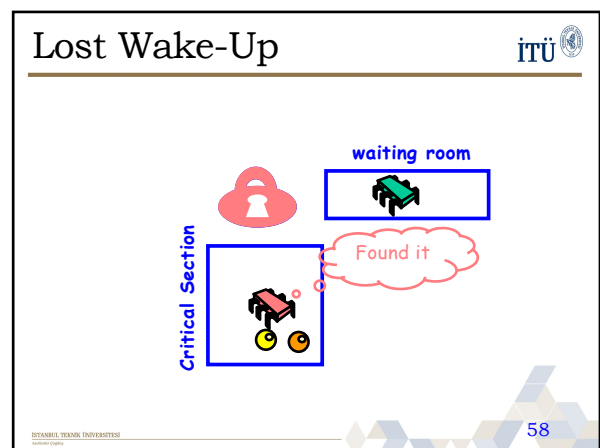
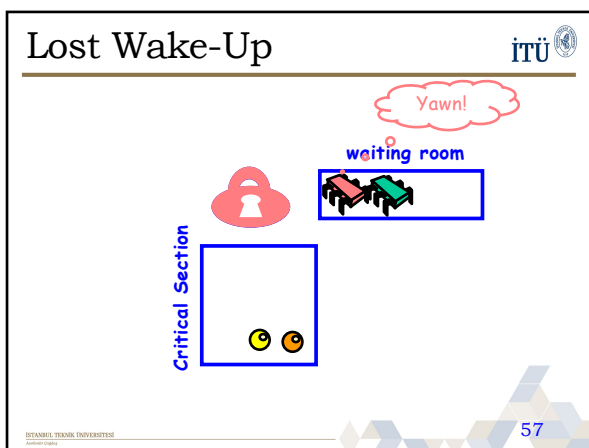
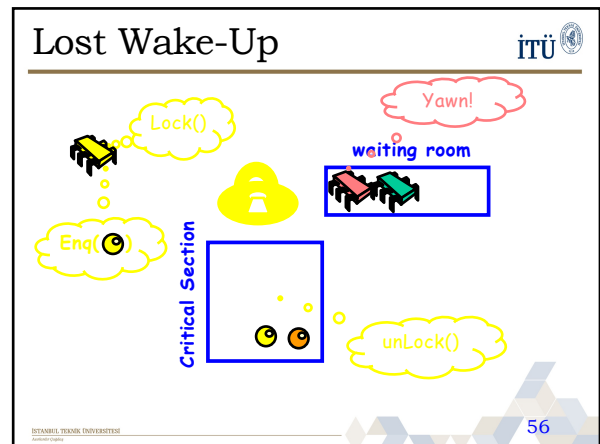
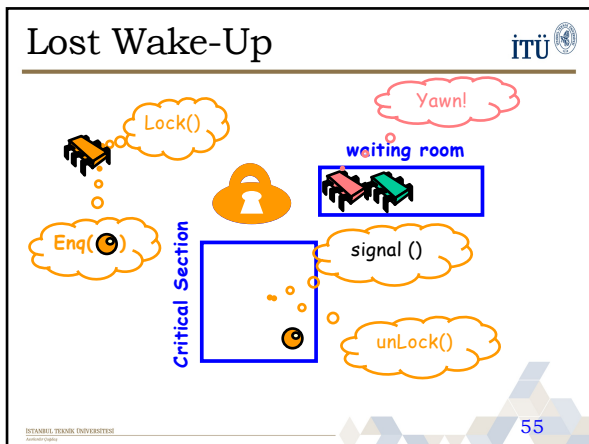
Unsuccessful Deq



48

İSTANBUL TEKNİK ÜNİVERSİTESİ





Java Synchronized Methods

```
public class Queue<T> {
    int head = 0, tail = 0;
    T[QSIZE] items;

    public synchronized T deq() {
        while (tail - head == 0)
            this.wait();
        T result = items[head % QSIZE]; head++;
        this.notifyAll();
        return result;
    }
}
```

ISTANBUL TEKNİK ÜNİVERSİTESİ

61

Java Synchronized Methods

```
public class Queue<T> {
    int head = 0, tail = 0;
    T[QSIZE] items;

    public synchronized T deq() {
        while (tail - head == 0)
            this.wait();
        T result = items[head % QSIZE]; head++;
        this.notifyAll();
        return result;
    }
}
```

Each object has an implicit lock with an implicit condition

ISTANBUL TEKNİK ÜNİVERSİTESİ

62

Java Synchronized Methods

```
public class Queue<T> {
    int head = 0, tail = 0;
    T[QSIZE] items;

    public synchronized T deq() {
        while (tail - head == 0)
            this.wait();
        T result = items[head % QSIZE]; head++;
        this.notifyAll();
        return result;
    }
}
```

Lock on entry, unlock on return

ISTANBUL TEKNİK ÜNİVERSİTESİ

63

Java Synchronized Methods

```
public class Queue<T> {
    int head = 0, tail = 0;
    T[QSIZE] items;

    public synchronized T deq() {
        while (tail - head == 0)
            this.wait();
        T result = items[head % QSIZE]; head++;
        this.notifyAll();
        return result;
    }
}
```

Wait on implicit condition

ISTANBUL TEKNİK ÜNİVERSİTESİ

64

Java Synchronized Methods

```
public class Queue<T> {
    int head = 0, tail = 0;
    T[QSIZE] items;

    public synchronized T deq() {
        while (tail - head == 0)
            this.wait();
        T result = items[head % QSIZE]; head++;
        this.notifyAll();
        return result;
    }
}
```

Signal all threads waiting on condition

ISTANBUL TEKNİK ÜNİVERSİTESİ

65

(Pop!) The Bounded Queue

```
public class BoundedQueue<T> {
    ReentrantLock enqLock, deqLock;
    Condition notEmptyCondition, notFullCondition;
    AtomicInteger permits;
    Node head;
    Node tail;
    int capacity;
    enqLock = new ReentrantLock();
    notFullCondition = enqLock.newCondition();
    deqLock = new ReentrantLock();
    notEmptyCondition = deqLock.newCondition();
}
```

ISTANBUL TEKNİK ÜNİVERSİTESİ

66

Bounded Queue Fields



```
public class BoundedQueue<T> {
    ReentrantLock enqLock, deqLock;
    Condition notEmptyCondition, notFullCondition;
    AtomicInteger permits;
    Node head;
    Node tail;
    int capacity;
    enqLock = new ReentrantLock();
    notFullCondition = enqLock.newCondition();
    deqLock = new ReentrantLock();
    notEmptyCondition = deqLock.newCondition();
}
```

Enq & deq locks

ISTANBUL TEKNİK ÜNİVERSİTESİ

67

Bounded Queue Fields



```
public class BoundedQueue<T> {
    ReentrantLock enqLock, deqLock;
    Condition notEmptyCondition, notFullCondition;
    AtomicInteger permits;
    Node head;
    Node tail;
    int capacity;
    enqLock = new ReentrantLock();
    notFullCondition = enqLock.newCondition();
    deqLock = new ReentrantLock();
    notEmptyCondition = deqLock.newCondition();
}
```

Enq lock's associated condition

ISTANBUL TEKNİK ÜNİVERSİTESİ

68

Bounded Queue Fields



```
public class BoundedQueue<T> {
    ReentrantLock enqLock, deqLock;
    Condition notEmptyCondition, notFullCondition;
    AtomicInteger permits;
    Node head;
    Node tail;
    int capacity;
    enqLock = new ReentrantLock();
    notFullCondition = enqLock.newCondition();
    deqLock = new ReentrantLock();
    notEmptyCondition = deqLock.newCondition();
}
```

Num permits: 0 to capacity

ISTANBUL TEKNİK ÜNİVERSİTESİ

69

Bounded Queue Fields



```
public class BoundedQueue<T> {
    ReentrantLock enqLock, deqLock;
    Condition notEmptyCondition, notFullCondition;
    AtomicInteger permits;
    Node head;
    Node tail;
    int capacity;
    enqLock = new ReentrantLock();
    notFullCondition = enqLock.newCondition();
    deqLock = new ReentrantLock();
    notEmptyCondition = deqLock.newCondition();
}
```

Head and Tail

ISTANBUL TEKNİK ÜNİVERSİTESİ

70

Enq Method Part One



```
public void enq(T x) {
    boolean mustWakeDequeueers = false;
    enqLock.lock();
    try {
        while (permits.get() == 0)
            notFullCondition.await();
        Node e = new Node(x);
        tail.next = e;
        tail = e;
        if (permits.getAndDecrement() == capacity)
            mustWakeDequeueers = true;
    } finally {
        enqLock.unlock();
    }
    ...
}
```

ISTANBUL TEKNİK ÜNİVERSİTESİ

71

Enq Method Part One



```
public void enq(T x) {
    boolean mustWakeDequeueers = false;
    enqLock.lock();
    try {
        while (permits.get() == 0)
            notFullCondition.await();
        Node e = new Node(x);
        tail.next = e;
        tail = e;
        if (permits.getAndDecrement() == capacity)
            mustWakeDequeueers = true;
    } finally {
        enqLock.unlock();
    }
    ...
}
```

Lock and unlock enq lock

ISTANBUL TEKNİK ÜNİVERSİTESİ

72

Enq Method Part One



```
public void enq(T x) {
    boolean mustWakeDequeueers = false;
    enqLock.lock();
    try {
        while (permits.get() == 0)
            notFullCondition.await();
        Node e = new Node(x);
        tail.next = e;
        tail = e;
        if (permits.getAndDecrement() == capacity)
            mustWakeDequeueers = true;
    } finally {
        enqLock.unlock();
    }
    ...
}
```

If queue is full, patiently await further instructions ...

73

Enq Method Part One



```
public void enq(T x) {
    boolean mustWakeDequeueers = false;
    enqLock.lock();
    try {
        while (permits.get() == 0)
            notFullCondition.await();
        Node e = new Node(x);
        tail.next = e;
        tail = e;
        if (permits.getAndDecrement() == capacity)
            mustWakeDequeueers = true;
    } finally {
        enqLock.unlock();
    }
    ...
}
```

Add new node

75

Be Afraid



```
public void enq(T x) {
    boolean mustWakeDequeueers = false;
    enqLock.lock();
    try {
        while (permits.get() == 0)
            notFullCondition.await();
        Node e = new Node(x);
        tail.next = e;
        tail = e;
        if (permits.getAndDecrement() == capacity)
            mustWakeDequeueers = true;
    } finally {
        enqLock.unlock();
    }
    ...
}
```

How do we know the permits field won't change?

74

Enq Method Part One



```
public void enq(T x) {
    boolean mustWakeDequeueers = false;
    enqLock.lock();
    try {
        while (permits.get() == 0)
            notFullCondition.await();
        Node e = new Node(x);
        tail.next = e;
        tail = e;
        if (permits.getAndDecrement() == capacity)
            mustWakeDequeueers = true;
    } finally {
        enqLock.unlock();
    }
    ...
}
```

If queue was empty, wake frustrated dequeueers

76

Enq Method Part Deux



```
public void enq(T x) {
    ...
    if (mustWakeDequeueers) {
        deqLock.lock();
        try {
            notEmptyCondition.signalAll();
        } finally {
            deqLock.unlock();
        }
    }
}
```

77

Enq Method Part Deux



```
public void enq(T x) {
    ...
    if (mustWakeDequeueers) {
        deqLock.lock();
        try {
            notEmptyCondition.signalAll();
        } finally {
            deqLock.unlock();
        }
    }
}
```

Are there dequeueers to be signaled?

78

Enq Method Part Deux



```
public void enq(T x) {
    ...
    if (mustWakeDequeuers) {
        deqLock.lock();
        try {
            notEmptyCondition.signalAll();
        } finally {
            deqLock.unlock();
        }
    }
}
```

Lock and unlock
deq lock

İSTANBUL TEKNİK ÜNİVERSİTESİ

79

Enq Method Part Deux



Signal dequeuers that
queue no longer empty

```
deqLock.lock();
try {
    notEmptyCondition.signalAll();
} finally {
    deqLock.unlock();
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

80

The Enq() & Deq() Methods



- Share no locks
 - That's good
- But do share an atomic counter
 - Accessed on every method call
 - That's not so good
- Can we alleviate this bottleneck?

İSTANBUL TEKNİK ÜNİVERSİTESİ

81

Split the Counter



- The enq() method
 - Decrements only
 - Cares only if value is zero
- The deq() method
 - Increments only
 - Cares only if value is capacity

İSTANBUL TEKNİK ÜNİVERSİTESİ

82

Split Counter




- Enqueuer decrements enqSidePermits
- Dequeuer increments deqSidePermits
- When enqueuer runs out
 - Locks deqLock
 - Transfers permits
- Intermittent synchronization
 - Not with each method call
 - Need both locks! (careful ...)

İSTANBUL TEKNİK ÜNİVERSİTESİ

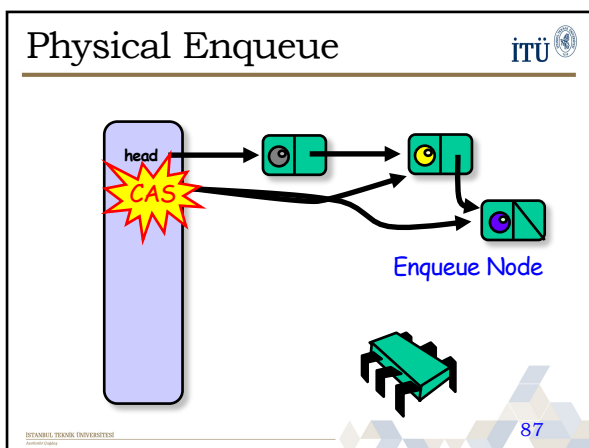
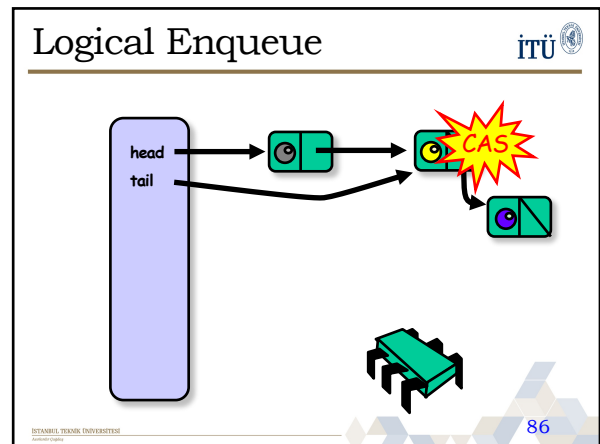
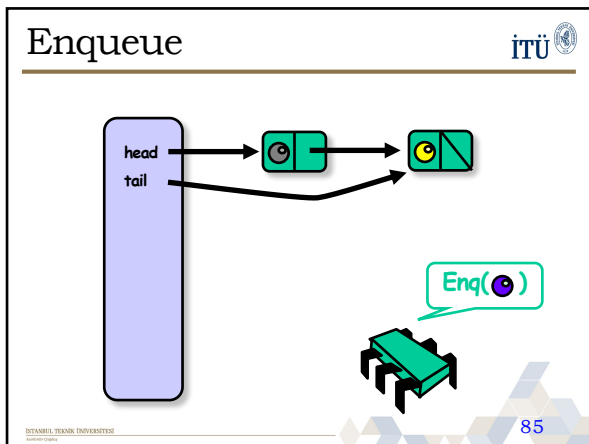
83

A Lock-Free Queue



İSTANBUL TEKNİK ÜNİVERSİTESİ

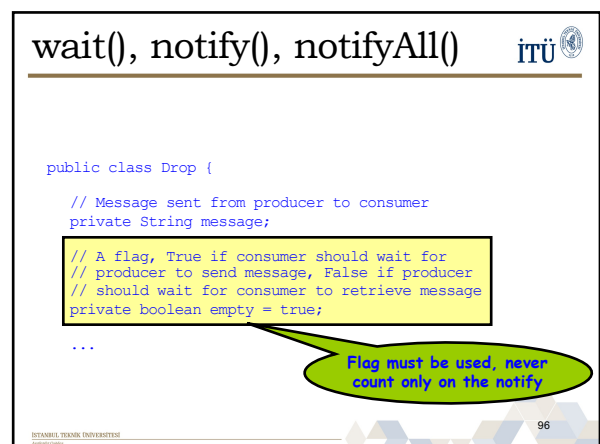
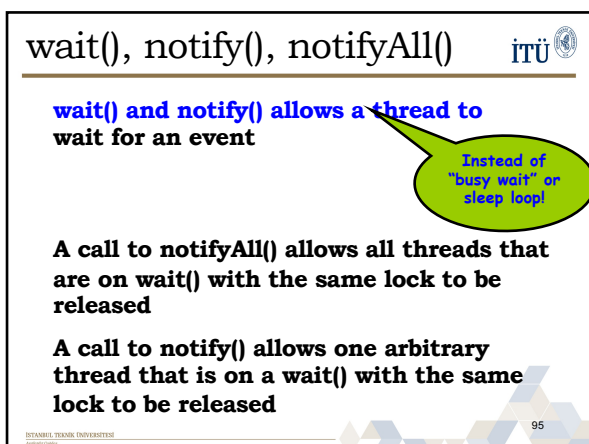
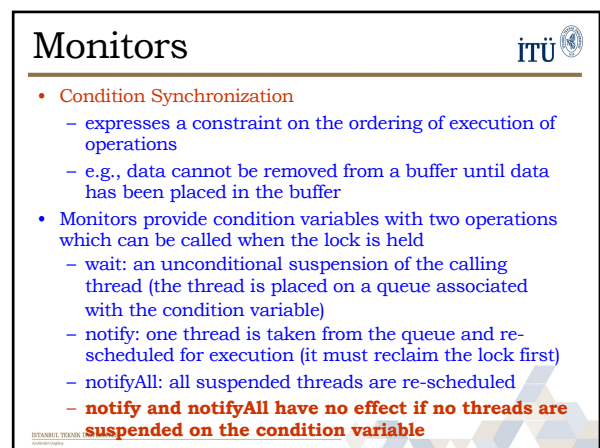
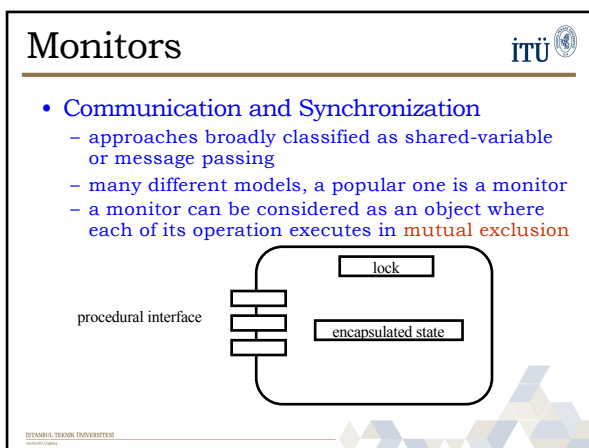

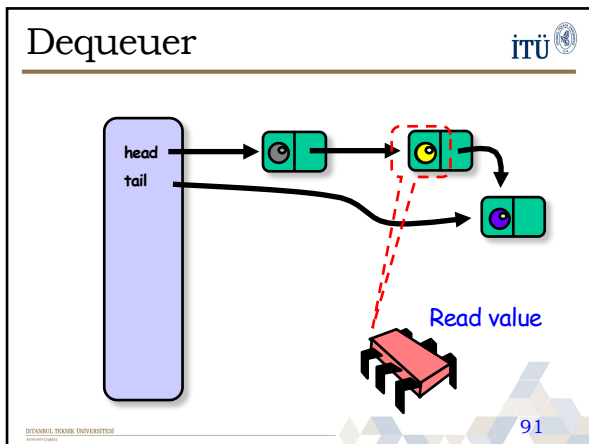
84



- ## Enqueue
- These two steps are not atomic
 - The tail field refers to either
 - Actual last Node (good)
 - Penultimate Node (not so good)
 - Be prepared!
- 88

- ## Enqueue
- What do you do if you find
 - A trailing tail?
 - Stop and help fix it
 - If tail node has non-null next field
 - CAS the queue's tail field to tail.next
- 89

- ## When CASs Fail
- During logical enqueue
 - Abandon hope, restart
 - Still lock-free
 - During physical enqueue
 - Ignore it
- 90



wait(), notify(), notifyAll()



```
public class Drop {  
    ...  
    public synchronized String take() {  
        // Wait until message is available  
        while (empty) {  
            // we do nothing on InterruptedException  
            // since the while condition is checked anyhow  
            try { wait(); } catch (InterruptedException e) {}  
        }  
        // Toggle status and notify on the status change  
        empty = true;  
        notifyAll();  
        return message;  
    }  
    ...  
}
```

ISTANBUL TEKNİK ÜNİVERSİTESİ

97

wait(), notify(), notifyAll()



```
public class Drop {  
    ...  
    public synchronized void put(String message) {  
        // Wait until message has been retrieved  
        while (!empty) {  
            // we do nothing on InterruptedException  
            // since the while condition is checked anyhow  
            try { wait(); } catch (InterruptedException e) {}  
        }  
        // Toggle status, store message and notify consumer  
        empty = false;  
        this.message = message;  
        notifyAll();  
    }  
    ...  
}
```

ISTANBUL TEKNİK ÜNİVERSİTESİ

98

The new wait() and notify()



- Remember **Lock.newCondition();** ?
- Condition factors out the Object monitor methods (**wait**, **notify** and **notifyAll**)
 - into distinct objects to give the effect of having multiple wait-sets per object
 - can combine them with the use of arbitrary **Lock** implementations.
- Where a **Lock** replaces the use of **synchronized** methods and statements, a **Condition** replaces the use of the Object monitor methods.

ISTANBUL TEKNİK ÜNİVERSİTESİ

Conditions



- Allows more than one wait condition per object
 - Even for built-in locks, via Locks utility class
- Allows much simpler implementation of

```
interface Condition {  
    void await() throws IE;  
    void awaitUninterruptibly();  
    long awaitNanos(long nanos) throws IE;  
    boolean awaitUntil(Date deadline) throws IE;  
    void signal();  
    void signalAll();  
}
```

ISTANBUL TEKNİK ÜNİVERSİTESİ

Concurrent Stack




- Methods
 - push(x)
 - pop()
- Last-in, First-out (LIFO) order
- Lock-Free!

ISTANBUL TEKNİK ÜNİVERSİTESİ

101


Empty Stack



ISTANBUL TEKNİK ÜNİVERSİTESİ

102


Push



İSTANBUL TEKNİK ÜNİVERSİTESİ

103


Push



İSTANBUL TEKNİK ÜNİVERSİTESİ

104


Push



İSTANBUL TEKNİK ÜNİVERSİTESİ

105


Push



İSTANBUL TEKNİK ÜNİVERSİTESİ

106


Push



İSTANBUL TEKNİK ÜNİVERSİTESİ

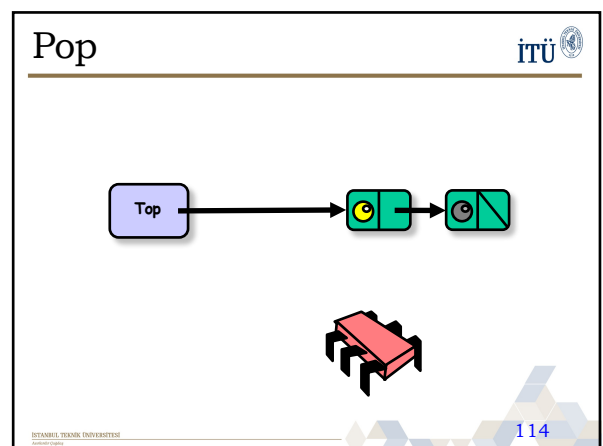
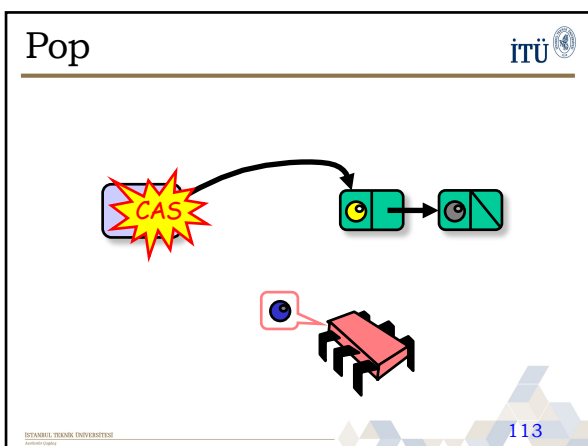
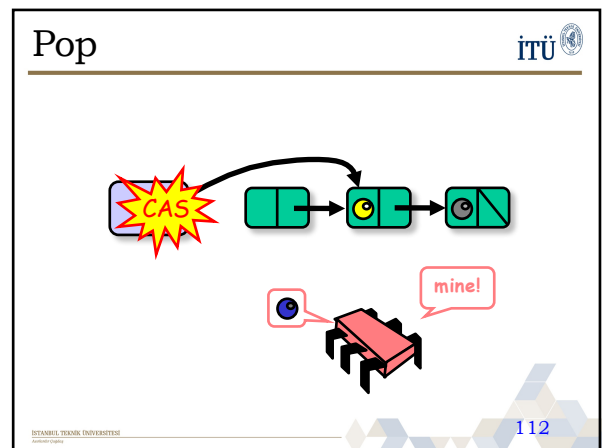
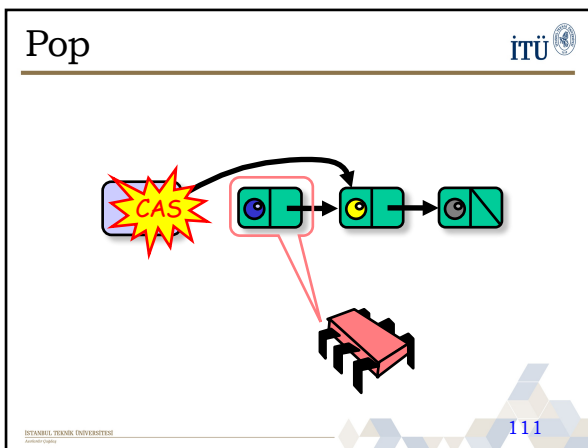
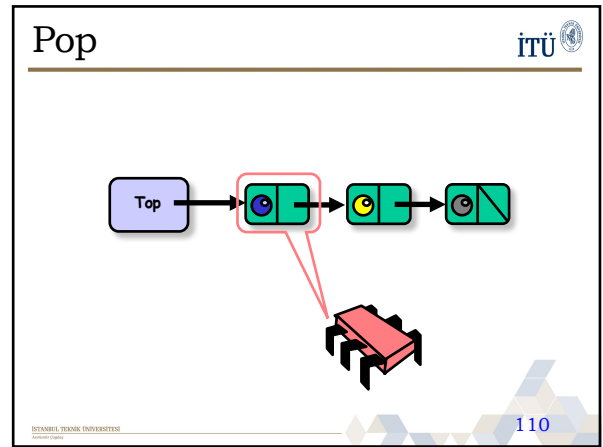
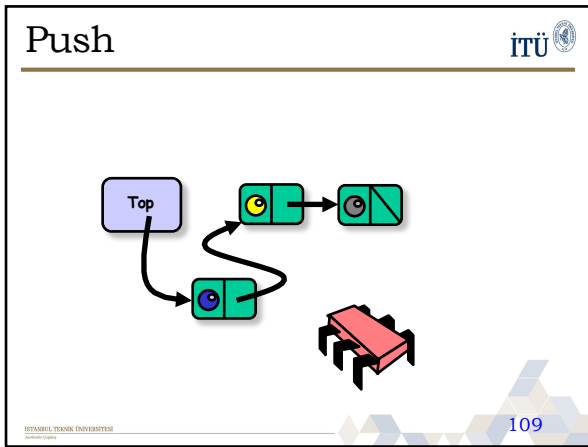
107

Push



İSTANBUL TEKNİK ÜNİVERSİTESİ

108



Lock-free Stack



```
public class LockFreeStack {
    private AtomicReference top = new AtomicReference(null);

    public boolean tryPush(Node node){
        Node oldTop = top.get();
        node.next = oldTop;
        return(top.compareAndSet(oldTop, node))
    }

    public void push(T value) {
        Node node = new Node(value);
        while (true) {
            if (tryPush(node)) {
                return;
            }
            else backoff.backoff();
        }
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

115

Lock-free Stack



- Good
 - No locking
- Bad
 - Without GC, fear ABA
 - Without backoff, huge contention at top
 - In any case, no parallelism

İSTANBUL TEKNİK ÜNİVERSİTESİ

116

Big Question



- Are stacks inherently sequential?
- Reasons why
 - Every pop() call fights for top item
- Reasons why not
 - Stay tuned ...

İSTANBUL TEKNİK ÜNİVERSİTESİ

117

Elimination-Backoff Stack

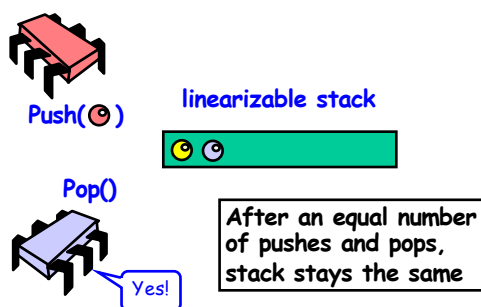


- How to
 - “turn contention into parallelism”
- Replace familiar
 - exponential backoff
- With alternative
 - elimination-backoff

İSTANBUL TEKNİK ÜNİVERSİTESİ

118


Observation



İSTANBUL TEKNİK ÜNİVERSİTESİ

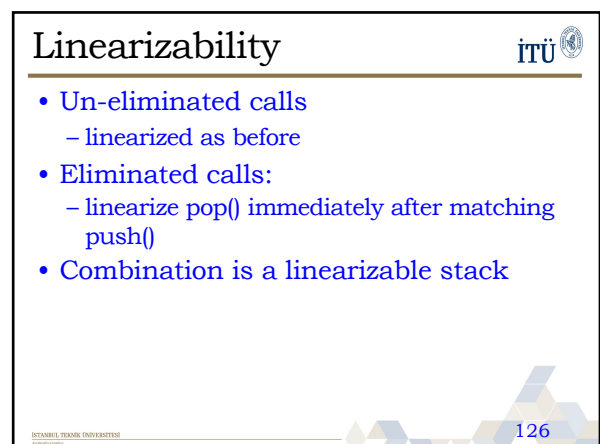
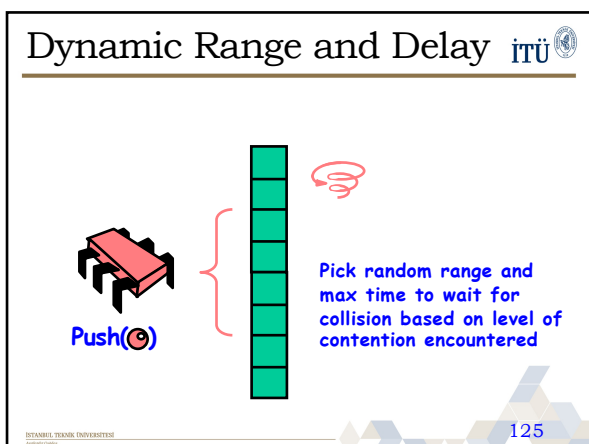
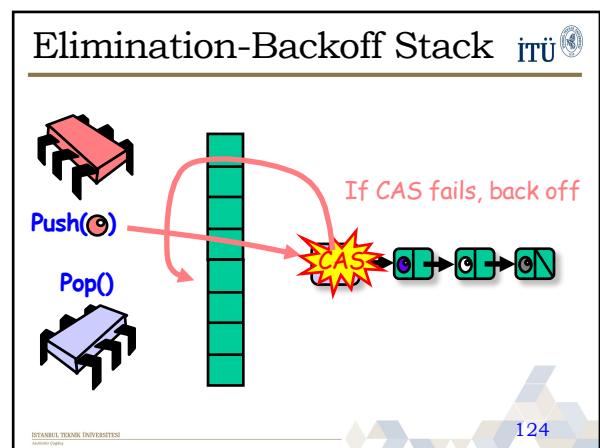
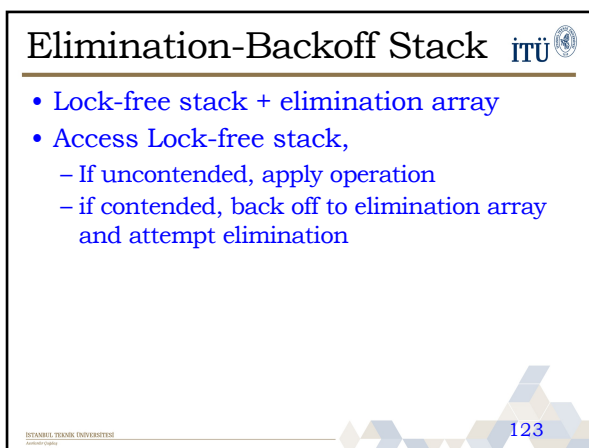
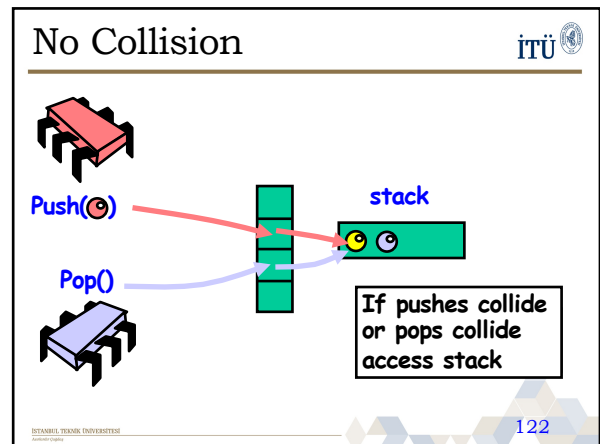
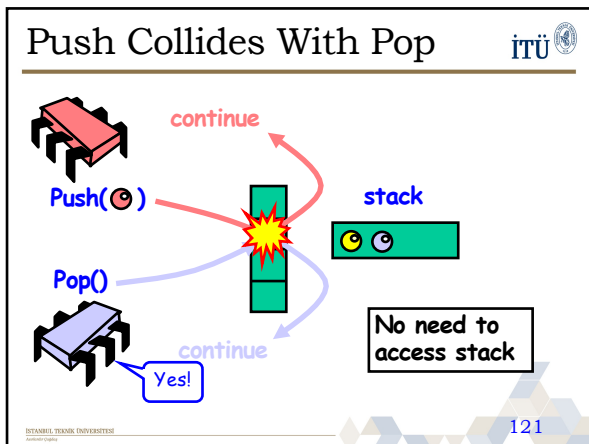
119

Idea: Elimination Array



İSTANBUL TEKNİK ÜNİVERSİTESİ

120



Backoff Has Dual Effect



- Elimination introduces parallelism
- Backoff onto array cuts contention on lock-free stack
- Elimination in array cuts down total number of threads ever accessing lock-free stack

İSTANBUL TEKNİK ÜNİVERSİTESİ

127

Elimination Array



```
public class EliminationArray {
    private static final int duration = ...;
    private static final int timeUnit = ...;
    Exchanger<T>[] exchanger;
    public EliminationArray(int capacity) {
        exchanger = new Exchanger[capacity];
        for (int i = 0; i < capacity; i++)
            exchanger[i] = new Exchanger<T>();
        ...
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

128

A Lock-Free Exchanger



```
public class Exchanger<T> {
    AtomicStampedReference<T> slot
    = new AtomicStampedReference<T>(null, 0);
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

129

Extracting Reference & Stamp



```
public T get(int[] stampHolder);
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

130

Exchanger Status



```
enum Status {EMPTY, WAITING, BUSY};
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

131

The Exchange



```
public T Exchange(T myItem, long nanos)
    throws TimeoutException {
    long timeBound = System.nanoTime() + nanos;
    int[] stampHolder = {EMPTY};
    while (true) {
        if (System.nanoTime() > timeBound)
            throw new TimeoutException();
        T herItem = slot.get(stampHolder);
        int stamp = stampHolder[0];
        switch(stamp) {
            case EMPTY: ... // slot is free
            case WAITING: ... // someone waiting for me
            case BUSY: ... // others exchanging
        }
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

132

Lock-free Exchanger

İTÜ

133

Lock-free Exchanger

İTÜ

In search of partner ...

134

Lock-free Exchanger

İTÜ

Still waiting ...

Try to exchange item and set state to BUSY

135

Lock-free Exchanger

İTÜ

Partner showed up, take item and reset to EMPTY

item stamp/state

136

Lock-free Exchanger

İTÜ

Partner showed up, take item and reset to EMPTY

item stamp/state

137

Exchanger State EMPTY

İTÜ

```

case EMPTY: // slot is free
if (slot.compareAndSet(herItem, myItem, EMPTY, WAITING)) {
while (System.nanoTime() < timeBound){
herItem = slot.get(stampHolder);
if (stampHolder[0] == BUSY) {
slot.set(null, EMPTY);
return herItem;
}
}
if (slot.compareAndSet(myItem, null, WAITING, EMPTY)){
throw new TimeoutException();
}
else {
herItem = slot.get(stampHolder);
slot.set(null, EMPTY);
return herItem;
}
}
break;

```

138

States WAITING and BUSY

```
case WAITING: // someone waiting for me
    if (slot.compareAndSet(herItem, myItem, WAITING, BUSY))
        return herItem;
break;
case BUSY: // others in middle of exchanging
break;
default: // impossible
break;
}}}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

139

The Exchanger Slot

- Exchanger is lock-free
- Because the only way an exchange can fail is if others repeatedly succeeded or no-one showed up
- The slot we need does not require symmetric exchange

İSTANBUL TEKNİK ÜNİVERSİTESİ

140

Elimination Array

```
public class EliminationArray {
    ...
    public T visit(T value, int Range) throws TimeoutException {
        int slot = random.nextInt(Range);
        int nanodur = convertToNanos(duration, TimeUnit);
        return (exchanger[slot].exchange(value, nanodur));
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

141

Elimination Stack Push

```
public void push(T value) {
    ...
    while (true) {
        if (tryPush(node)) {
            return;
        } else try {
            T otherValue = eliminationArray.visit(value, policy.Range);
            if (otherValue == null) {
                return;
            }
        }
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

142

Elimination Stack Pop

```
public T pop() {
    ...
    while (true) {
        if (tryPop()) {
            return returnNode.value;
        } else {
            try {
                T otherValue =
                    eliminationArray.visit(null, policy.Range);
                if (otherValue != null) {
                    return otherValue;
                }
            }
        }
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ

143

Exchangers

- A synchronization point at which two threads can exchange objects
- Can be thought of as a CyclicBarrier with a count of two plus allowing an exchange of state at the barrier
- On calling the exchange() method the thread provides some object as input
- The exchange() methods returns the object entered as input by the second thread to the calling thread
- Useful for example in the case where one thread is filling a buffer (filler thread) and the other emptying (emptying thread)
- On calling exchange() the filler thread is passed an empty buffer
- The emptying thread obtain the newly full buffer from the filler thread

İSTANBUL TEKNİK ÜNİVERSİTESİ

Concurrent Collections



The Concurrency framework provides implementation of several commonly used collections classes optimized for concurrent operations

```
public interface Queue<E> extends Collection<E>:
```

A collection class that hold elements prior to processing, generally orders element in a FIFO manner (however can also be LIFO etc).

Supports the following functionality

- Offer method inserts an element if possible else return false
- remove() and poll() return and remove the head of the queue
- element() and peek() return but do not remove the head of the queue

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Concurrent Collections



The Concurrency framework queue implementations

- CopyOnWriteArrayList: A thread-safe variant of ArrayList
- ConcurrentLinkedQueue : An unbounded thread-safe queue based on linked nodes.
- *BlockingQueue* : A Queue that additionally supports operations that wait for the queue to become non-empty when retrieving an element, and wait for space to become available in the queue when storing an element.

Summary of BlockingQueue methods

	Throws exception	Special value	Blocks	Times out
Insert	add(e)	offer(e)	put(e)	offer(e, time, unit)
Remove	remove()	poll()	take()	poll(time, unit)
Examine	element()	peek()	not applicable	not applicable

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Concurrent Collections



The Concurrency framework queue implementations

- *BlockingQueue*
 - ArrayBlockingQueue: A bounded blocking queue backed by an array.
 - DelayQueue: An element can only be taken when its delay has expired. Holds elements of *Delayed* instances.
 - LinkedBlockingDeque: An optionally-bounded blocking deque based on linked nodes.
 - LinkedTransferQueue: When transfer is used producing threads may wait until consumed.
 - PriorityBlockingQueue: An unbounded blocking queue that uses comparators for prioritizing.
 - SynchronousQueue: A blocking queue in which each insert operation must wait for a corresponding remove operation by another thread, and vice versa.

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Concurrent Collections



```
public interface ConcurrentMap<K,V> extends Map<K,V>:
```

An extension to the Map interface that provides support for concurrency

- Supports concurrent putIfAbsent, remove and replace methods
- putIfAbsent(key, value): If the specified key is not associated with a value associate it with a value. Performed atomically
- remove(key, value): Removes entry for a key if associated with value. Performed atomically
- remove(key, oldvalue, newvalue): Replaces entry for a key if associated with oldvalue with newvalue. Performed atomically

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Concurrent Collections



The Concurrency framework queue implementations

- *ConcurrentMap*
 - ConcurrentHashMap: A hash table supporting full concurrency of retrievals and high expected concurrency for updates.
 - ConcurrentNavigableMap: A ConcurrentMap supporting NavigableMap operations, and recursively so for its navigable sub-maps
 - ConcurrentSkipListMap: A scalable concurrent ConcurrentNavigableMap implementation.

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Concurrent Collections



```
public class ConsumerThread implements Runnable{
    private static int capacity ;
    private BlockingQueue<Integer> intqueue;
    public consumerThread(BlockingQueue<Integer> queue, int cap){
        capacity = cap;
        this.intqueue = queue;
    }
    public void run(){
        int num;
        for(int i = 0; i<capacity; i++){
            try {
                num = intqueue.take();
                if(num == -1) break;
                System.out.println("Square of " + num + " : " + num*num);
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}
```

İSTANBUL TEKNİK ÜNİVERSİTESİ
İTÜ

Concurrent Collections



```
public class ProducerThread implements Runnable {
    private static int capacity;
    private BlockingQueue<Integer> intqueue;
    public producerThread(BlockingQueue<Integer> queue, int cap) {
        capacity = cap;
        this.intqueue = queue;
    }
    public void run() {
        for(int i = 0; i < capacity-1; i++) {
            try {
                intqueue.put(i);
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
        try {
            intqueue.put(-1);
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
    }
}
```

Concurrent Collections



```
public class Tester {
    public static void main(String [] args) {
        BlockingQueue<Integer> queue = new
        ArrayBlockingQueue<Integer>(100);

        ConsumerThread consumer = new ConsumerThread(queue, 100);

        ProducerThread producer = new ProducerThread(queue, 100);

        new Thread(consumer).start();

        new Thread(producer).start();
    }
}
```

Concurrent Collections



That was an example of using blocking queues to implement producer consumer relationships

- The producer class fills in a queue of integers
- The consumer class pulls integers of this queue and finds the square
- If the queue is empty a take() blocks, if full a put() blocks
- We could extend this to use a thread pool for the producers

Summary



- We saw both lock-based and lock-free implementations of
- queues and stacks
- Don't be quick to declare a data structure inherently sequential
 - Linearizable stack is not inherently sequential
- ABA is a real problem, pay attention

This work is licensed under a [Creative Commons Attribution-ShareAlike 3.0 license](https://creativecommons.org/licenses/by-sa/3.0/).

- **You are free:**
 - to **Share** — to copy, distribute and transmit the work
 - to **Remix** — to adapt the work
- **Under the following conditions:**
 - **Attribution.** You must attribute the work to "The Art of Multiprocessor Programming" (but not in any way that suggests that the authors endorse you or your use of the work).
 - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
 - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.