İTÜ

Shared Counters and Parallelism

Companion slides for
The Art of Multiprocessor Programming
by Maurice Herlihy & Nir Shavit

A Shared Pool

İTÜ

```
public interface Pool {
  public void put(Object x);
  public Object remove();
}
```

Unordered set of objects

- Put
  - Inserts object
  - blocks if full
- Remove
  - Removes & returns an object
  - blocks if empty

Simple Locking Implementation

İTÜ

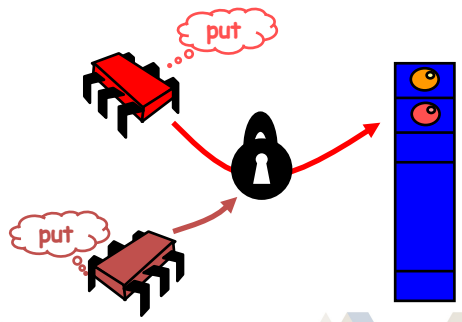

Simple Locking Implementation

İTÜ

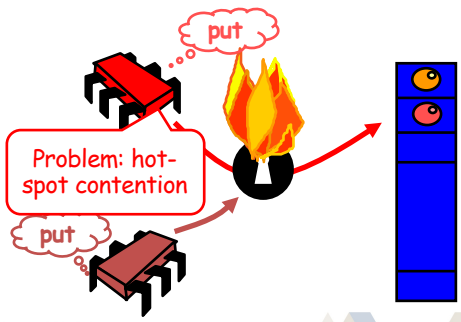

Simple Locking Implementation

İTÜ

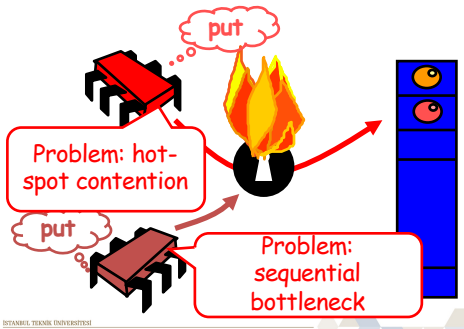

Simple Locking Implementation

İTÜ

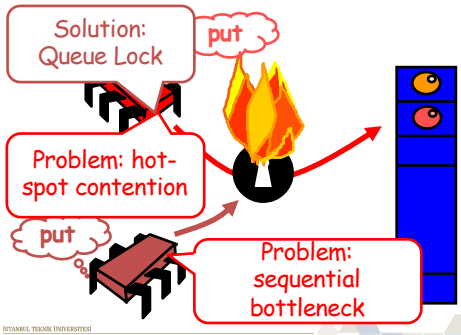

Simple Locking Implementation

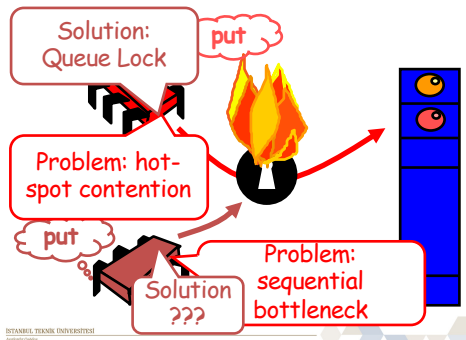

Shared Counter

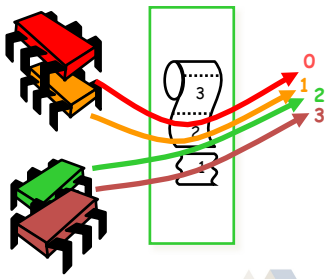

Shared Counter

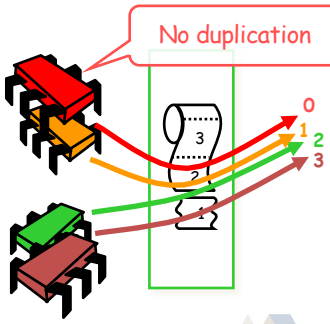

Shared Counter

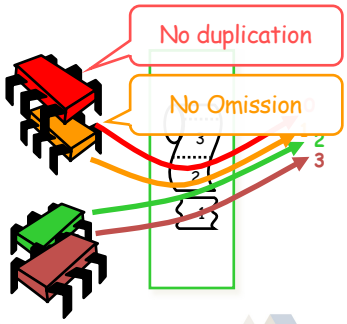

Shared Counter

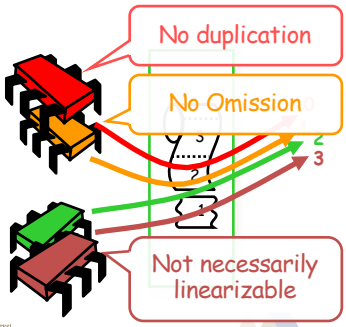

Shared Counters

- Can we build a shared counter with
  - Low memory contention, and
  - Real parallelism?
- Locking
  - Can use queue locks to reduce contention
  - No help with parallelism issue …

Software Combining Tree

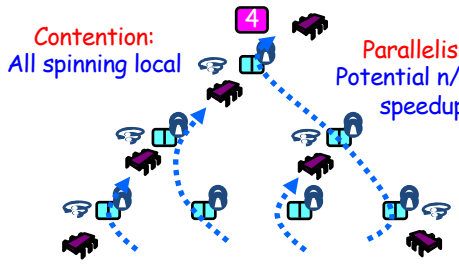

Combining Trees

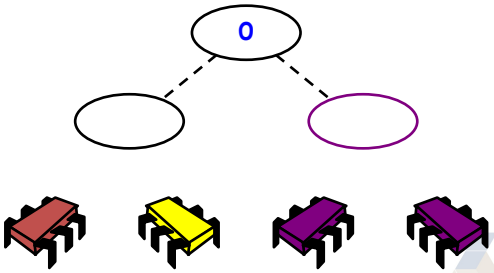

Combining Trees

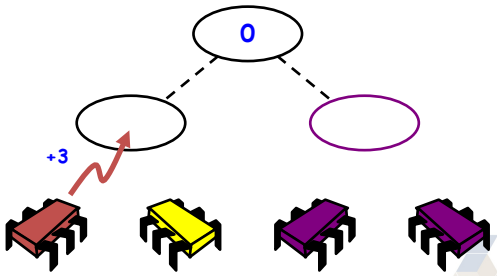

Combining Trees

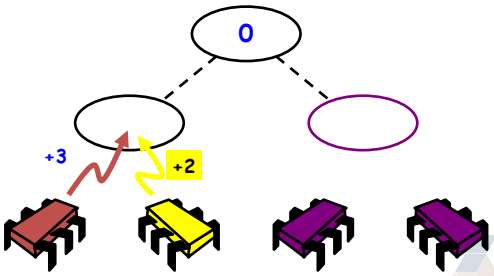

Combining Trees

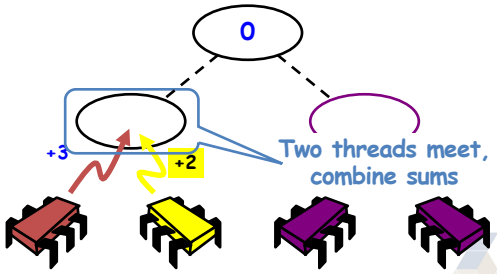

Combining Trees

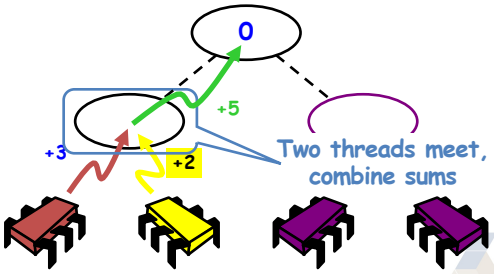

Combining Trees

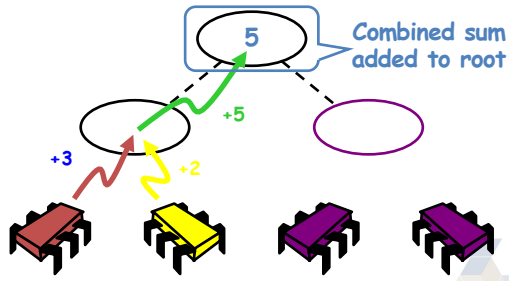

Combining Trees

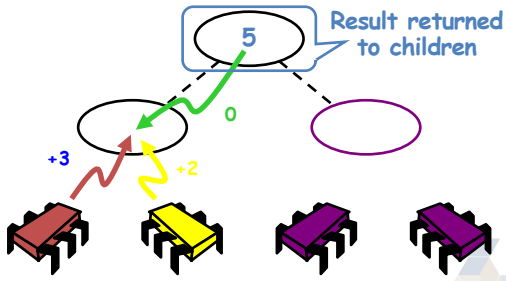

Combining Trees

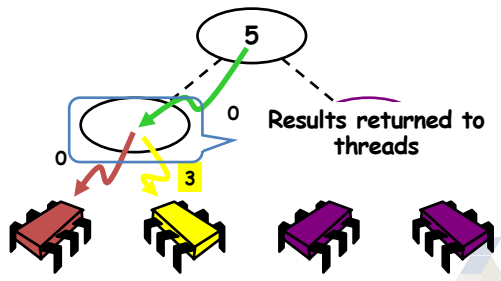

Devil in the Details

- What if
  - threads don't arrive at the same time?
- Wait for a partner to show up?
  - How long to wait?
  - Waiting times add up …
- Instead
  - Use multi-phase algorithm
  - Try to wait in parallel …

Combining Status

```
enum CStatus{
 IDLE, FIRST, SECOND, DONE, ROOT};
```


getAndIncrement

```
public int getAndIncrement() {
  Stack<Node> stack = new Stack<Node>();
  Node myLeaf = leaf[ThreadID.get()/2];

  // precombining phase
  // combining phase
  // operation phase
  // distribution phase
}
```

Node Synchronization

- Short-term
  - Synchronized methods
  - Consistency during method call
- Long-term
  - Boolean locked field
  - Consistency across calls

Phases

- Precombining
  - Set up combining rendez-vous
- Combining
  - Collect and combine operations
- Operation
  - Hand off to higher thread
- Distribution
  - Distribute results to waiting threads

Precombining Phase

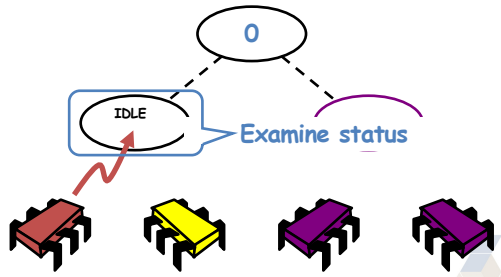

Precombining Phase

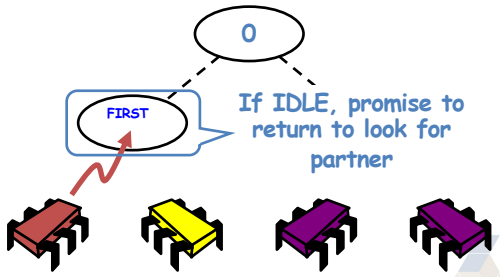

Precombining Phase

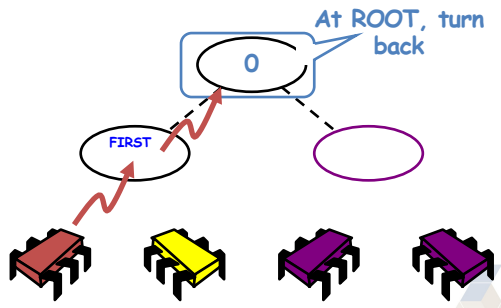

Precombining Phase

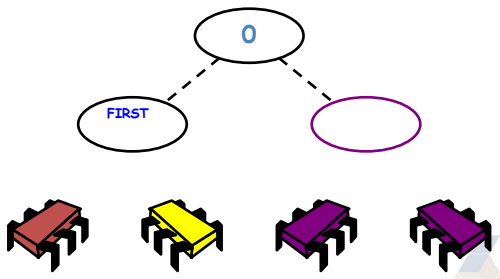

Precombining Phase

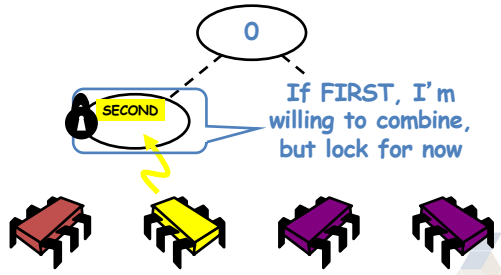

Code

- Tree class
  - In charge of navigation
- Node class
  - Combining state
  - Synchronization state
  - Bookkeeping

Precombining Navigation

```
Node node = myLeaf;
while (node.precombine()) {
  node = node.parent;
  }
Node stop = node;
```


Precombining Navigation

```
Node node = myLeaf;
while (node.precombine()) {
  node = node.parent;
  }
Node stop = node;
```

Start at leaf

Precombining Navigation

```
Node node = myLeaf;
while (node.precombine()) {
  node = node.parent;
  }
Node stop = node;
```

Move up while instructed to do so

Precombining Navigation

```
Node node = myLeaf;
while (node.precombine()) {
  node = node.parent;
  }
Node stop = node;
```

Remember where we stopped

Precombining Node

```
synchronized boolean precombine() {
 while (locked) wait();
 switch (cStatus) {
  case IDLE: cStatus = CStatus.FIRST;
             return true;
  case FIRST: locked = true;
              cStatus = CStatus.SECOND;
              return false;
  case ROOT: return false;
  default: throw new PanicException()
  }
}
```


Precombining Node

```
synchronized boolean precombine() {
 while (locked) wait();
 switch (cStatus) {
  case IDLE: cStatus = CStatus.FIRST;
             return true;
  case FIRST: locked = true;
              cStatus = CStatus.SECOND;
              return false;
  case ROOT: return false;
  default: throw new PanicException()
  }
}
```

Short-term synchronization

Synchronization

```
synchronized boolean precombine() {
 while (locked) wait();
 switch (cStatus) {
  case IDLE: cStatus = CStatus.FIRST;
             return true;
  case FIRST: locked = true;
              cStatus = CStatus.SECOND;
              return false;
  case ROOT: return false;
  default: throw new PanicException()
  }
}
```

Wait while node is locked

Precombining Node

```
synchronized boolean precombine() {
 while (locked) wait();
 switch (cStatus) {
  case IDLE: cStatus = CStatus.FIRST;
             return true;
  case FIRST: locked = true;
              cStatus = CStatus.SECOND;
              return false;
  case ROOT: return false;
  default: throw new PanicException()
  }
}
```

Check combining status

Node was IDLE

```
synchronized boolean precombine() {
 while (locked) {wait();}
 switch (cStatus) {
  case IDLE: cStatus = CStatus.FIRST;
             return true;
  case FIRST: locked = true;
              cStatus = CStatus.SECOND;
              return false;
  case ROOT: return false;
  default: throw new PanicException()
```

I will return to look for combining value

Precombining Node

```
synchronized boolean precombine() {
 while (locked) {wait();}
 switch (cStatus) {
  case IDLE: cStatus = CStatus.FIRST;
             return true;
  case FIRST: locked = true;
              cStatus = CStatus.SECOND;
              return false;
  case ROOT: return false;
  default: throw new PanicException()
  }
}
```

Continue up the tree

I'm the 2nd Thread

```
synchronized boolean precombine() {
 while (locked) {wait();}
 switch (cStatus) {
  case IDLE: cStatus = CStatus.FIRST;
             return true;
  case FIRST: locked = true;
              cStatus = CStatus.SECOND;
              return false;
  case ROOT: return false;
```

If 1st thread has promised to return, lock node so it won't leave without me

Precombining Node

```
synchronized boolean precombine() {
 while (locked) {wait();}
 switch (cStatus) {
  case IDLE: cStatus = CStatus.FIRST;
             return true;
  case FIRST: locked = true;
              cStatus = CStatus.SECOND;
              return false;
  case ROOT: return false;
  default: throw new
  }
}
```

Prepare to deposit 2nd value

Precombining Node

End of phase 1, don't continue up tree

```
) {
Y) {wait();}
  case IDLE: cStatus = CStatus.FIRST;
             return true;
  case FIRST: locked = true;
              cStatus = CStatus.SECOND;
              return false;
  case ROOT: return false;
  default: throw new PanicException()
  }
}
```


Node is the Root

If root, phase 1 ends, don't continue up tree

```
) {
Y) {wait();}
  case IDLE: cStatus = CStatus.FIRST;
             return true;
  case FIRST: locked = true;
              cStatus = CStatus.SECOND;
              return false;
  case ROOT: return false;
  default: throw new PanicException()
  }
}
```


Precombining Node

Always check for unexpected values!

```
synchronized boolean phase1() {
 while (locked) {w
 switch (cStatus)
  case IDLE: cStatus = CStatus.FIRST;
             return true;
  case FIRST: locked = true;
              cStatus = CStatus.SECOND;
              return false;
  case ROOT: return false;
  default: throw new PanicException()
  }
}
```


Combining Phase

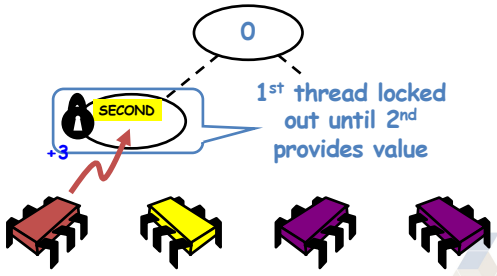

Combining Phase

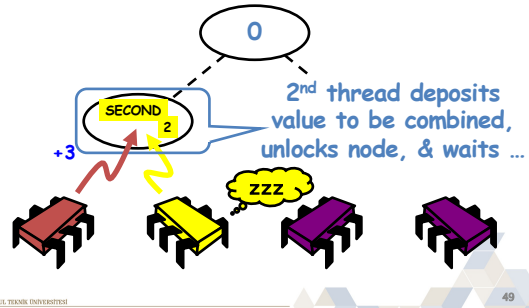

Combining Phase

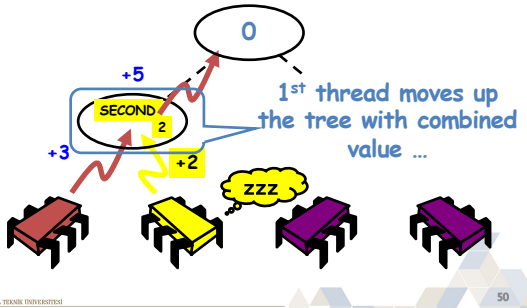

Combining (reloaded)

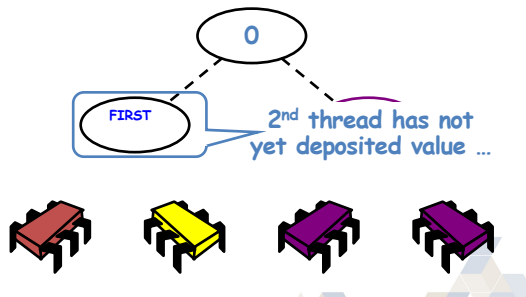

Combining (reloaded)

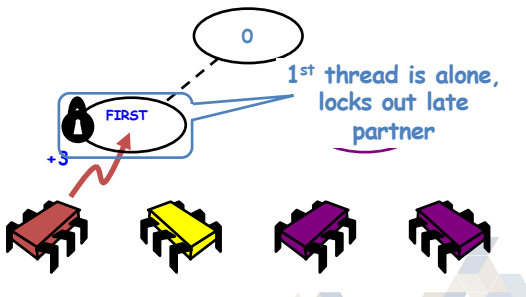

Combining (reloaded)

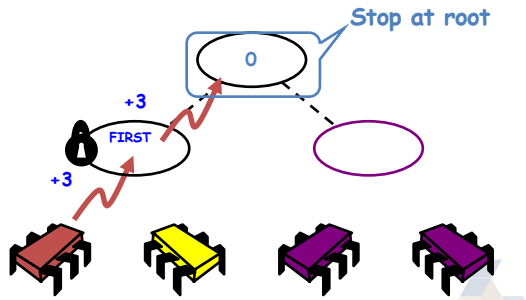

Combining (reloaded)

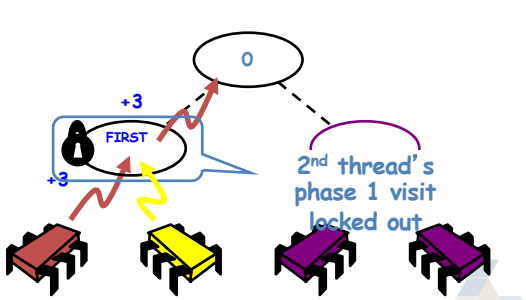

Combining Navigation

```
node = myLeaf;
int combined = 1;
while (node != stop) {
  combined = node.combine(combined);
  stack.push(node);
  node = node.parent;
  }
```


Combining Navigation

```
node = myLeaf;
int combined = 1;
while (node != stop) {
  combined = node.combine(combined);
  stack.push(node);
  node = node.parent;
  }
```

Start at leaf

Combining Navigation

```
node = myLeaf;
int combined = 1;
while (node != stop) {
  combined = node.combine(combined);
  stack.push(node);
  node = node.parent;
  }
```

Add 1

Combining Navigation

```
node = myLeaf;
int combined = 1;
while (node != stop) {
  combined = node.combine(combined);
  stack.push(node);
  node = node.parent;
  }
```

Revisit nodes visited in phase 1

Combining Navigation

```
node = myLeaf;
int combined = 1;
while (node != stop) {
  combined = node.combine(combined);
  stack.push(node);
  node = node.parent;
  }
```

Accumulate combined values, if any

Combining Navigation

```
node = myLeaf;
int combined = 1;
while (node != stop) {
  combined = node.combine(combined);
  stack.push(node);
  node = node.parent;
  }
```

We will retraverse path in reverse order …

Combining Navigation

```
node = myLeaf;
int combined = 1;
while (node != stop) {
  combined = node.combine(combined);
  stack.push(node);
  node = node.parent;
  }
```

Move up the tree

Combining Phase Node

```
synchronized int combine(int combined) {
  while (locked) wait();
  locked = true;
  firstValue = combined;
  switch (cStatus) {
    case FIRST:
      return firstValue;
    case SECOND:
      return firstValue + secondValue;
    default: …
    }
  }
```


Combining Phase Node

```
synchronized int combine(int combined) {
  while (locked) wait();
  locked = true;
  firstValue = combined;
  switch (cStatus) {
    case FIRST:
      return firstValue;
    case SECOND:
      return firstValue + secondValue;
    default: …
    }
  }
```

Wait until node is unlocked

Combining Phase Node

```
synchronized int combine(int combined) {
  while (locked) wait();
  locked = true;
  firstValue = combined;
  switch (cStatus) {
    case FIRST:
      return firstValue;
    case SECOND:
      return firstValue + secondValue;
    default: …
    }
  }
```

Lock out late attempts to combine

Combining Phase Node

```
synchronized int combine(int combined) {
  while (locked) wait();
  locked = true;
  firstValue = combined;
  switch (cStatus) {
    case FIRST:
      return firstValue;
    case SECOND:
      return firstValue + secondValue;
    default: …
    }
  }
```

Remember our contribution

Combining Phase Node

```
synchronized int combine(int combined) {
  while (locked) wait();
  locked = true;
  firstValue = combined;
  switch (cStatus) {
    case FIRST:
      return firstValue;
    case SECOND:
      return firstValue + secondValue;
    default: …
    }
  }
```

Check status

Combining Phase Node

```
synchronized int combine(int combined) {
  while (locked) wait();
  locked = true;
  firstValue = combined;
  switch (cStatus) {
    case FIRST:
      return firstValue;
    case SECOND:
      return firstValue + secondValue;
    default: …
    }
  }
```

1st thread is alone

Combining Node

```
synchronized int combine(int combined) {
  while (locked) wait();
  locked = true;
  firstValue = combined;
  switch (cStatus) {
    case FIRST:
      return firstValue;
    case SECOND:
      return firstValue + secondValue;
    default: …
    }
  }
```

Combine with 2nd thread

Operation Phase

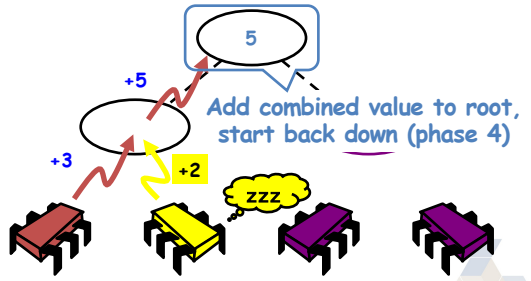

Operation Phase (reloaded)

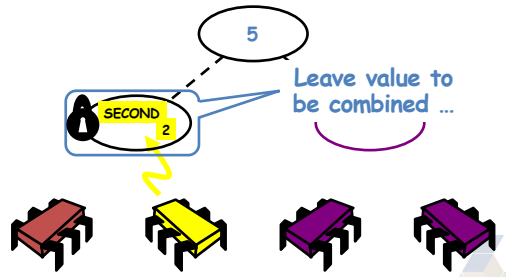

Operation Phase (reloaded)

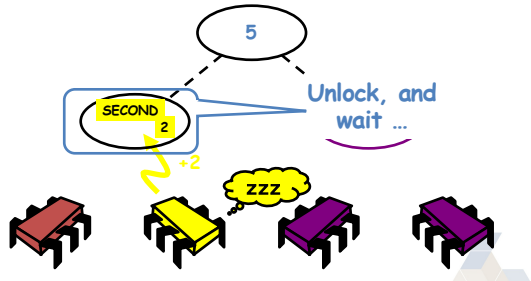

Operation Phase Navigation

```
prior = stop.op(combined);
```

Operation Phase Navigation

```
prior = stop.op(combined);
```

Get result of combining

Operation Phase Node

```
synchronized int op(int combined) {
  switch (cStatus) {
    case ROOT: int oldValue = result;
      result += combined;
      return oldValue;
    case SECOND: secondValue = combined;
      locked = false; notifyAll();
      while (cStatus != CStatus.DONE) wait();
      locked = false; notifyAll();
      cStatus = CStatus.IDLE;
      return result;
    default: …
```


At Root

```
synchronized int op(int combined) {
  switch (cStatus) {
    case ROOT: int oldValue = result;
      result += combined;
      return oldValue;
    case SECOND: secondValue = combined;
      locked = false; notifyAll();
      while (cStatus != CStatus.DONE) wait();
      locked = false; notifyAll();
      cStatus = CStatus.IDLE;
      return result;
    default: …
```

Add sum to root, return prior value

Intermediate Node

```
synchronized int op(int combined) {
  switch (cStatus) {
    case ROOT: int oldValue = result;
      result += combined;
      return oldValue;
    case SECOND: secondValue = combined;
      locked = false; notifyAll();
      while (cStatus != CStatus.DONE) wait();
      locked = false; notifyAll();
      cStatus = CStatus.IDLE;
      return result;
    default: …
```

Deposit value for later combining …

Intermediate Node

```
synchronized int op(int combined) {
  switch (cStatus) {
    case ROOT: int oldValue = result;
      result += combined;
      return oldValue;
    case SECOND: secondValue = combined;
      locked = false; notifyAll();
      while (cStatus != CStatus.DONE) wait();
      locked = false; notifyAll();
      cStatus = CStatus.IDLE;
      return result;
    default: …
```

Unlock node, notify 1st thread

Intermediate Node

```
synchronized int op(int combined) {
  switch (cStatus) {
    case ROOT: int oldValue
      result += combined;
      return oldValue;
    case SECOND: secondValue = combined;
      locked = false; notifyAll();
      while (cStatus != CStatus.DONE) wait();
      locked = false; notifyAll();
      cStatus = CStatus.IDLE;
      return result;
    default: …
```

Wait for 1st thread to deliver results

Intermediate Node

```
synchronized int op(int combined) {
  switch (cStatus) {
    case ROOT: int oldValue
      result += combined;
      return oldValue;
    case SECOND: secondValue = combined;
      locked = false; notifyAll();
      while (cStatus != CStatus.DONE) wait();
      locked = false; notifyAll();
      cStatus = CStatus.IDLE;
      return result;
    default: …
```

Unlock node & return

Distribution Phase

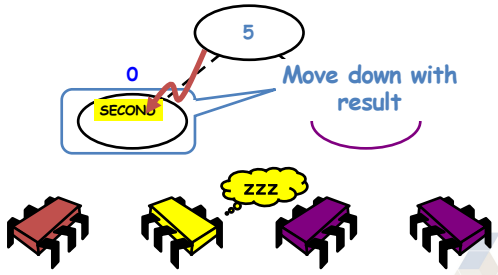

Distribution Phase

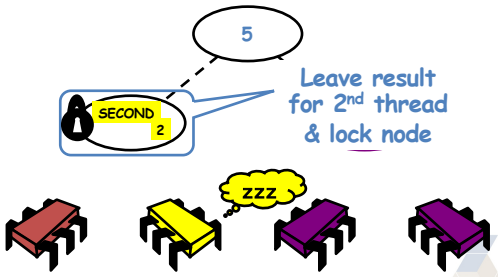

Distribution Phase

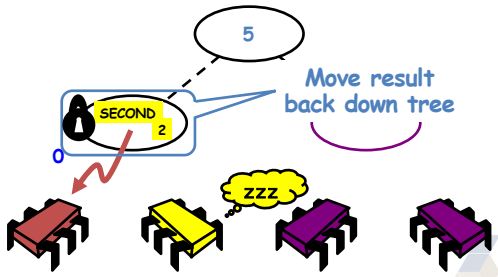

Distribution Phase

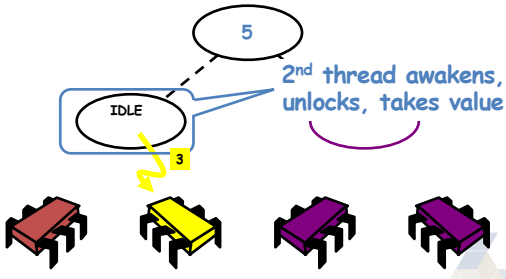

Distribution Phase Navigation

```
while (!stack.empty()) {
  node = stack.pop();
  node.distribute(prior);
  }
return prior;
```

Distribution Phase Navigation

```
while (!stack.empty()) {
  node = stack.pop();
  node.distribute(prior);
  }
return prior;
```

Traverse path in reverse order

Distribution Phase Navigation

```
while (!stack.empty()) {
  node = stack.pop();
  node.distribute(prior);
  }
return prior;
```

Distribute results to waiting 2nd threads

Distribution Phase Navigation

```
while (!stack.empty()) {
  node = stack.pop();
  node.distribute(prior);
  }
return prior;
```

Return result to caller

Distribution Phase

```
synchronized void distribute(int prior) {
   switch (cStatus) {
     case FIRST:
       cStatus = CStatus.IDLE;
       locked = false; notifyAll();
       return;
     case SECOND:
       result = prior + firstValue;
       cStatus = CStatus.DONE; notifyAll();
       return;
     default: …
```


Distribution Phase

```
synchronized void distribute(int prior) {
   switch (cStatus) {
     case FIRST:
       cStatus = CStatus.IDLE;
       locked = false; notifyAll();
       return;
     case SECOND:
       result = prior + firstValue;
       cStatus = CStatus.DONE; notifyAll();
       return;
     default: …
```

No combining, unlock node & reset

Distribution Phase

```
synchronized void distribute(int prior) {
   switch (cStatus) {
     case FIRST:
       cStatus = CStatu
       locked = false;
       return;
     case SECOND:
       result = prior + firstValue;
       cStatus = CStatus.DONE; notifyAll();
       return;
     default: …
```

Notify 2nd thread that result is available

Bad News: High Latency

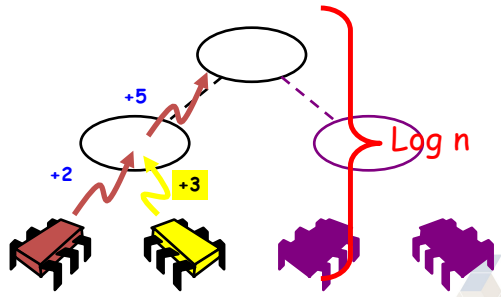

Good News: Real Parallelism

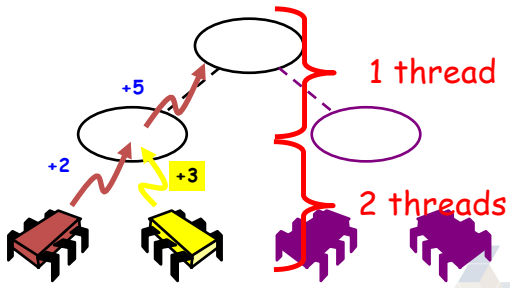

Throughput Puzzles

- Ideal circumstances
  - All n threads move together, combine
  - n increments in O(log n) time
- Worst circumstances
  - All n threads slightly skewed, locked out
  - n increments in O(n · log n) time

Index Distribution Benchmark

```
void indexBench(int iters, int work) {
 while (int i < iters) {
  i = r.getAndIncrement();
  Thread.sleep(random() % work);
 }}
```


Index Distribution Benchmark

```
void indexBench(int iters, int work) {
 while (int i < iters) {
  i = r.getAndIncrement();
  Thread.sleep(random() % work);
 }}
```

How many iterations

Index Distribution Benchmark

```
void indexBench(int iters, int work) {
 while (int i < iters) {
  i = r.getAndIncrement();
  Thread.sleep(random() % work);
 }}
```

Expected time between incrementing counter

Index Distribution Benchmark

```
void indexBench(int iters, int work) {
 while (int i < iters) {
  i = r.getAndIncrement();
  Thread.sleep(random() % work);
 }}
```

Take a number

Index Distribution Benchmark

```
void indexBench(int iters, int work) {
 while (int i < iters) {
  i = r.getAndIncrement();
  Thread.sleep(random() % work);
 }}
```

Pretend to work
(more work, less concurrency)

Performance Benchmarks

- Alewife
  - NUMA architecture
  - Simulated

MIT - ALEWIFE

- **Throughput:**
  - average number of **inc** operations in 1 million cycle period.
- **Latency:**
  - average number of simulator cycles per **inc** operation.

Performance

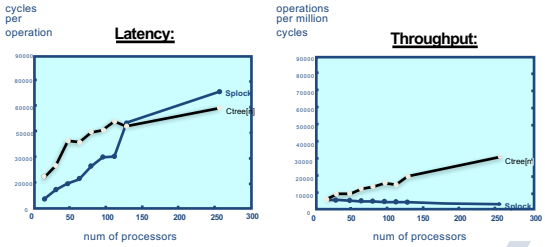

The Combining Paradigm

- Implements any RMW operation
- When tree is loaded
  - Takes 2 log n steps
  - for n requests
- Very sensitive to load fluctuations:
  - if the arrival rates drop
  - the combining rates drop
  - overall performance deteriorates!

Better to Wait Longer

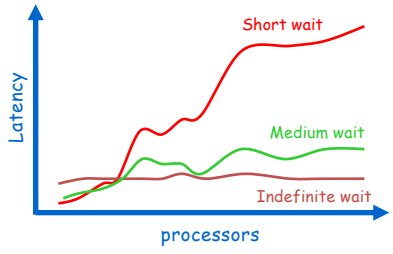

Conclusions

- Combining Trees
  - Work well under high contention
  - Sensitive to load fluctuations
  - Can be used for getAndMumble() ops
- Next
  - Counting networks
  - A different approach …

A Balancer

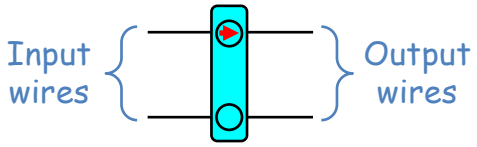

Tokens Traverse Balancers

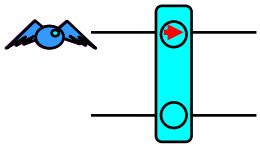

- Token i enters on any wire
- leaves on wire i mod (fan-out)

Tokens Traverse Balancers

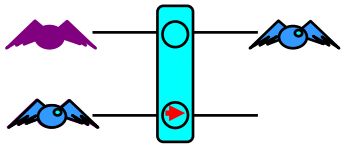

Tokens Traverse Balancers

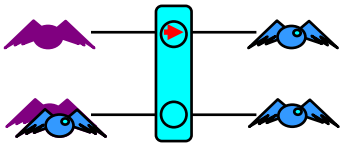

Tokens Traverse Balancers

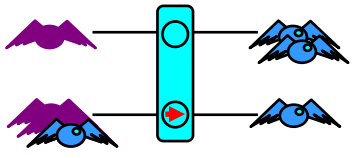

Tokens Traverse Balancers

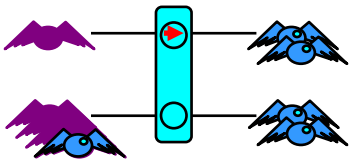

Tokens Traverse Balancers

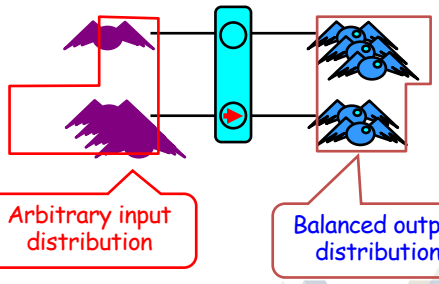

Smoothing Network

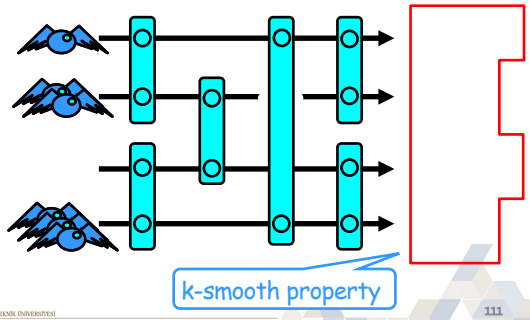

Counting Network

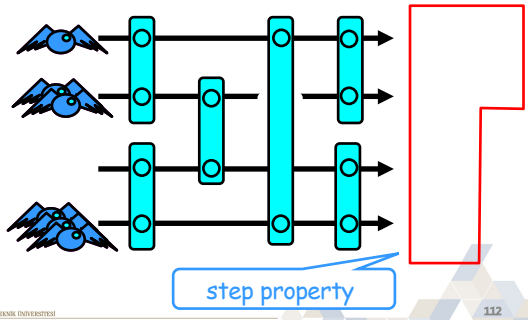

Counting Networks Count!

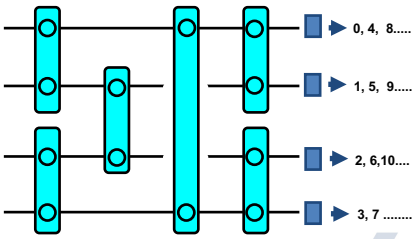

Bitonic[4]

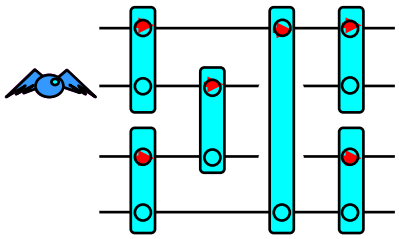

Bitonic[4]

İTÜ

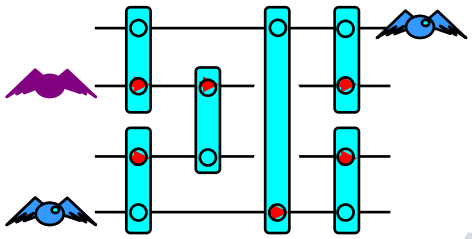

İSTANBUL TEKNİK ÜNİVERSİTESİ 115

Bitonic[4]

İTÜ

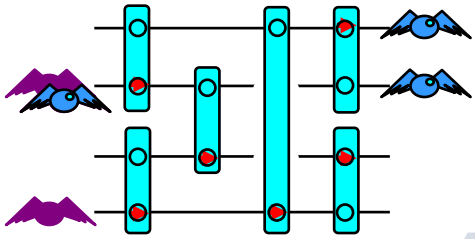

İSTANBUL TEKNİK ÜNİVERSİTESİ 116

Bitonic[4]

İTÜ

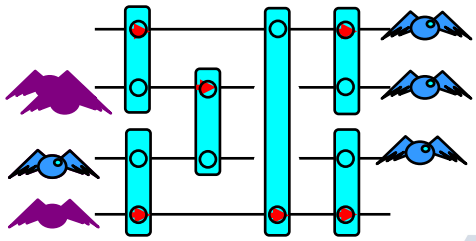

İSTANBUL TEKNİK ÜNİVERSİTESİ 117

Bitonic[4]

İTÜ

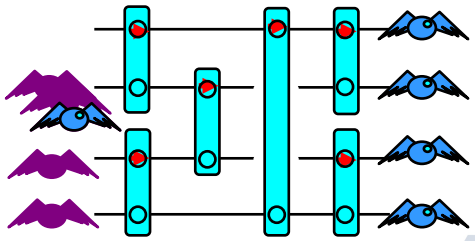

İSTANBUL TEKNİK ÜNİVERSİTESİ 118

Bitonic[4]

İTÜ

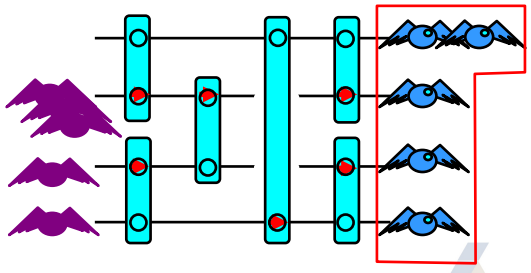

İSTANBUL TEKNİK ÜNİVERSİTESİ 119

Counting Networks

İTÜ

- Good for counting number of tokens
- low contention
- no sequential bottleneck
- high throughput
- practical networks depth

$\log^2 n$

İSTANBUL TEKNİK ÜNİVERSİTESİ 120

Bitonic[k] is not Linearizable

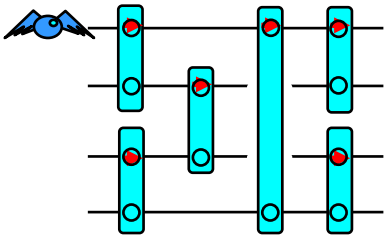

Bitonic[k] is not Linearizable

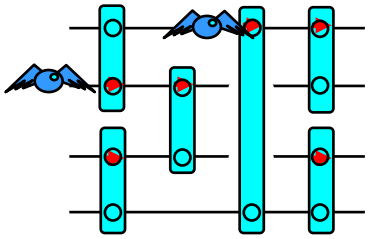

Bitonic[k] is not Linearizable

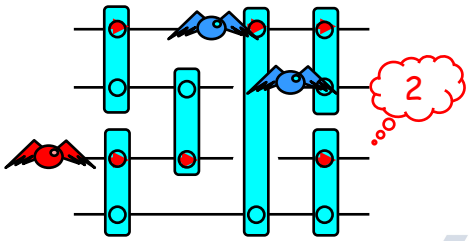

Bitonic[k] is not Linearizable

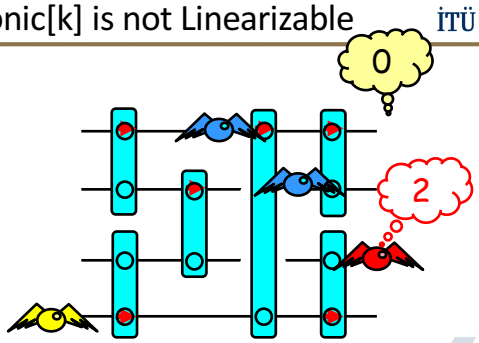

Bitonic[k] is not Linearizable

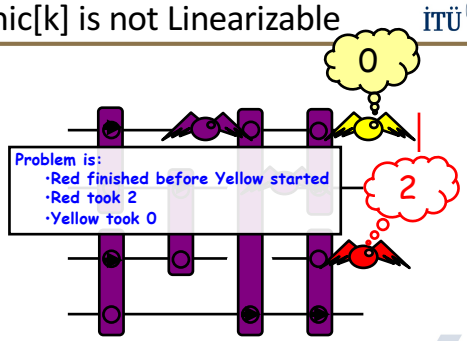

Shared Memory Implementation

```
class Balancer {
 boolean toggle;
 Balancer[] next;

synchronized boolean flip() {
 boolean oldValue = this.toggle;
 this.toggle = !this.toggle;
 return oldValue;
}
```

Shared Memory Implementation

```
Balancer traverse (Balancer b) {
 while(!b.isLeaf()) {
  if (toggle)
    b = b.next[0]
  else
    b = b.next[1]
  boolean toggle = b.flip();
  return b;
}
```


Bitonic[2k] Schematic

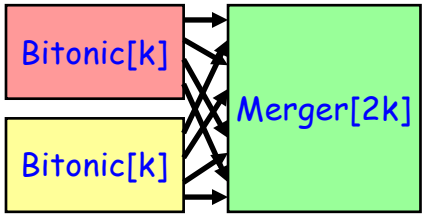

Bitonic[2k] Layout

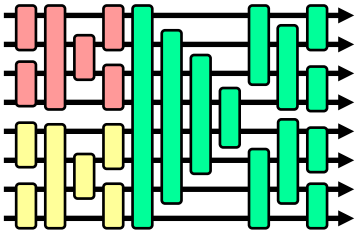

Unfolded Bitonic Network

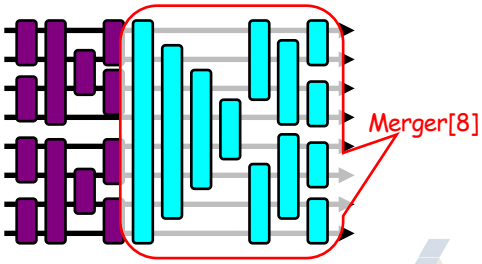

Unfolded Bitonic Network

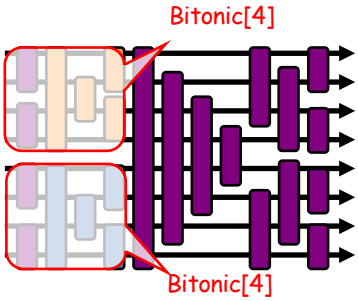

Unfolded Bitonic Network

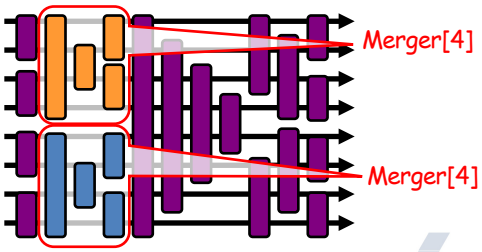

Unfolded Bitonic Network

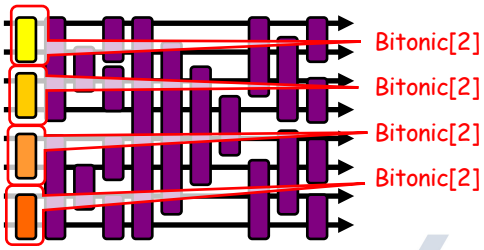

Bitonic[k] Depth

- Width k
- Depth is $(\log_2 k)(\log_2 k + 1)/2$

Network Depth

- Each block[k] has depth $\log_2 k$
- Need $\log_2 k$ blocks
- Grand total of $(\log_2 k)^2$

Index Distribution Benchmark

```
void indexBench(int iters, int work) {
 while (int i = 0 < iters) {
  i = fetch&inc();
  Thread.sleep(random() % work);
 }
}
```


Performance (Simulated)

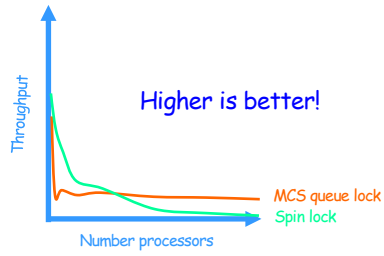

* All graphs taken from Herlihy,Lim,Shavit, copyright ACM.

Performance (Simulated)

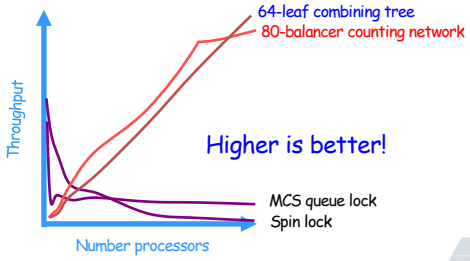

* All graphs taken from Herlihy,Lim,Shavit, copyright ACM.

Performance (Simulated)

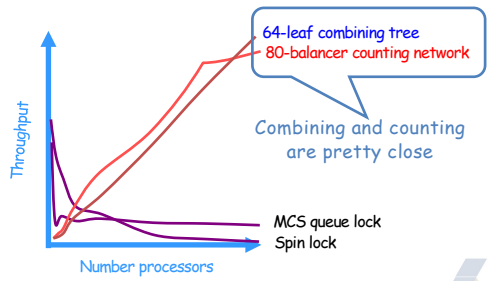

* All graphs taken from Herlihy,Lim,Shavit, copyright ACM.

Performance (Simulated)

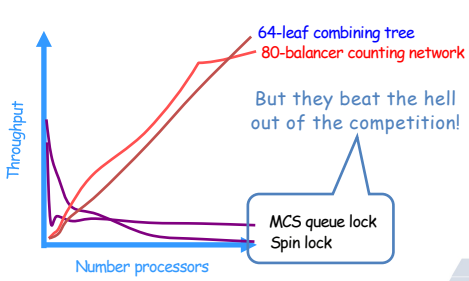

* All graphs taken from Herlihy,Lim,Shavit, copyright ACM.

Saturation and Performance

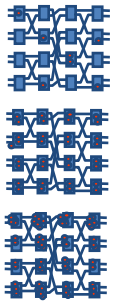

Undersaturated P < w log w

Optimal performance

Saturated P = w log w

Oversaturated P > w log w

Throughput vs. Size

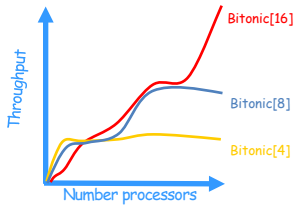

What About

- Decrements
- Adding arbitrary values

Anti-Tokens

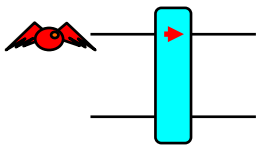

Tokens & Anti-Tokens Cancel

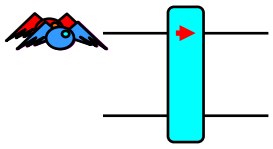

Tokens & Anti-Tokens Cancel

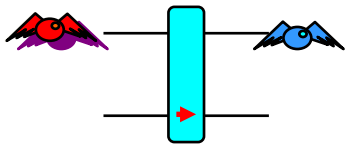

Tokens & Anti-Tokens Cancel

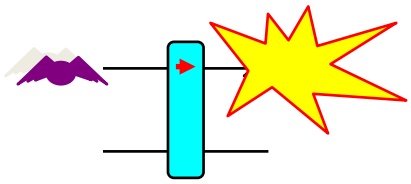

Tokens & Anti-Tokens Cancel

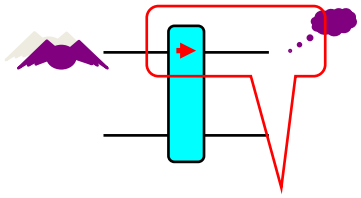

Tokens vs Antitokens

- Tokens
  - read balancer
  - flip
  - proceed
- Antitokens
  - flip balancer
  - read
  - proceed

Pumping Lemma

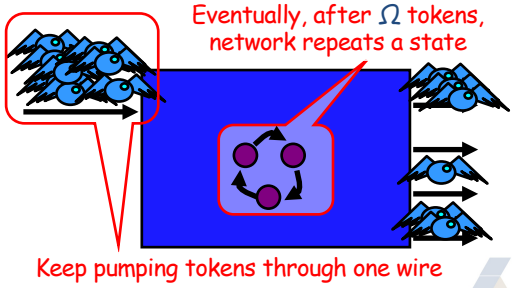

Anti-Token Effect